

# 2INC0 - Operating Systems

## Introduction to Operating Systems

Dr. Geoffrey Nelissen



Interconnected  
Resource-aware  
Intelligent Systems

**TU/e**

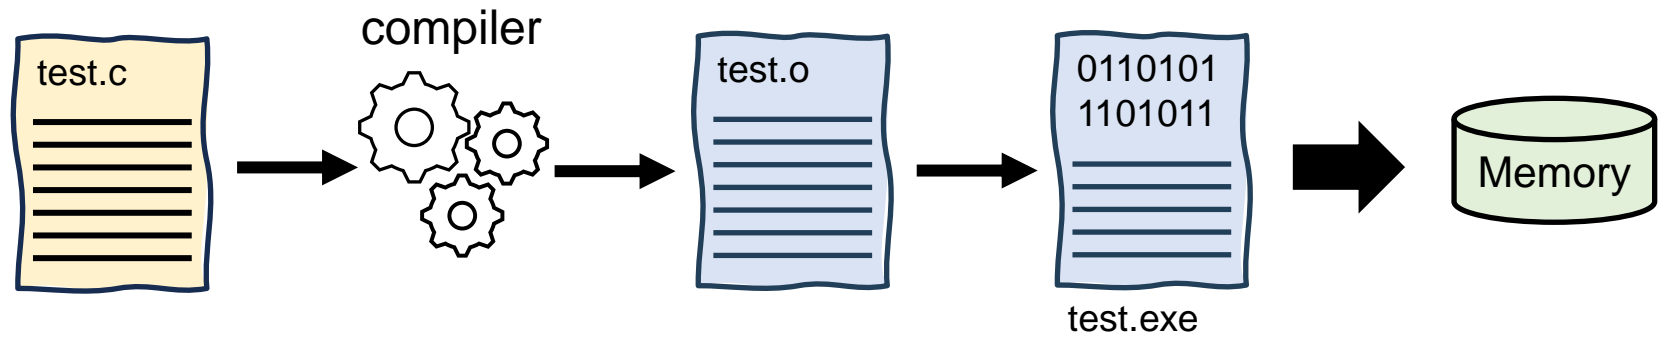
Technische Universiteit  
**Eindhoven**  
University of Technology

Where innovation starts

- **Introduction to operating systems** (lecture 1)
- **Processes, threads and scheduling** (lecture 2)
- **Concurrency and synchronization**
  - atomicity and interference (lecture 3)
  - action synchronization (lecture 4)
  - condition synchronization (lecture 5)
  - deadlock (lecture 6)
- **Memory management** (lecture 7 ad 8)
- **Input/output** (lectures 9 and 10)
  - general issues
  - file systems

- Course overview
- **OS: place in the computer system**
- Motivation and OS tasks
- Extra-functional requirements

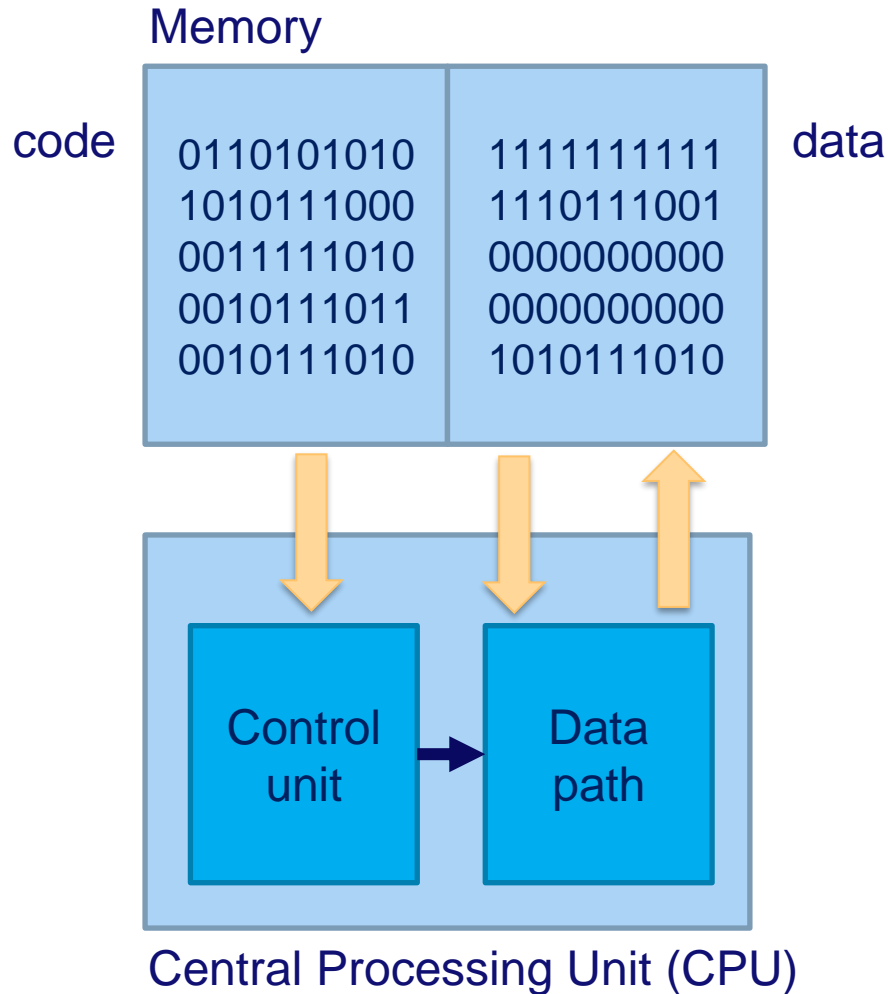
# What are the steps to execute a program?




Executable file is **platform dependent** (ISA, cpu arch, ...)

# How is code executed by your computer?


# How is code executed by your computer?



# How is code executed by your computer?




# How is code executed by your computer?






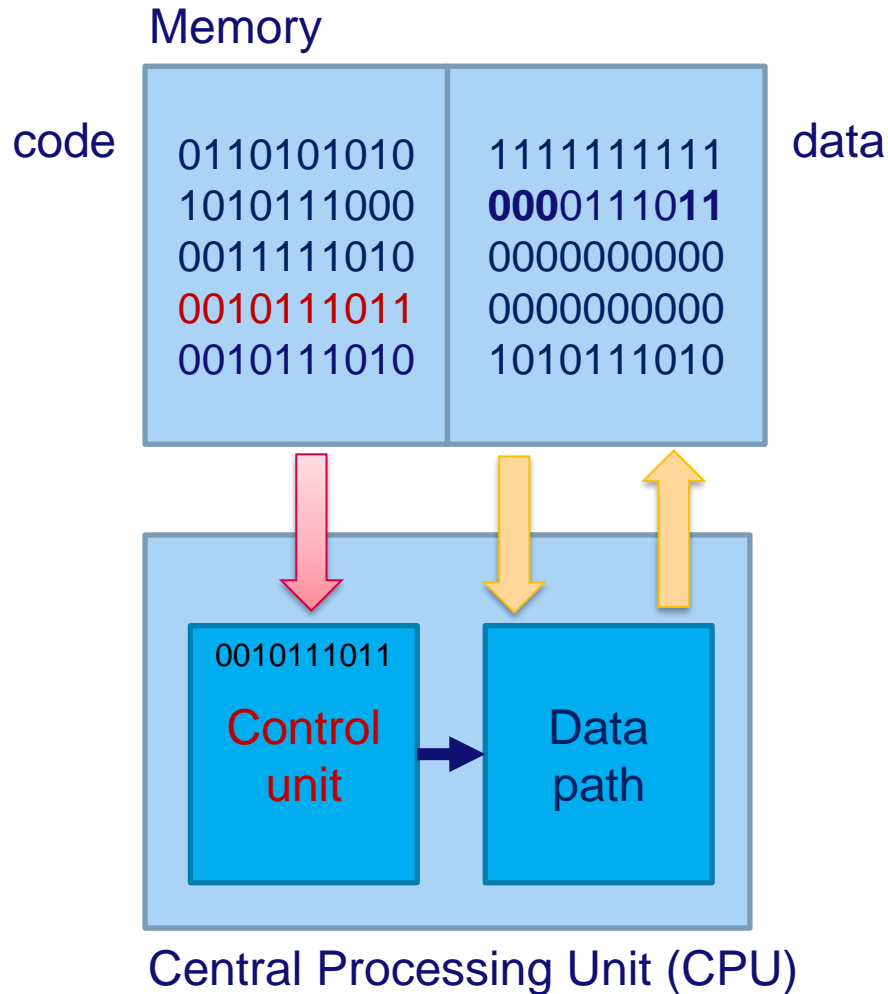
# How is code executed by your computer?



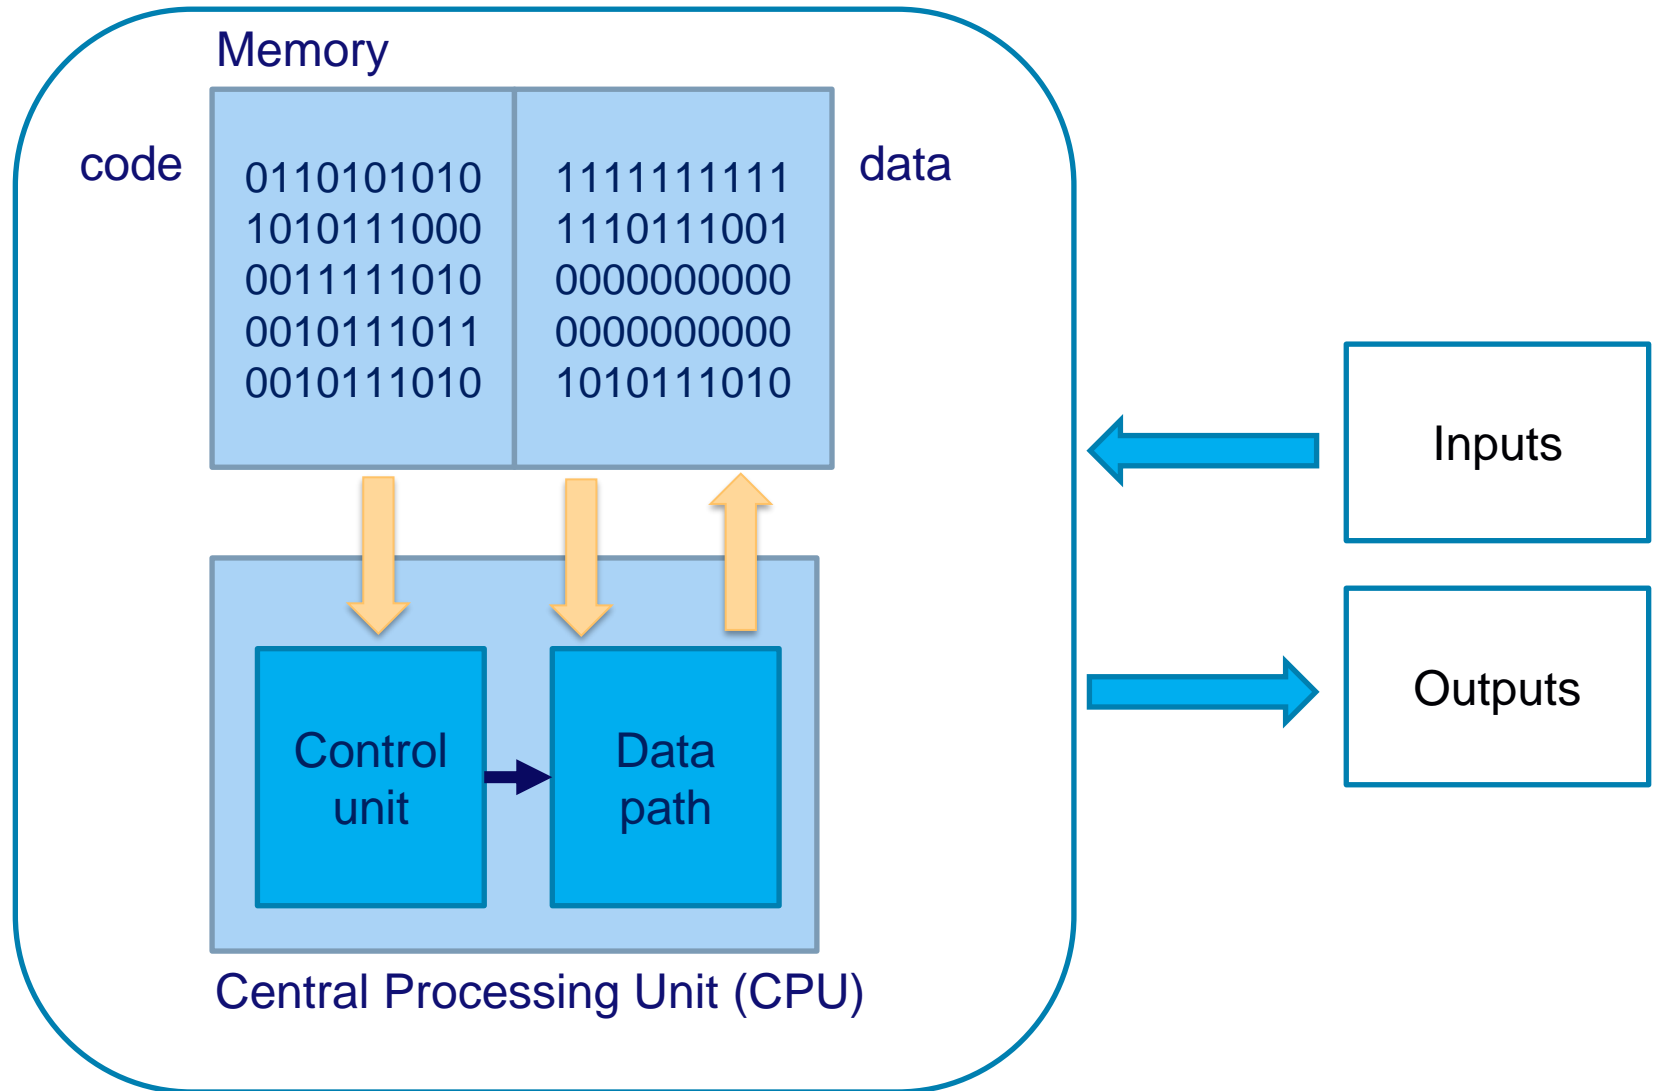
# How is code executed by your computer?



# How is code executed by your computer?

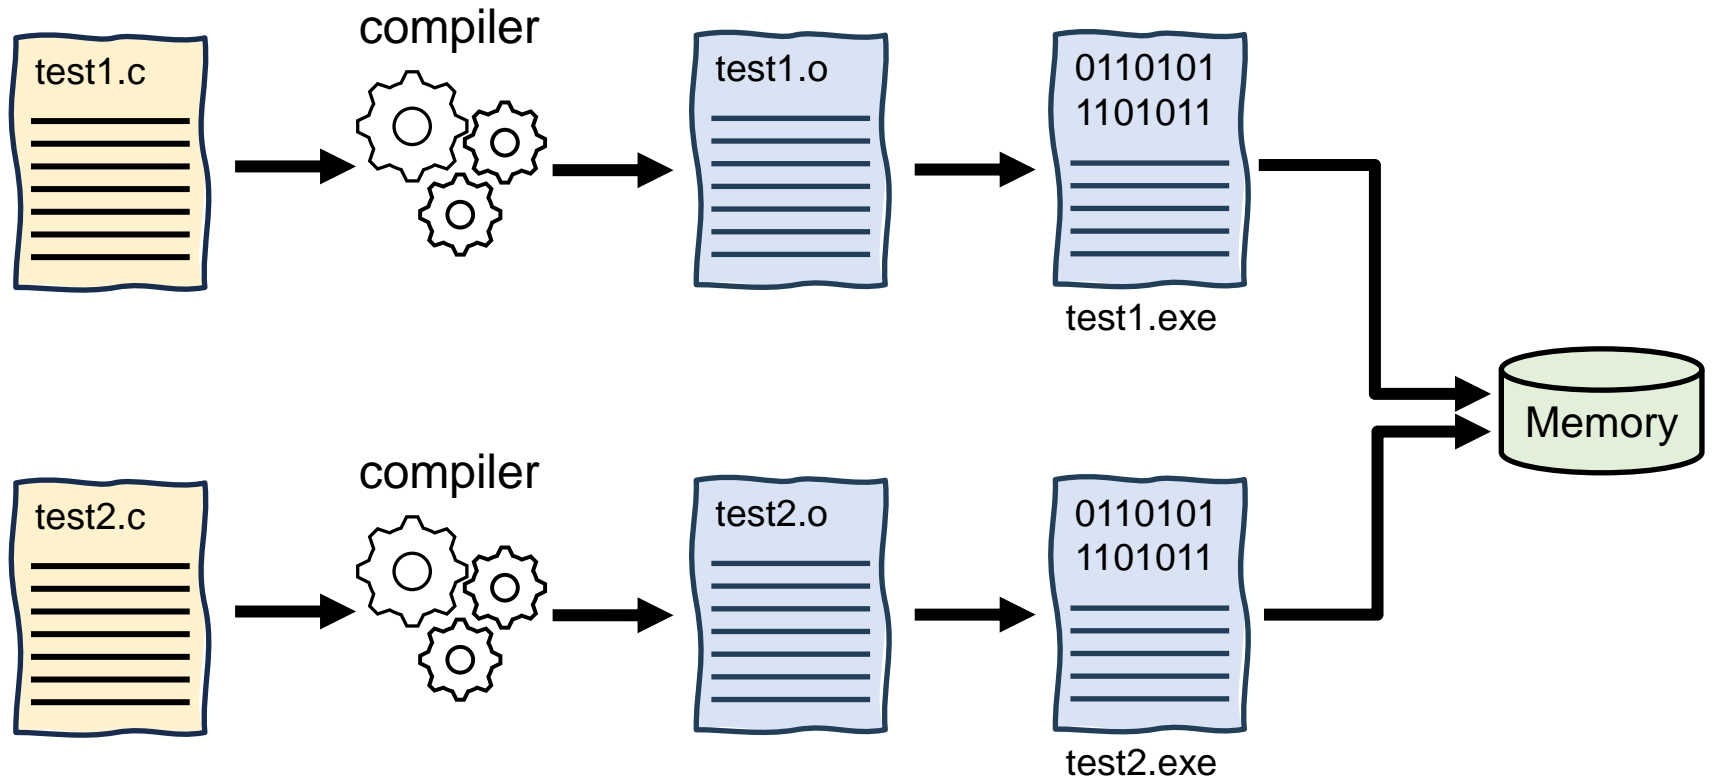


# How is code executed by your computer?



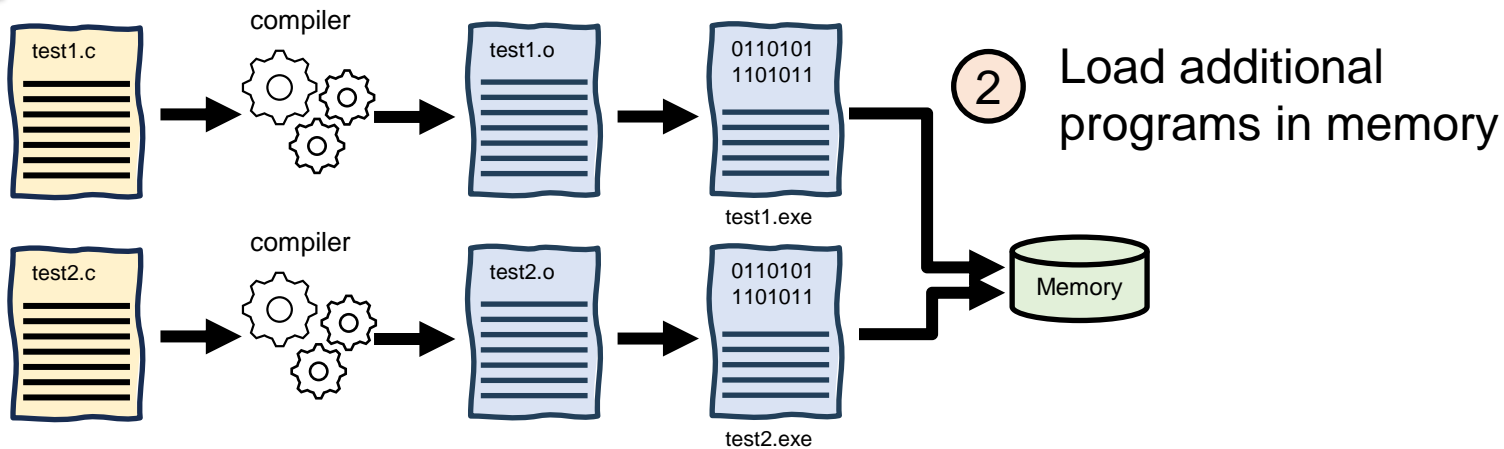
Why do we need an operating system?

# What are the steps to execute a program?



What executes?

# What are the steps to execute a program?



Main program that always runs on the processor from system startup and **manages execution of other programs**

① When system starts, start executing the operating system



③ Request to execute `test1.exe` (using system call)

Operating System

④ Operating system starts execution of `test1.exe`

Processor

Any other reason we might want to use  
an operating system?





## Users

- humans
- machines



## Software

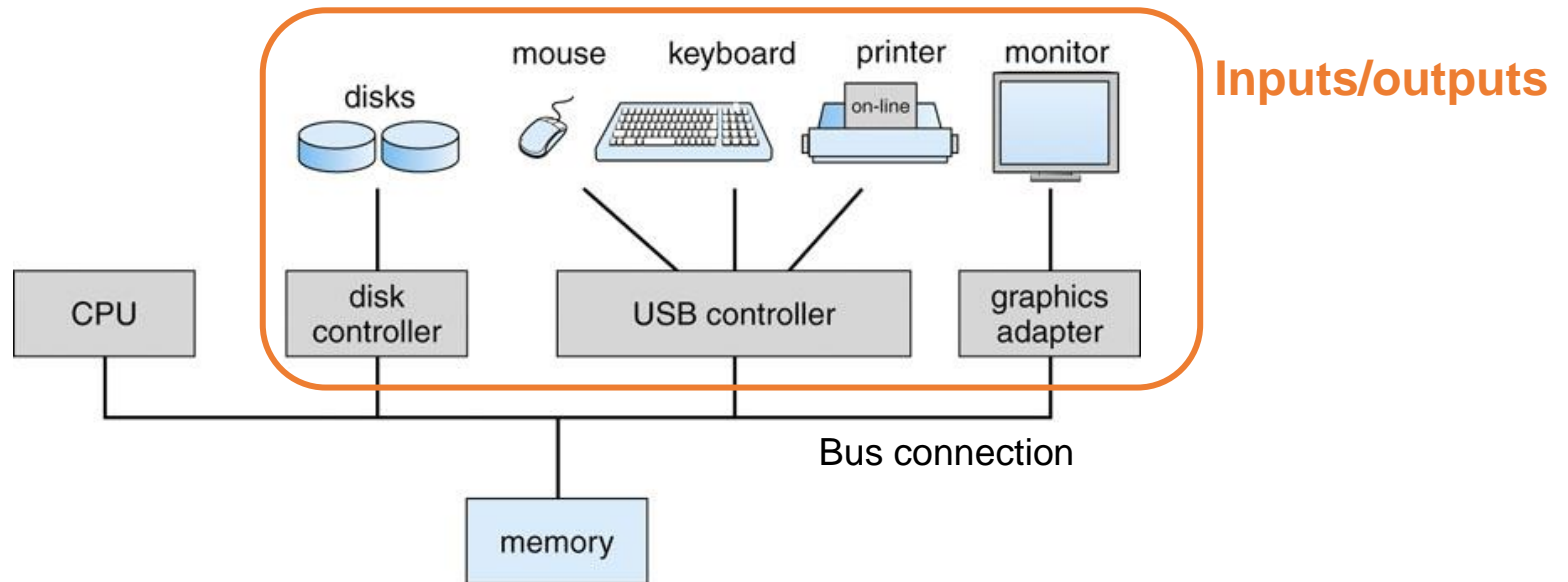
- application programs
- OS, system programs



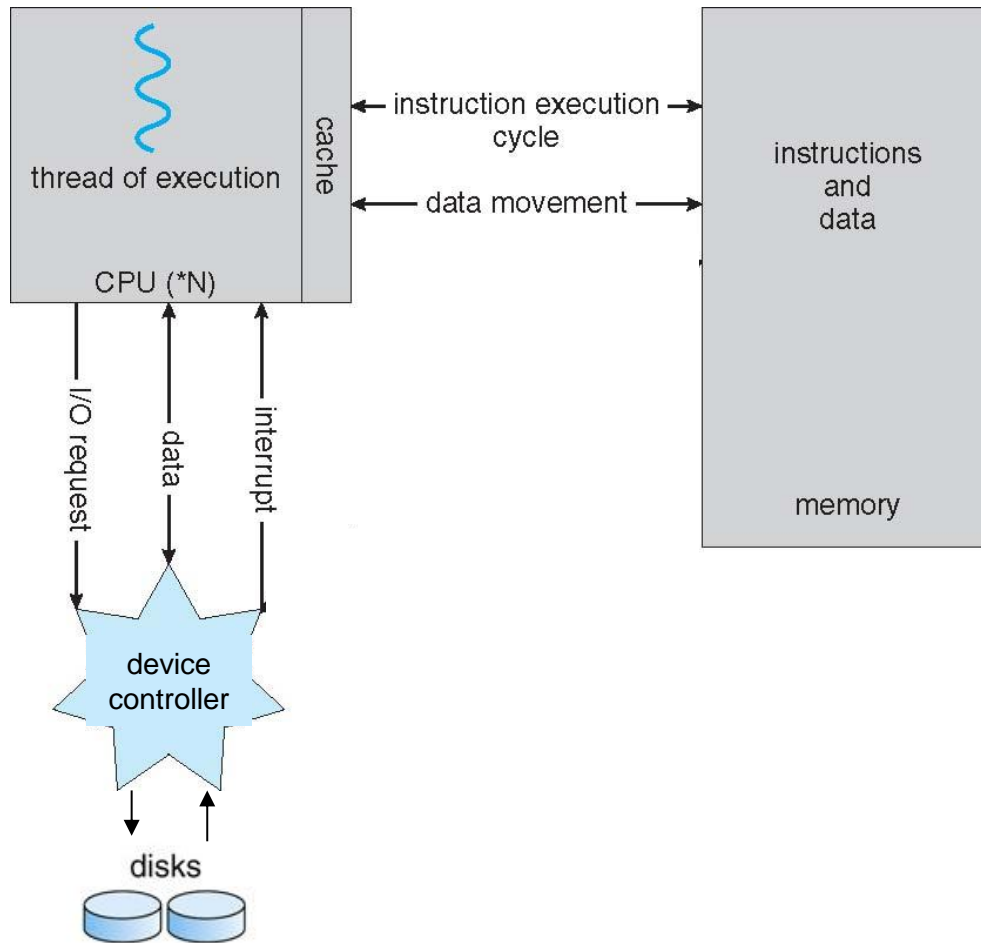
## Hardware

- CPU, memory (storage), battery
- input/output devices

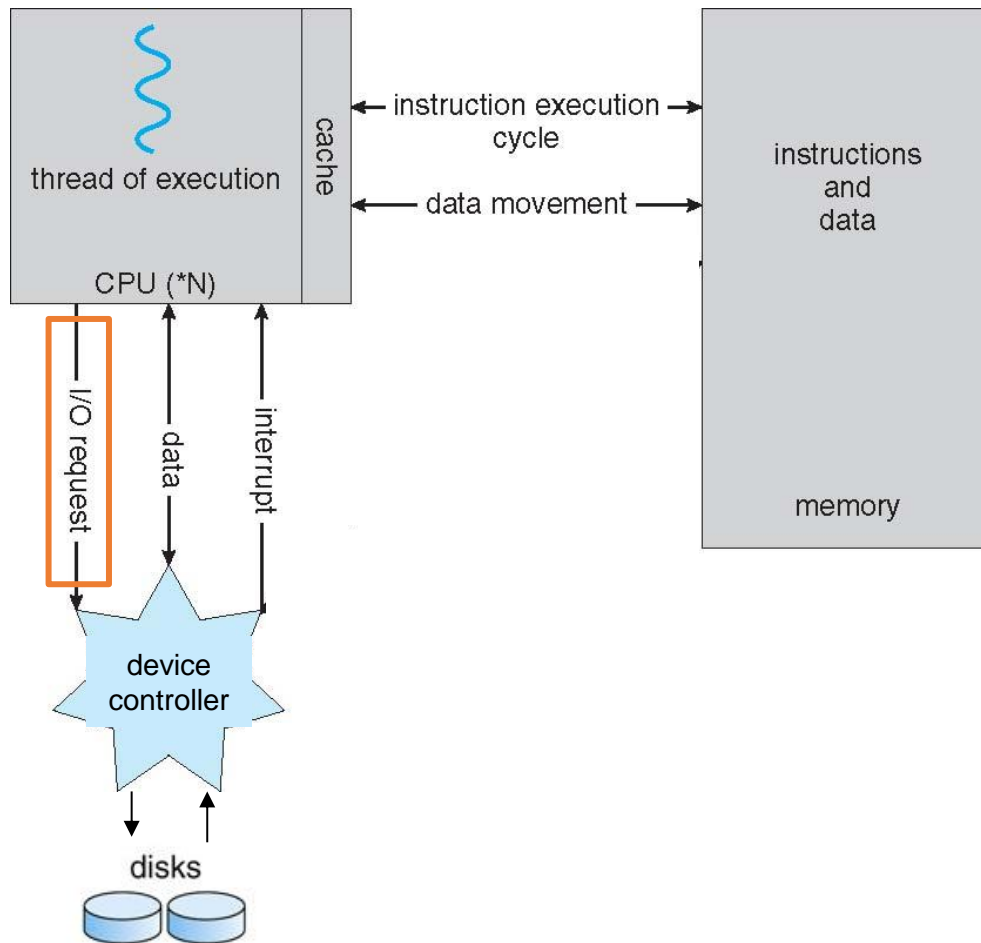
Images from Techstereo



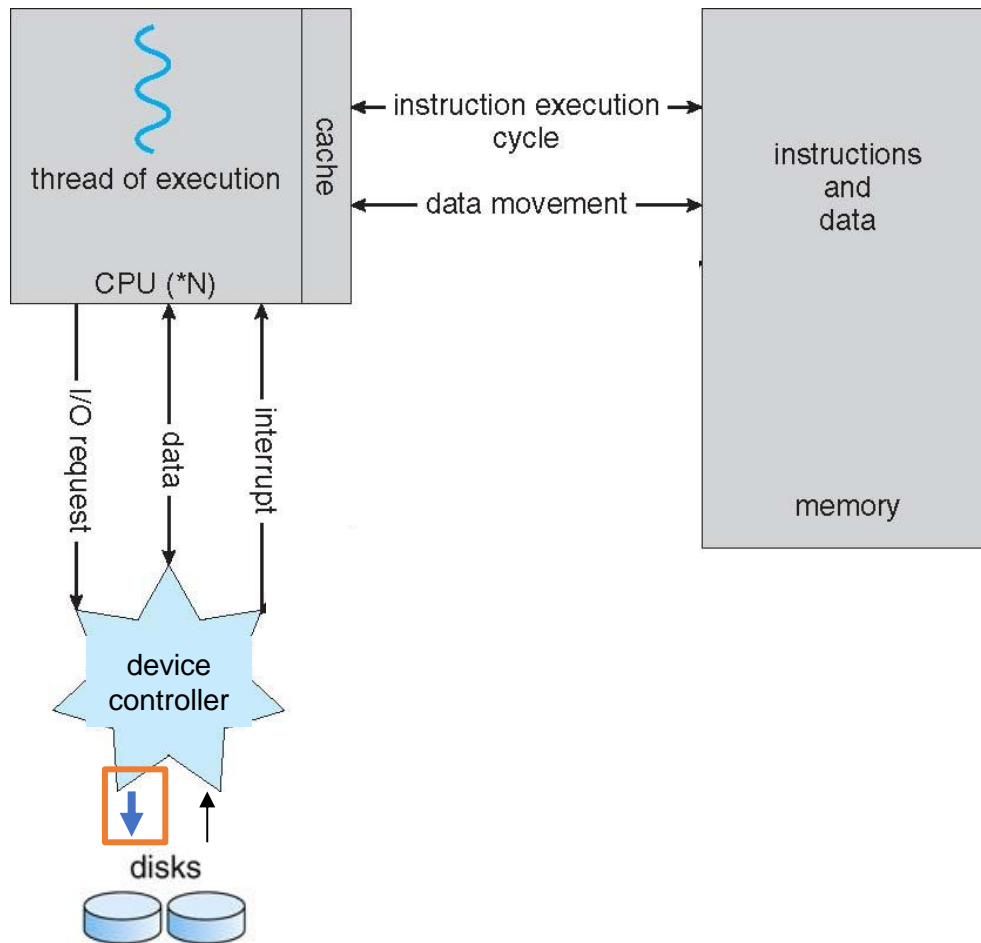
# How does a computer work?



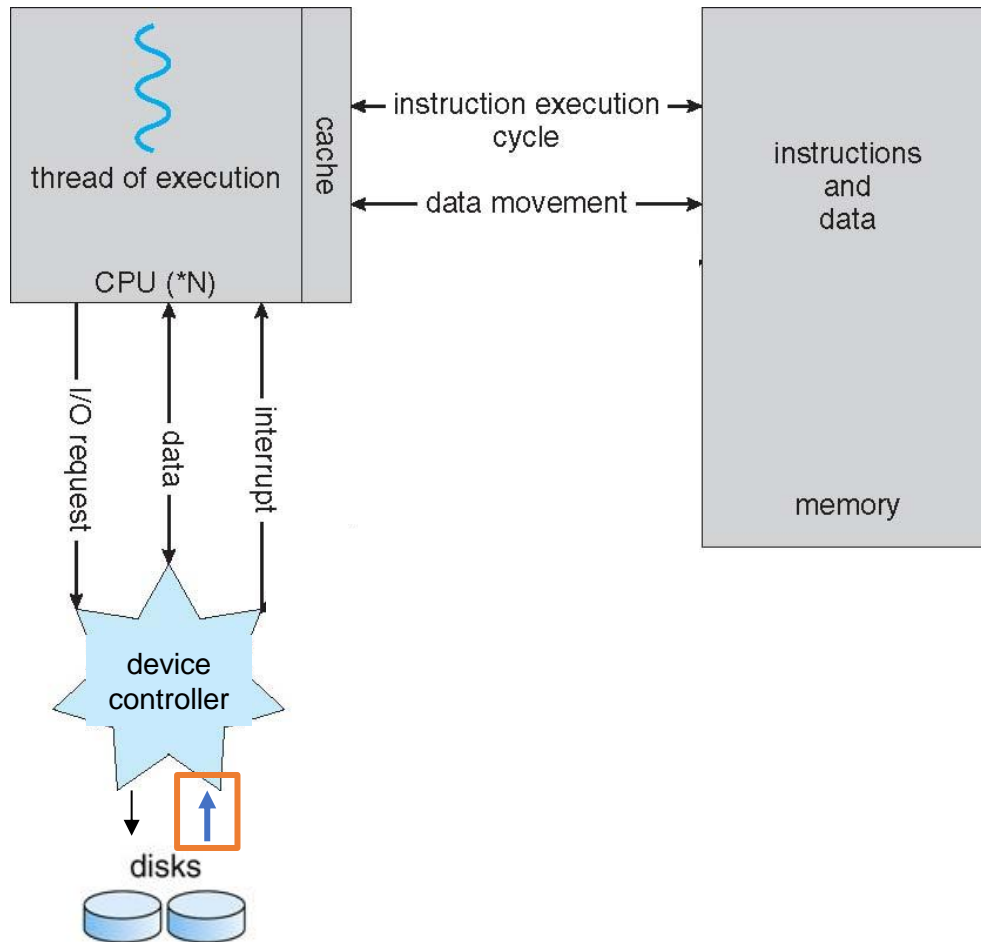
- **Example: access data on a hard drive:**



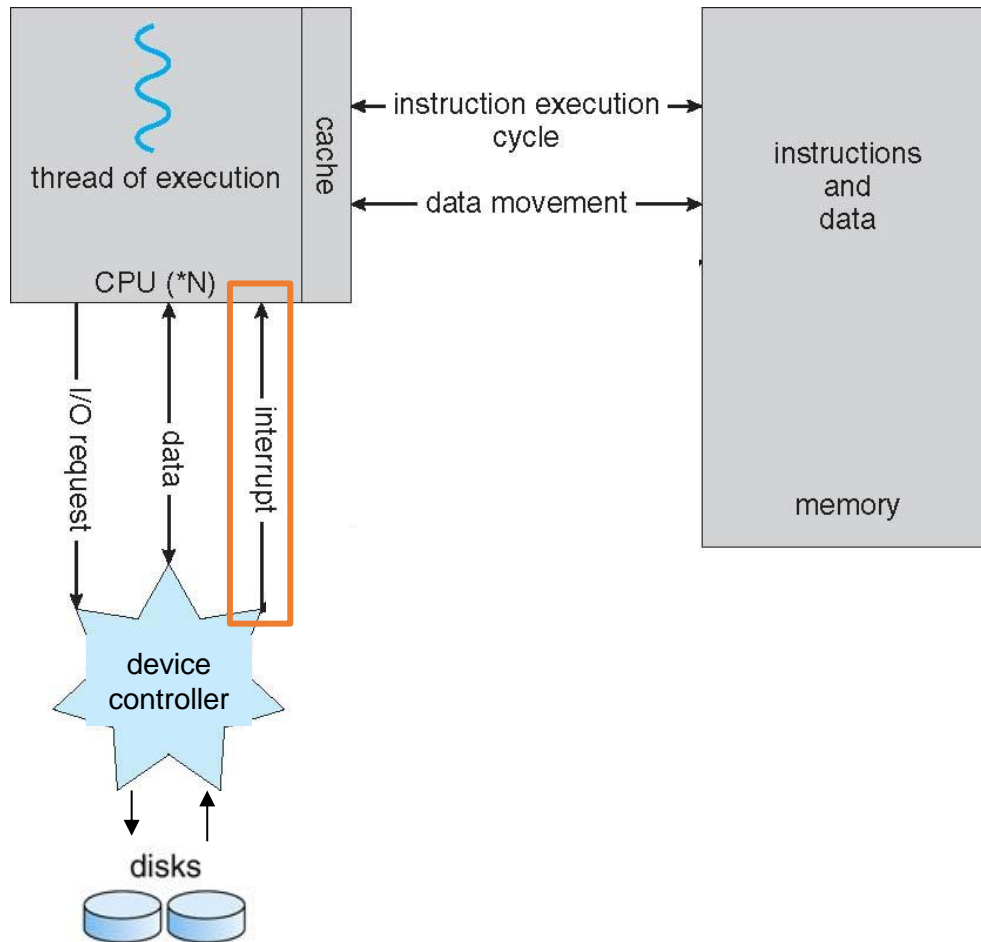
- **Example: access data on a hard drive:**
  1. **CPU tells the device controller what data it wants to read** (CPU writes command + address in **device controller registers**)



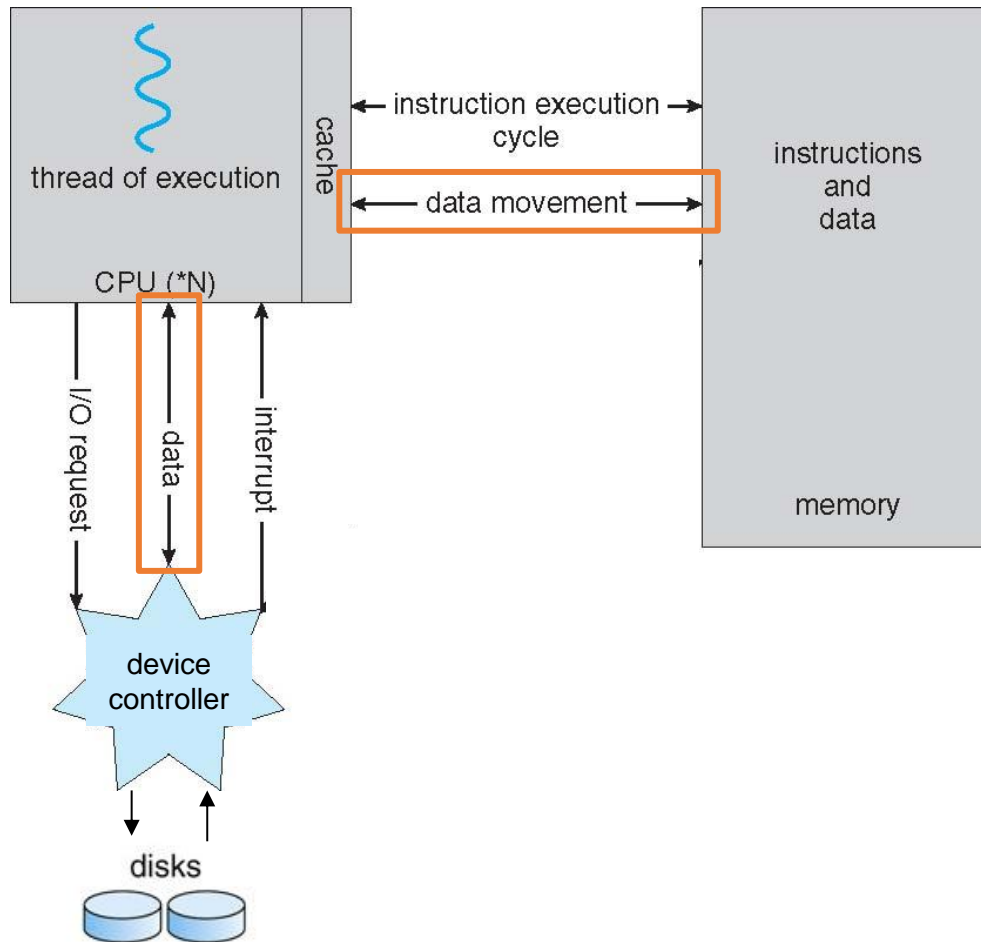
- **Example: access data on a hard drive:**
  1. **CPU tells the device controller what data it wants to read** (CPU writes command + address in **device controller registers**)
  2. **Device controller accesses the hard drive**



- **Example: access data on a hard drive:**
  1. **CPU tells the device controller what data it wants to read** (CPU writes command + address in **device controller registers**)
  2. **Device controller accesses the hard drive**
  3. **Hard drive recovers data** on disc and writes it in the **device controller buffer**

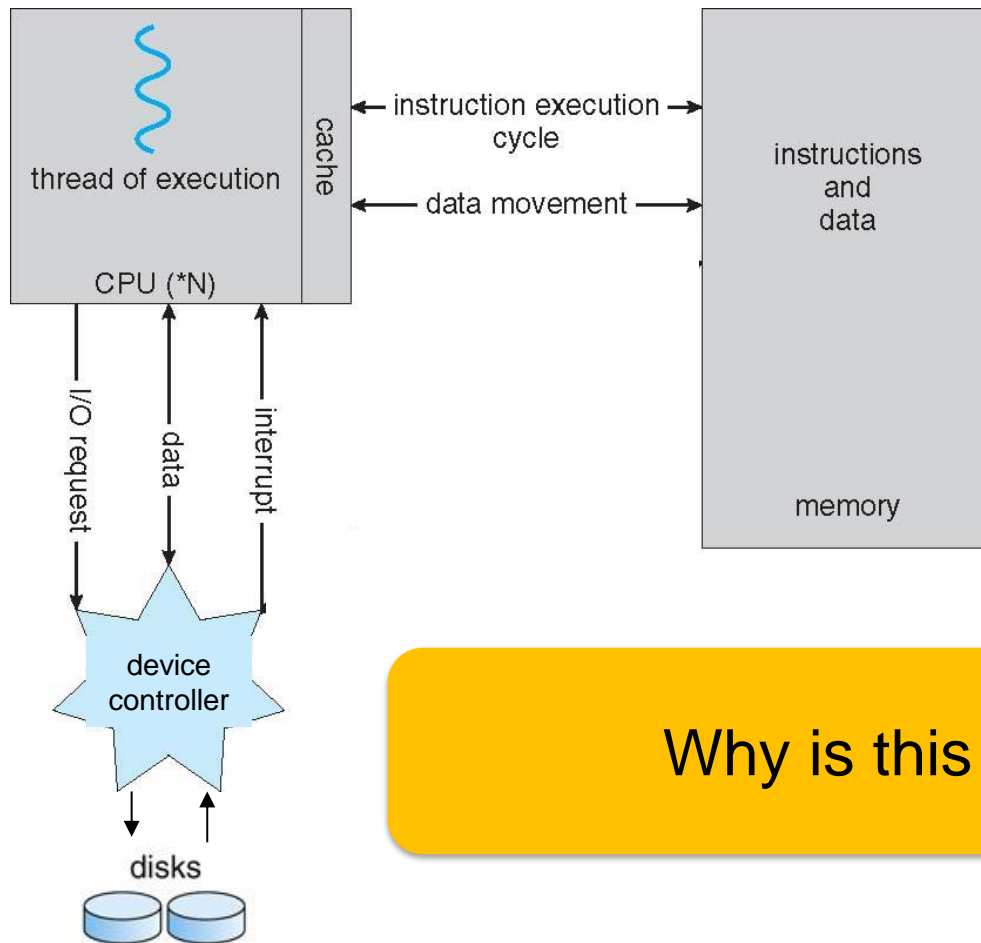


- **Example: access data on a hard drive:**
  1. **CPU tells the device controller what data it wants to read** (CPU writes command + address in **device controller registers**)
  2. **Device controller accesses the hard drive**
  3. **Hard drive recovers data** on disc and writes it in the **device controller buffer**
  4. **Device controller tells the CPU that the data is ready** using an **interrupt**



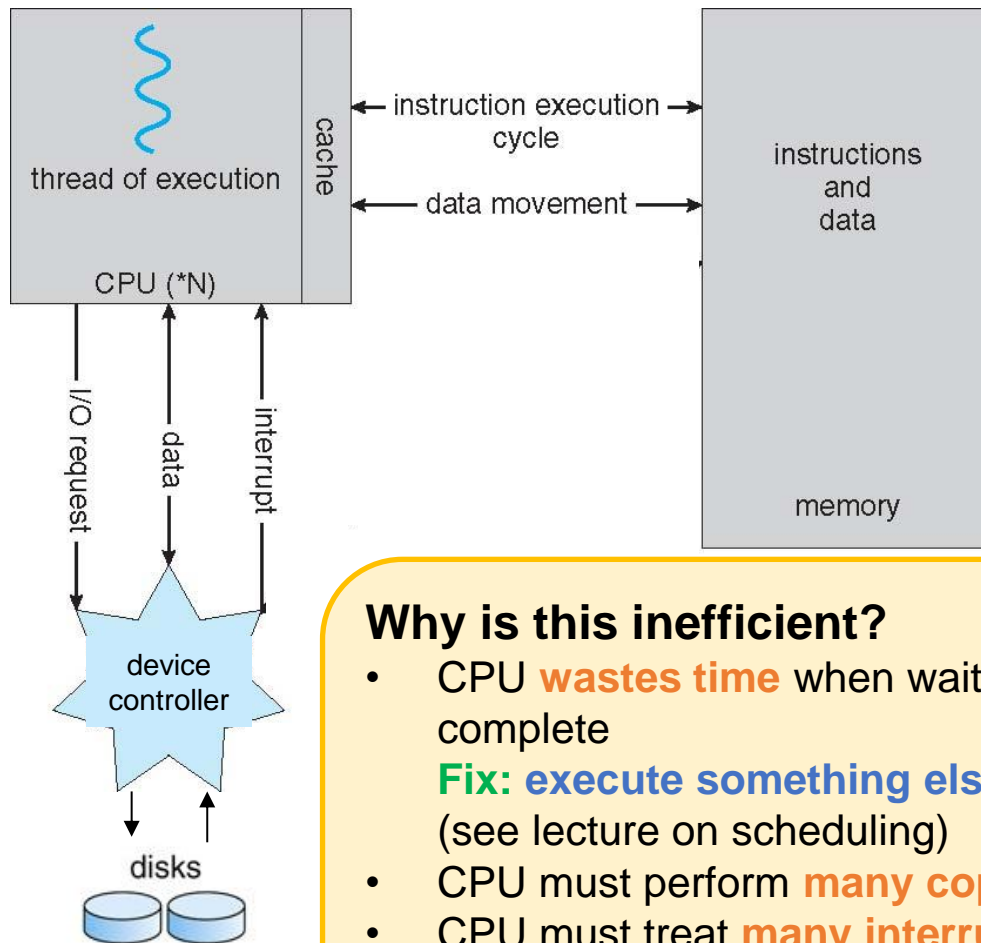
- **Example: access data on a hard drive:**
  1. **CPU tells the device controller what data it wants to read** (CPU writes command + address in **device controller registers**)
  2. **Device controller accesses the hard drive**
  3. **Hard drive recovers data** on disc and writes it in the **device controller buffer**
  4. **Device controller tells the CPU that the data is ready** using an **interrupt**
  5. **CPU reads data** in the device controller and **copies it in memory**





- **Example: access data on a hard drive:**
  1. **CPU tells the device controller what data it wants to read** (CPU writes command + address in **device controller registers**)
  2. **Device controller accesses the hard drive**
  3. **Hard drive recovers data** on disc and writes it in the **device controller buffer**
  4. **Device controller tells the CPU that the data is ready** using an **interrupt**
  5. **CPU reads data** in the device controller and **copies it in memory**
  6. Return to 1 to read next data

Why is this inefficient?

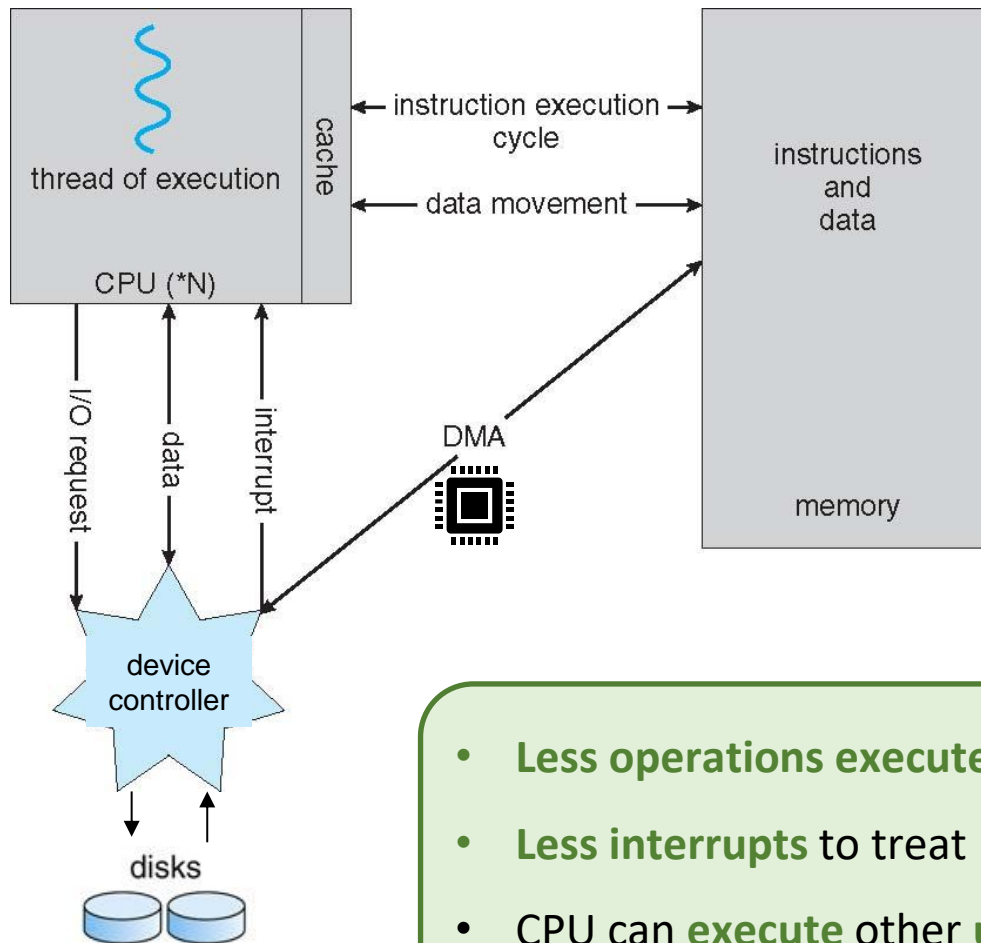


- **Example: access data on a hard drive:**
  1. **CPU tells the device controller what data it wants to read** (CPU writes command + address in **device controller registers**)
  2. **Device controller accesses the hard drive**
  3. **Hard drive recovers data** on disc and writes it in the **device controller buffer**
  4. **Device controller tells the CPU that the data is ready** using an **interrupt**
  5. **CPU reads data** in the device controller and **copies it in memory**
  6. Return to 1 to read next data

## Why is this inefficient?

- CPU **wastes time** when waiting for the I/O operation to complete  
**Fix: execute something else** on the CPU while waiting (see lecture on scheduling)
- CPU must perform **many copy operations**
- CPU must treat **many interrupts**  
**Fix: use direct memory access (DMA)**

Imagine we must move 1000 data




- **Example: access data on a hard drive using a DMA:**

1. **CPU tells the DMA what block of data it wants to read** (CPU writes command + **start address** + **size** of block of data it wants to read in **DMA registers**)
2. **DMA configure the I/O device controller for every data that must be read**
3. **Device controller accesses the hard drive**
4. **Hard drive recovers data** on disc and writes it in the **device controller buffer**
5. **DMA reads data** in the device controller and **copies it in memory**
6. **Once all data are transferred** in the memory, **DMA tells the CPU the I/O transfer is completed using an interrupt**

- **Less operations executed** by the CPU
- **Less interrupts** to treat
- CPU can **execute** other **useful code in parallel**

# How does a computer work?



Someone must **decide what to execute** while waiting for the I/O

Someone must provide code to **configure** the **HW devices** and **treat** the **interrupts**

Someone must ensure that the **program state remains consistent** when switching between executing tasks

- HW operates at a very **low abstraction level** compared to user programs
  - A mediator is needed in between
- HW resources are **shared among running programs** (in time and in space)
  - A burden that must be taken away from the programmer
- HW can be expensive
  - **Use of HW resources needs to be optimized**

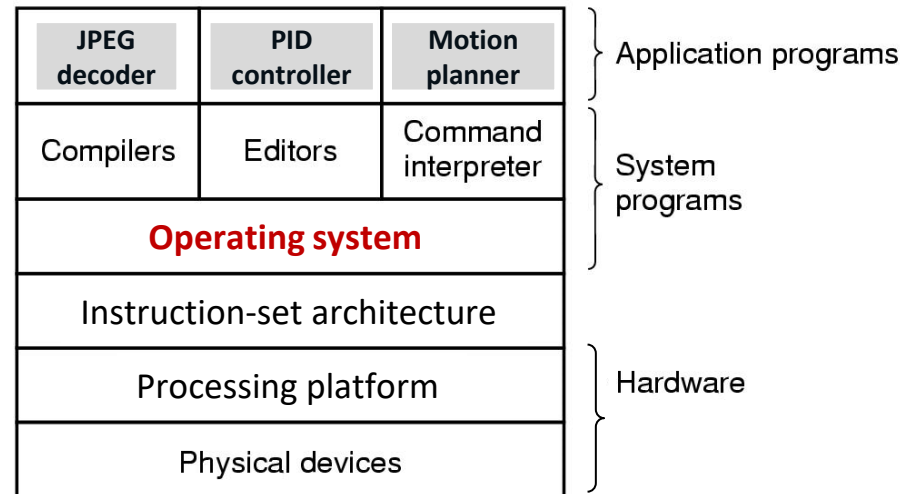
# Why do we need an operating system?

An operating system is a piece of software that acts as an intermediary between **users/applications** and **computer hardware**.

It provides a level of abstraction that **hides the gory details of the HW architecture from applications**.

It **manages** the **sharing of HW resources** among applications and users

It **optimizes performance** through **resource management**




# Without operating systems ...

In a “**bare-metal implementation**”, application developers must

- Know hardware platform details to be able to write applications
- Manage the memory and address space (very complex)
- Manage I/O operations
- Manage communications
- Manage the file system

Operating systems turn ugly (and complex) hardware into **beautiful abstractions**




Beautiful interface

Ugly (and complex) interface

Source of the figure: Tanenbaum, Modern Operating Systems 3 e, (c) 2008 Prentice-Hall, Inc.

Application code

Hardware



Application development becomes **time consuming** and **error prone**

**Low portability** (the code cannot be easily used on other HW platforms)

It also becomes very **hard to use** **third-party code and libraries**

# What does an operating system provide?

## Process management

- Creation and deletion
- Scheduling, suspension, and resumption
- Synchronization, inter-process communication

## Memory management

- Allocate and deallocate memory space as requested
- Efficient utilization when the memory resource is heavily contended
- Keep track of which parts of memory are currently being used and by whom

## I/O management

- A buffer-caching system
- A general device-driver interface
- Drivers for specific hardware devices

## File management

- Manipulation of files and directories
- Map files onto (nonvolatile) secondary storage -disks
- Free space management and storage allocation
- Disk scheduling

## Error detection

- Debugging facilities

## Protection

- concurrent processes should not interfere with each other (memory)

## Security

- Against other processes and outsiders

## Accounting

- Using timers
  - CPU cycles,
  - main memory usage
  - I/O devices usage
  - ...



# What does an operating system provide?

## Process management

- Creation and deletion
- Scheduling, suspension, and resumption
- Synchronization, inter-process communication

## Memory management

- Allocate and deallocate memory space as requested
- Efficient utilization when the memory resource is heavily contended
- Keep track of which parts of memory are currently being used and by whom

## I/O management

- A buffer-caching system
- A general device-driver interface
- Drivers for specific hardware devices

## File management

- Manipulation of files and directories
- Map files onto (nonvolatile) secondary storage -disks
- Free space management and storage allocation
- Disk scheduling

## Error detection

- Debugging facilities

## Protection

- concurrent processes should not interfere with each other (memory)

## Security

- Against other processes and outsiders

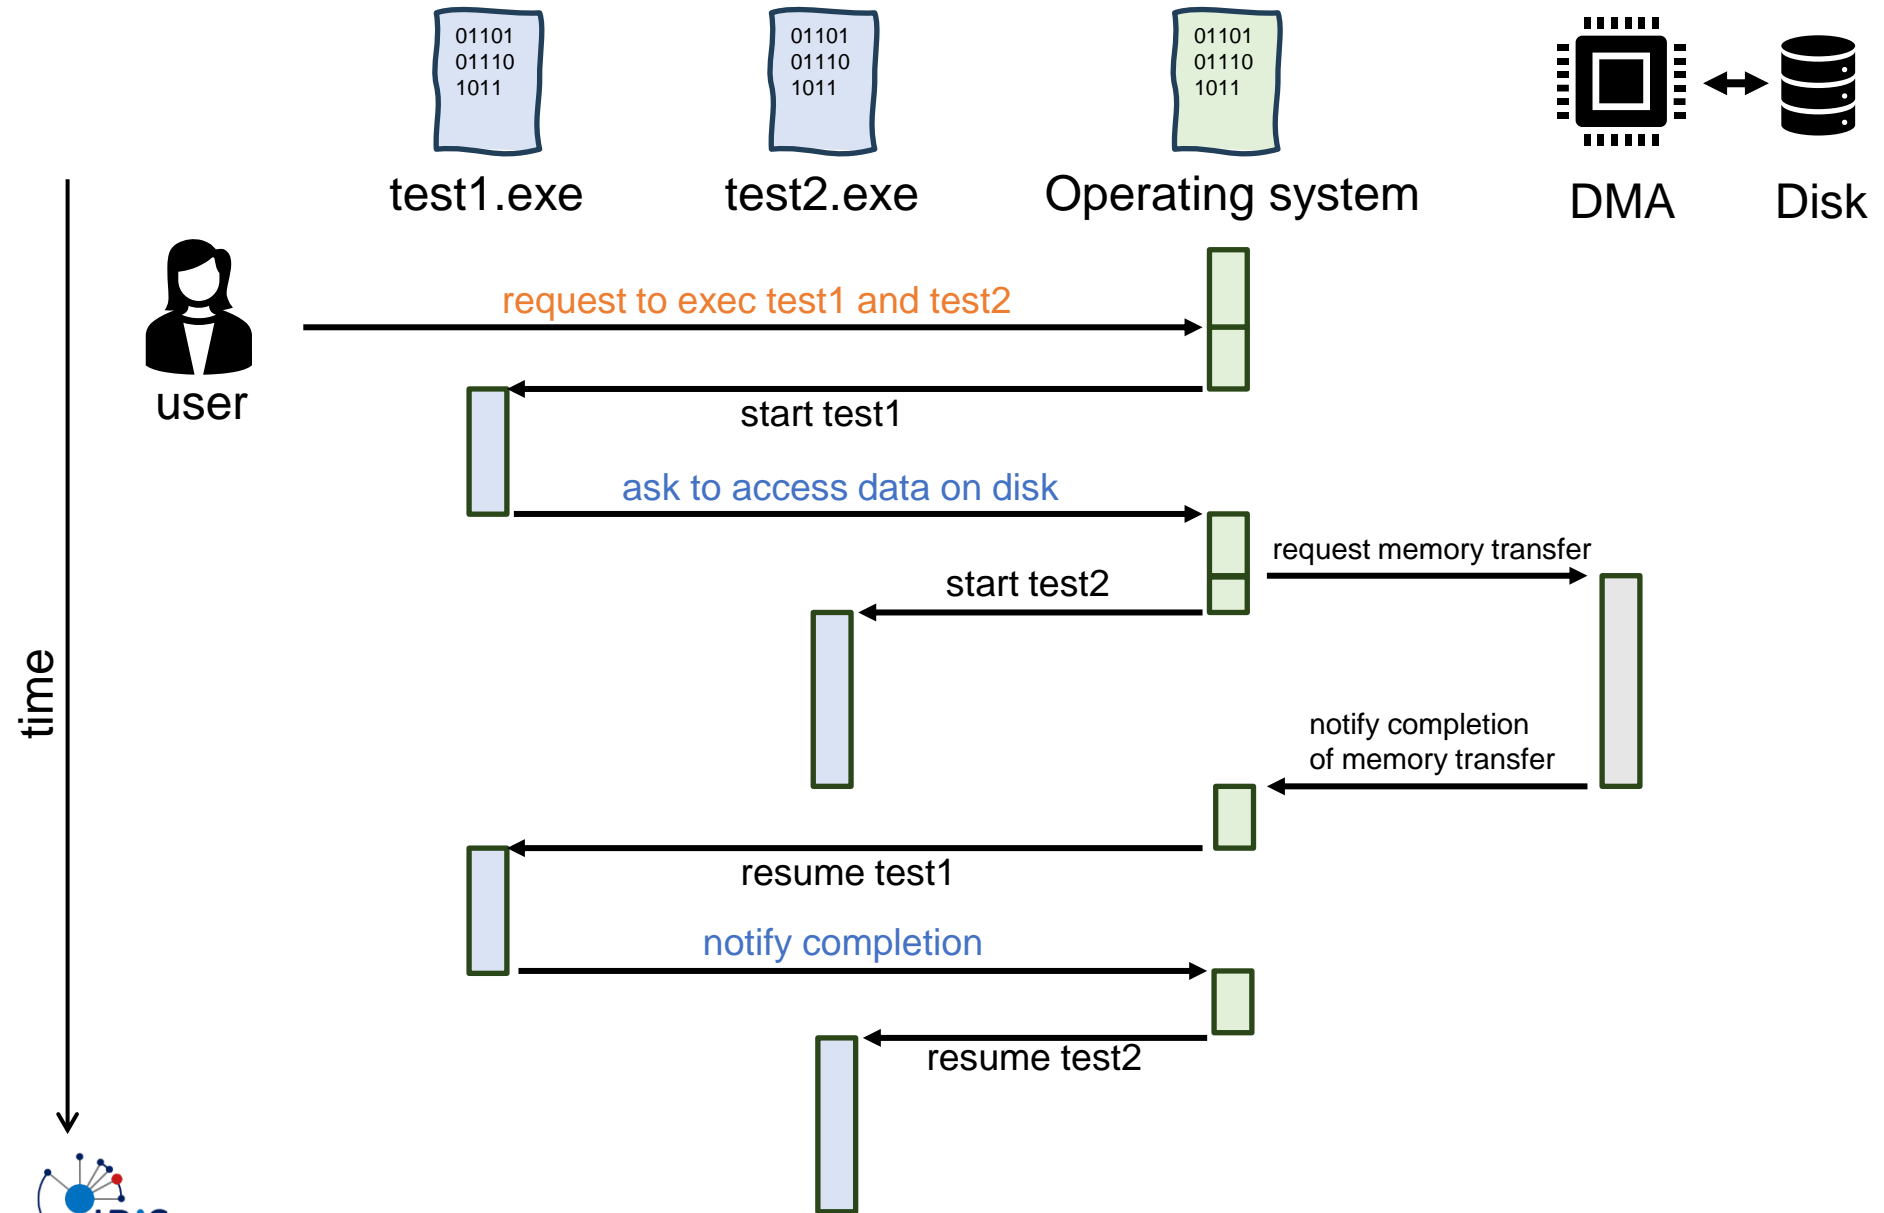
## Accounting

- Using timers
- CPU cycles,
  - main memory usage
  - I/O devices usage
  - ...

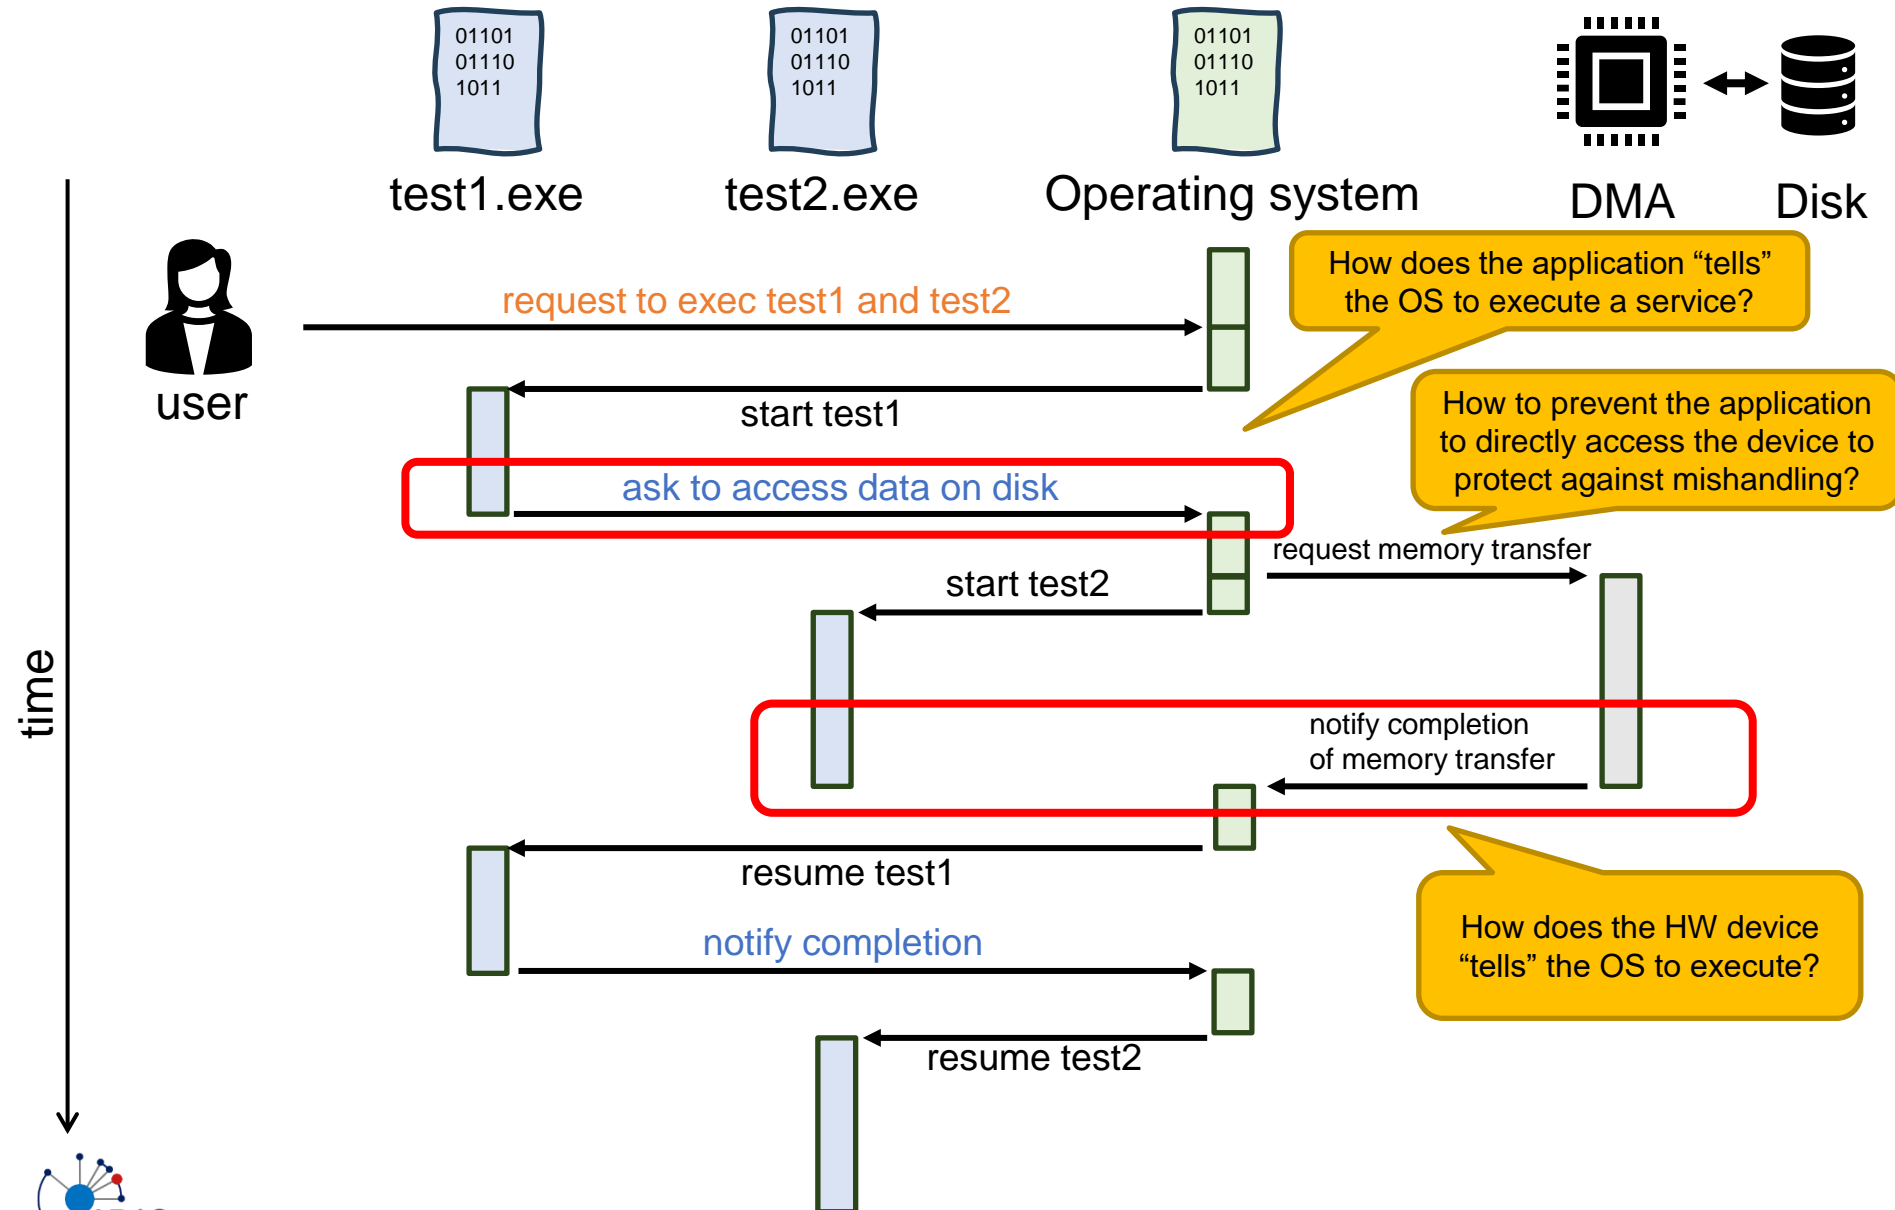
Covered in this course

# How applications/external devices interact with an OS?

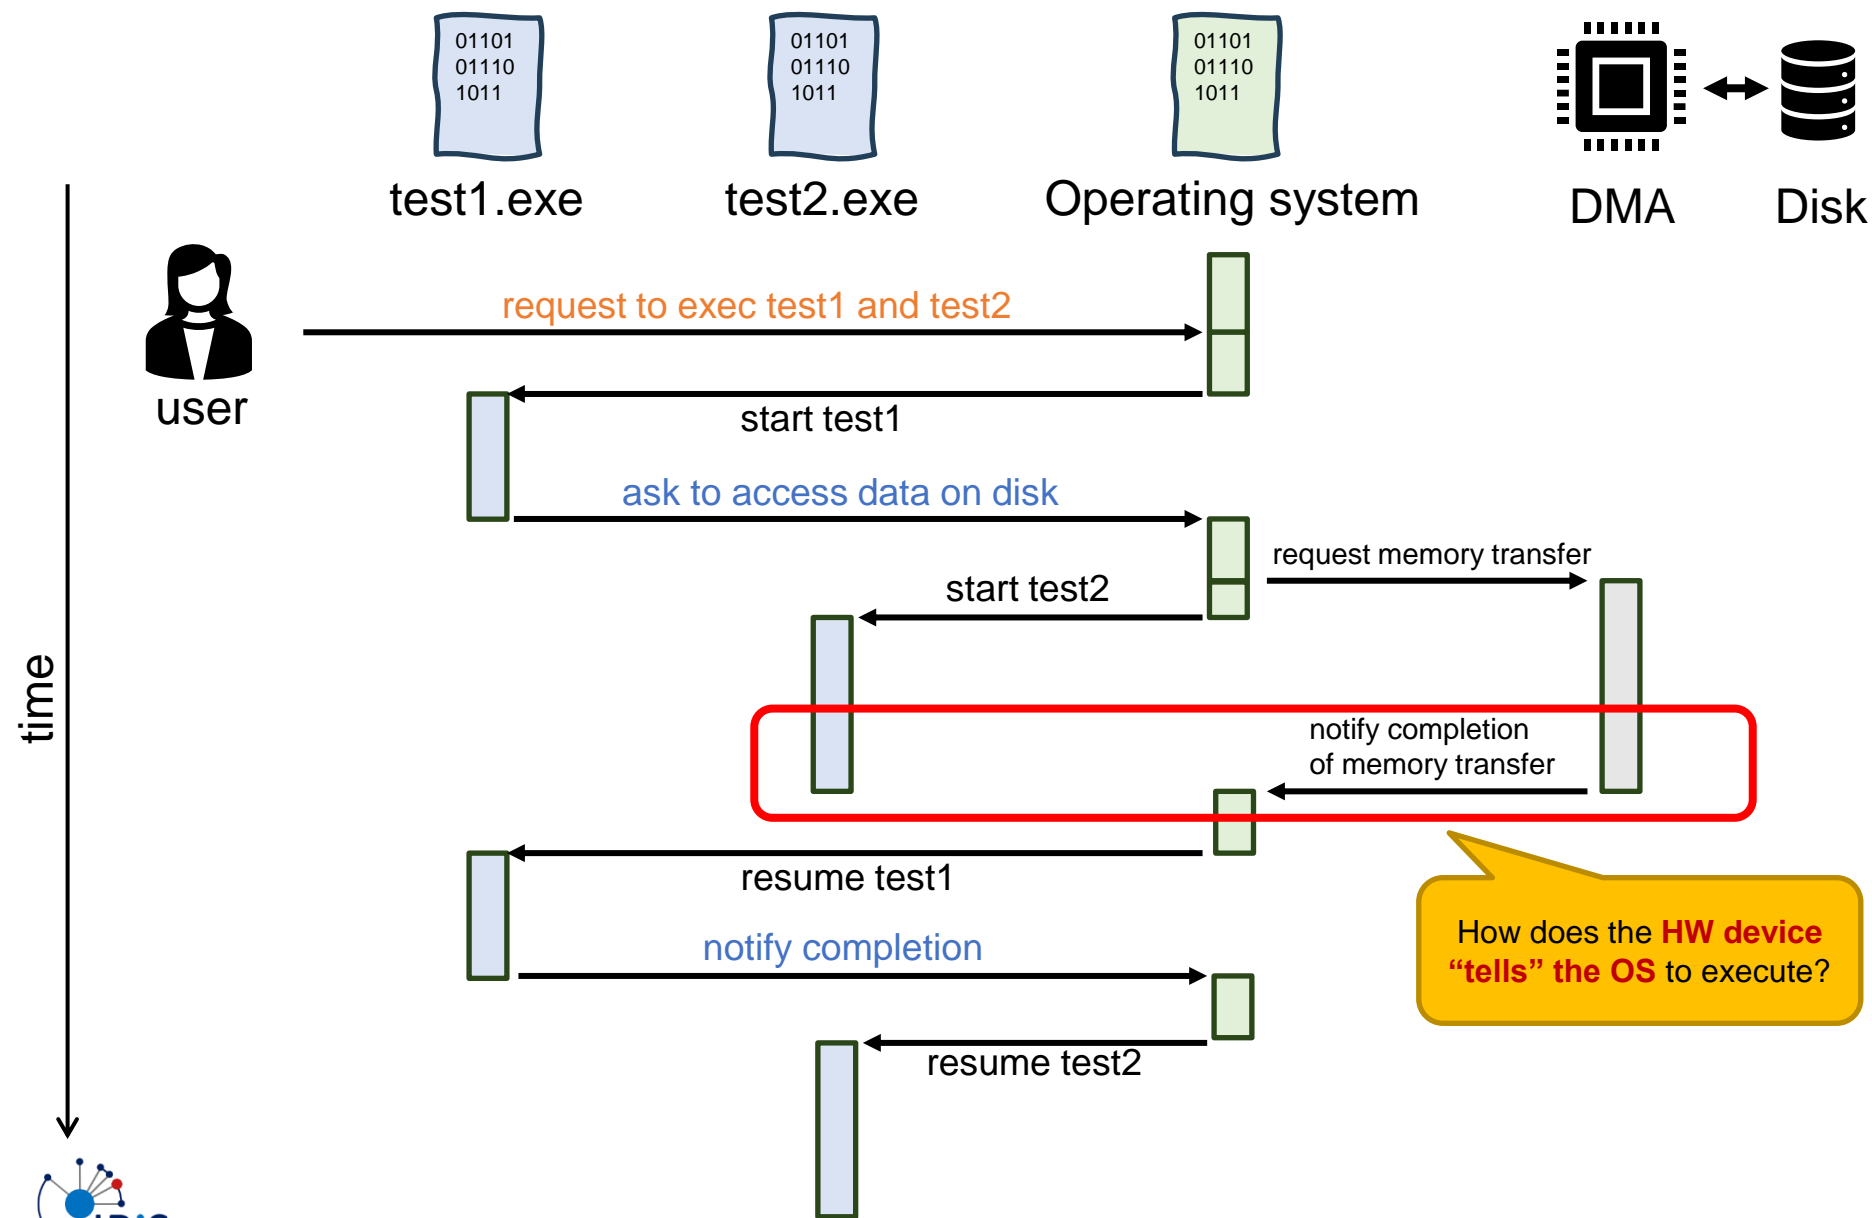
# How OS and application/devices interact?



# How OS and application/devices interact?

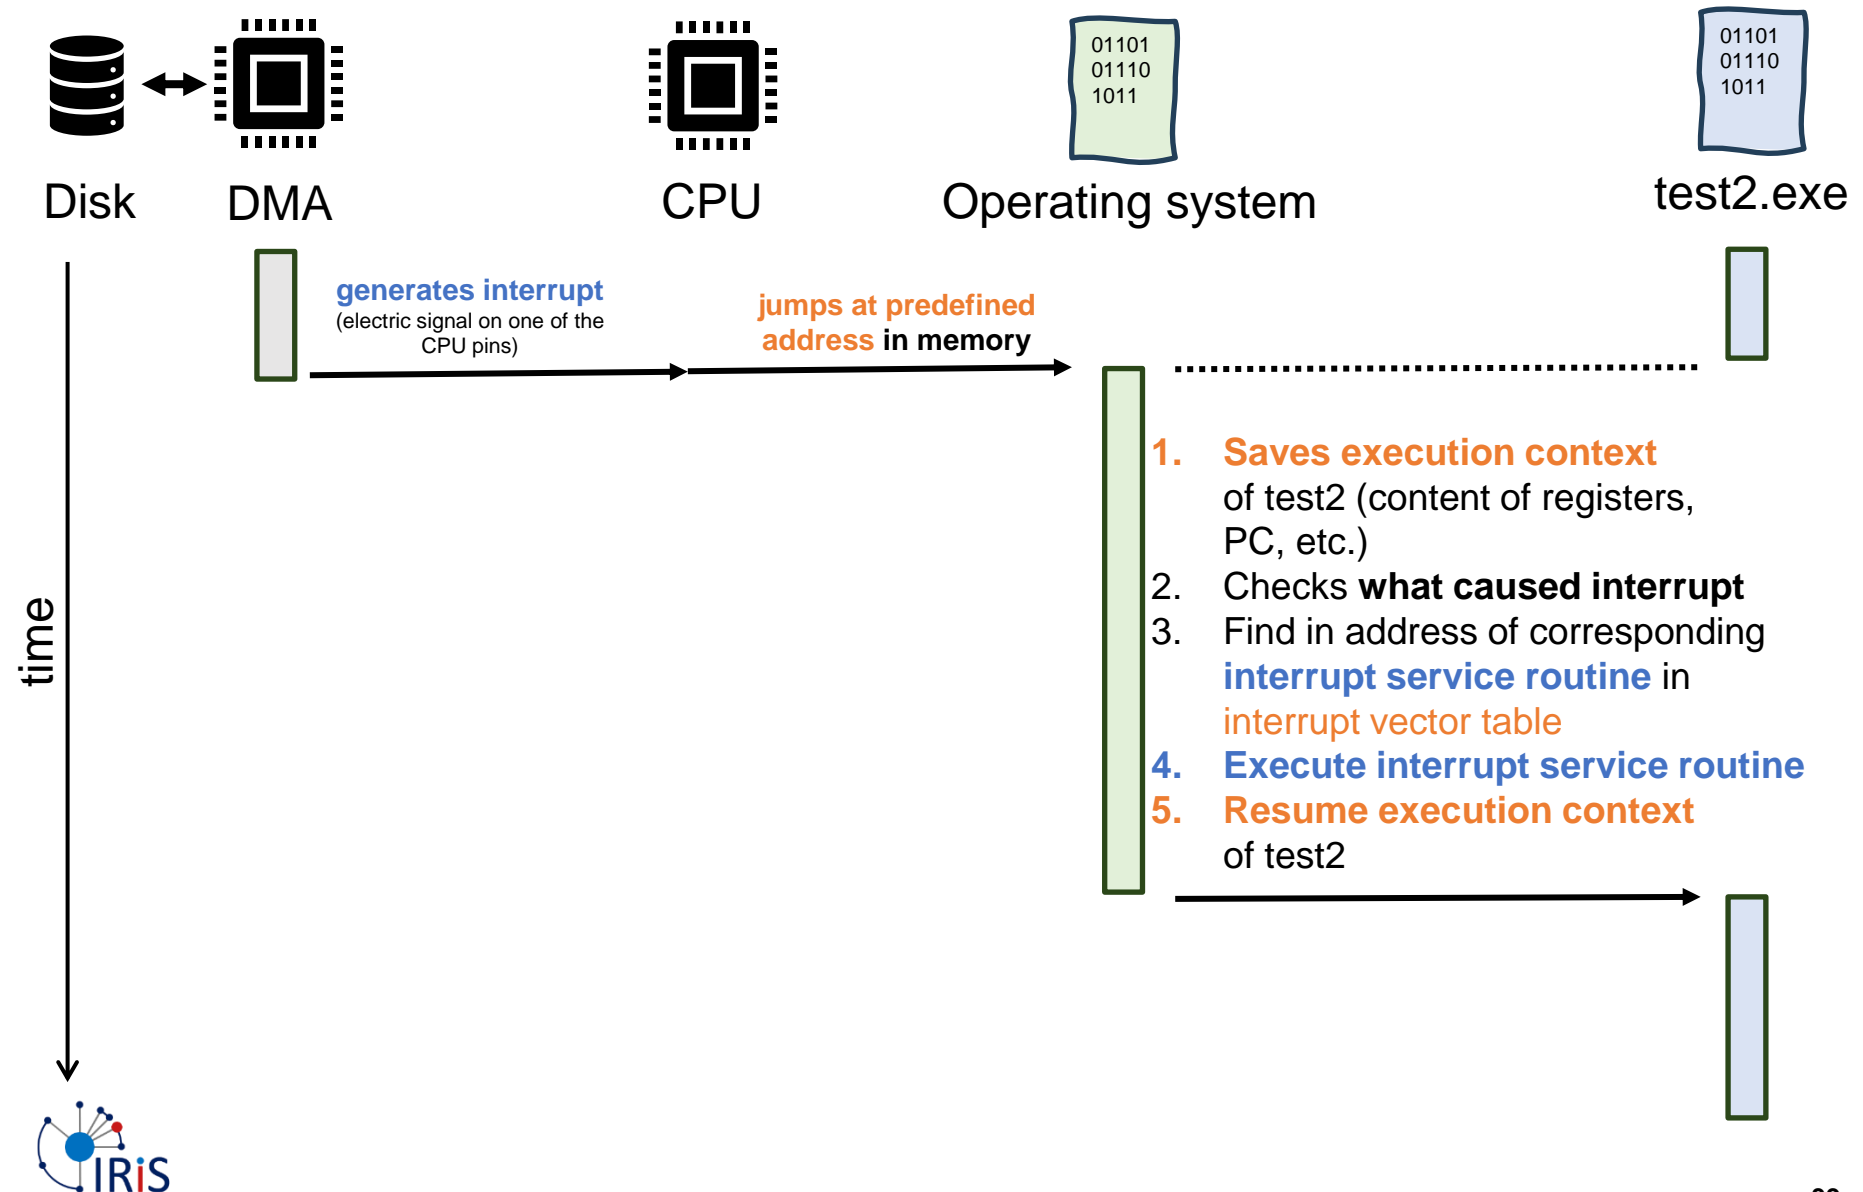


# How OS and application/devices interact?

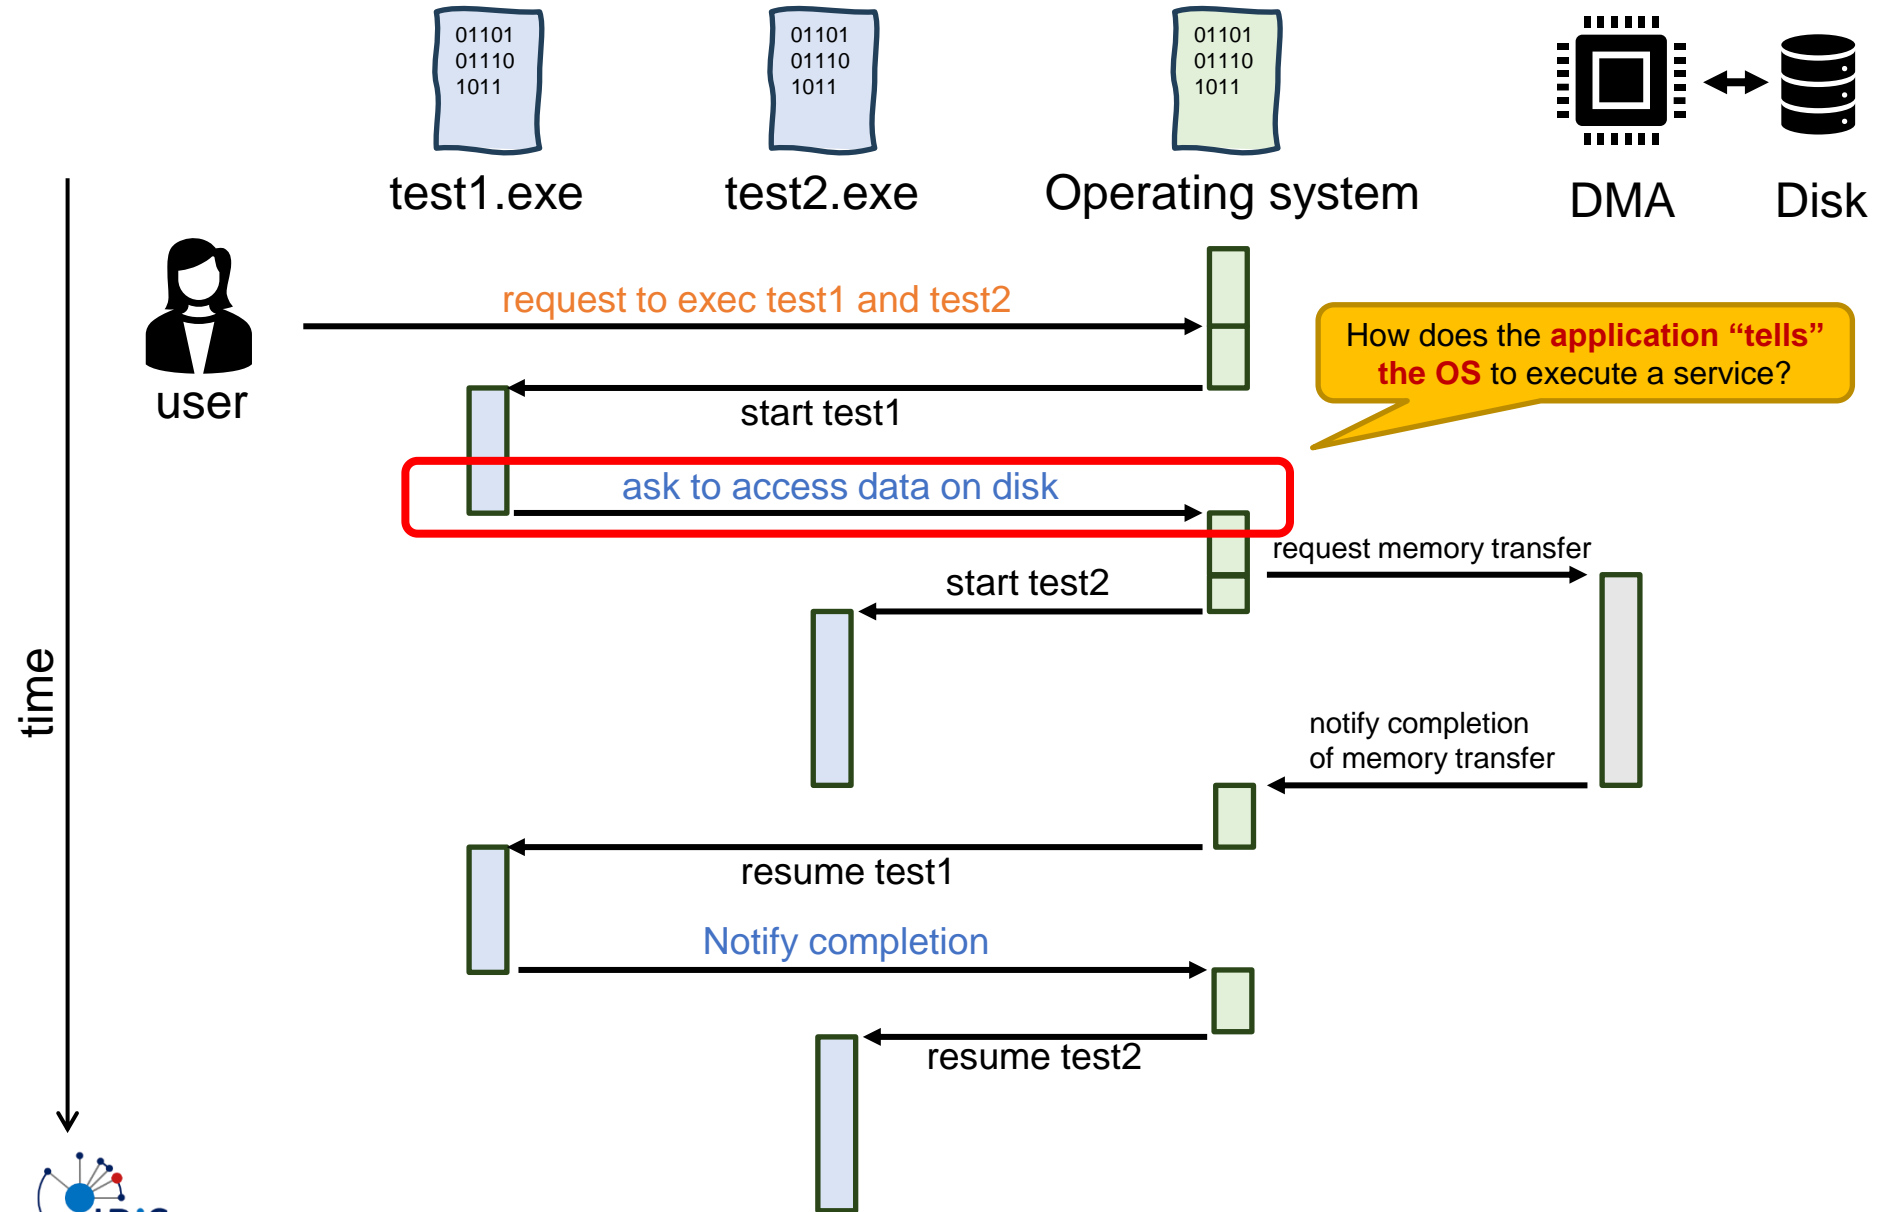


# Communication with devices: OSs are interrupt driven

# Communication with devices: OSs are interrupt driven

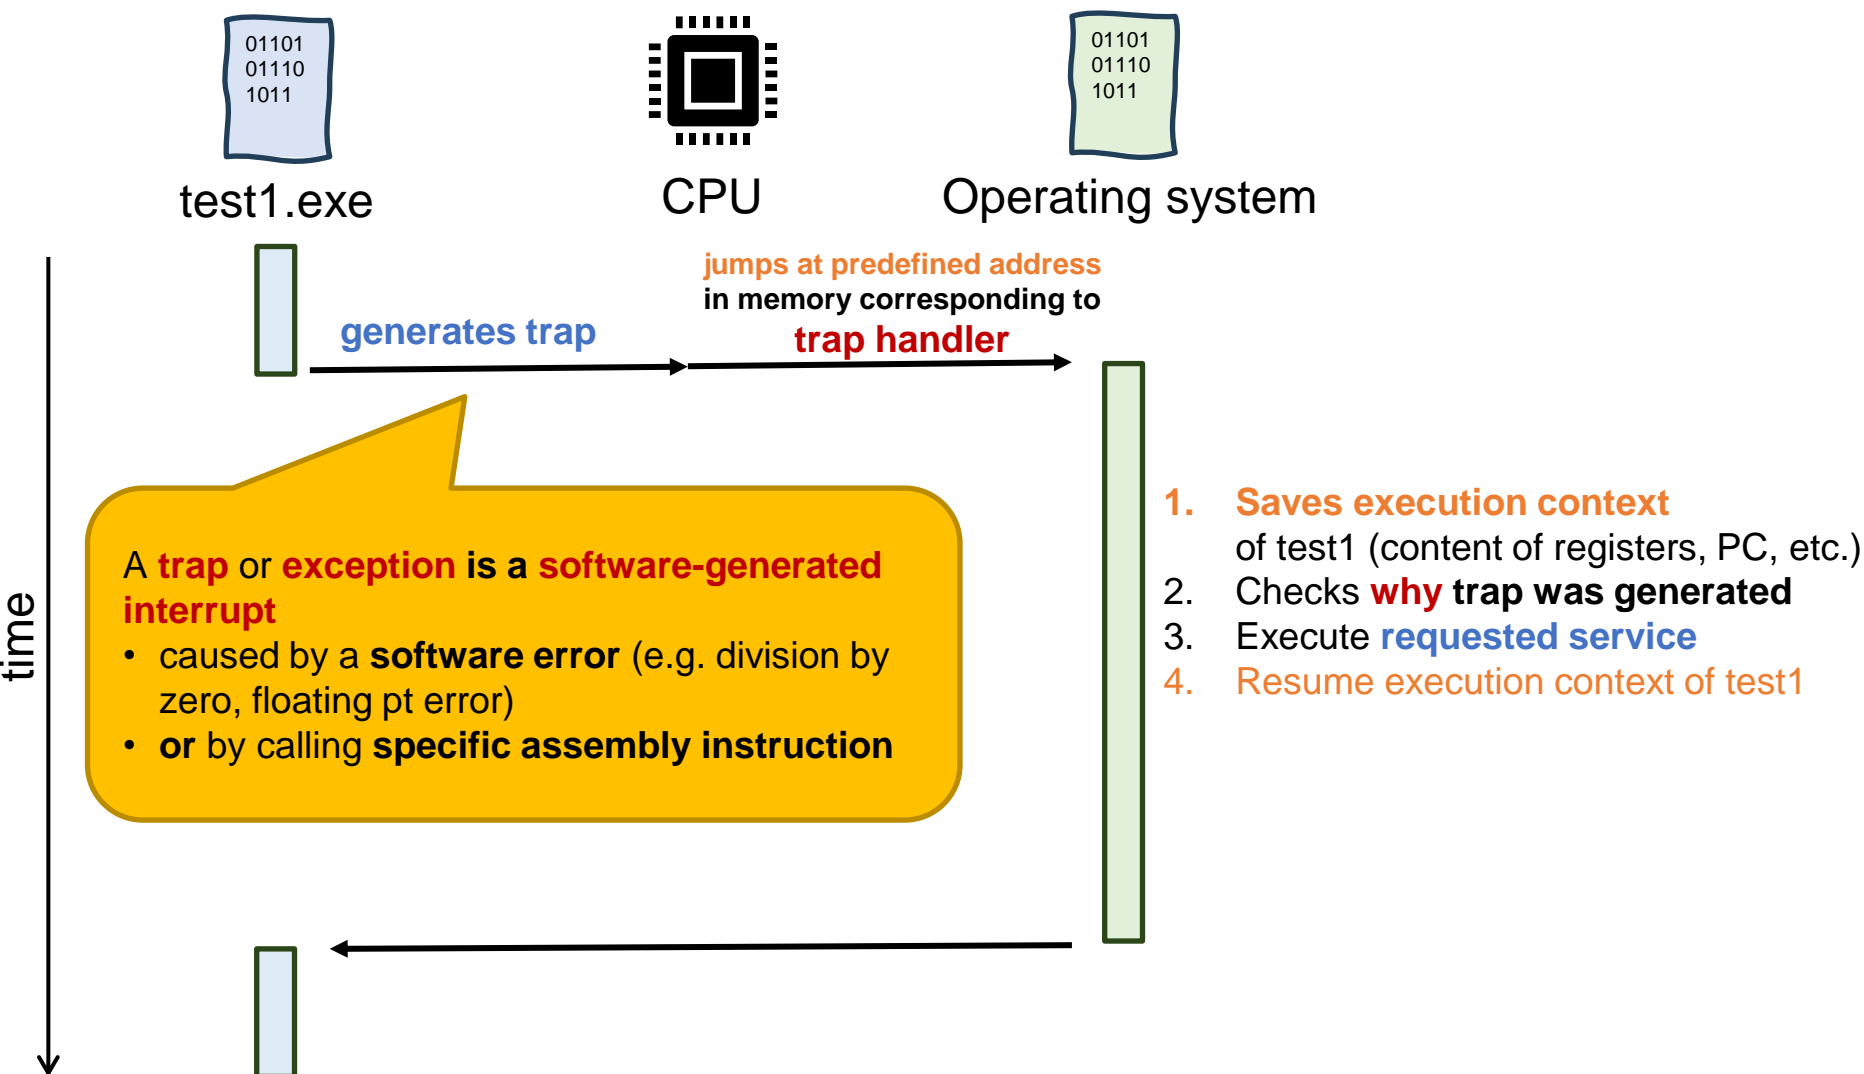


# How OS and application/devices interact?

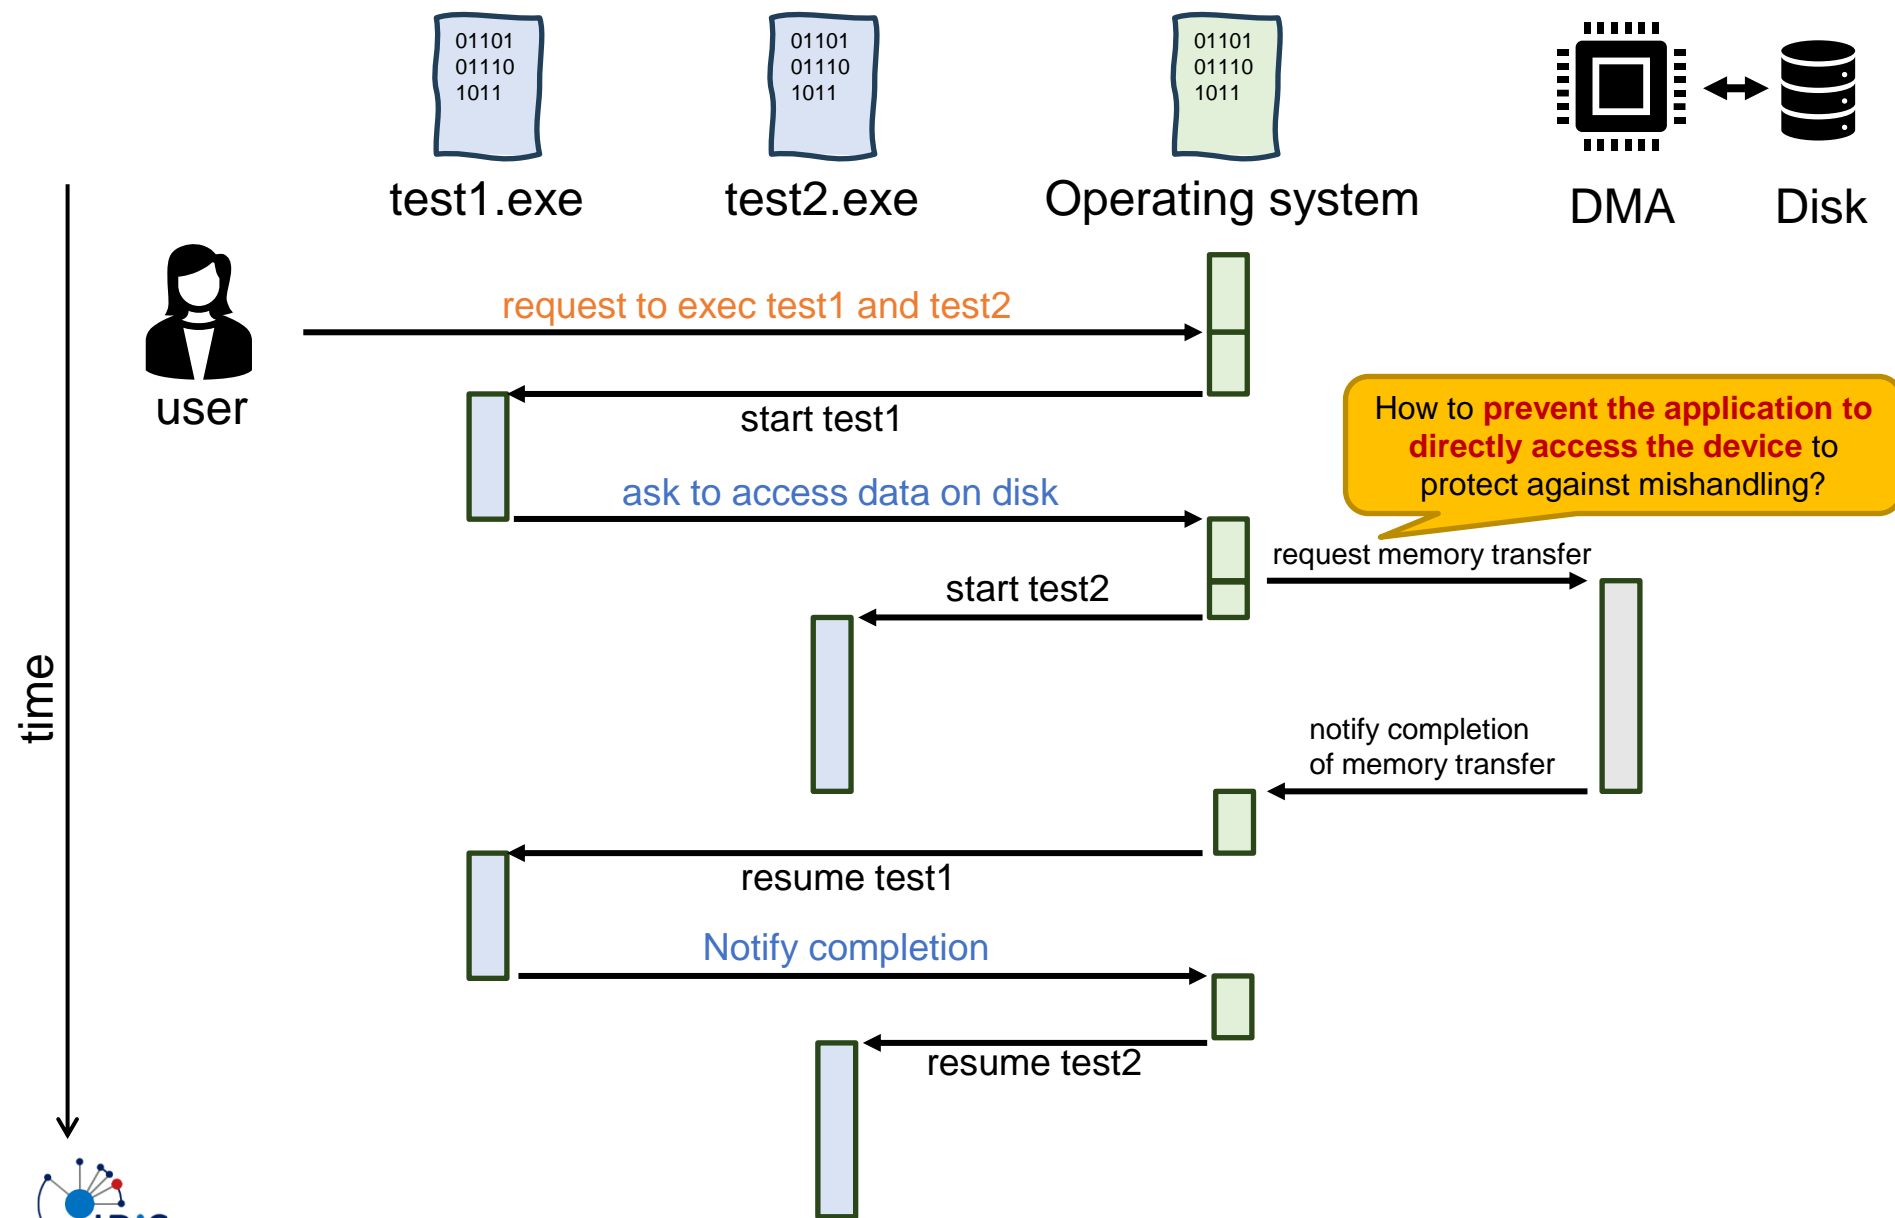




# Communication between applications and OS: OS's are interrupt driven

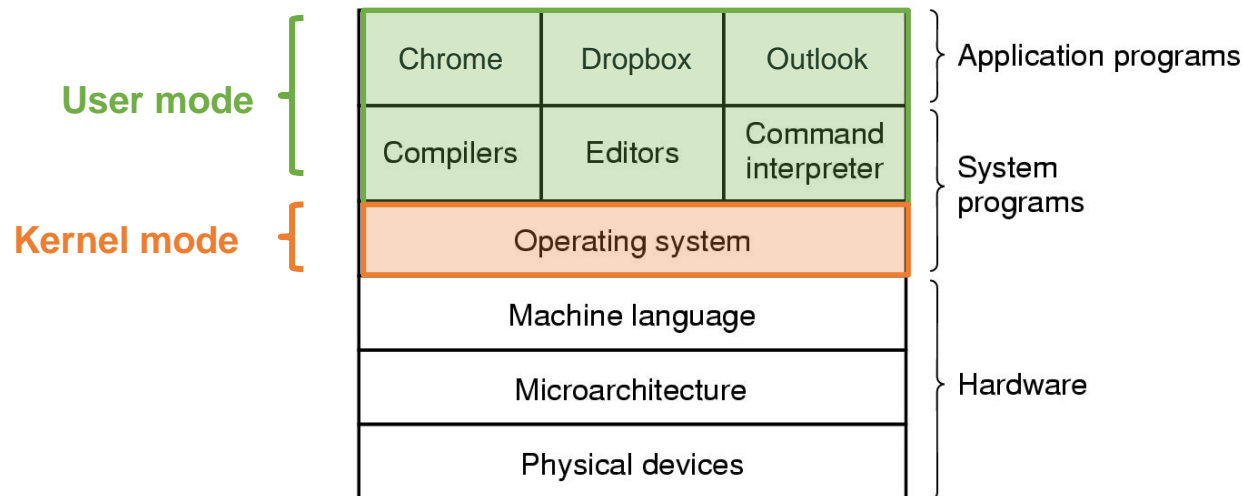


# How OS and application/devices interact?



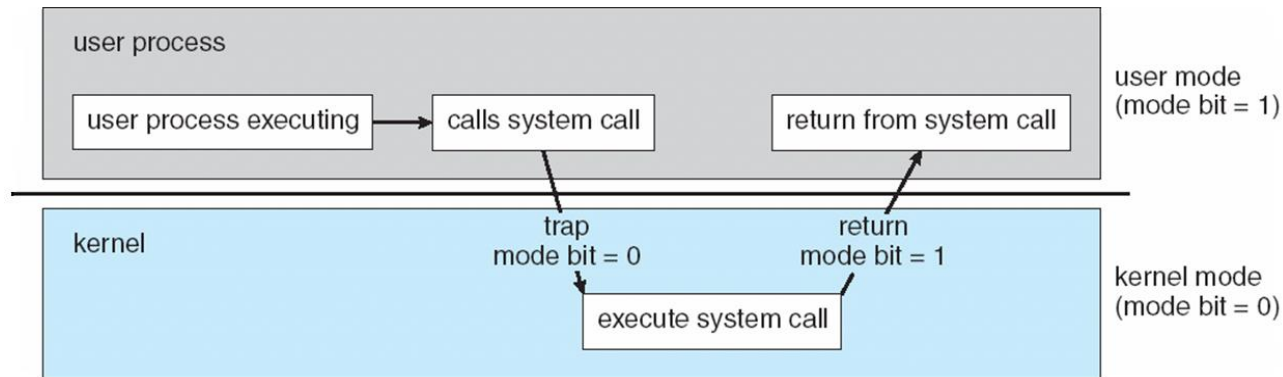
# OS: Dual-mode operation

- **Different privileges** required **for different types of code**
  - **User code** is executed in **user mode**
  - **Operating system** is executed in **kernel mode**
  - Most instructions can be executed in **user mode**, but not all
    - For instance, in user mode, the CPU cannot access memory locations reserved for OS
  - Some “**privileged**” instructions are only executable in **kernel mode**



Dual-mode allows the OS to **protect** itself and **other system components**

- **Mode bit** provided by hardware
  - Provides ability to **distinguish** when system is **running user code or kernel code**
  - **A system call changes the execution mode** to kernel mode, returning from the system call resets the mode to user mode



## Example:

**read data from file using** *read(fd, buffer, nbytes)*

Reads *nbytes* bytes from *fd*, storing it in *buffer*:

1. **push parameters** needed by '*read(.,.,.)*' (pointer to *fd*, buffer address, number of bytes) **on the stack, in registers or in pre-defined place in memory**  
**write a code identifying '*read(.,.,.)*' in a register** so that the trap handler knows that it must **call the code for *read***
2. **generate a trap**: switch mode & call trap handler
3. Trap handler calls and execute the *read(.,.,.)* function handler
4. Return in user mode and **give control back to the caller**

# Example: read data from file using *read(fd, buffer, nbytes)*

Reads *nbytes* bytes from *fd*, storing it in *buffer*:

User mode

1. push parameters needed by '*read(.,.,.)*' (pointer to *fd*, buffer address, number of bytes) on the stack, in registers or in pre-defined place in memory
2. write a code identifying '*read(.,.,.)*' in a register so that the trap handler knows that it must call the code for *read*

3. generate a trap: switch mode & call trap handler

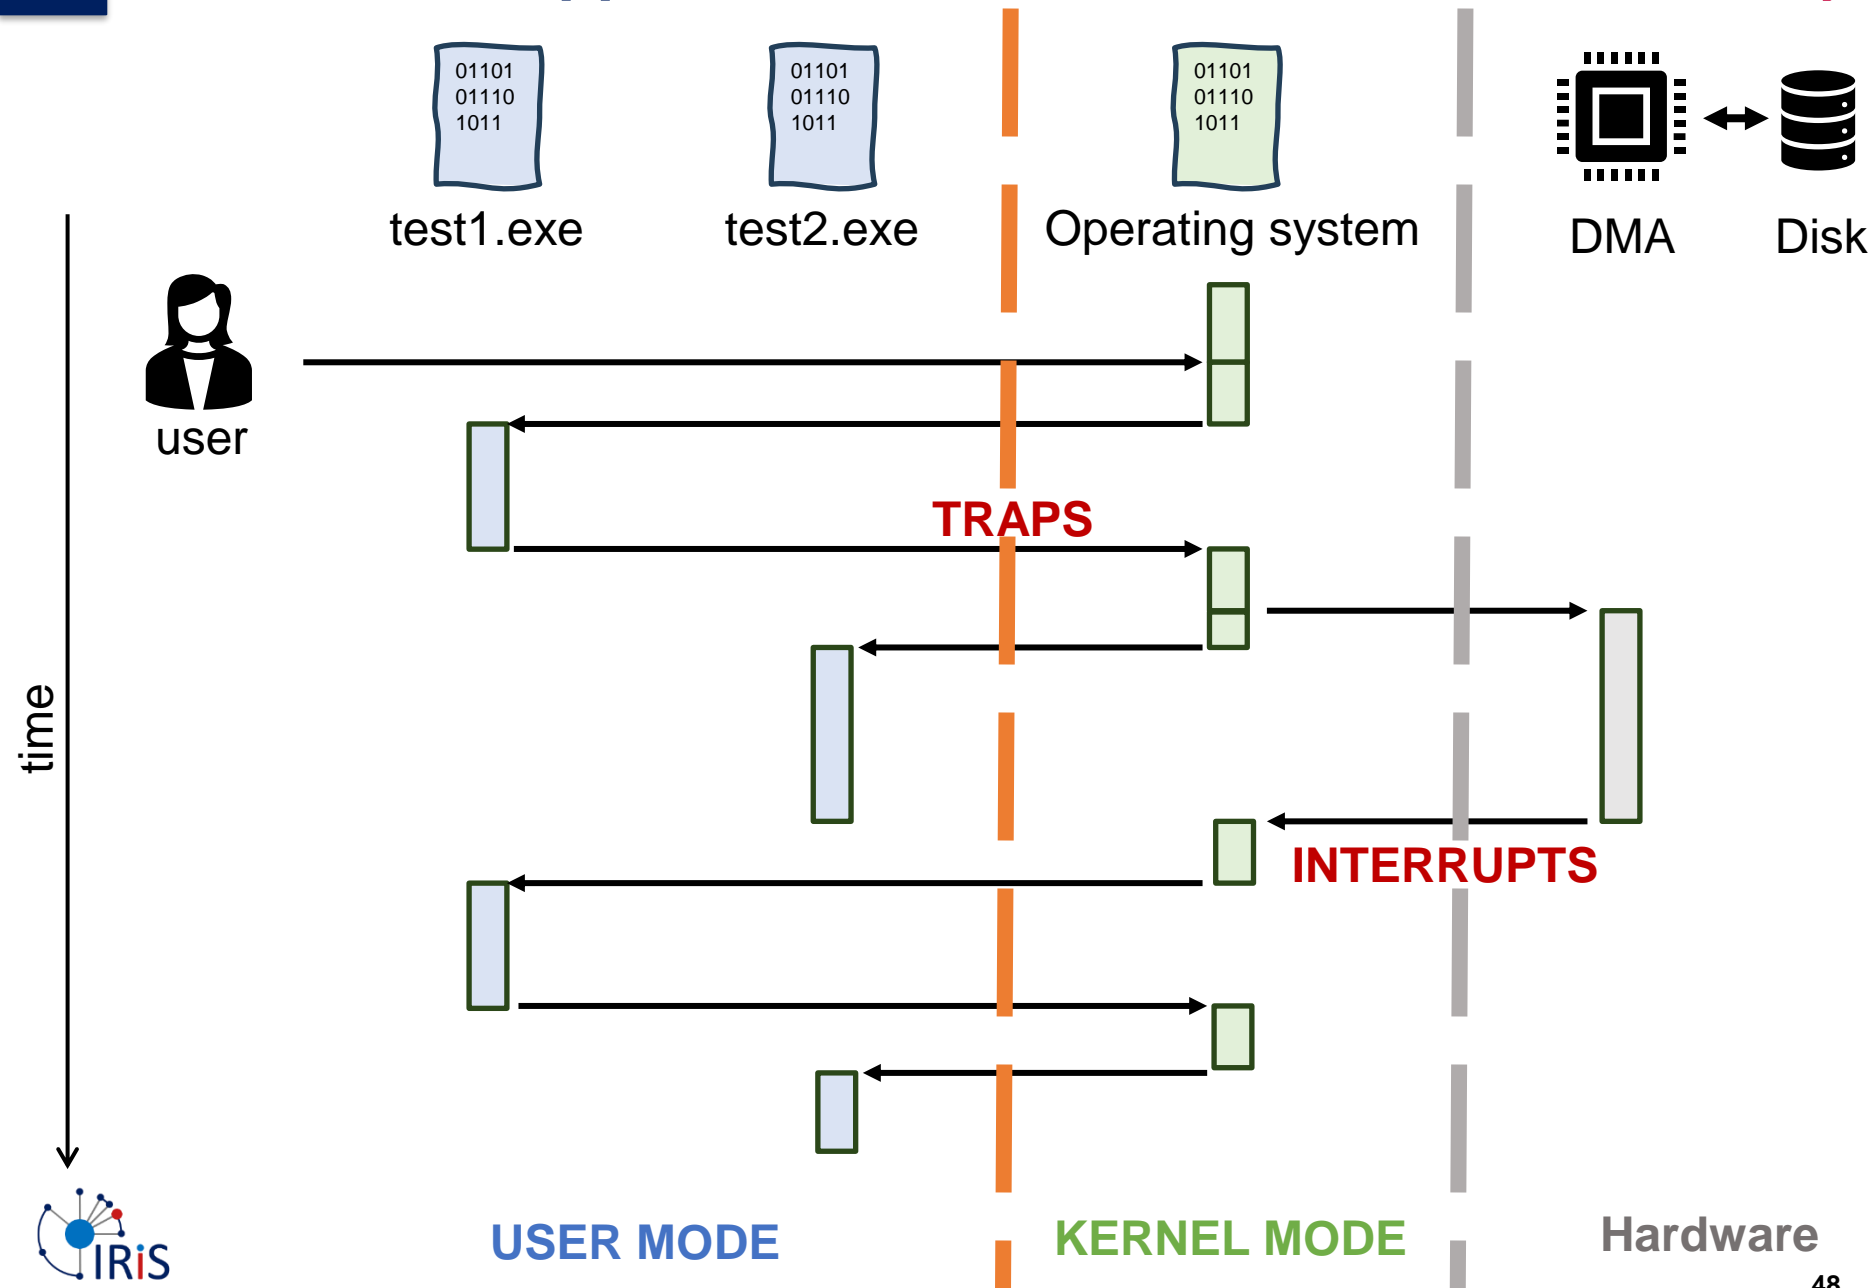
Kernel mode

4. Trap handler calls and execute *read(.,.,.)* function handler

5. Return in user mode and give control back to the caller

User mode

# How OS and application/devices interact?







- Course overview
- OS: place in the computer system
- **Motivation & OS tasks**
- Extra-functional requirements

- Deal with diversity
- Transparency
- Virtualization
- Support for shared functionality
- Portability
- Support for concurrency


- Many system implementations
  - **diverse CPUs** (Instruction Set Architectures)
  - **diverse architectures and organization**



Multicore system




Clustered system




Distributed system


- **Hide details** with respect to a given issue.
- Examples:
  - **processor architecture** (Instruction Set Architecture - ISA)
    - mechanism: use **compiler**
  - physical **memory size**
    - mechanism: **virtual memory** i.e., present linear memory model that is larger than physical
  - **Location** of programme and data **in physical memory**
    - mechanism: **indirection** using **logical memory address**




Reality




OS layer



What the programmer sees

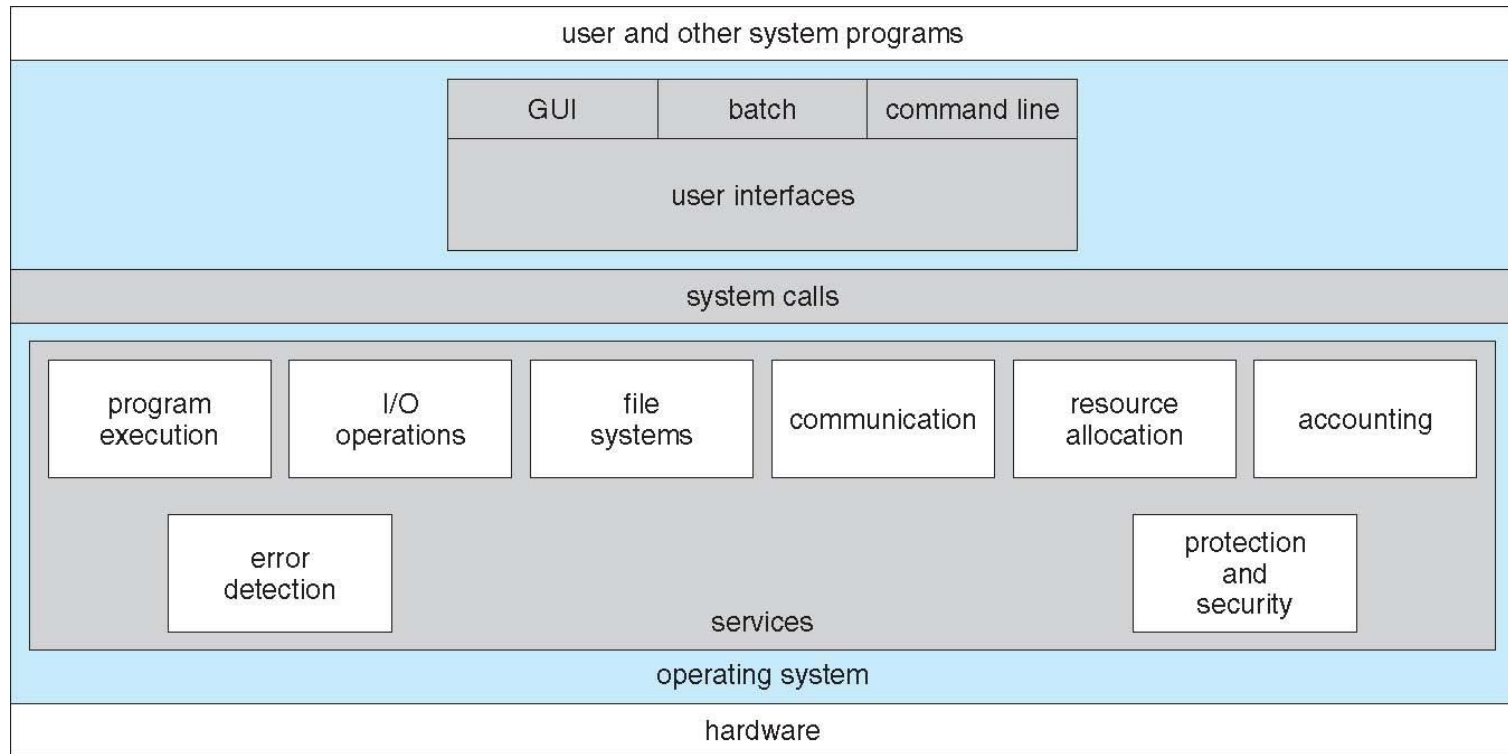


- Provides a **simple, abstract, logical model** of the system
  - virtual memory, virtual CPU, virtual disk
  - **Program can be developed as if they were the sole user of the HW resources**
- Current systems tend to **virtualize the entire hardware**
  - i.e., it yields a **virtual machine**



- OS task: provide *functionality common to most programs*
  - introduce **well-defined *abstractions*** of *concepts*
    - files and file systems instead of disk blocks
    - read from or write to I/O devices as if writing in a file instead of manipulating registers in device controller
    - exceptions and traps rather than ‘something goes wrong’
    - linear memory rather than memory blocks, pages and disk space
- provide **system calls** to access the functions provided by the OS

- **System calls:** Programming interface to the services provided by the OS,...
- ...typically written in a high-level language (C or C++)



- **Types of system calls:**
  - **Process management**
    - create, destroy, communication, synchronization, ...
  - **File management**
    - open, close, read, write, ...
  - **Memory management**
    - allocation, free, share memory between processes, ...
  - **Device management**
    - access control, open, attach, send/receive, ....
  - **Communications**
    - setup communications, exchange messages, ....
  - **Miscellaneous**
    - timers, inspect system resources, ...



# OS Motivation: Support for shared functionality

## API vs. System calls

- System calls are mostly accessed by programs via a high-level **Application Programming Interface (API)**
  - [Win32](#) API for Windows, [POSIX](#) API for POSIX-based systems (including virtually all versions of UNIX, Linux, and Mac OS X),
  - [Java API](#) for the Java Virtual Machine (JVM),
  - [Cocoa Touch](#) for iOS, Java based [Android API](#) for Android (runs on top of Android Runtime – previously on top of Dalvik VM).


### Advantage of an API over system calls

- API works at a [higher abstraction](#) closer to programs need → **easier** to work with
- API calls may [combine many system calls](#) to implement higher level tasks (e.g., copy a file) → **less calls** needed
- a program written using an API can [compile/run on any system supporting the API](#) → **portable**

### Disadvantage of an API to system calls

- Adds one more layer → **increased overhead**

- a program written using an OS or an API can **compile/run on any system supporting that OS or API**
- Protect investments in application software
  - Reduces cost to **support several platforms**
  - Reduces cost to **upgrade/change** the execution platform
  - **Simplifies integration** of several application components
- OS gives a ***unified machine view*** to applications, effectively defining a ***virtual machine***



- **The machine must be *shared*** between multiple activities ('tasks', 'processes') and multiple users
- OS:
  - realizes **concurrency transparency** (virtualization)
    - each task virtually has a (virtual) machine of its own
  - **manages and protects** tasks from each others
    - Limits resource usage between tasks *and* between users
    - e.g., processor, memory, files, i/o equipment, ....
  - **ensures efficiency** by scheduling, e.g.,
    - to avoid starvation;
    - to enforce timing requirements;
    - to reduce power consumption;
    - ...

- **Course overview**
- **OS: place in the computer system**
- **Motivation and OS tasks**
- **Extra-functional requirements**

- **Extensibility**

- support for adding application-specific (domain-specific) functionalities

- **Scalability**

- wide range of environments, functionalities, machines
- Scale to large-enough number of tasks, cpus, clusters, ...

- **Dependability**

- robust, correct, safe and secure
- How dependable? → depends... on the application domain

- **Real-time OS**

- predictability
  - *known performance* instead of *high performance* (all resources)
  - *Known bounds* on the maximum *latencies* (response times) of API calls
- support for dealing with real-time control
  - *predictable scheduling policies*
  - *explicit control over resources*
  - *real-time facilities: clocks and timers*
- stringent dependability

- **Embedded OS**

- small footprint (e.g., leave all superfluous parts out)
- low system requirements (e.g. processor speed, energy)
- stringent dependability

## MOST POPULAR REAL-TIME OPERATING SYSTEMS

(2020)

- Deos (DDC-I) Compliant with the **safety standard** DO-178 required in **avionics** (spin-out from Honeywell)
- embOS (SEGGER) Used in automotive industry (it is free for non-commercial use)
- FreeRTOS (owned by Amazon since 2017) It is the second most popular/used operating system for embedded systems (according to a study by [www.embedded.com](https://www.embedded.com), 2019)
- Integrity (Green Hills Software)
- Keil RTX (ARM) Only supports ARM-Cortex M
- LynxOS (Lynx Software Technologies)
- MQX (Philips NXP / Freescale) It is also the kernel of QNX
- Nucleus (Mentor Graphics)
- Neutrino (BlackBerry)
- PikeOS (Sysgo) Small company that is now part of Thales
- SafeRTOS (Wittenstein)
- ThreadX (Microsoft Express Logic)
- µC/OS (Micrium) Priority-based preemptive real-time kernel for microcontrollers
- VxWorks (Wind River)
- Zephyr (Linux Foundation)

microkernel

Source: <https://www.embedded.com/2019-embedded-markets-study-reflects-emerging-technologies-continued-c-c-dominance/>

Most Used	World	Americas	Europe	Asia
Embedded Linux	31%	32%	31%	<b>26%</b>
FreeRTOS	27%	25%	24%	<b>37%</b>
Android	14%	12%	<b>10%</b>	<b>26%</b>

- Operating system
  - is an **interface** between HW and applications
  - **abstracts the HW** by providing high level services
  - **manages concurrent access** to HW devices
  - **optimizes efficiency**
- The operating system **protects itself and other components** using **two modes of executions** (user vs kernel mode)
- **HW** devices **communicate** using **interrupts**
- **SW** applications **use system calls** to perform **privileged operations**



- *Abstraction*

- provide useful generic concepts/functionalities
- .... to handle complexity (transparency)

- *Virtualization*

- provide the same abstract model for a wide range of systems
- each process/user sees single machine, linear memory, ...

- *Resource management*

- resource sharing
- optimize performance
- access control for protection and security
- accounting and usage limit enforcement

- Friday at 8h45
- How to implement concurrency in an OS
  - Quiz
  - Exercises
- Prepare!
  - 11 pages to read in the reference book
  - 25 minutes of video