

Operating Systems (2INC0)

Process (02)

Active program execution

Dr. Geoffrey Nelissen

Courtesy of Dr. Tanir Ozcelebi



Interconnected
Resource-aware
Intelligent Systems



Technische Universiteit
Eindhoven
University of Technology

Where innovation starts

Process concept

- A program is a passive entity
- A process is a ***program in execution***
 - Several processes can run the same program
 - User processes run in user mode

Process Control Block (PCB) – inside kernel

A process is identified by an ID (number) and a pointer to a PCB in kernel space

PCB contains:


- Process state (see next slide)
- Process number (ID)
- Program counter (needed for context switching)
- CPU registers (needed for context switching)
- Memory management info (values of base, limit registers)
- I/O status information (list of open files, I/O devices)
- Scheduling information (priority, scheduling parameters)
- Accounting information (e.g. CPU time used)
- and more...

PCB

process state
process number
program counter
registers
memory limits
list of open files
...

Process states

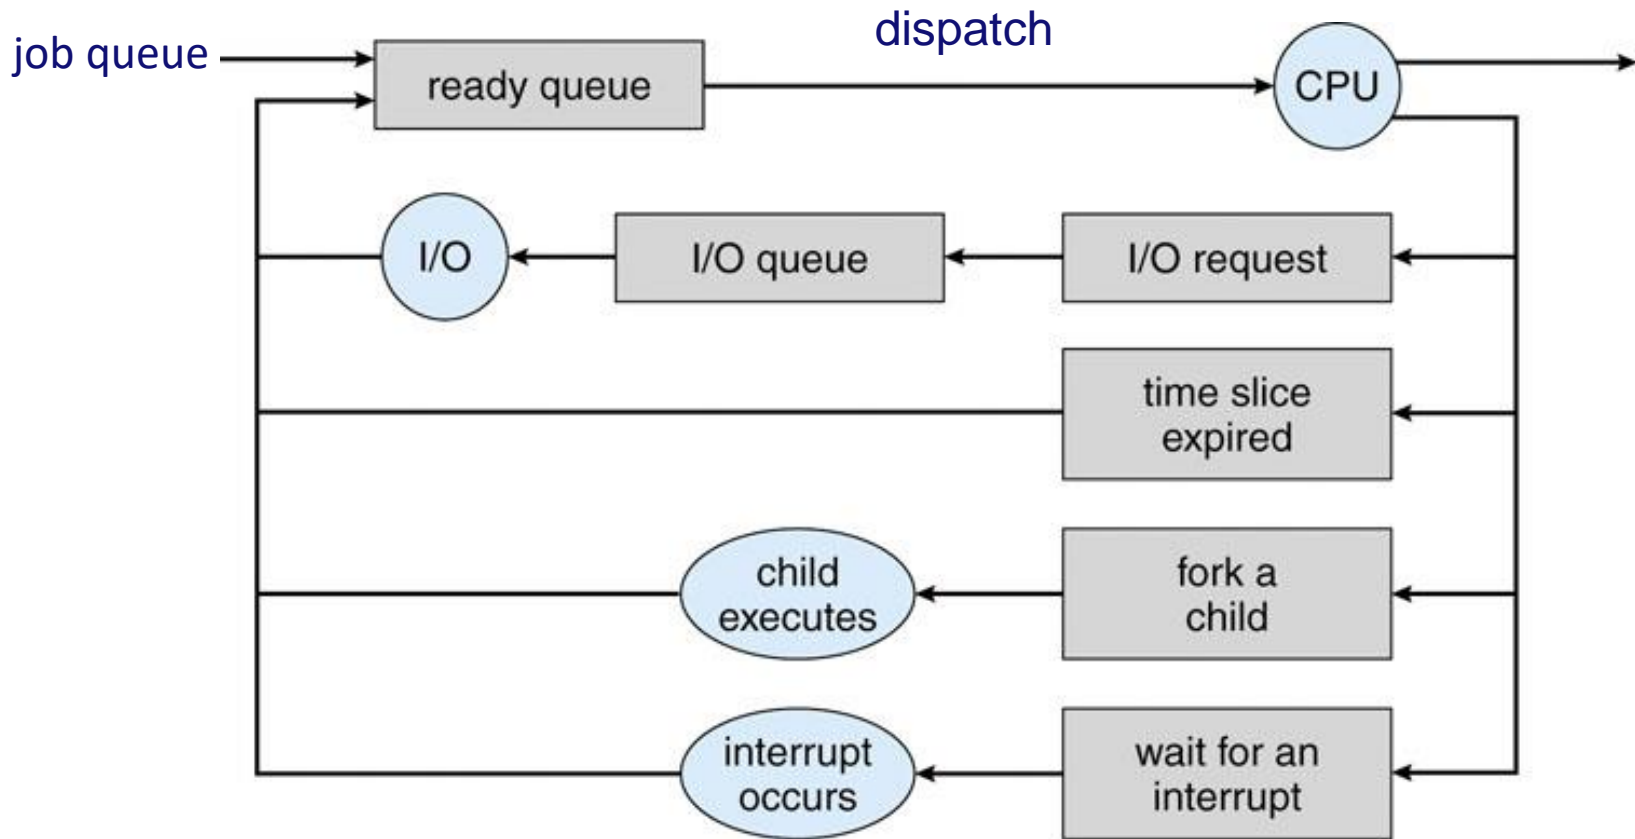
- **new**: process is being created
- **ready**: process is waiting to be assigned to a CPU (before **scheduler dispatch**)
- **running**: instructions are being executed
- **waiting**: process is waiting for some event to occur
- **terminated**: process has finished execution




Only one process
can run on a CPU
at any time instant!
→ **must schedule**

Process scheduling

Ready queue and device (I/O) queues


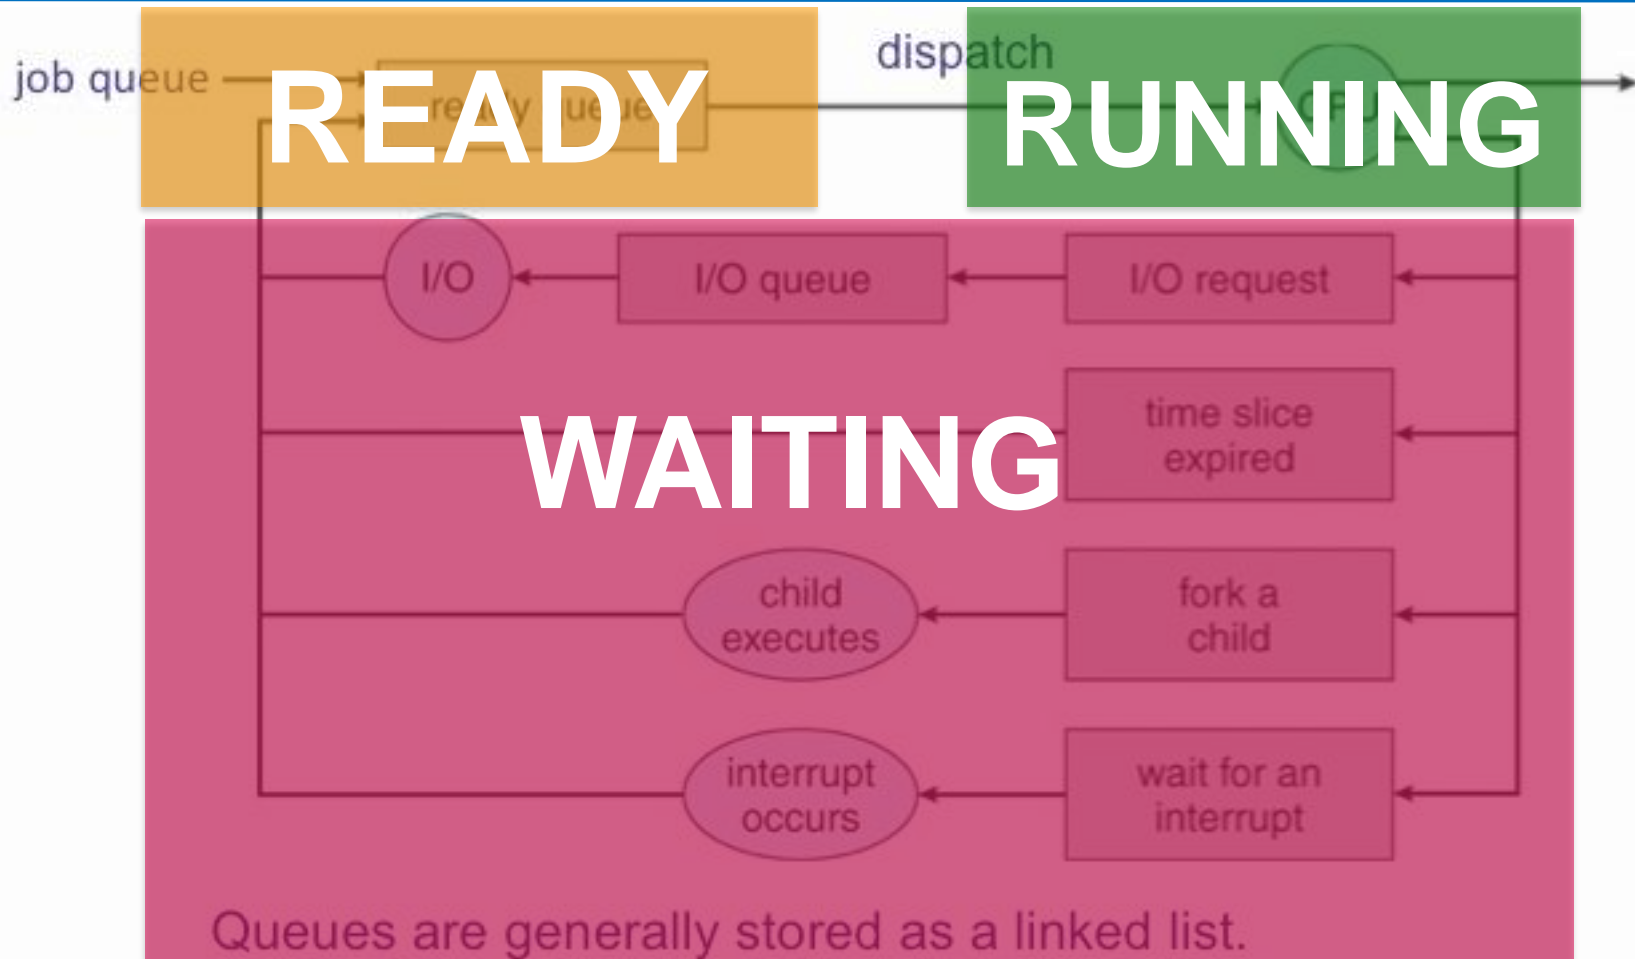


Queues are generally stored as linked lists.



Process scheduling

Ready queue and device (I/O) queues



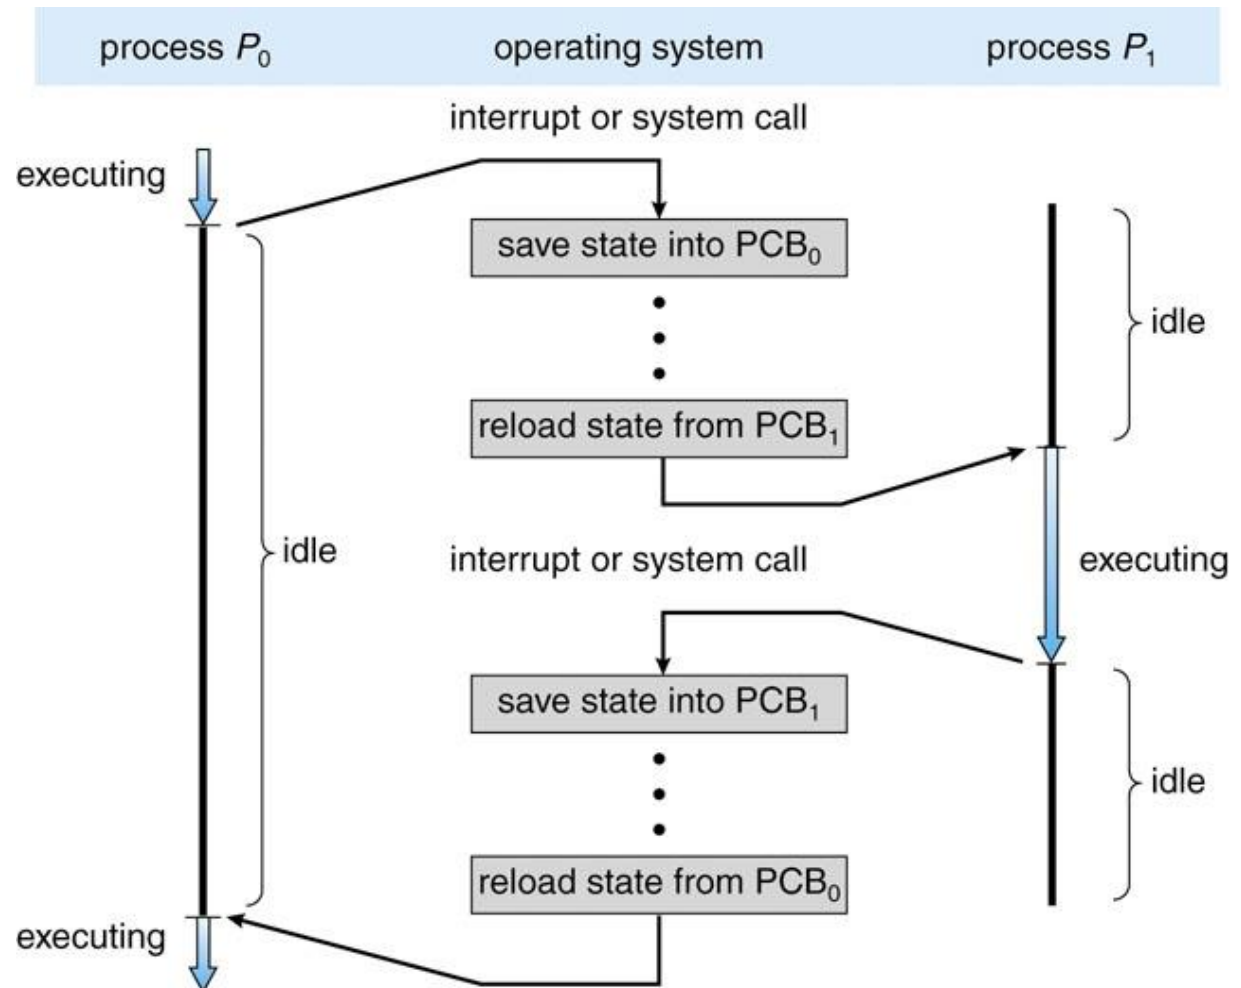
Context switch

Context switch:

Saving the state of a process whose execution is to be suspended, and reloading this state when the process is to be resumed.

Overhead:

(typically) takes a few usecs. → depends on HW



Operating Systems (2INC0)

Process (02)

Process creation/termination

Dr. Geoffrey Nelissen

Courtesy of Dr. Tanir Ozcelebi



Interconnected
Resource-aware
Intelligent Systems

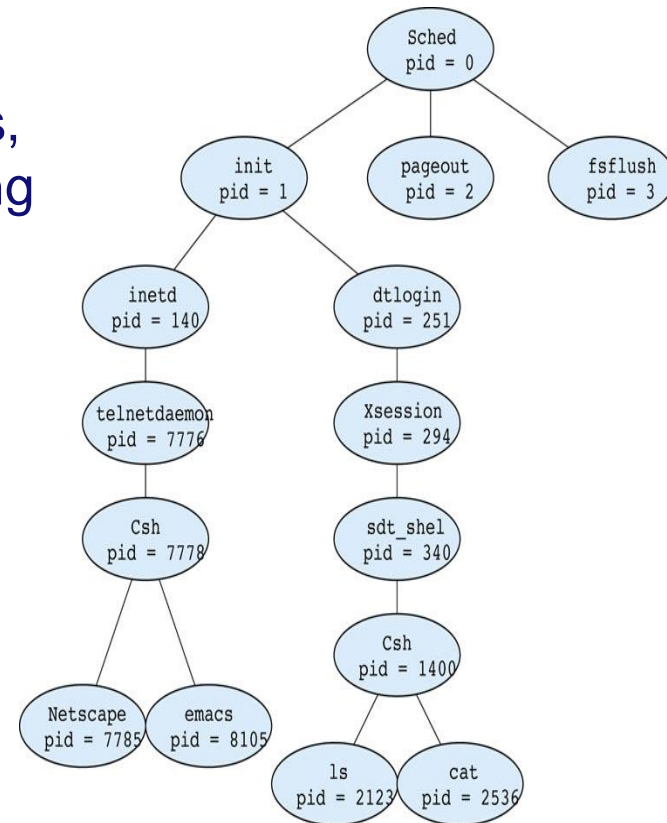
TU/e

Technische Universiteit
Eindhoven
University of Technology

Where innovation starts

Process Creation

- Starting point is an already running process
- **Parent** process **creates children** processes, which, in turn create other processes, forming a **tree of processes**
 - new processes given process identifiers (pid)
- Resource sharing options
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources



Process Creation (cont'd)

- Execution options
 - The parent and children **execute concurrently**
 - The parent **waits** until children terminate
- Address space options
 - The child is a **copy** of the parent (e.g., default in Linux)
 - The child has a **new program** loaded into its address space (e.g., Windows)

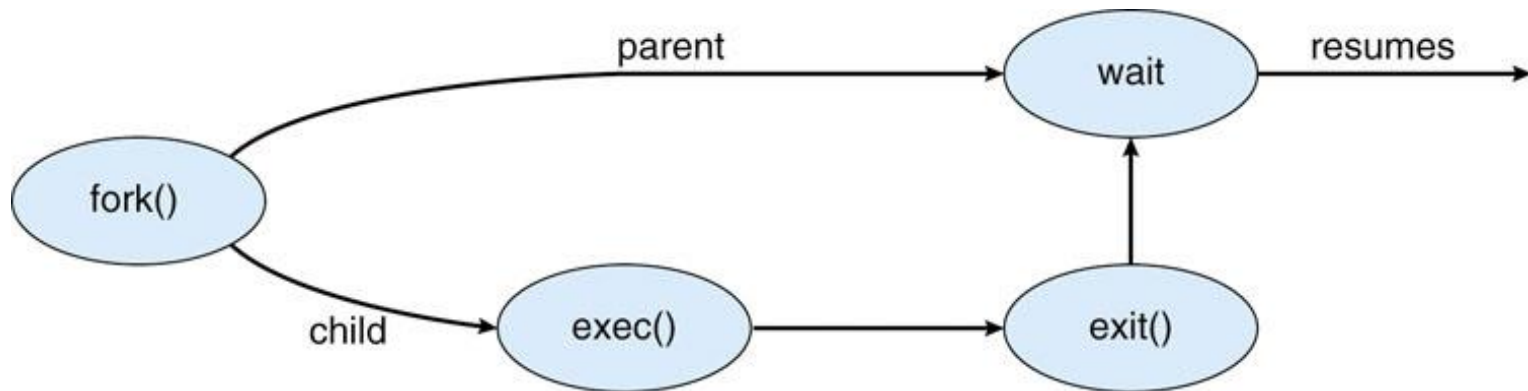
Example:

Using the Portable Operating System Interface (POSIX) API

- **IEEE standard on the API** (it's not an OS!)
 - Goal: reduce **portability** effort for applications
 - Many operating systems use the POSIX API (or a subset).
- Implementations may vary ...
- ... but a system supporting POSIX must provide
 - a host language and compiler (often: C)
 - programming interface definition files (e.g., C-header files)
 - programming interface implementation binary or code (e.g., C-libraries)
 - a run-time system (a platform: OS or the like)

POSIX processes (1003.1)


- POSIX 1003.1 API: *fork()* – *exec()* – *exit()* – *wait()*




- *fork()* function **creates** new process (duplicate address space of the parent in another location)
- *exec()* function used after a fork to write over the process' memory space with a **new program**
- *wait()* function may be issued by the parent to **wait** until the execution of the **child is finished**.

Before and after a *fork()*

Parent process about to perform *child = fork()*



Parent and child process after *fork()*
child is a literal copy of parent



variable *child* of child process equals 0

variable *child* in parent points to child process

Process Termination

- Process may ask the operating system to delete itself (**exit**)...
 - return status value from child to parent (via wait)
 - the process' resources are taken away by OS
- ... or parent may terminate execution of children processes (**abort**)
 - The child has **exceeded allocated resources**
 - The task assigned to child is **no longer required**
 - If the **parent is exiting**
 - Some operating systems do not allow the child to continue if its parent terminates, in which case all children are terminated - **cascading termination**
 - Some operating systems **attach orphans to a 'grandfather' process** (e.g. *init*)

Termination of children: fragment

- Use `exit(status)` to terminate
- In some OSs, the parent **must wait** for children **to free children's resources**
- Functions `wait()`, `waitpid()`
 - `wait()` blocks until any child exit
 - `waitpid()` blocks until a specific child changes its state
- Alternative (not shown here): use asynchronous notification signals

parent:

```
...
pid_t child, terminated; int status;
...
/* blocking wait */
while (child != wait (&status)) /* nothing */;

/* or polling wait */
terminated = (pid_t) 0;
while (terminated != child) {
    terminated = waitpid (child, &status,
        WNOHANG);
    /* other useful activities */
}
/* both cases: status == 23 */
```

child:

```
..... exit (23);
```

Operating Systems (2INC0)

Process (02)

Interprocess communication

Dr. Geoffrey Nelissen

Courtesy of Dr. Tanir Ozcelebi



Interconnected
Resource-aware
Intelligent Systems


TU/e

Technische Universiteit
Eindhoven
University of Technology

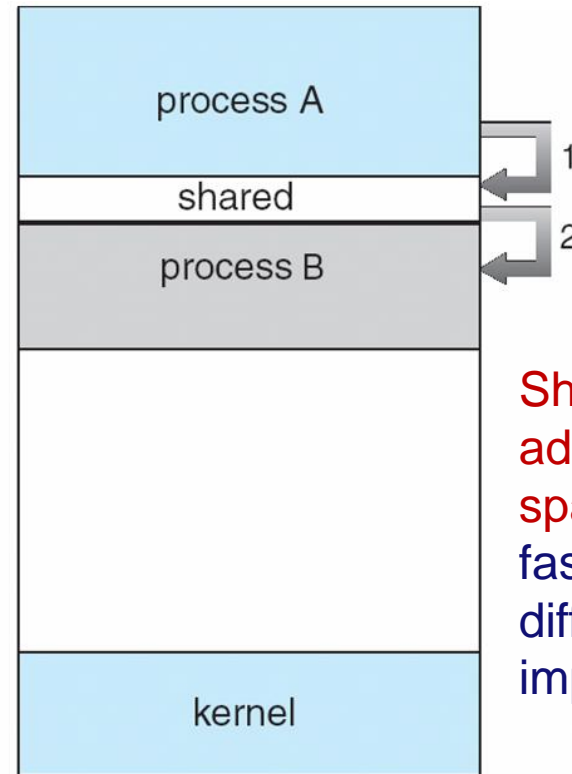
Where innovation starts

Cooperating processes need interprocess communication (IPC)

No shared address space:
useful for exchanging small amounts of data



(a)
Message passing



(b)
Shared memory

Shared address space:
faster but difficult to implement

Message passing between processes

- The **easiest** solution especially for **distributed systems**
- The **producer(s) send(s)** messages through a pipe or in a mailbox
 - Can be blocking or non-blocking
- The **consumer(s) read(s)** messages
 - Can be blocking or non-blocking
- If **both send and receive are blocking** then we call it a *rendez-vous*
 - communicating processes synchronize on the transmission
- A temporary queue (**buffer**) can be associated to the communication
 - Allows for **asynchronous executions** and/or applications with **multiple producers/consumers**
 - The producer may or may not be blocked when the buffer is full
 - If non-blocking, the consumers see the **most recent data** and may **miss old ones**(e.g., control systems)

Shared memory between processes

- A memory segment is reserved in the kernel memory
- It is **added to the space addressable** in user mode
- *No kernel involvement is required upon subsequent reading and writing*

