

2INC0 - Operating Systems

Processes

Dr. Geoffrey Nelissen



Interconnected
Resource-aware
Intelligent Systems

TU/e

Technische Universiteit
Eindhoven
University of Technology

Where innovation starts

- **Introduction to operating systems** (lecture 1)
- **Processes, threads and scheduling** (lectures 2 and 3)
- **Concurrency and synchronization**
 - atomicity and interference (lecture 4)
 - action synchronization (lecture 5)
 - condition synchronization (lecture 6)
 - deadlock (lecture 7)
- **File systems** (lectures 8)
- **Memory management** (lectures 9 and 10)
- **Input/output** (lecture 11)

- **First practical assignment** will be **released today**.
 - **Deadline in 2.5 weeks**
 - There is a question in the assignment asking to **show that the program may reach a deadlock**. This **topic is covered next Friday**.
 - Possibility to get a **bonus point** by submitting a second solution avoiding the deadlock situation.
- **First homework** assignment will be **released today**.
Deadline Sunday of next week.
 - We redesigned the homework based on previous years students' feedback.
 - Homework requires to solve exercises similar to those done in today's lecture.
- Check that you are **registered to a group** for the practical assignments (groups of 3 students)

- **Reminder of lecture 1 (quiz)**
- **Processes**

Reminder of lecture 1

- Connect to www.menti.com
- Code **7553 5681**

Which of the following is (are) motivation(s) for the existence of operating systems?

1. Improve portability
2. Manage concurrency
3. Reduce execution overhead
4. Provide a simplified view of the execution platform
5. Provide support for shared functionality
6. Specify a unified hardware specification
7. Accelerate memory and I/O accesses

Question 1

Which of the following is (are) motivation(s) for the existence of operating systems?

1. Improves portability
2. Manages concurrency
3. Reduces execution overhead
4. Provides a simplified view of the execution platform
5. Provides support for shared functionality
6. Specifies a unified hardware specification
7. Accelerates memory and I/O accesses

(green = correct, red = incorrect)

Question 2

Which of the following is/are true about DMA (Direct Memory Access)

1. A DMA is a memory access protocol
2. A DMA virtualizes memory
3. A DMA is part of the operating system
4. A DMA replaces I/O controllers
5. A DMA is a hardware component

Question 2

Which of the following is/are true about DMA (Direct Memory Access)

1. A DMA **is a memory access protocol**
2. A DMA **virtualizes memory**
3. A DMA **is part of the operating system**
4. A DMA **replaces I/O controllers**
5. A DMA **is a hardware component**

(green = correct, red = incorrect)

Question 3

Which of the following is/are true?

1. A DMA increases the work done by the CPU
2. A DMA may impact negatively the execution time of a program

Question 3

Which of the following is/are true?

1. A DMA increases the work done by the CPU
2. A DMA may impact negatively the execution time of a program

Question 5

Which of the following is/are true about interrupts and traps?

1. Upon an interrupt, the CPU enters into kernel mode.
2. An interrupt is a signal from the CPU to a device controller, informing it that the CPU is ready to perform an I/O.
3. Traps are used to transfer data from a file to a buffer.
4. System calls are implemented using traps.

(green = correct, red = incorrect)

Question 5

Which of the following is/are true about interrupts and traps?

1. Upon an interrupt, the CPU enters into kernel mode.
2. An interrupt is a signal from the CPU to a device controller, informing it that the CPU is ready to perform an I/O.
3. Traps are used to transfer data from a file to a buffer.
4. System calls are implemented using traps.

(green = correct, red = incorrect)



- Reminder of lecture 1 (quiz)
- **Processes**

Process


- It is a **program in execution**.
- It has a **context of execution**
- A process owns resources (has memory and can own other resources such as files, etc.)

Thread

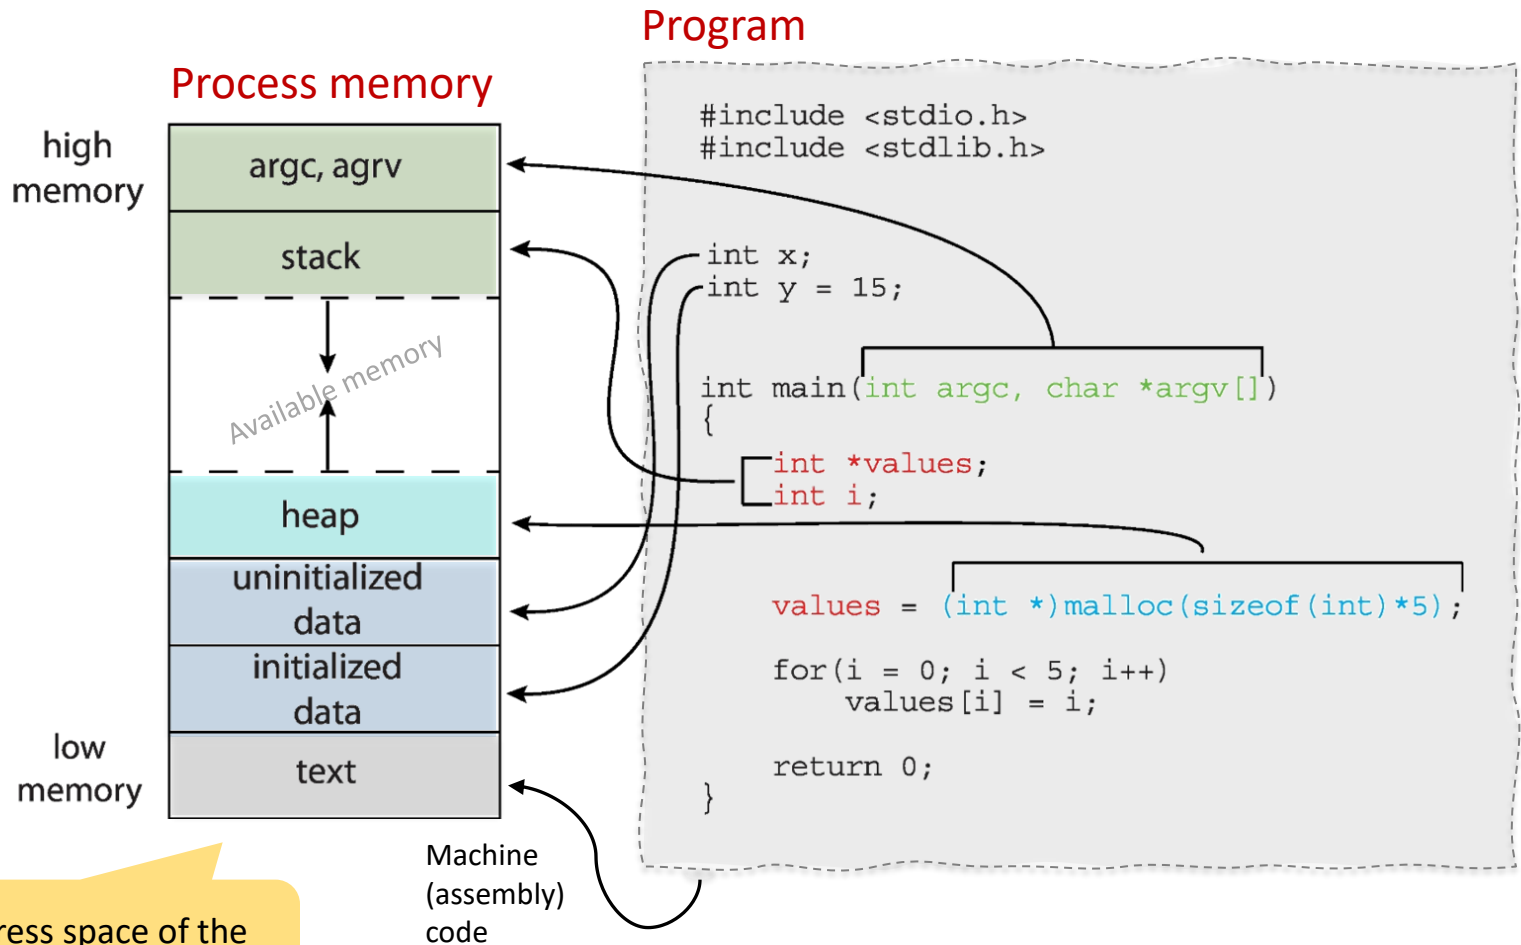
- A **dispatchable unit of work within a process**
- Threads within a process share code and data segments (i.e., **share memory address space**)



- A process is *a program in execution*
 - Several processes can run the same program
 - They all have a different context of execution
- A process is defined by:
 - Text (program code)
 - Stack
 - e.g. local variables, function parameters, return addresses (their sizes known at compile time).
 - heap
 - used to store **data whose size may be unknown** before runtime, e.g. used by a program that reads a file whose size is unknown or uses malloc.
 - data (global variables whose size are known)
 - Information about the current (latest) state of execution

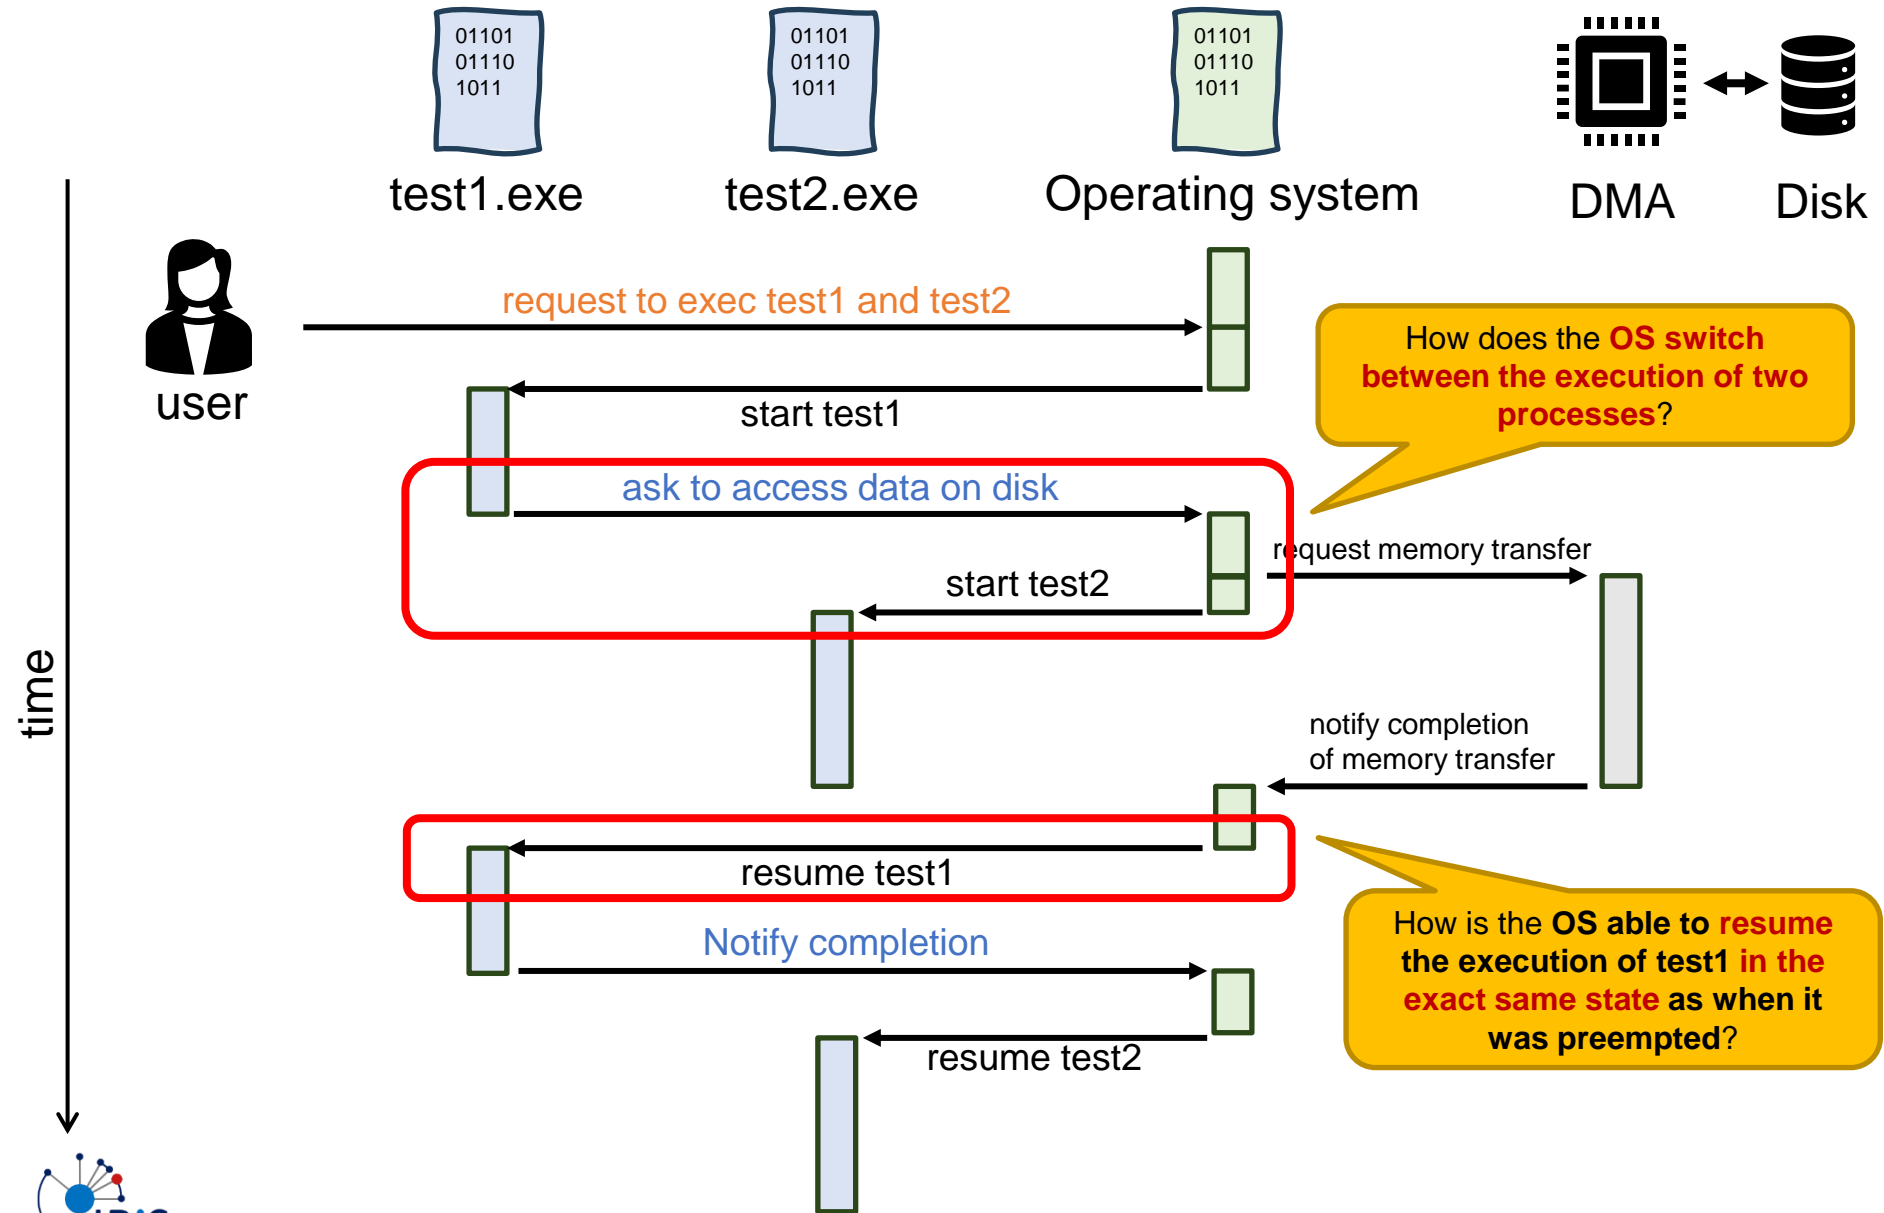


Memory layout of a program in a process



Memory address space of the process running this program

Remember this example?



Process control block (PCB)

The **PCB records** information associated with a **process context of execution**

- **Process state** – running, waiting, etc.
- **Program counter** – location of next instruction to execute
- **CPU registers** – contents of all process-centric registers
- **CPU scheduling information** – priorities, scheduling queue pointers
- **Memory-management information** – memory allocated to the process
- **Accounting information** – CPU used, clock time elapsed since start, time limits
- **I/O status information** – I/O devices allocated to process, list of open files

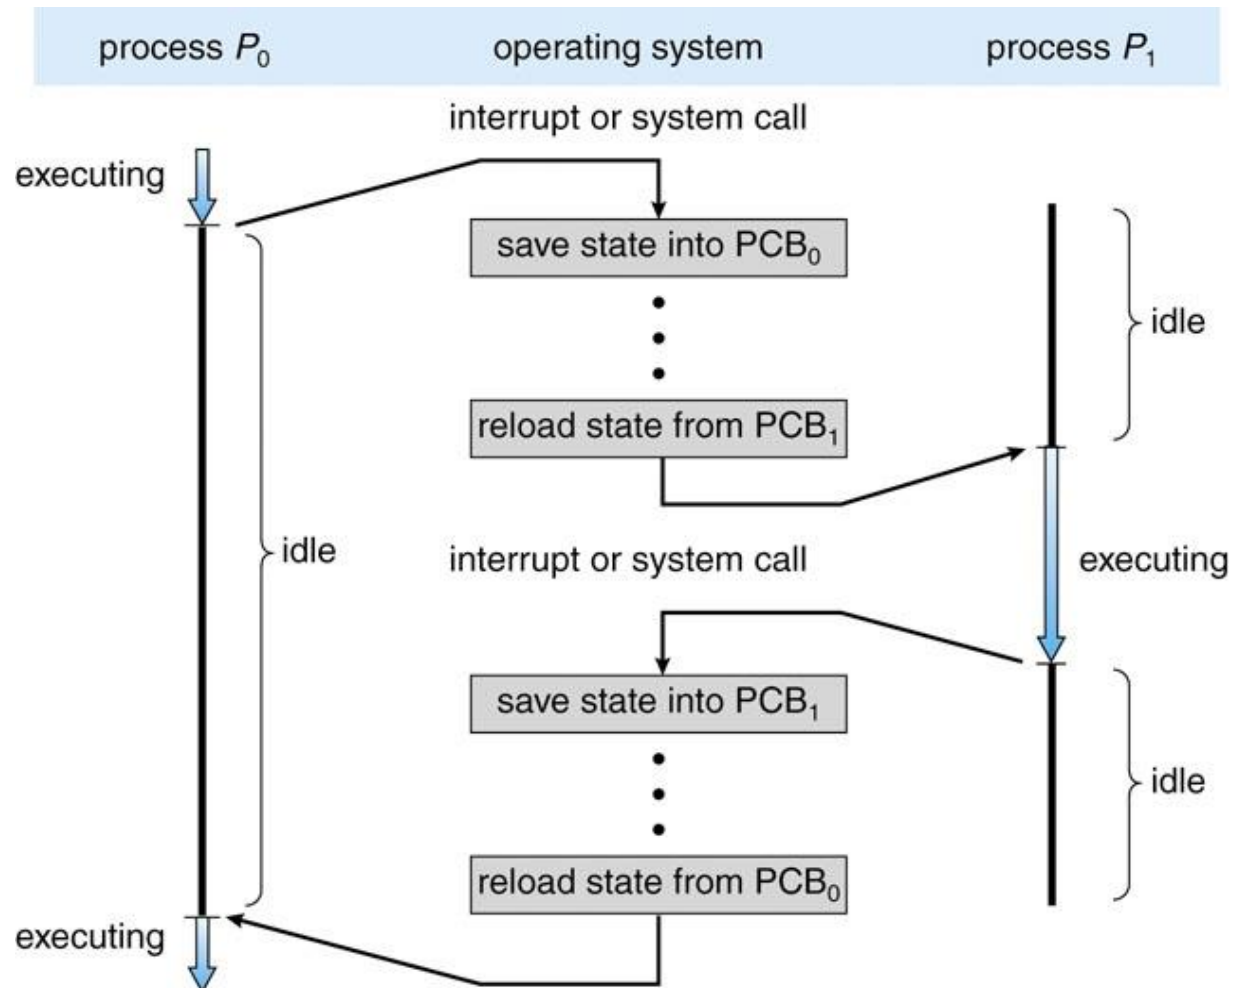
process state
process number
program counter
registers
memory limits
list of open files
• • •

Context switch:

Saving the state of a process whose execution is to be suspended, and reloading this state when the process is to be resumed.

Overhead:

(typically) takes a few usecs. → depends on HW



How can we use processes?

- **Parent** process creates **children** processes, which, in turn creates other processes, forming a **tree** of processes
- Generally, process identified and managed via a **process identifier (pid)**

When a process creates children, it can share its memory or resources (such as files) with them.

Resource sharing options

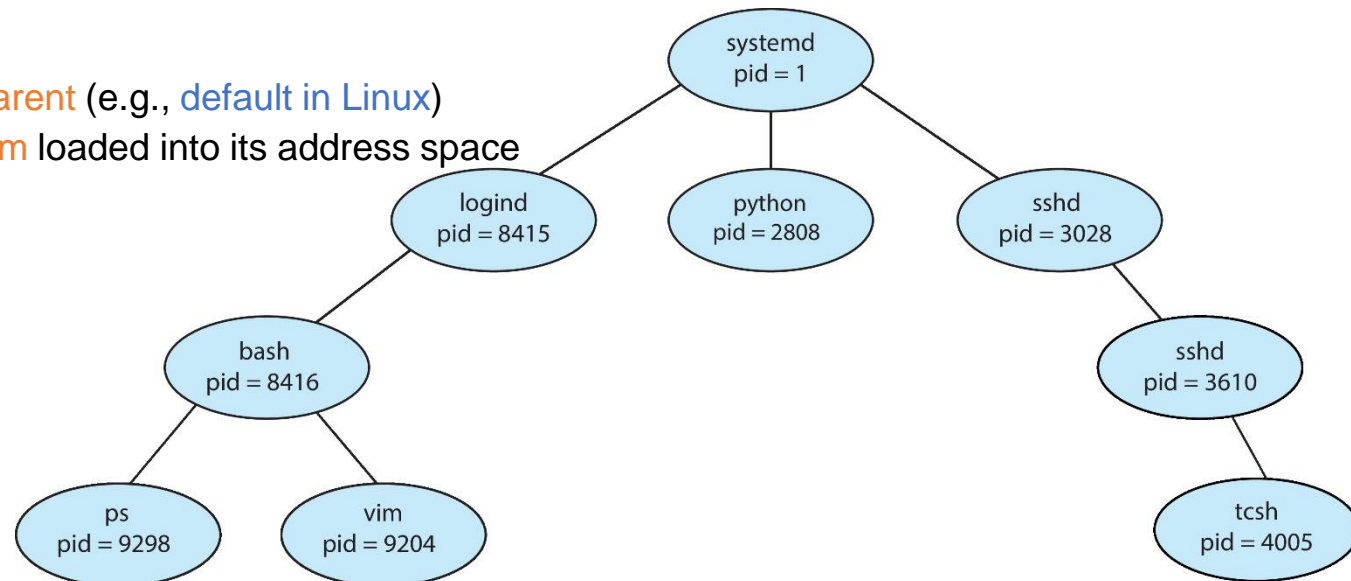
- Parent and children share all resources
- Children share subset of parent's resources
- Parent and child share no resources

Execution options

- Parent and children execute concurrently
- Parent waits until children **terminate**

Address space options

- The **child** is a **copy** of the **parent** (e.g., **default in Linux**)
- The child has a **new program** loaded into its address space (e.g., **Windows**)



Example from Linux

- IEEE standard on the API (POSIX is not an OS)
 - Goal: reduce portability effort for applications
 - Many operating systems use POSIX API.
- A system supporting POSIX provides
 - a host language and compiler (often: C)
 - programming interface definition files (e.g., C-header files)
 - programming interface implementation binary or code (e.g., C-libraries)
 - a run-time system (a platform: OS or the like)

We will use the POSIX API to create and terminate processes and threads in this lecture

Example: create new processes



```
pid_t child;
.....
child = fork();
if (child < 0) /* error occurred */ {
    perror ("fork"); exit (-1); }

if (child == 0) /* the child */ {
    execlp ("/bin/lis", "lis", arg0, arg1, ..., NULL);
    /* this place is reached only in case of error */
    perror ("execlp"); exit (-1);
}
else /* the parent; child == process id of child */ {
    /* do whatever you want, e.g., just return from
    this routine */
    int status;
    wait (&status);
}
.....
```


Example: create new processes

fork() creates an identical copy of the caller in a separate memory space;



Only the return value stored in variable '*child*' differs between the two: **0 for the child**, and **child-pid** for the parent

```
pid_t child;
.....
child = fork();
if (child < 0) /* error occurred */ {
    perror ("fork"); exit (-1); }

if (child == 0) /* the child */ {
    execlp ("/bin/ls", "ls", arg0, arg1, ..., NULL);
    /* this place is reached only in case of error */
    perror ("execlp"); exit (-1);
}
else /* the parent; child == process id of child */ {
    /* do whatever you want, e.g., just return from
    this routine */
    int status;
    wait (&status);

}
.....
```

Example: create new processes

fork() creates an identical copy of the caller in a separate memory space;



Only the return value stored in variable '*child*' differs between the two: **0 for the child**, and **child-pid** for the parent

```
pid_t child;
.....
child = fork();
if (child < 0) /* error occurred */ {
    perror ("fork"); exit (-1); }

if (child == 0) /* the child */ {
    execlp ("/bin/ls", "ls", arg0, arg1, ..., NULL);
    /* this place is reached only in case of error */
    perror ("execlp"); exit (-1);
}
else /* the parent; child == process id of child */ {
    /* do whatever you want, e.g., just return from
    this routine */
    int status;
    wait (&status);

}
.....
```

Example: create new processes

fork() creates an identical copy of the caller in a separate memory space;

Only the return value stored in variable '*child*' differs between the two: **0 for the child**, and **child-pid** for the parent



```
pid_t child;
.....
child = fork();
if (child < 0) /* error occurred */ {
    perror ("fork"); exit (-1); }
if (child == 0) /* the child */ {
    execlp ("/bin/ls", "ls", arg0, arg1, ..., NULL);
    /* this place is reached only in case of error */
    perror ("execlp"); exit (-1);
}
else /* the parent; child == process id of child */ {
    /* do whatever you want, e.g., just return from
    this routine */
    int status;
    wait (&status);
}
.....
```

Example: create new processes

fork() creates an identical copy of the caller in a separate memory space;

Only the return value stored in variable '*child*' differs between the two: **0 for the child**, and **child-pid** for the parent

execvp() **overwrites** the calling process with its argument (an executable programme).

This means that code following **execvp** is never reached.

perror() is a generic error printing routine



```
pid_t child;
.....
child = fork();
if (child < 0) /* error occurred */ {
    perror ("fork"); exit (-1); }

if (child == 0) /* the child */ {
    execvp ("/bin/ls", "ls", arg0, arg1, ..., NULL);
    /* this place is reached only in case of error */
    perror ("execvp"); exit (-1);
}
else /* the parent; child == process id of child */ {
    /* do whatever you want, e.g., just return from
    this routine */
    int status;
    wait (&status);
}
.....
```

Example: create new processes

fork() creates an identical copy of the caller in a separate memory space;

Only the return value stored in variable '*child*' differs between the two: **0** for the child, and **child-pid** for the parent

execvp() overwrites the calling process with its argument (an executable programme).

This means that code following **execvp** is never reached.

perror() is a generic error printing routine



```
pid_t child;
.....
child = fork();
if (child < 0) /* error occurred */ {
    perror ("fork"); exit (-1); }

if (child == 0) /* the child */ {
    execvp ("/bin/ls", "ls", arg0, arg1, ..., NULL);
    /* this place is reached only in case of error */
    perror("execvp"); exit(-1);
}

else /* the parent; child == process id of child */ {
    /* do whatever you want, e.g., just return from
    this routine */
    int status;
    wait (&status);
}
.....
```



Multi-shadow-clone jutsu
Naruto

- Every POSIX child process is an **exact copy of the parent process at its creation**
- Child processes become **independent after creation**

- Use **exit(status)** to terminate
- Functions **wait()**, **waitpid()**
 - **wait()** blocks until one of the children exit
 - **waitpid()** blocks until a specific child changes its state

parent:

```
...  
pid_t  child, terminated;  
int status;
```

```
/* blocking wait */  
while (child != wait (&status))    /* nothing */;
```

Or

```
/* or polling wait */  
terminated = (pid_t) 0;  
while (terminated != child) {  
    terminated = waitpid (child, &status, WNOHANG);  
    /* other useful activities */  
    do_something()  
}
```

```
/* both cases: status == 23 */
```

child:

```
..... exit (23);
```

More examples: <https://www.geeksforgeeks.org/wait-system-call-c/>

- Use **exit(status)** to terminate
- Functions **wait()**, **waitpid()**
 - **wait()** blocks until one of the children exit
 - **waitpid()** blocks until a specific child changes its state

parent:

```
...  
pid_t  child, terminated;  
int status;
```

```
/* blocking wait */  
while (child != wait (&status))  /* nothing */;
```

Or

```
/* or polling wait */  
terminated = (pid_t) 0;  
while (terminated != child) {  
    terminated = waitpid (child, &status, WNOHANG);  
    /* other useful activities */  
    do_something()  
}
```

```
/* both cases: status == 23 */
```

This parameter is used to continue without hanging if the child "Child" did not finish yet.

child:

```
..... exit (23);
```

More examples: <https://www.geeksforgeeks.org/wait-system-call-c/>

- Process may ask the operating system to delete itself (**exit**)...
 - **exit()** returns status value from child to parent (if the parent was “**waiting**”)
 - After termination, the process’ **resources** are taken away by OS
- ... or **parent** may terminate execution of **children** processes (**abort / kill**)
 - The child has **exceeded allocated resources**
 - The task assigned to child is **no longer required**
 - If the **parent is exiting**
 - Some operating systems do not allow the child to continue if its parent terminates, in which case all children are terminated - **cascading termination**
 - Some operating systems **attach the orphan children to a ‘grandfather’ process** (e.g. *init*)
- If no parent is **waiting** (did not invoke **wait()**) the child process is a **zombie**
- If parent **terminated** without invoking **wait()**, the child process is an **orphan**

Exercises

```
#include ...
```

```
#include ...
```

```
int main() {  
    pid_t smith;  
    smith = fork( );  
    if (smith == 0) { smith = fork( );}  
    if (smith > 0) { smith = fork( );}  
    /* do something useful */  
    return 0;  
}
```

How many processes are created,
including the initial parent process?

Answer: 5

```
#include ...
```

```
#include ...
```

```
int main() {
```

```
    pid_t smith;
```

```
    smith = fork( );
```

```
    if (smith == 0) { smith = fork( );}
```

```
    if (smith > 0) { smith = fork( );}
```

```
    /* do something useful */
```


```
    return 0;
```

```
}
```

```
#include ...
```

```
#include ...
```

```
int main() {  
    pid_t smith;  
    smith = fork( );  
    if (smith == 0) { smith = fork( );}  
    if (smith > 0) { smith = fork( );}  
    /* do something useful */  
    return 0;  
}
```



```
#include ...
```

```
#include ...
```

```
int main() {
```

```
    pid_t smith;
```

```
    smith = fork( );
```


```
    if (smith == 0) { smith = fork( );}
```

```
    if (smith > 0) { smith = fork( );}
```

```
    /* do something useful */
```

```
    return 0;
```

```
}
```



```
#include ...
```

```
#include ...
```

```
int main() {
```

```
    pid_t smith;
```

```
    smith = fork( );
```

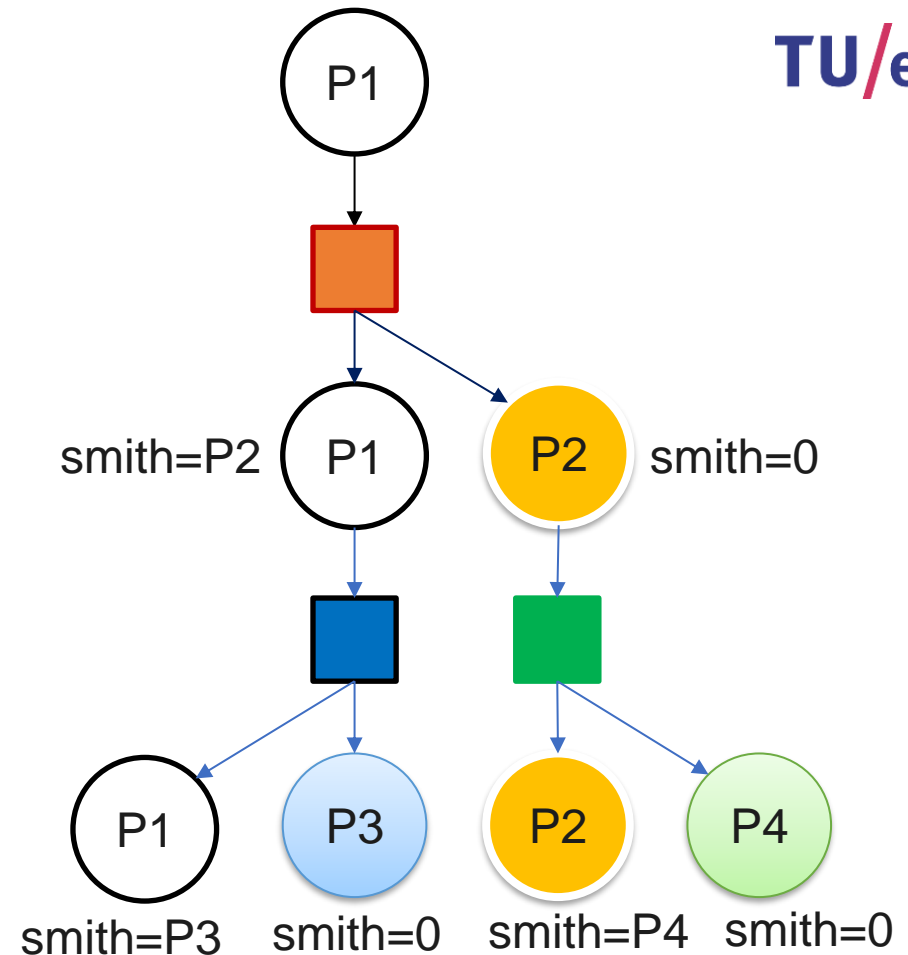
```
    if (smith == 0) { smith = fork( );}
```

```
    if (smith > 0) { smith = fork( );}
```

```
    /* do something useful */
```

```
    return 0;
```

```
}
```



```
#include ...
```

```
#include ...
```

```
int main() {
```

```
    pid_t smith;
```

```
    smith = fork( );
```

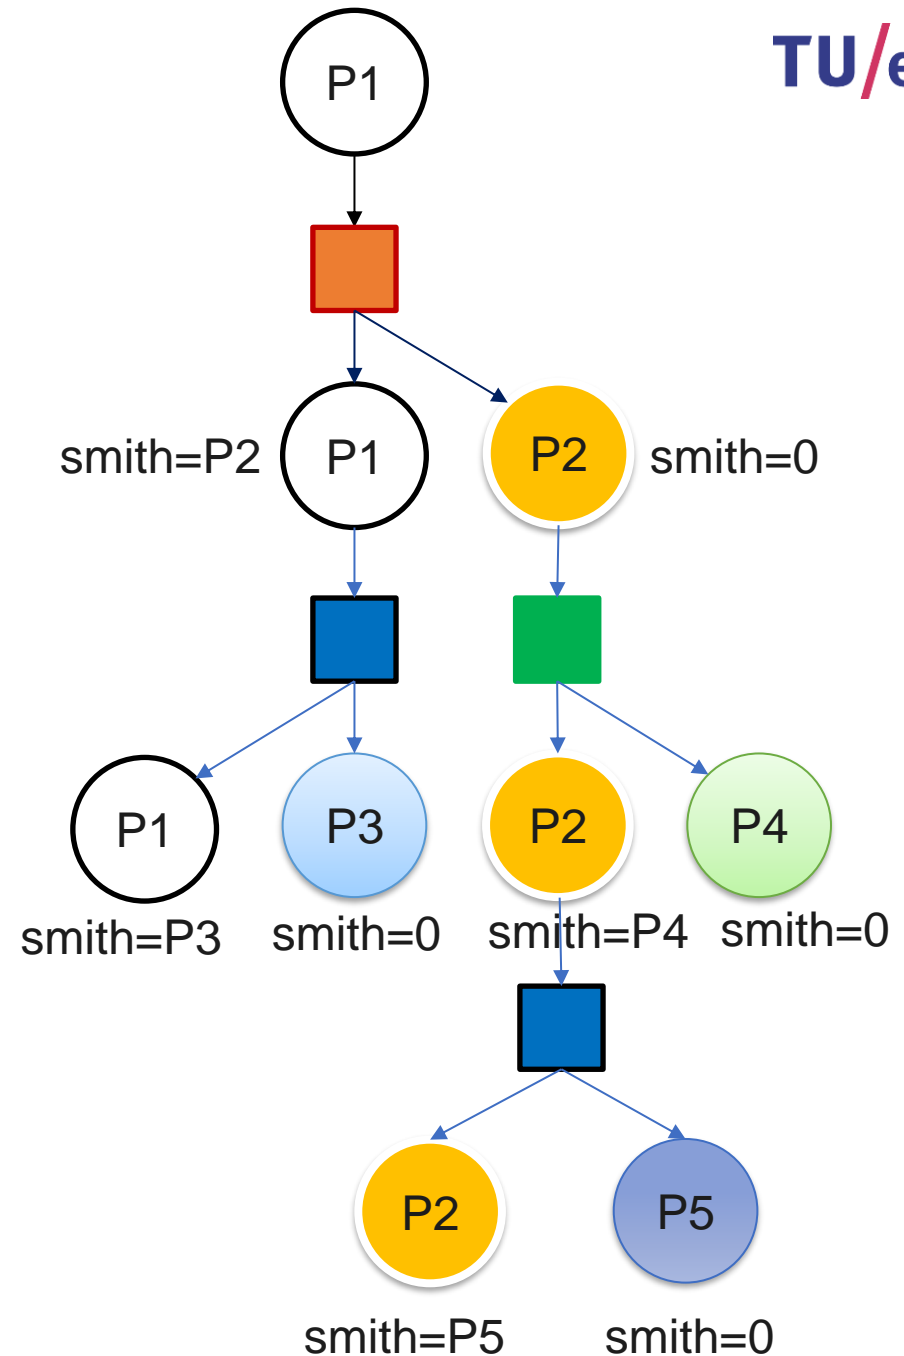
```
    if (smith == 0) { smith = fork( );}
```

```
    if (smith > 0) { smith = fork( );}
```

```
    /* do something useful */
```

```
    return 0;
```

```
}
```




```
#include ...  
#include ...  
  
int main() {  
    pid_t smith;  
    smith = fork( );  
    if (smith == 0) { smith = fork( );}  
    if (smith > 0) { smith = fork( );}  
    /* do something useful */  
    return 0;  
}
```

There are **two issues** in the code. Can you spot them?

```
int i = 0;
int main() {
    pid_t id;
    id = fork( );
    if (id == 0) {
        exit(0);
    }
    id = fork();
    if (id > 0) { fork(); }
    i++;
    /* do something useful */
    return 0;
}
```

How many **processes** are created
(including the initial parent process)?

How many processes **execute** the
“**i++**” instruction?

What is the **final value** of the
variable **i** in the last process that
reaches “return 0”?

Process execution

```
int i = 0;
int main() {
    pid_t id;
    id = fork( );
    if (id == 0) {
        exit(0);
    }
    id = fork();
    if (id > 0) { fork(); }
    i++;
    /* do something useful */
    return 0;
}
```

How many **processes** are created
(including the initial parent process)?

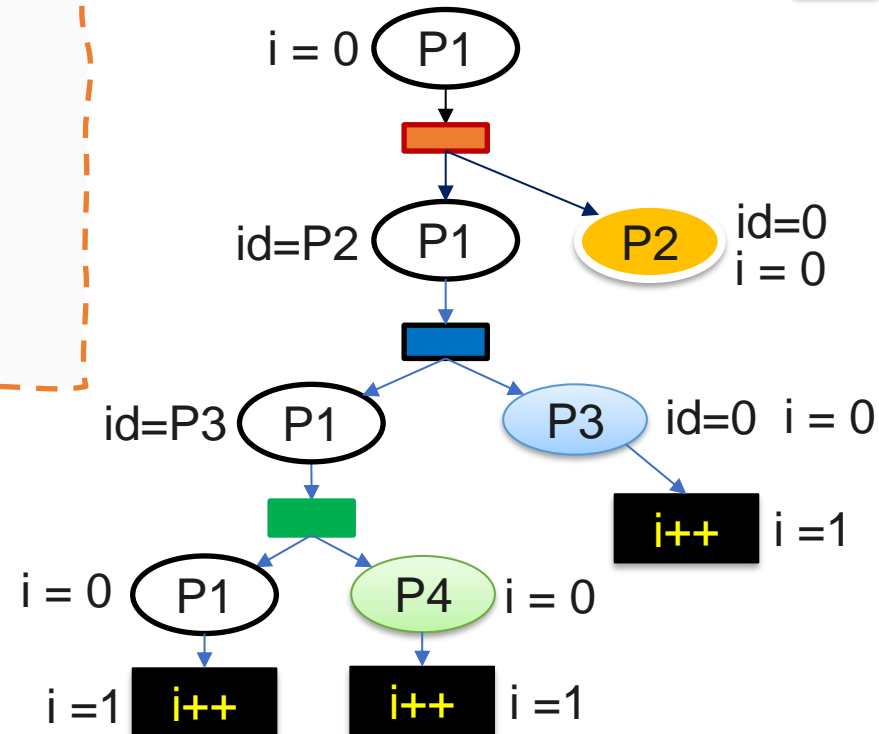
4

How many processes **execute the**
“**i++**” instruction?

3

What is the **final value** of the
variable **i** in the last process that
reaches “return 0”?

1



Remember: the content of the preparatory material is also part of the exam
(e.g., process state, inter-process communication, etc.)

- Ingredients of concurrent execution: processes and threads
 - Process
 - a **program in execution**.
 - has a **context of execution**
 - **owns resources**
 - Thread
 - dispatchable **unit of work within a process**
 - **shares code and data**
- Creation/deletion of process

- Lecture on threads and scheduling
- **Preparation:**
 - Reference book (~11 pages)