

# 2INC0 - Operating Systems

## Atomicity and interference

Geoffrey Nelissen



Interconnected  
Resource-aware  
Intelligent Systems

**TU/e**

Technische Universiteit  
**Eindhoven**  
University of Technology

Where innovation starts

- First **homework**: **deadline on Sunday**
- New homework released every weekend
- Topics covered by the **final exam**:
  - **preparatory material**
  - **lecture content**
  - **additional exercises**

- **Introduction to operating systems** (lecture 1)
- **Processes, threads and scheduling** (lectures 2 and 3)
- **Concurrency and synchronization**
  - **atomicity and interference** (lecture 4)
  - action synchronization (lecture 5)
  - condition synchronization (lecture 6)
  - deadlock (lecture 7)
- **File systems** (lecture 8)
- **Memory management** (lectures 9 and 10)
- **Input/output** (lecture 11)

- **Reminder of lectures 2-3**
- **Atomicity, Traces and concurrency**
- **Synchronization**
- **Peterson's algorithm**
- **Synchronization with mutexes**

- Let  $x$  be a **global** variable initialized to 0.
- Let Task1 and Task2 be POSIX **processes** executing the same program concurrently.
- Task2 is the child of Task1.
- What is the final value of  $x$  in Task1 and Task2?

Task1:  $[x := x+1;]$

||

Task2:  $[x := x+3;]$

Answer:

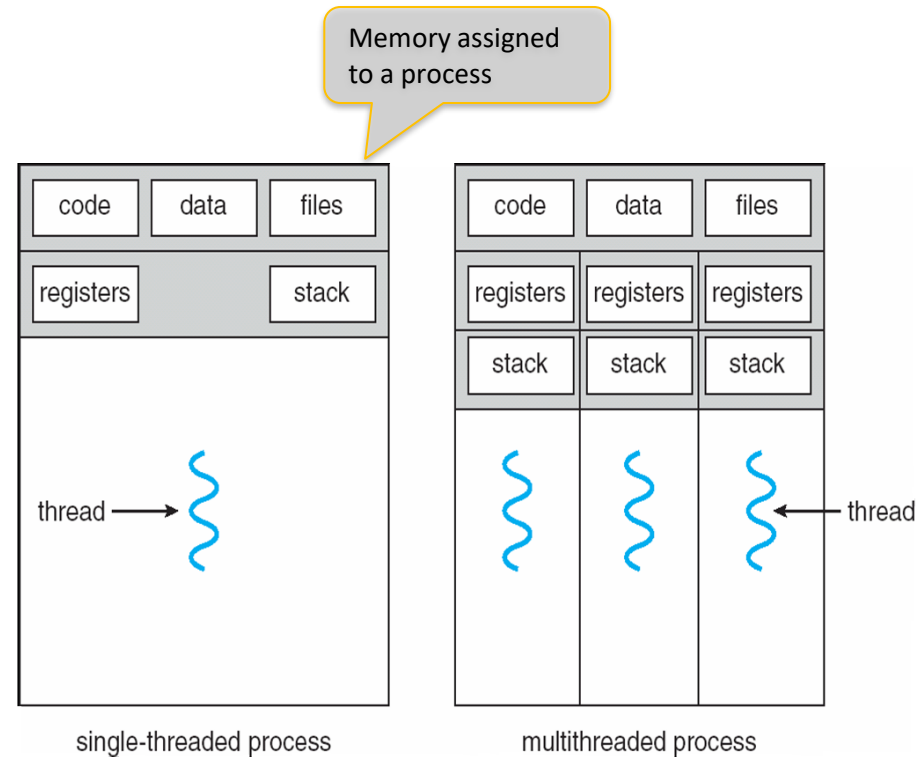
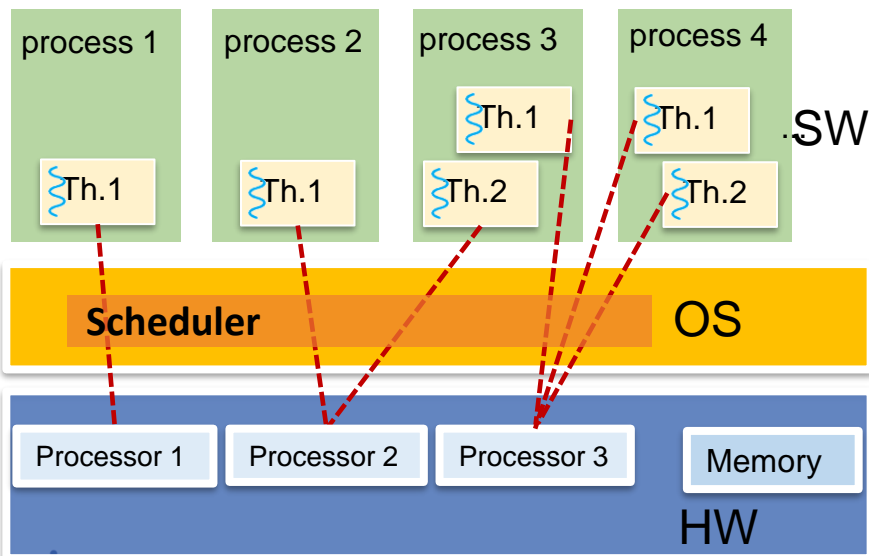
- 1 in Task1
- 3 in Task2

## Process

- It is a **program in execution**.
- It has a **context of execution**
- A process owns resources (has memory and can own other resources such as files, etc.)

## Thread

- A **dispatchable unit of work within a process**
- Threads within a process share code and data segments (i.e., **share memory address space**)



- Let  $x$  be a **global** variable initialized to 0.
- Let Task1 and Task2 be **threads** executing the same program concurrently.
- What is the final value of  $x$  in Task1 and Task2?

Task1:  $[x := x+1;]$

||

Task2:  $[x := x+3;]$

Answer:  
 $x = 1$  or  $3$  or  $4$   
for both Task1  
and Task2

- Reminder of lecture 2+3
- **Atomicity, Traces and concurrency**
- **Synchronization**



## Atomic action

- **Indivisible** step in the evolution of an executed program (outcome only depends on system state when the action starts and not on what runs concurrently during the execution of the action)

## Trace

## Concurrent trace

**How can we know if a statement is atomic?**

## Atomic action

- **Indivisible** step in the evolution of an executed program (outcome only depends on system state when the action starts and not on what runs concurrently during the execution of the action)

## Trace

## Concurrent trace


### Single reference rule:


a statement (expression) may be regarded as **atomic** if it makes **at most one reference to a shared variable**.

- If  $x$  and  $y$  are **shared variables** and  $z$  is a **private variable**, which of the following statements can be **considered as atomic**

1.  $x := x + y;$  

2.  $x := y * z;$  

3.  $z := z + 1;$  

4.  $z := z + z * y;$  

5.  $y++;$  

**Note:** if we write  $\langle x := x + z \rangle$   
(placing the statement between  $\langle \dots \rangle$ )  
then the statement should be  
considered to be **atomic**

## Atomic action

- **Indivisible** step in the evolution of an executed program

## Trace

- **Sequence of atomic actions** in the execution of a program
- Typically, a sequence of **assignments** and **tests**

## Concurrent trace

- Trace made of an **interleaving** of atomic actions of two or more tasks

### Single reference rule:

a statement (expression) may be regarded as **atomic** if it makes **at most one reference to a shared variable**.

- Let  $x$  be a **global** variable **initialized to 0**.
- Let Task1 and Task2 be **threads** executing concurrently.
- What is the final value of  $x$ ?

Task1:  $[x := x+1;]$

||

Task2:  $[x := x+1;]$

We use two additional variables, **reg1** and **reg2** to represent the **content of the internal processor registers** used by Task1 and Task2, respectively.

Possible trace 1:  $(\text{reg2}:=x)(\text{reg2} := \text{reg2}+1)(x := \text{reg2}) \{x=1\} (\text{reg1}:=x) \{\text{reg1}=1\} (\text{reg1} := \text{reg1}+1)(x := \text{reg1}) \{x=2\}$

Possible trace 2:  $(\text{reg1}:=x) \{\text{reg1}=0\} (\text{reg2}:=x) \{\text{reg2}=0\} (\text{reg2} := \text{reg2}+1) \{\text{reg2}=1\} (x := \text{reg2}) \{x=1\}$   
 $(\text{reg1} := \text{reg1}+1) \{\text{reg1}=1\} (x := \text{reg1}) \{x=1\}$

...

**Note:** for a **complete proof**, we would need to **write all possible concurrent traces** and show they all lead to  $x=1$  or  $x=2$

Answer:  
 $x = 1$  or  $2$

1. Form a group with the people around you.
2. Together, solve the problem that will be shown on the next slide.
  - You have 10 minutes.



Anime: One piece  
Character: Luffy

```
#include <stdio.h>
#include <pthread.h>

int y;

void Count_10 ()
{
    for (int i = 0; i < 10; i ++) {
        y = y+1;
    }
}

void main()
{
    pthread_t thread_id;

    y = 0;
    pthread_create (&thread_id, NULL, Count_10, NULL);
    Count_10 ();
    pthread_join (thread_id, NULL);
}
```

- The statement ' $y=y+1$ ' is *not* atomic.
- What are the possible final values of  $y$ ?
- **Use traces** to prove your answer

# Exercise on concurrency

```
#include <stdio.h>
#include <pthread.h>
```

```
int y;
```

```
void Count_10 ()
```

```
{
  for (int i = 0; i < 10; i++) {
    y = y+1;
  }
}
```

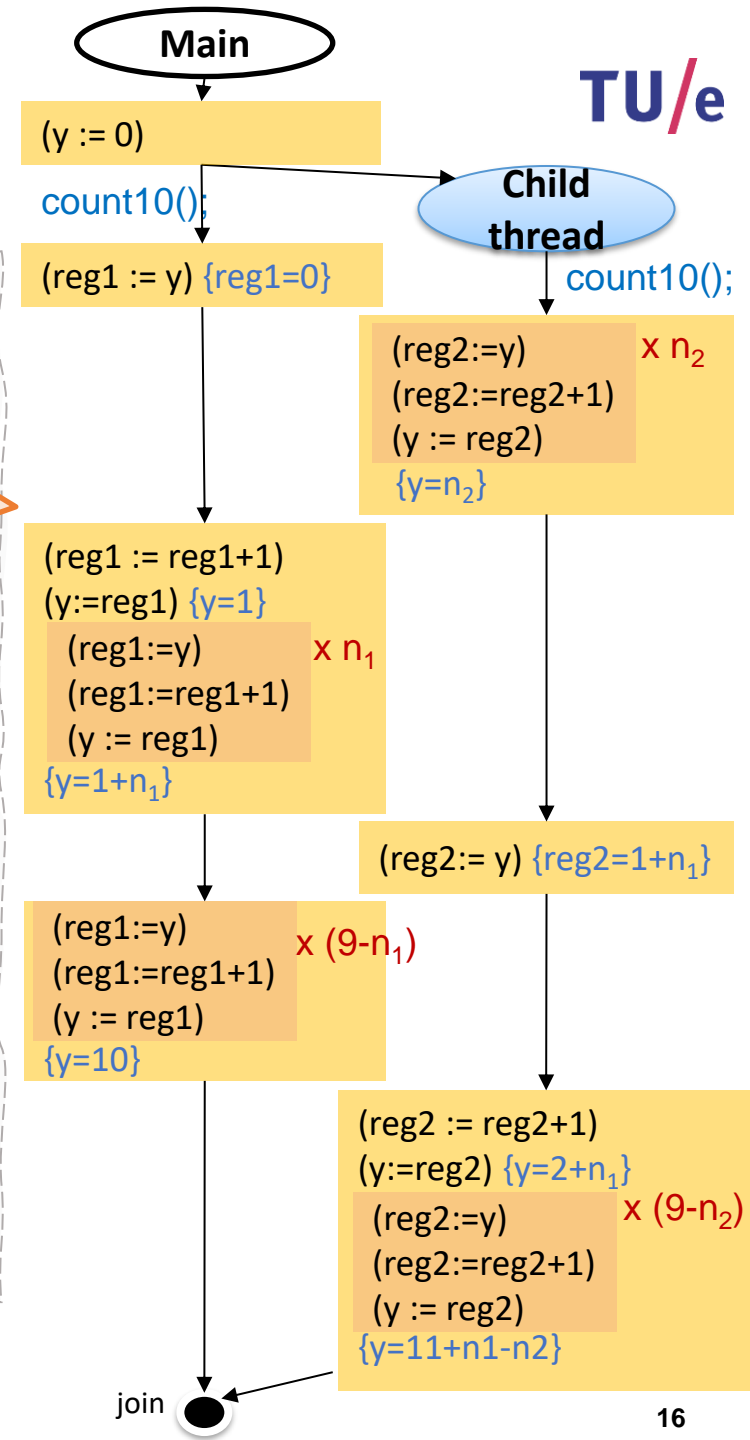
```
void main()
```

```
{
  pthread_t thread_id;

  y = 0;
  pthread_create (&thread_id, NULL, Count_10, NULL);
  Count_10 ();
  pthread_join (thread_id, NULL);
}
```

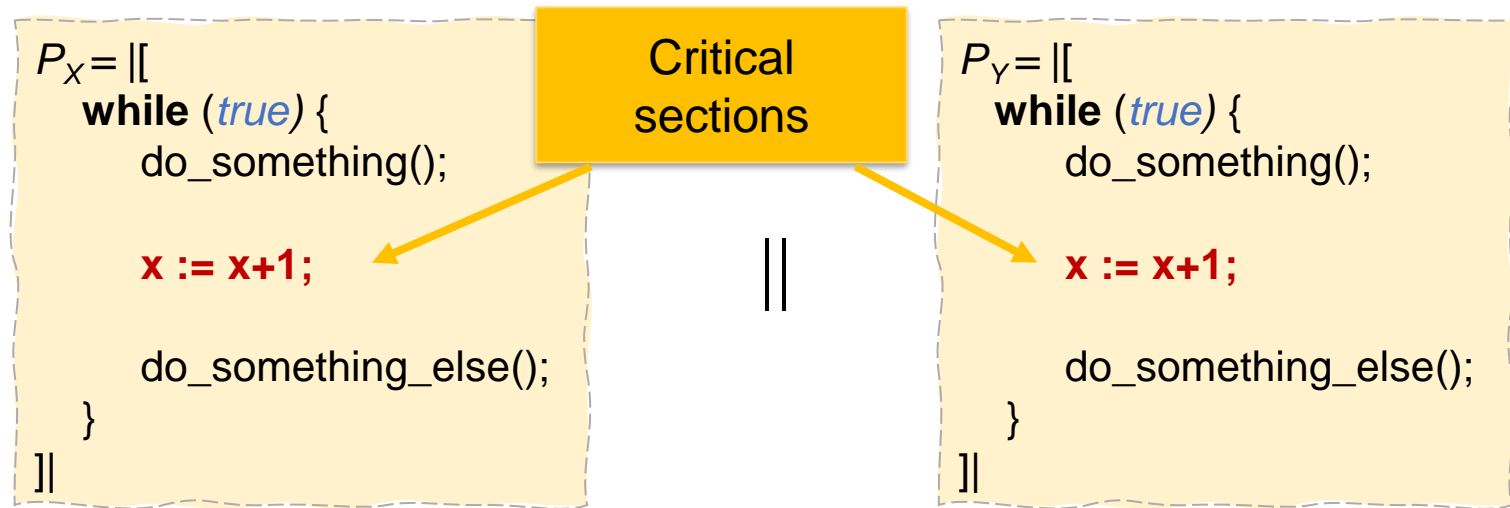
In this trace,  $n_1$  and  $n_2$  can take any value between 0 and 9

**Answer:** y is equal to any number between 2 and 20





- Reminder of lectures 2+3
- Atomicity, Traces and concurrency
- **Synchronization**
- **Peterson's algorithm**
- **Synchronization with mutexes**



## Our synchronization goal:

How can we coordinate the execution of the given programs such that **no two tasks execute the statement  $x:=x+1$  at the same time**?

**Mutual exclusion**

**Synchronization refers to *ordering the execution***, i.e., **forbidding certain traces** (more complex synchronization goals addressed in the next two lectures)


- Introduce global **boolean variables**  $bX$  and  $bY$  to record when we are in the critical section
- Initially:  $bX=false$  and  $bY=false$

```
PX = [[  
  while (true) {  
    do_something();  
     $bX := true$ ;  
     $x := x+1$ ;  
     $bX := false$ ;  
    do_something_else();  
  }  
]]
```

||

```
PY = [[  
  while (true) {  
    do_something();  
     $bY := true$ ;  
     $x := x+1$ ;  
     $bY := false$ ;  
    do_something_else();  
  }  
]]
```

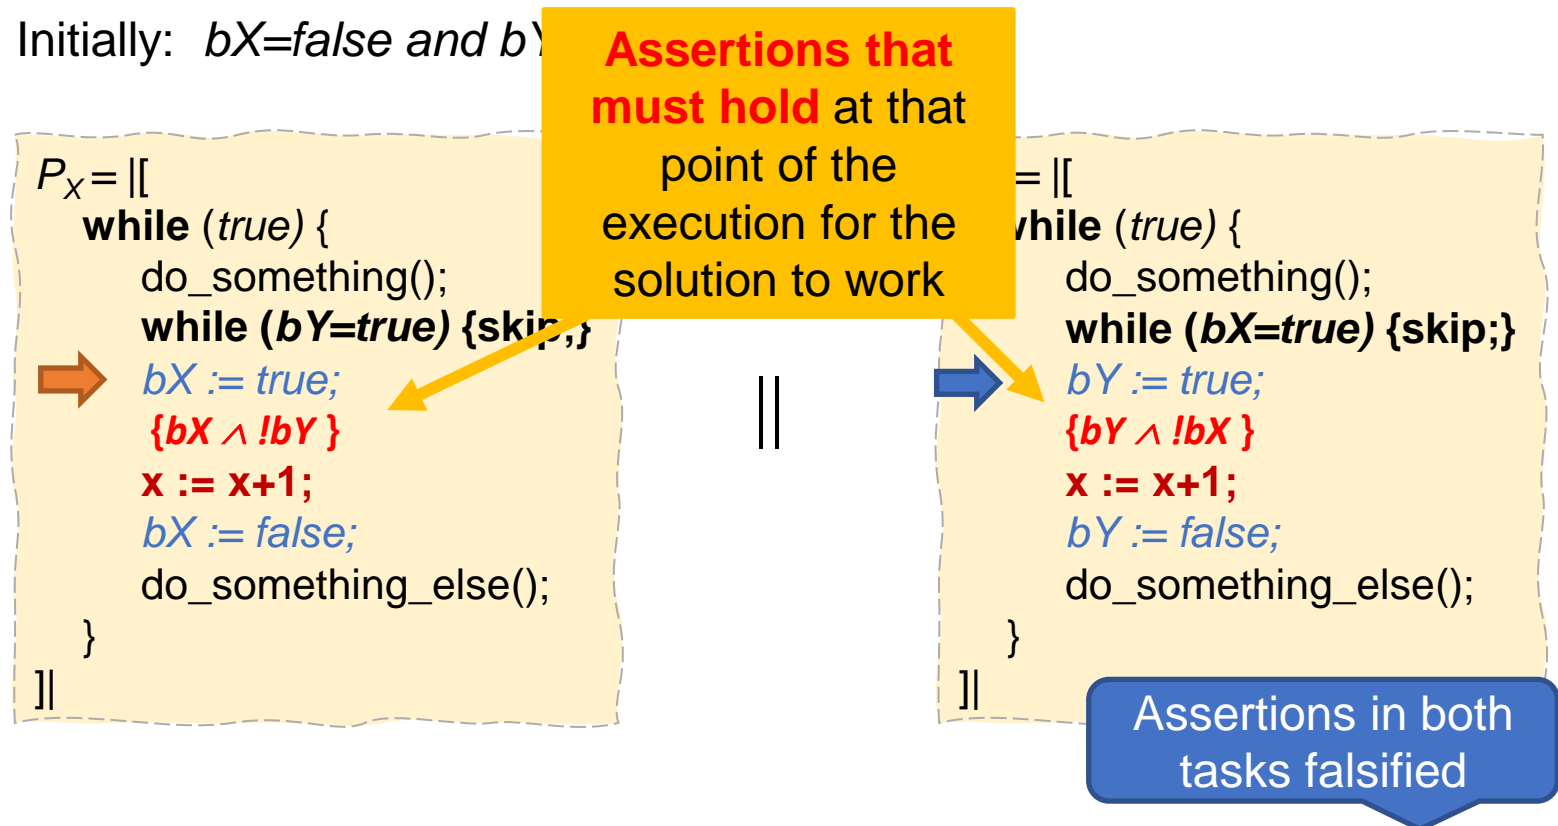
- Introduce global **boolean variables**  $bX$  and  $bY$  to record when we are in the critical section
- Initially:  $bX=false$  and  $bY=false$



Does this work? Why?

- Introduce global **boolean variables**  $bX$  and  $bY$  to record when we are in the critical section

- Initially:  $bX=false$  and  $bY=false$



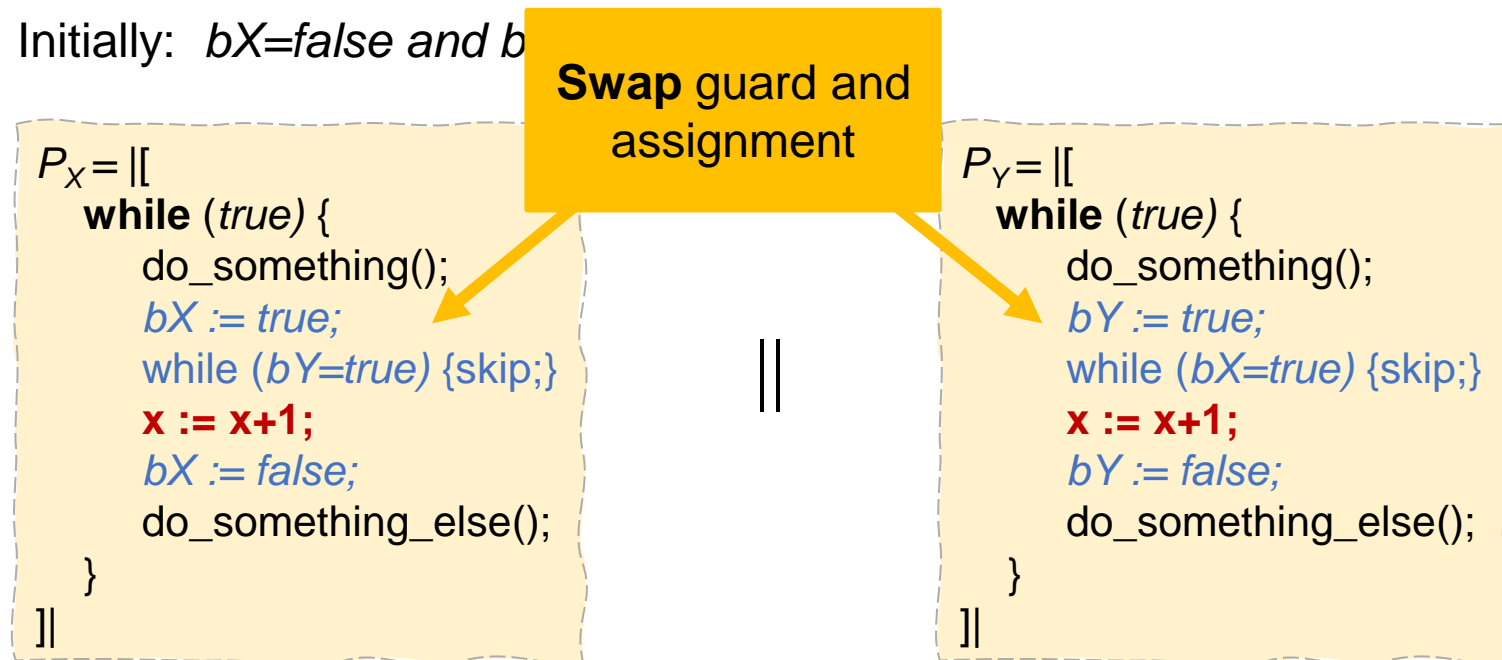
**Trace:**  $\{!bY \wedge !bX\} ( !(bY=true) ) ( !(bX=true) ) (bY:=true) (bX:=true) \{bY \wedge bX\}$

Let's try something else...

# Swap the order of the Boolean assignment and guard...

- Instead of recording when we are in the critical section,  $bX$  and  $bY$  could **record the interest** in entering the critical section

- Initially:  $bX=false$  and  $bY=false$



Does this work? Why?

# Change the order...

- Instead of recording when we are in the critical section,  $bX$  and  $bY$  could **record the interest** in entering the critical section
- Initially:  $bX=false$  and  $bY=false$

```
PX = [[  
  while (true) {  
    do_something();  
    →  $bX := true$ ;  
    while ( $bY=true$ ) {skip;}  
     $x := x+1$ ;  
     $bX := false$ ;  
    do_something_else();  
  }  
]]
```

||

```
PY = [[  
  while (true) {  
    do_something();  
    →  $bY := true$ ;  
    while ( $bX=true$ ) {skip;}  
     $x := x+1$ ;  
     $bY := false$ ;  
    do_something_else();  
  }  
]]
```

No task can proceed:  
**Deadlock**

**Trace:**  $(bX:=true)(bY:=true) \{bY \wedge bX\}$



Let's try something else...

- Rather than trying to **obtain** access, the task **give** the access away using a shared variable  $t$  that tells who can go next
- Initially:  $t = X$  or  $t = Y$

```
PX = [[  
  while (true) {  
    do_something();  
    t := Y;  
    while( t != X){ skip;}  
    x := x+1;  
    do_something_else();  
  }  
]]
```

||

```
PY = [[  
  while (true) {  
    do_something();  
    t := X;  
    while( t != Y){ skip;}  
    x := x+1;  
    do_something_else();  
  }  
]]
```

Does this work? Why?

- Rather than trying to **obtain** access, the task **give** the access away using a shared variable  $t$  that tells who can go next
- Initially:  $t = X$  or  $t = Y$

```
PX = [[  
  while (true) {  
    do_something();  
    t := Y;  
    while( t != X){ skip;}  
    x := x+1;  
    do_something_else();  
  }  
]]
```

||

```
PY = [[  
  while (true) {  
    do_something();  
    t := X;  
    while( t != Y){ skip;}  
    x := x+1;  
    do_something_else();  
  }  
]]
```

- **Functionally correct:** mutual exclusion is respected
- **No deadlock**
- **Waiting is not minimal:** a task must wait for the other to be able to enter the critical section

- **Functional correctness:**

- **satisfy** the given **specification** (e.g., mutual exclusion) -- an assertion that is needed for correctness must not be disturbed (by another process).

- **Minimal waiting: (“make progress”)**

- **waiting only when correctness is in danger.**

- **Absence of deadlock:**

- don't manoeuvre (part of) the system into a state such that **progress is no longer possible.**

- **Absence of livelock:**

- ensure **convergence towards a decision** in a synchronization protocol

- **Fairness in competition:**

- **(weak) eventually**, each contender should be **admitted to proceed.**
- **(strong)** we can put a **bound on the waiting time** of a contender.
- **absence of fairness:** leads to **starvation** of tasks.

- **Note:** the last four are only required when there is concurrency

- Reminder of lectures 2+3
- Atomicity, Traces and concurrency
- Synchronization
- **Peterson's algorithm**
- **Synchronization with mutexes**

**Combine the ideas** of the last two solutions

```
Px = [[  
  while (true) {  
    do_something();  
  
    x := x+1;  
  
    do_something_else();  
  }  
]]
```

||

```
Py = [[  
  while (true) {  
    do_something();  
  
    x := x+1;  
  
    do_something_else();  
  }  
]]
```

**Combine the ideas** of the last two solutions

- **take turns** in crowded circumstances (use a variable  $t$ )

```
PX = [  
  while (true) {  
    do_something();  
    t := Y;  
    while( t != X ) { skip; }  
    x := x+1;  
  
    do_something_else();  
  }  
]
```

||

```
PY = [  
  while (true) {  
    do_something();  
    t := X;  
    while( t != Y ) { skip; }  
    x := x+1;  
  
    do_something_else();  
  }  
]
```

**Combine the ideas** of the last two solutions

- **take turns** in crowded circumstances (use a variable  $t$ )
- **don't wait if there is no need** (look at the boolean  $bY$  or  $bX$ )

```
PX = [  
  while (true) {  
    do_something();  
    bX := true;  
    t := Y;  
    while( bY=true ∧ t!=X ){ skip;}  
    x := x+1;  
    bX := false;  
    do_something_else();  
  }  
]
```

||

```
PY = [  
  while (true) {  
    do_something();  
    bY := true;  
    t := X;  
    while( bX=true ∧ t!=Y ){ skip;}  
    x := x+1;  
    bY := false;  
    do_something_else();  
  }  
]
```



We want to show:

- **mutual exclusion**
- **absence of deadlock**

We can prove it

- using traces (tedious), or
- using **detailed annotation** of the program

- Proof by **contradiction**:
  - **Assume *there is a deadlock***, then...  
both tasks must be blocked **on the while loop**.  
→ The guards must hold for both tasks at the same time

```
PX = [[  
  while (true) {  
    do_something();  
    bX := true;  
    t := Y;  
    while( bY=true ∧ t!=X ){ skip;}  
    x := x+1;  
    bX := false;  
    do_something_else();  
  }  
]]
```

||

```
PY = [[  
  while (true) {  
    do_something();  
    bY := true;  
    t := X;  
    while( bX=true ∧ t!=Y ){ skip;}  
    x := x+1;  
    bY := false;  
    do_something_else();  
  }  
]]
```

- Proof by **contradiction**:
  - **Assume *there is a deadlock***, then...  
both tasks must be blocked **on the while loop**.
    - ➔ The guards must hold for both tasks at the same time
    - ➔ For  $P_X$ :  $bY \wedge t \neq X$   
and for  $P_Y$ :  $bX \wedge t \neq Y$
    - ➔  $bY \wedge t \neq X \wedge bX \wedge t \neq Y$   
  
 $= \{t \neq X \text{ and } t \neq Y \text{ cannot be true at the same time}\}$   
*false*
    - ➔ *We reached a contradiction*
    - ➔ *There is no deadlock*

- Proof by **contradiction**:
  - Assume **there is no mutual exclusion**, then...  
both tasks must be in their critical section **at the same time**.
  - ➔ **States of  $P_X$  and  $P_Y$  after their guards** must hold at the same time.

What are those states?

```
PX = [  
  while (true) {  
    do_something();  
    bX := true;  
    t := Y;  
    while( bY=true ∧ t!=X ){ skip;}  
    x := x+1;  
    bX := false;  
    do_something_else();  
  }  
]
```

||

```
PY = [  
  while (true) {  
    do_something();  
    bY := true;  
    t := X;  
    while( bX=true ∧ t!=Y ){ skip;}  
    x := x+1;  
    bY := false;  
    do_something_else();  
  }  
]
```

- Proof by **contradiction**:
  - Assume **there is no mutual exclusion**, then...  
both tasks must be in their critical section **at the same time**.
  - ➔ **States of  $P_X$  and  $P_Y$  after their guards** must hold at the same time.

What are those states?

```
PX = [  
  while (true) {  
    do_something();  
    bX := true;  
    (*)  
    t := Y;  
    while( bY=true ∧ t!=X ){ skip;}  
    { bX ∧ (t = X ∨ !bY ∨ PY at (**)) }  
    x := x+1;  
    bX := false;  
    do_something_else();  
  }  
]
```

```
PY = [  
  while (true) {  
    do_something();  
    bY := true;  
    (**)  
    t := X;  
    while( bX=true ∧ t!=Y ){ skip;}  
    { bY ∧ (t = Y ∨ !bX ∨ PX at (*)) }  
    x := x+1;  
    bY := false;  
    do_something_else();  
  }  
]
```

||

**Proving that those assertions are correct  
is out of the scope of the course**

- Proof by **contradiction**:
  - Assume **there is no mutual exclusion**, then...  
both tasks must be in their critical section **at the same time**.
  - **States of  $P_X$  and  $P_Y$  after their guards** must hold at the same time.

What are those states?

```
PX = [[  
  while (true) {  
    do_something();  
    <bX, auxX := true, true>;  
    <t, auxX := Y, false>;  
    while( bY=true ∧ t!=X ) { skip;}  
    { bX ∧ (t = X ∨ !bY ∨ auxY )}  
    x := x+1;  
    bX := false;  
    do_something_else();  
  }  
]]
```

||

```
PY = [[  
  while (true) {  
    do_something();  
    <bY, auxY := true, true>;  
    <t, auxY := X, false>;  
    while( bX=true ∧ t!=Y ) { skip;}  
    { bY ∧ (t = Y ∨ !bX ∨ auxX )}  
    x := x+1;  
    bY := false;  
    do_something_else();  
  }  
]]
```

- Proof by **contradiction**:
  - Assume **there is no mutual exclusion**, then...  
both tasks must be in their critical section at the same time.
  - States of  $P_X$  and  $P_Y$  after their guards must hold at the same time.

What are those states?

For  $P_X$ :  $\{ bX \wedge (t = X \vee !bY \vee auxY) \}$

For  $P_Y$ :  $\{ bY \wedge (t = Y \vee !bX \vee auxX) \}$

→  $\{ bX \wedge (t = X \vee !bY \vee auxY) \} \wedge \{ bY \wedge (t = Y \vee !bX \vee auxX) \}$

= {simplifying}

$bX \wedge bY \wedge (t = X \wedge t = Y)$

= {  $t$  cannot be equal to both  $X$  and  $Y$  }

**false**

- We reached a contradiction.
- We have mutual exclusion

Note:  $auxX$  and  $auxY$  cannot be true since  $P_X$  and  $P_Y$  are in their critical sections

- Can synchronize **only two tasks**
  - (extension for more tasks exists but more complex)
- Busy-wait
  - **wastes CPU cycles**
  - **does not work for single core** if the scheduling policy is **non-preemptive**

Why?

Operating systems kernels provide **semaphores** and **mutexes** instead



- **Reminder of lecture 2**
- **Atomicity, Traces and concurrency**
- **Synchronization**
- **Peterson's algorithm**
- **Synchronization with mutexes**

- **Special, two-state (or binary) semaphore:** *mutex*  
(see next lecture for details)
- A mutex *m* **must be initialized to 1**
- Two operations:
  - **Lock**(*m*):  $\langle \text{await}(m > 0); m := m - 1 \rangle$
  - **Unlock**(*m*):  $\langle m := m + 1 \rangle$

The **implementation** of *Lock* and *Unlock* must guarantee their **atomicity**

A semaphore is **always non-negative**.

Other names for *Lock* and *Unlock*: *wait / signal*, *wait / post*, *P / V*

```
mutex  $m := 1$ ;
```

```
 $P_x = [$   
  while (true) {  
    do_something();  
    lock( $m$ );  
     $x := x+1$ ;  
    unlock( $m$ );  
    do_something_else();  
  }  
 $]$ 
```

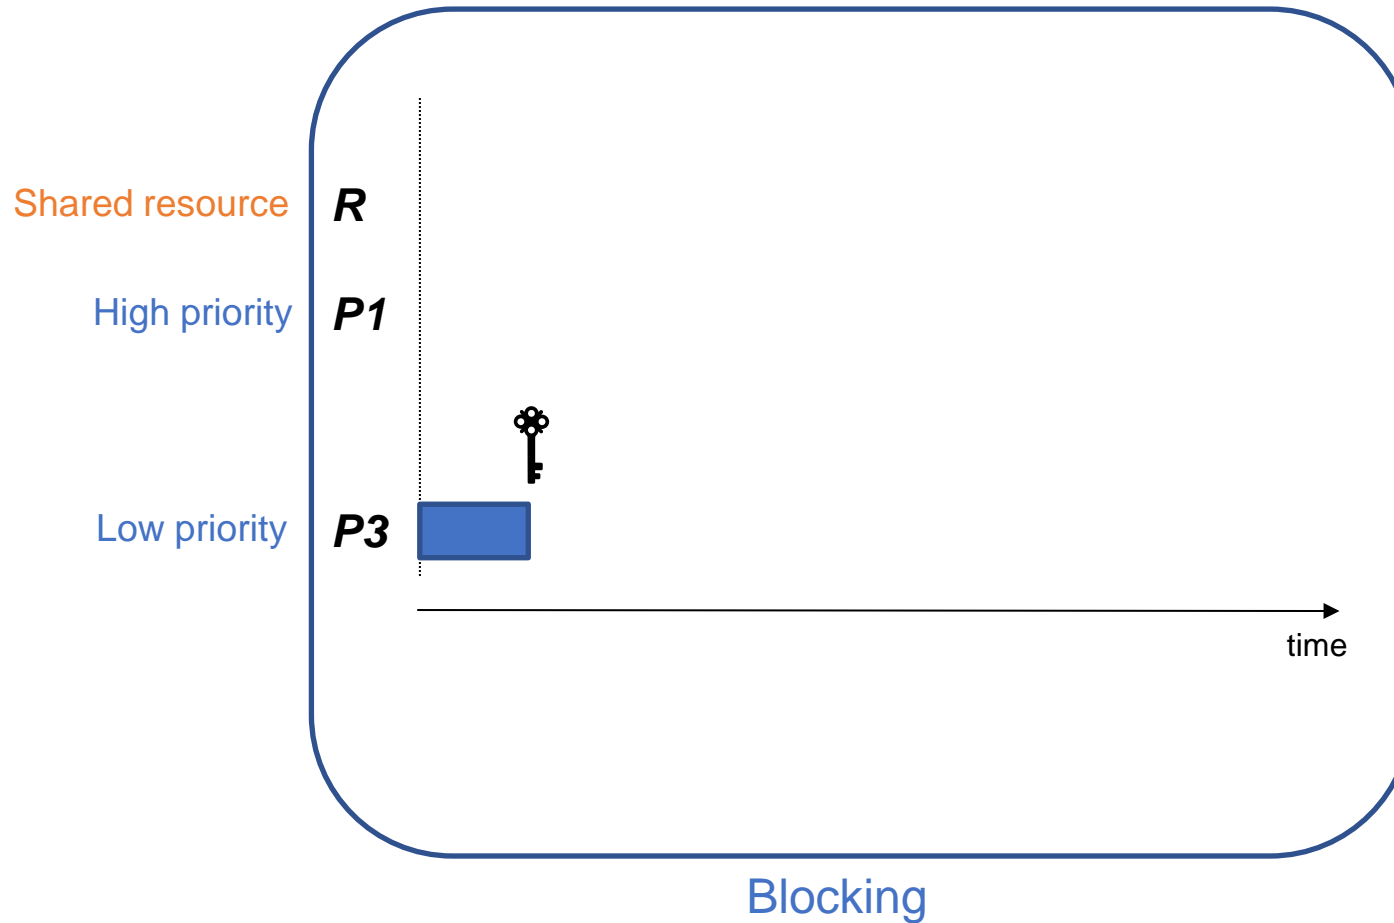
||

```
 $P_y = [$   
  while (true) {  
    do_something();  
    lock( $m$ );  
     $x := x+1$ ;  
    unlock( $m$ );  
    do_something_else();  
  }  
 $]$ 
```

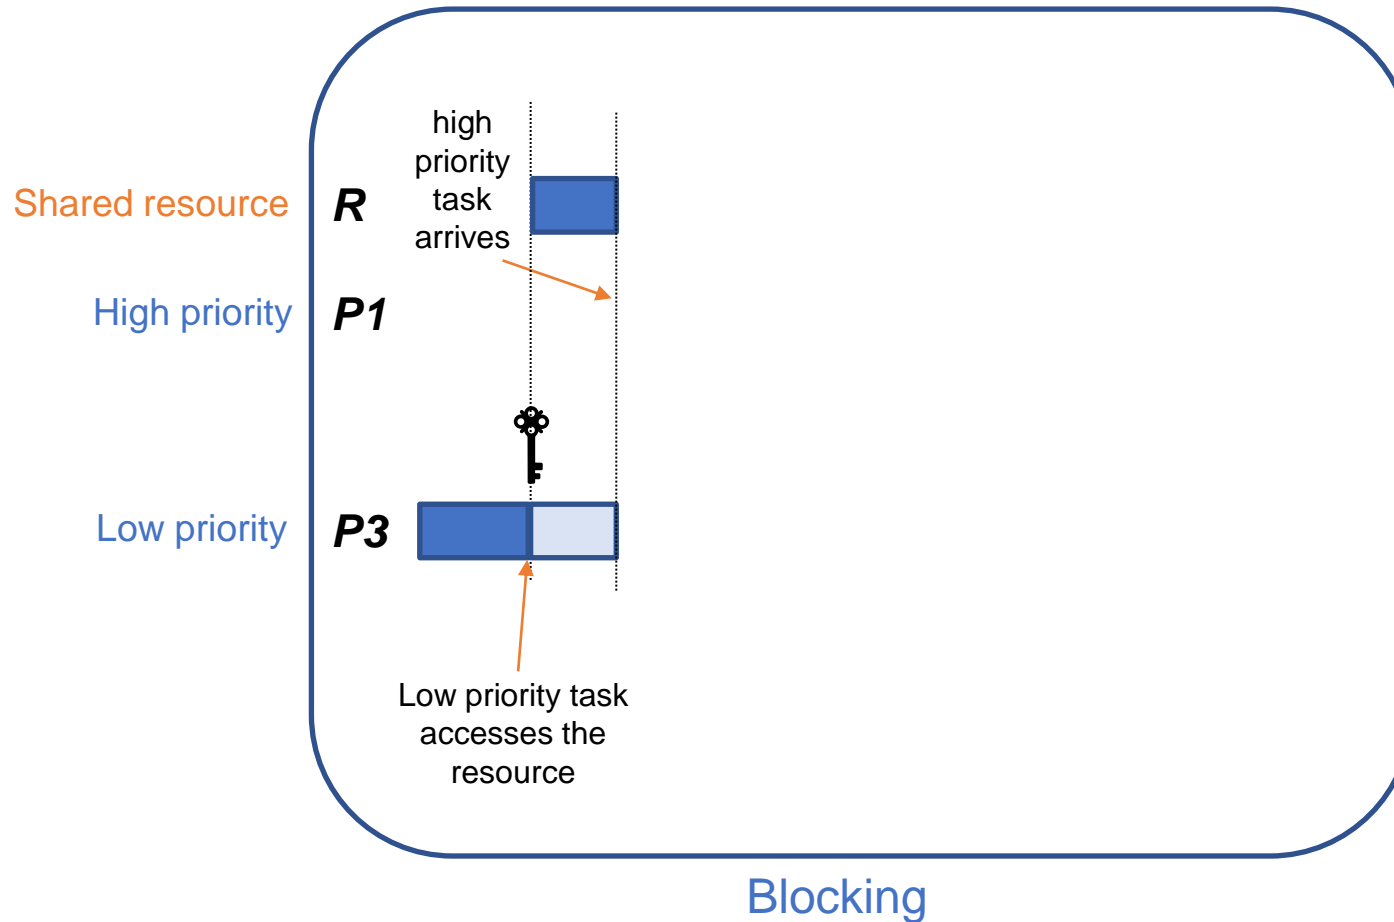
```
pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
/* static initialization, not always possible */
status = pthread_mutex_init (&m, attr); /* attr: NULL; should return 0 */
status = pthread_mutex_destroy (&m); /* should return 0 */
status = pthread_mutex_lock (&m); /* should return 0 */
status = pthread_mutex_trylock (&m); /* returns EBUSY if m is locked */
status = pthread_mutex_unlock (&m); /* should return 0 */
```

## Challenges with priority scheduling when protecting shared resource accesses

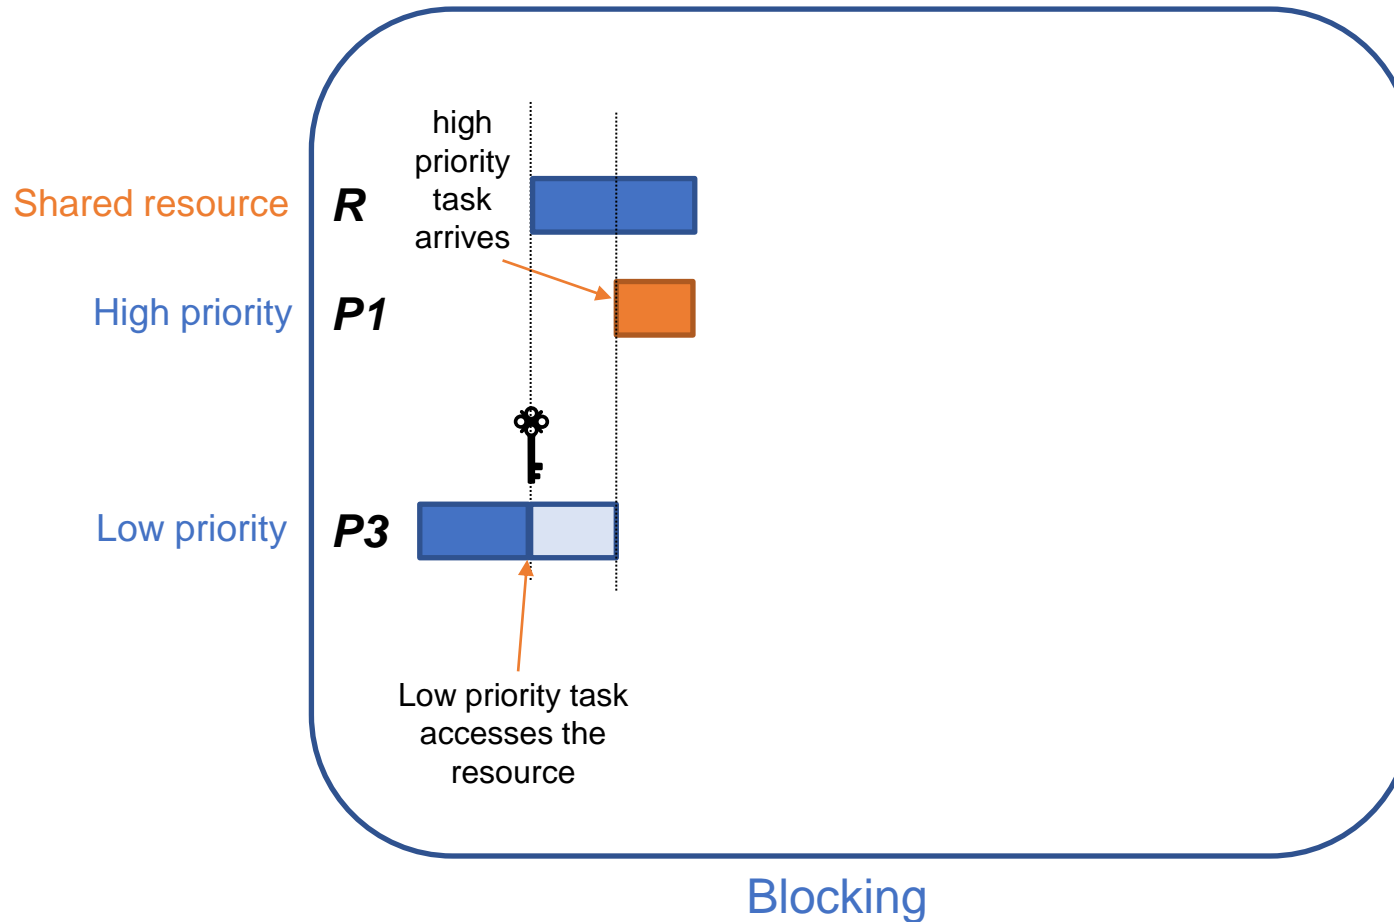
- Blocking:
  - A low priority task obtains a resource; a high priority task then waits on it



- Blocking:
  - A low priority task obtains a resource; a high priority task then waits on it

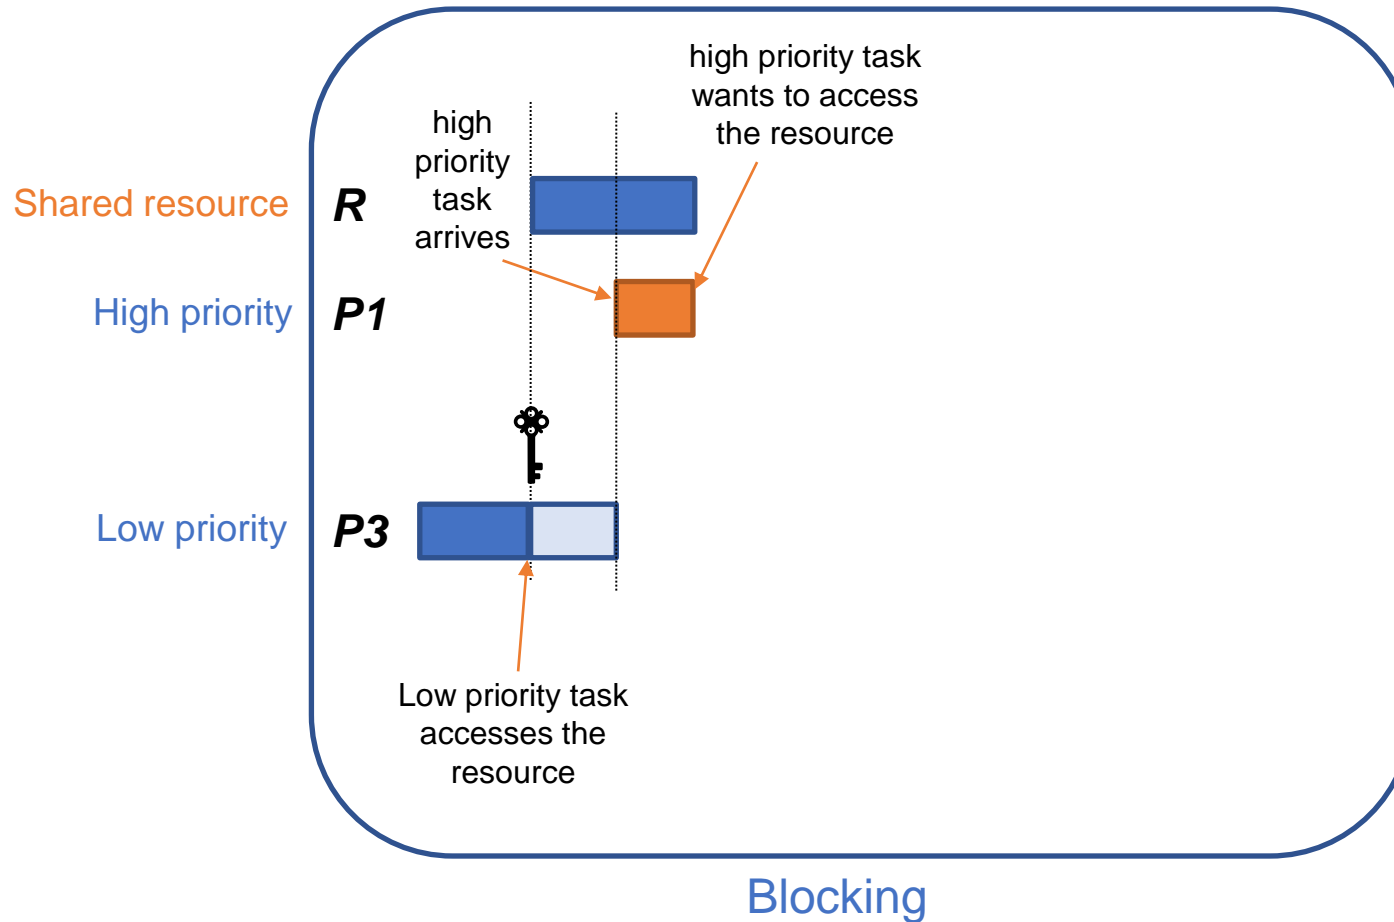


- Blocking:
  - A low priority task obtains a resource; a high priority task then waits on it

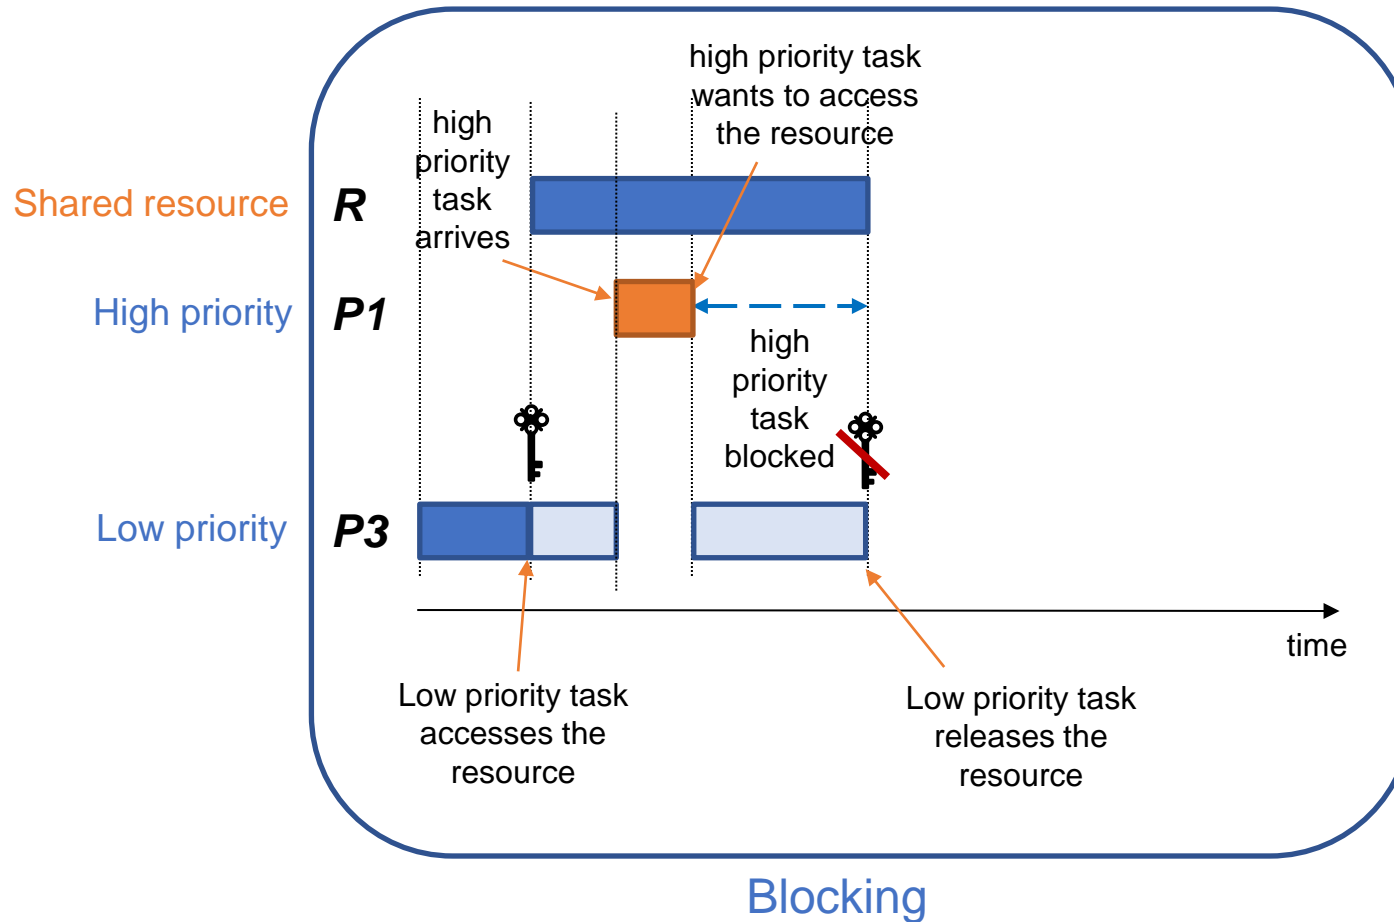




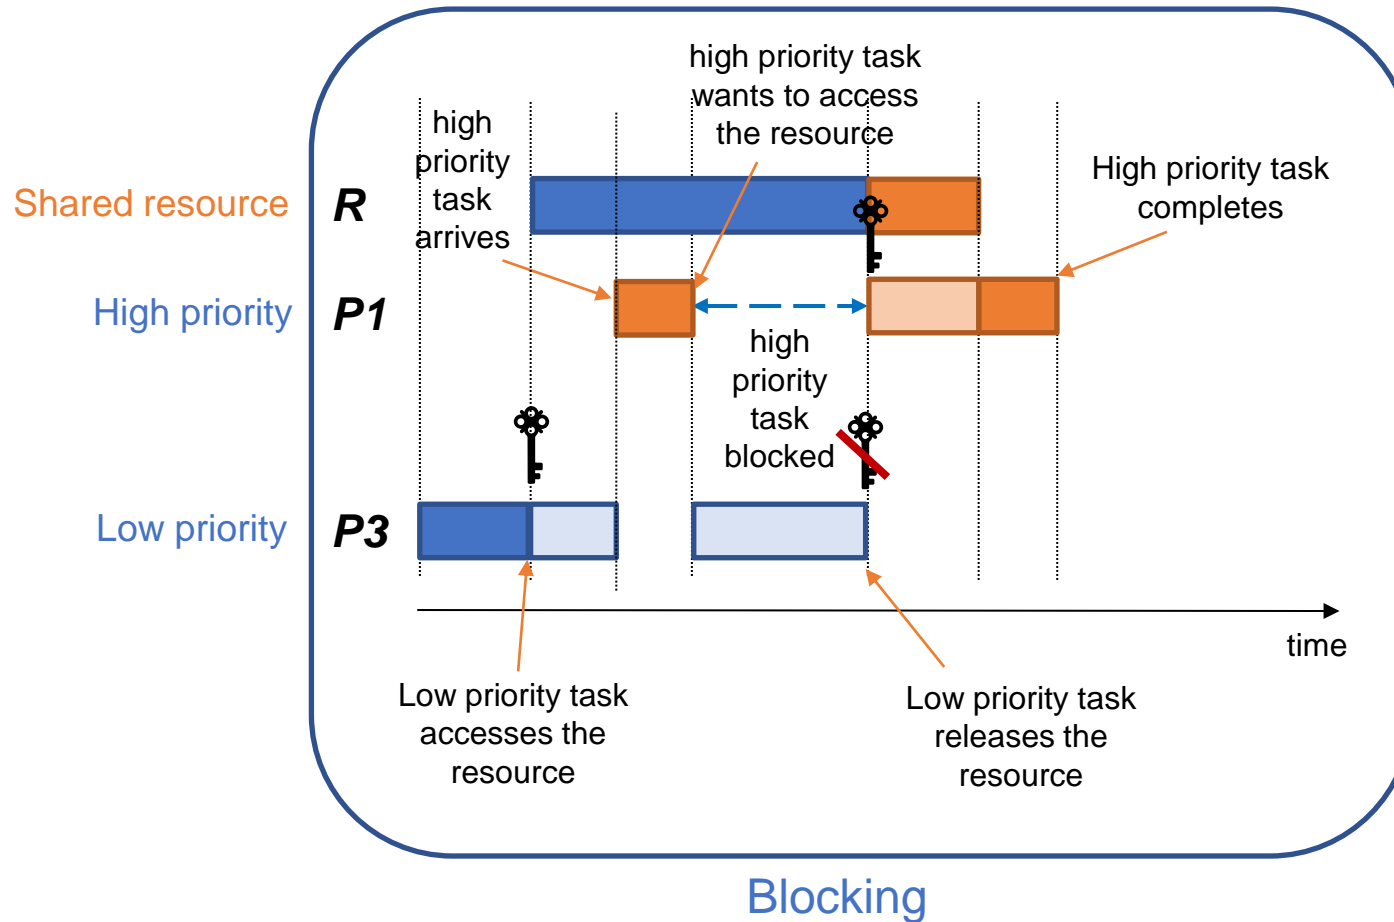
- Blocking:
  - A low priority task obtains a resource; a high priority task then waits on it



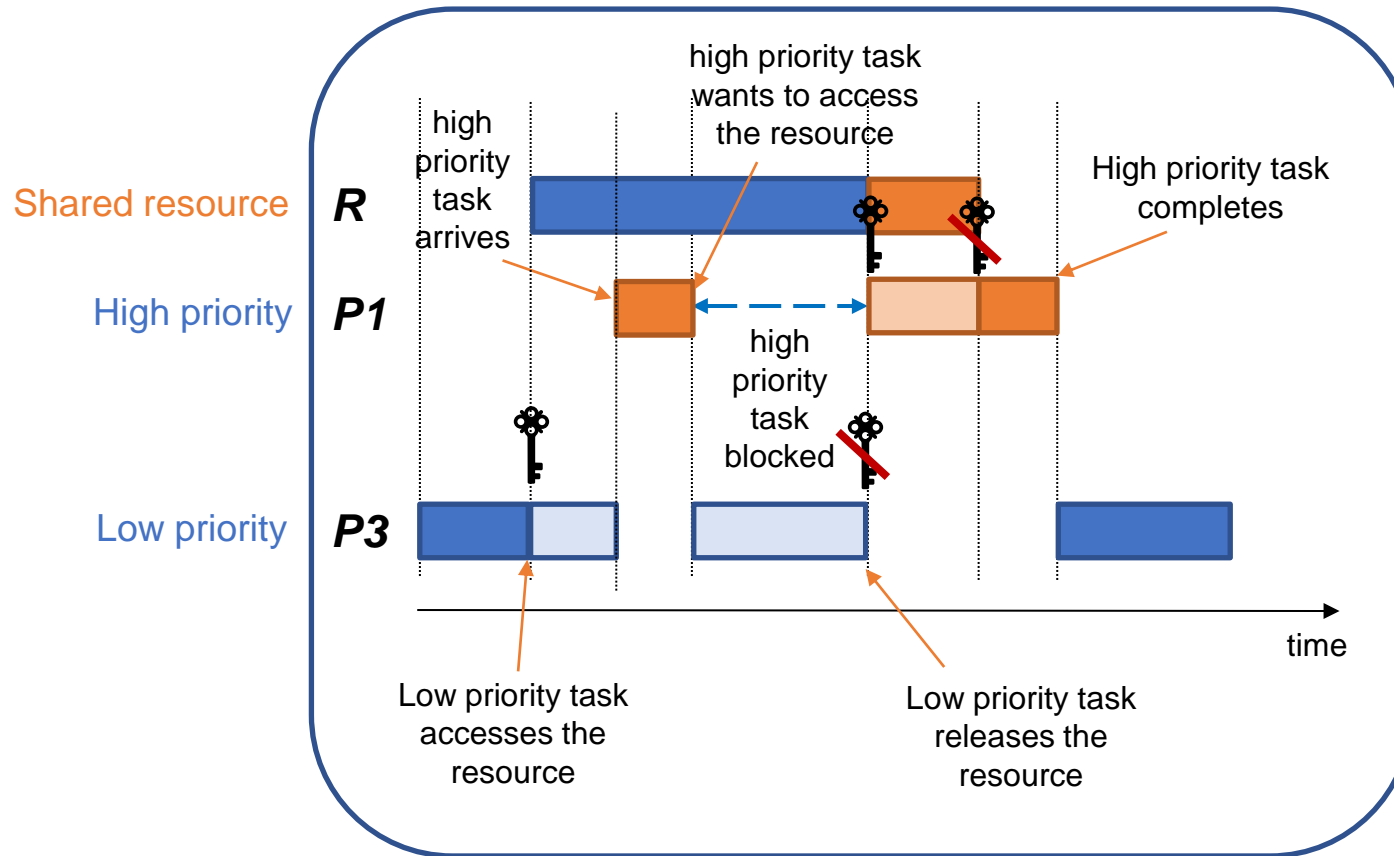
- Blocking:
  - A low priority task obtains a resource; a high priority task then waits on it



- Blocking:
  - A low priority task obtains a resource; a high priority task then waits on it




- **Blocking:**
  - A low priority task obtains a resource; a high priority task then waits on it



Blocking

- Mars Pathfinder

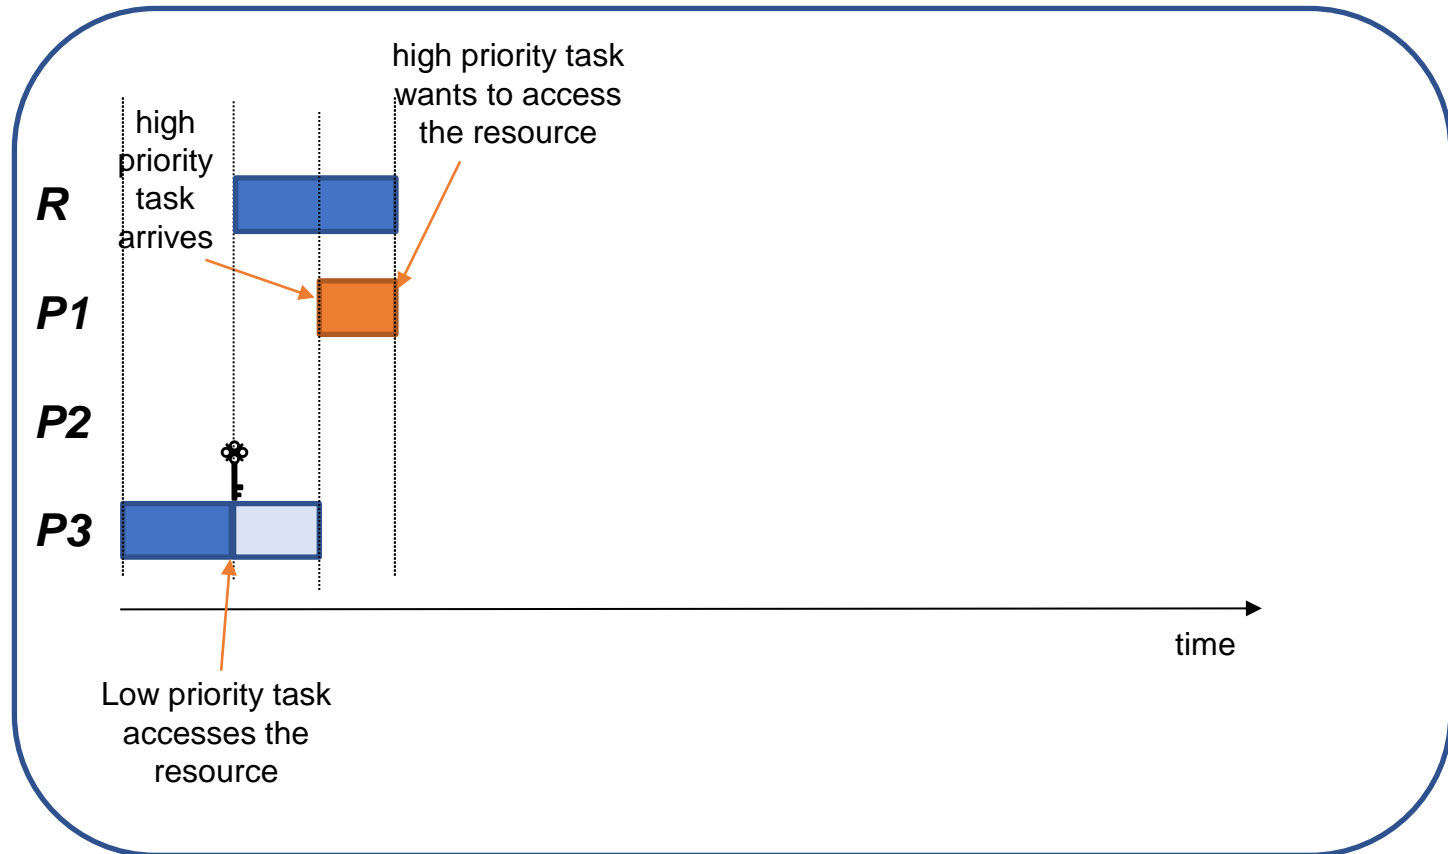


Copyright: NASA - <http://photojournal.jpl.nasa.gov/catalog/PIA01122>

**What really happened on Mars? Authoritative account:**

[https://web.archive.org/web/20150611035733/http://research.microsoft.com/en-us/um/people/mbj/mars\\_pathfinder/authoritative\\_account.html](https://web.archive.org/web/20150611035733/http://research.microsoft.com/en-us/um/people/mbj/mars_pathfinder/authoritative_account.html)

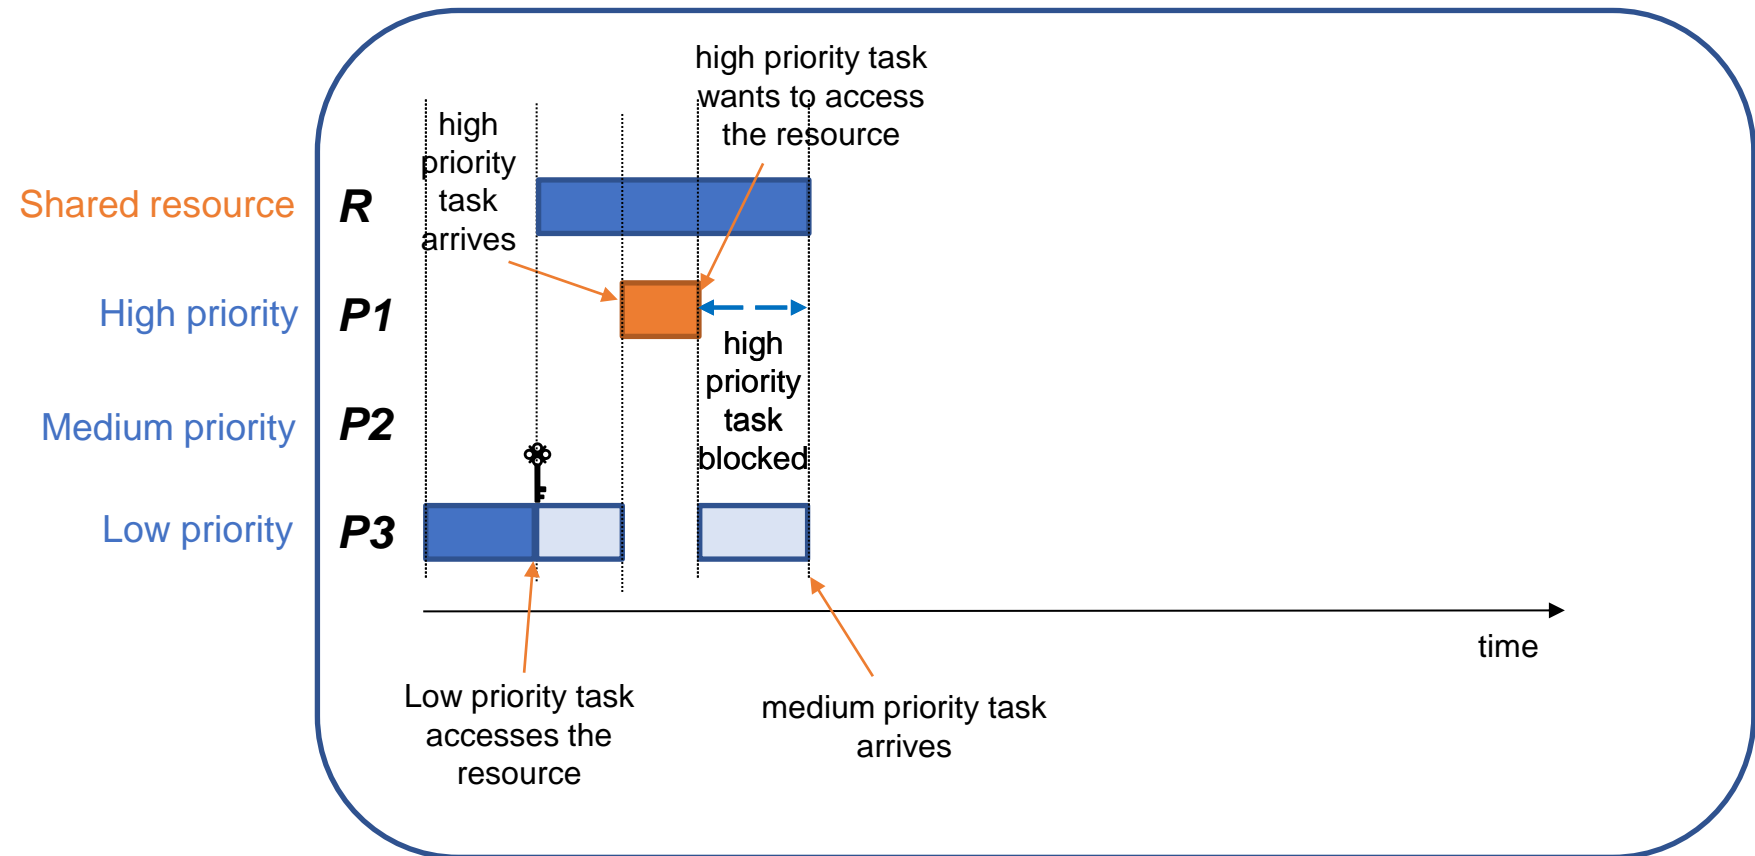
- Priority inversion:



Priority inversion

# Priority inversion

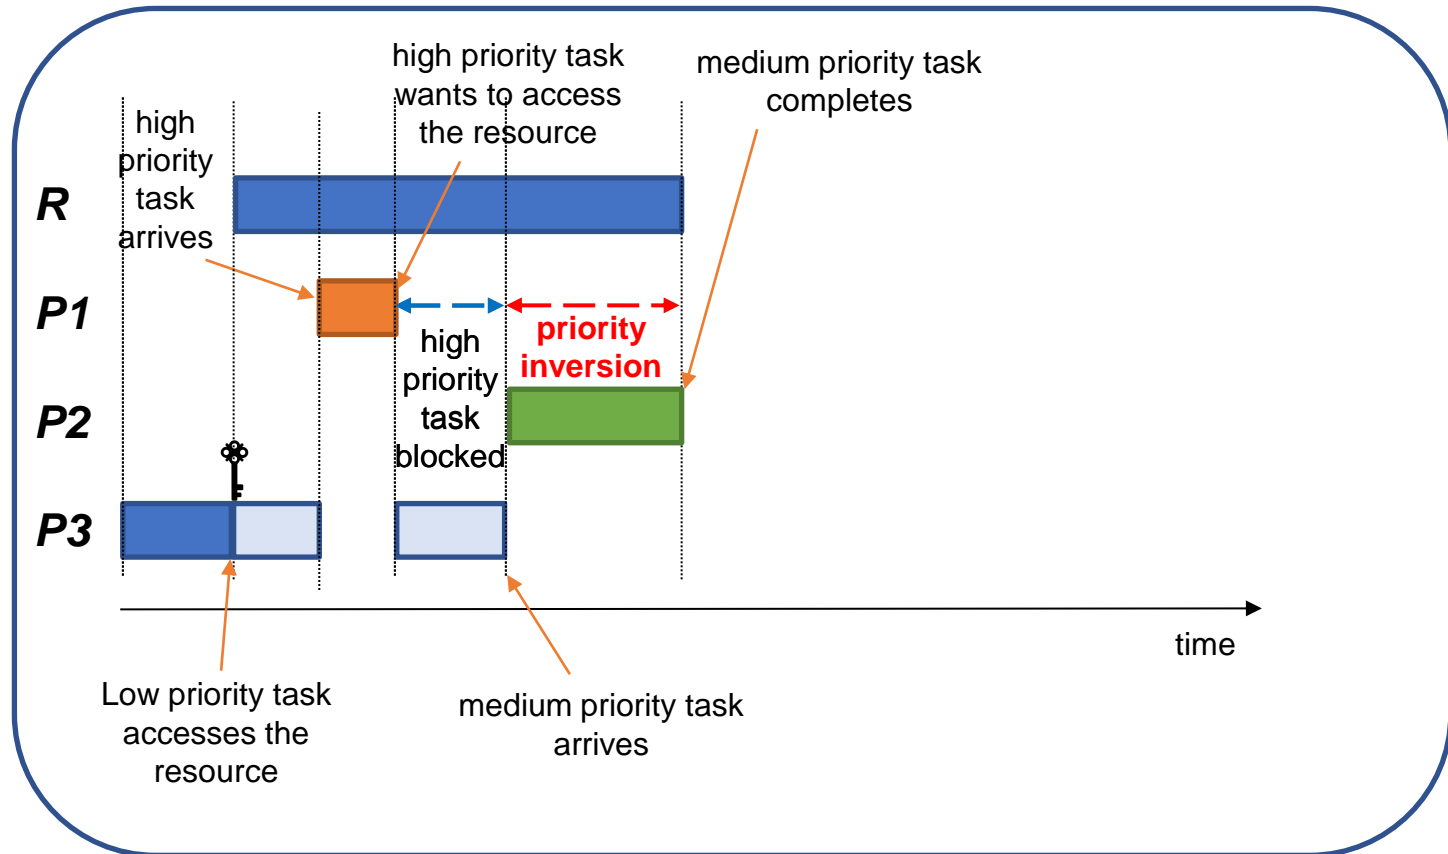
- Priority inversion:



Priority inversion

# Priority inversion

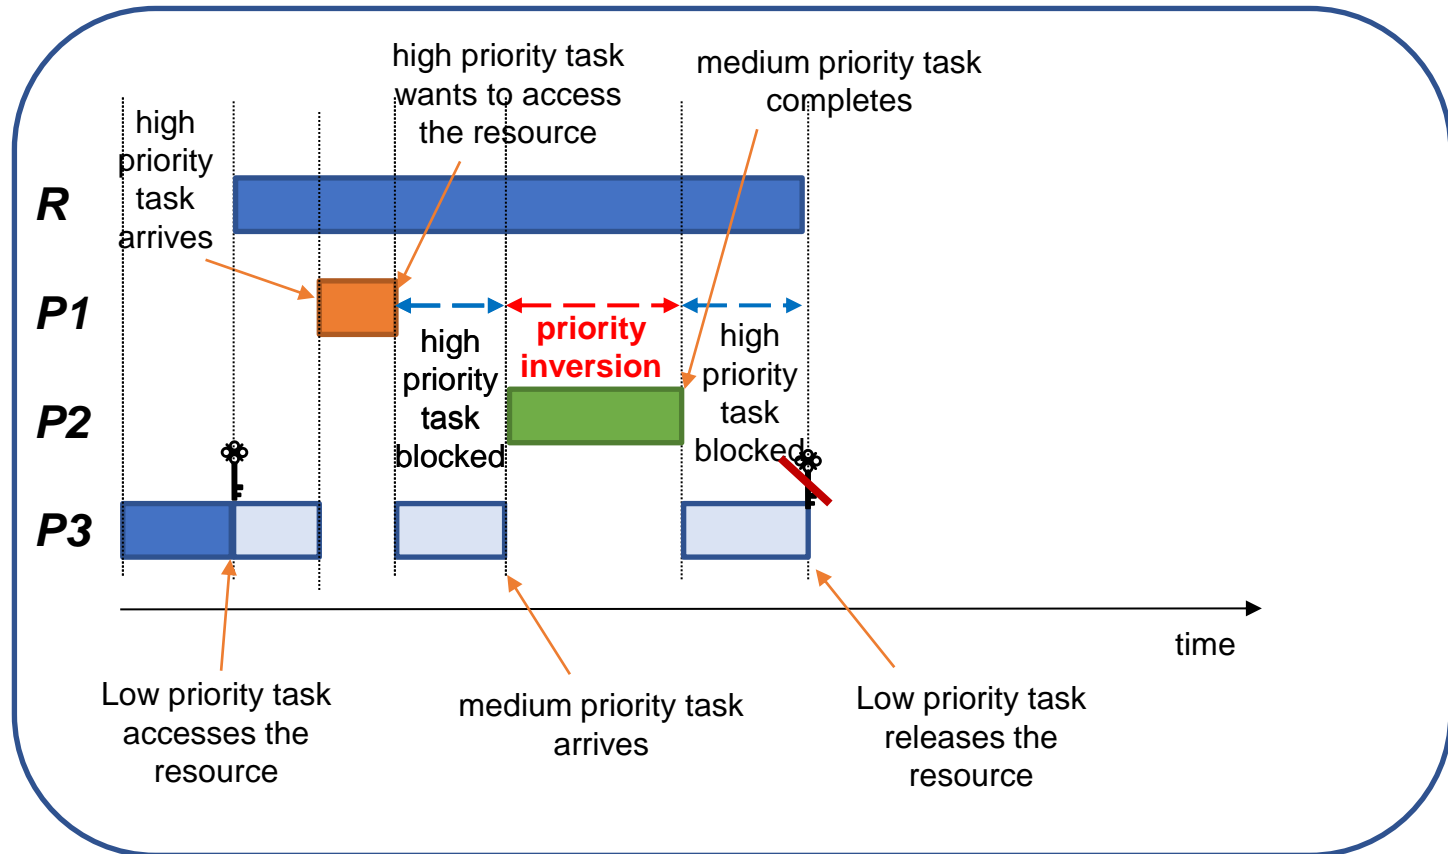
- Priority inversion:



Priority inversion



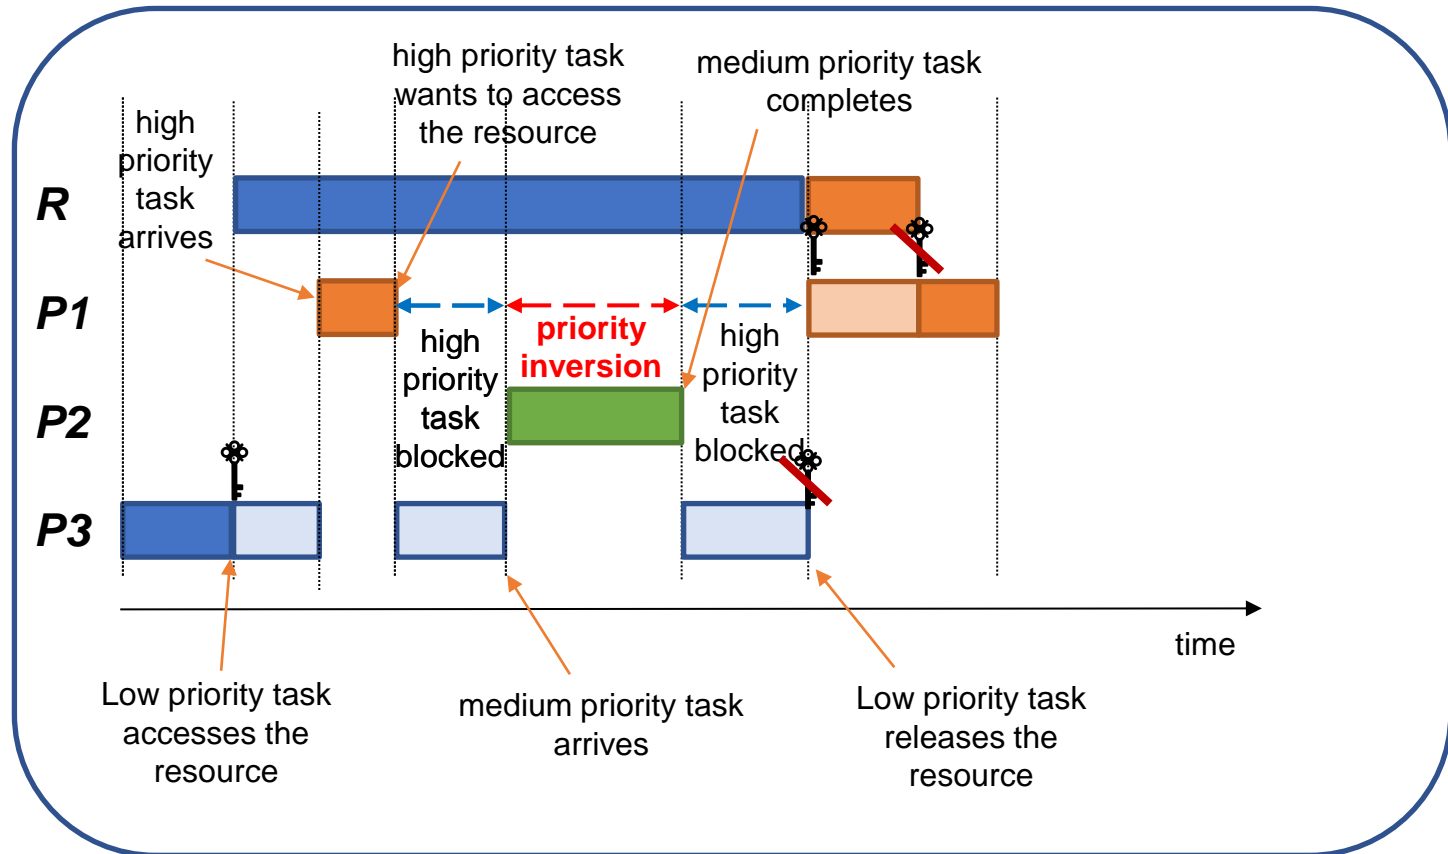
- Priority inversion:



Priority inversion

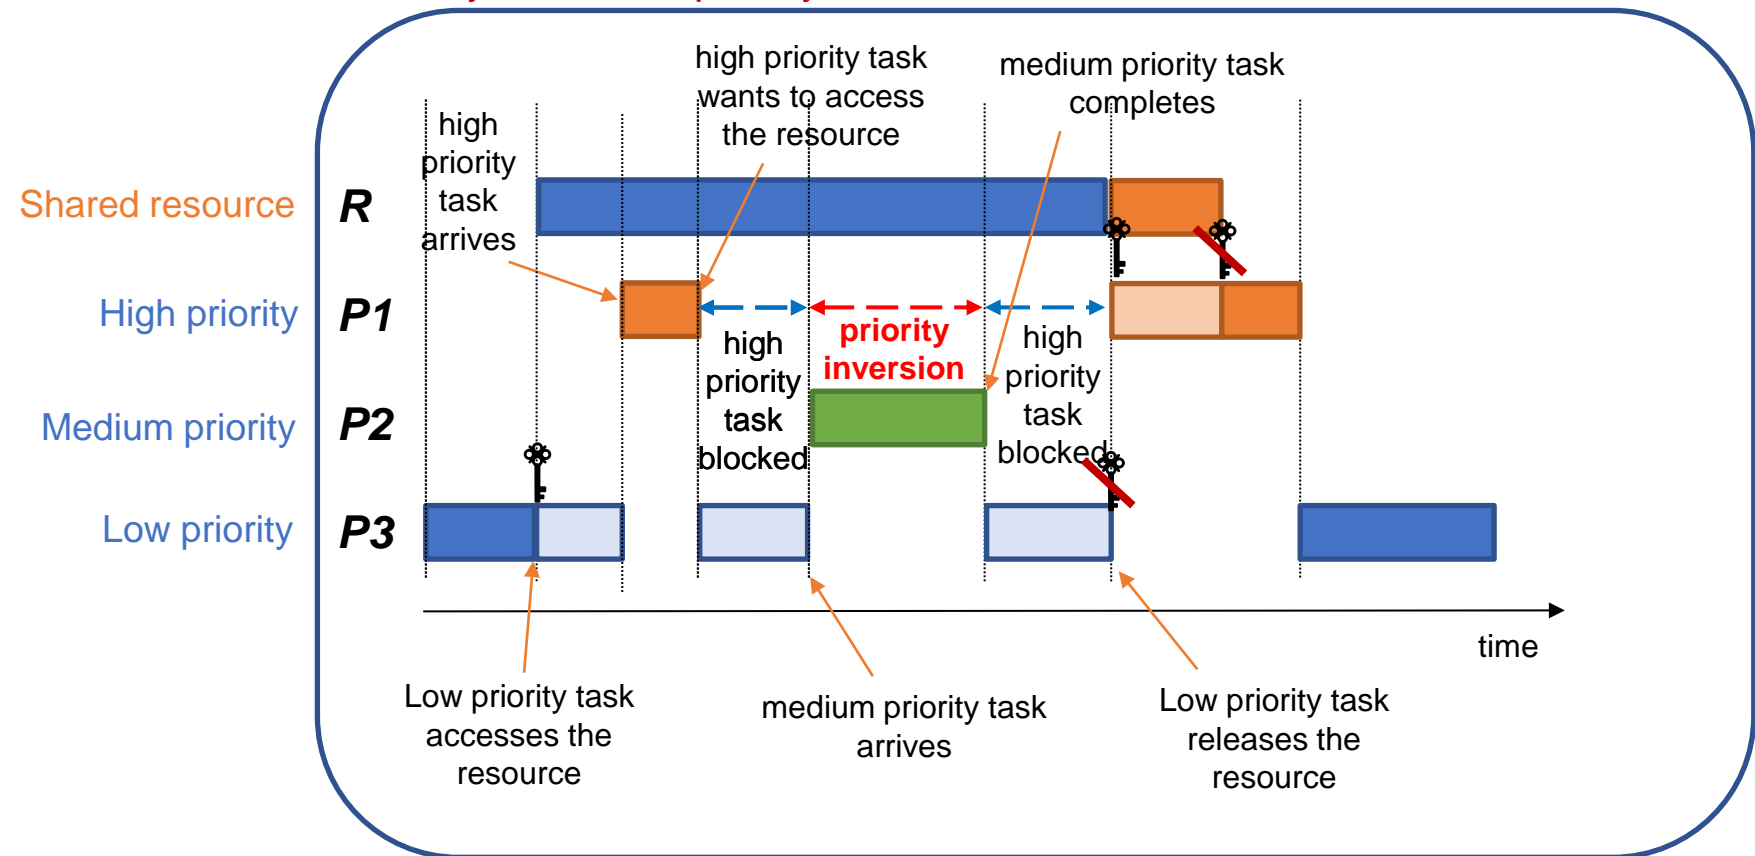
# Priority inversion

- Priority inversion:



Priority inversion

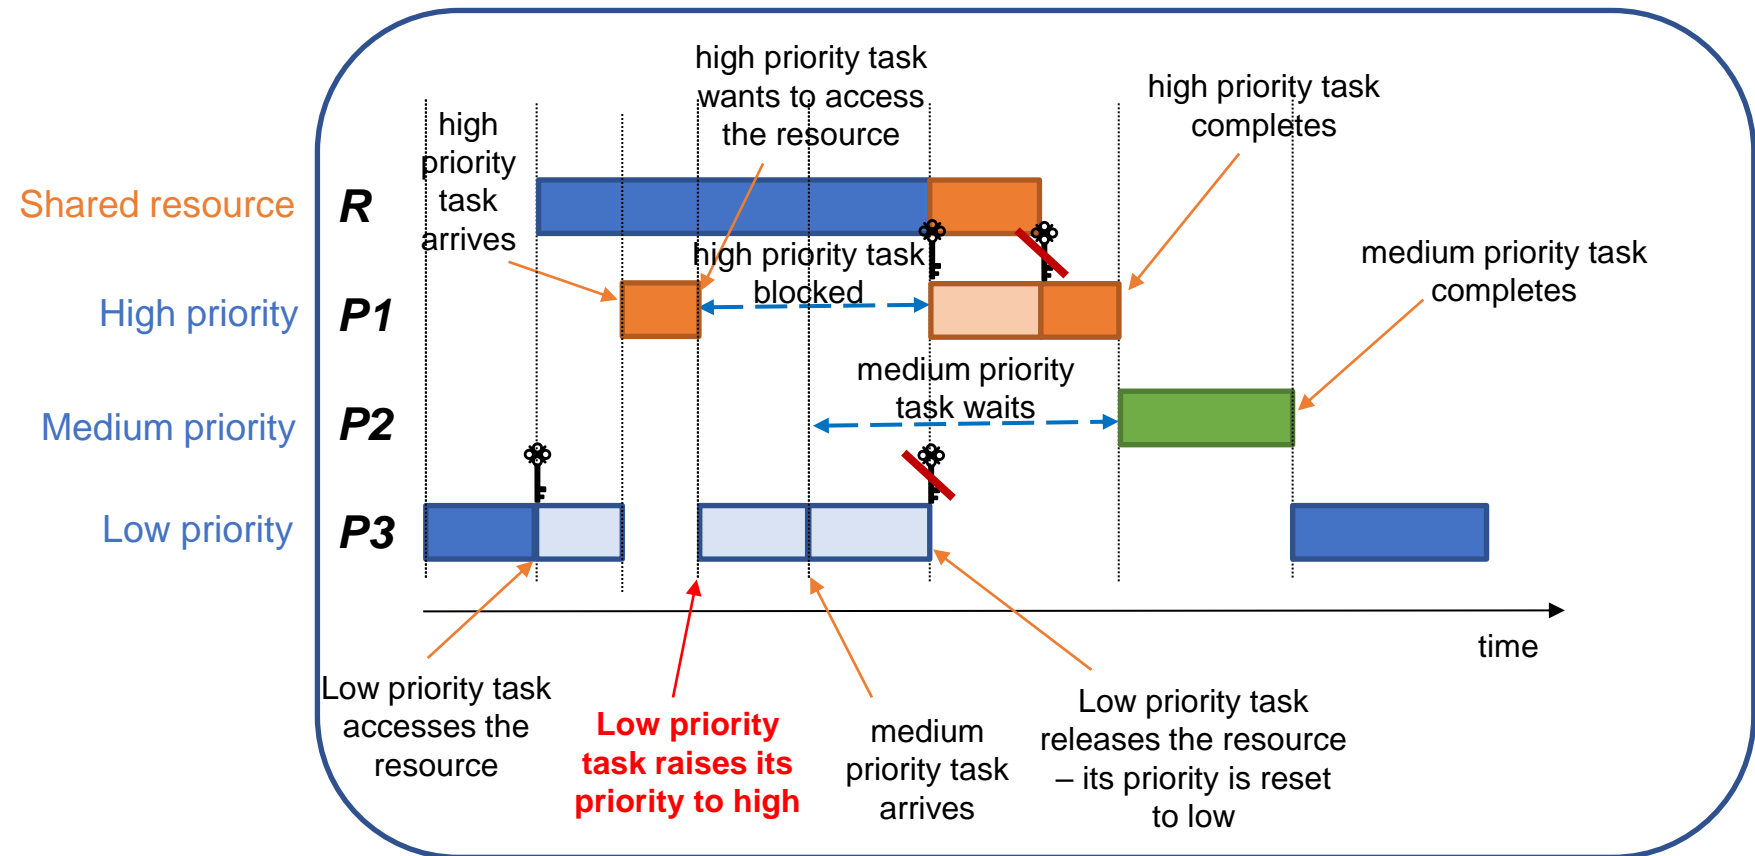
- Priority inversion:
  - A middle priority task pre-empt the low priority task that holds the resource
    - the **high priority task** now waits on the middle priority task and executes effectively at the low priority



Priority inversion

# A solution to priority inversion: priority inheritance protocol

- The **priority** of the task  $P$  using the resource is **dynamically adjusted** to be the maximum of **the priority of any other task that is blocked on the allocated resources** of  $P$  and its own priority
- ...middle priority tasks will wait now.



Priority inheritance

- Atomic actions and traces
- Issue with **concurrency** → **race conditions**
- Synchronization using **Peterson's algorithm**
- **5 correctness requirements**
- Proof for mutual exclusion and absence of deadlock using a **proof by contradiction**
- Preparation for next lecture (Wednesday)
  - 3 short videos (total: 20 minutes)
  - Topic: actions synchronization using semaphores