

Operating Systems (2INC0)

Concurrency and Atomicity (05) Introduction

Dr. Geoffrey Nelissen

Courtesy of Prof. Dr. Johan Lukkien and
Dr. Tanir Ozcelebi



Interconnected
Resource-aware
Intelligent Systems

TU/e


Technische Universiteit
Eindhoven
University of Technology

Where innovation starts

Starting point: the *sequential task*

Discrete sequence of states (e.g., observable in program code):

- needed: indivisible, atomic steps/actions between the states
- execution never observed to be half-way an atomic action



Execution: path through state-space

Initially:
 $StateP = 1$

Program:
 $StateP := 2;$
 $StateP := 3;$
 $StateP := 4;$
 $StateP := 5;$


In a task, the finest level of detail w.r.t. progress consists of *atomic actions*.

Atomic? (consider shared variables x, y and atomic memory operations)


- $x := 1$
 - `mov #1, r1; str r1, @x`
 - no 'internal' interference point, hence to be regarded as atomic, assuming a correct implementation of interrupt handling (note: actual execution not atomic)
- $x := y$
 - `mov @y, r1; str r1, @x`
 - 'internal' interference point: r1 may store an old copy of y for a long time while computations with y continue.
- $x := x+1$ (similar problem)
 - `mov @x, r1; inc r1; str r1, @x`
- **Single reference rule:** a **statement** (expression) in a programming language may be regarded as **atomic if at most one reference to a shared variable occurs** (we ignore here compiler optimizations).
- **Defined atomicity:** When we want to regard a non-atomic statement S as atomic, we write **< S >**, e.g. **< $x := x+1$ >**.
 - needs OS and/or hardware support...

What about $z := x; x := z+1$ (z private variable)?
→ 2 atomic statements
interference point visible in the program


Concurrent execution



- Execution path of a single task (above)
- Concurrent processes take a joint path through the joint state space.
- **Trace**: **sequence of atomic actions**, obtained by **interleaving** the execution of concurrent tasks
 - maintains the internal **execution order of the individual tasks**
- **many possible traces** in the joint state space (e.g. blue and red traces)



Traces of concurrent tasks




- In between any pair of instructions of one task, (part of) another task or collection of tasks can be executed, including the OS.
 - **Problem:** Old copies of shared variables can be stored in internal registers or in memory locations.
 - Example:
 - initially: $x=1, y=2$
 - program: $x := y \parallel y := x$
- $\xrightarrow{\text{red}} \text{mov } @x, r; \text{str } r, @y$
 $\xrightarrow{\text{blue}} \text{mov } @y, r; \text{str } r, @x$
- What are the possible traces? ... the final values?
 - 1) $\text{mov } @x, r; \text{str } r, @y; \text{mov } @y, r; \text{str } r, @x \rightarrow (1,1)$
 - 2) $\text{mov } @y, r; \text{str } r, @x; \text{mov } @x, r; \text{str } r, @y \rightarrow (2,2)$
 - ...
 - 6) $\text{mov } @x, r; \text{mov } @y, r; \text{str } r, @y; \text{str } r, @x \rightarrow (2,1)$
 - **Exercise:** Is the result (1,2) possible?
 - **Exercise:** Assuming all traces are equally likely, what is the probability of the result being (1,1)? How about (2,1)?

More on traces

- **Trace = sequence of (atomic) actions**;
It represents the possible steps of a program execution
 - **actions = assignments or tests**
- **Possible trace = trace in which all the tests yield true**
 - being possible depends on the initial program state
- Traces can be **finite or infinite**, and a program text has many traces

Example

- Consider two concurrent programs:



Example traces of the **left** program (mind the notation):

- $(x:=false)(x)$
 $(x:=false)(\neg x)(skip)(x)$
 $(x:=false)(\neg x)(skip)(\neg x)(skip)(x)$
.....
 - note: (x) and $(\neg x)$ denote tests
- For this program in isolation** (without the right task):
Only “the infinite trace” is possible, because (x) never yields *true*.

More on traces (2)

- The traces of a concurrent program are obtained by interleaving traces of all concurrent parts. Example: interleave $(x:=false)(x)$ with $(y:=E)(x:=true)$.
 - $(y:=E)(x:=true)(x:=false)(x)$
 - $(y:=E)(x:=false)(x:=true)(x) \rightarrow \text{possible}$
 - $(y:=E)(x:=false)(x)(x:=true)$
 - $(x:=false)(y:=E)(x:=true)(x) \rightarrow \text{possible}$
 - $(x:=false)(y:=E)(x)(x:=true)$
 - $(x:=false)(x)(y:=E)(x:=true)$
- Note: **Finite traces are now possible traces.**
 - For example, the **two traces in bold** are now possible, while $(x:=false)(x)$ was not a possible trace for single program in isolation

REMEMBER:

A trace is a possible one
if all its tests yield true

Summary

- **Shared variables:** accessible to several tasks (processes/threads)
- **Private variables:** accessible only to a single task
- **Atomic action:** finest grain of detail, indivisible
 - typically, assignments and tests in a program
 - **sufficient at program level: single reference to shared variable in statement**
 - ignoring possible optimization and reordering by compiler/processor
- **Concurrent execution:** interleaving of atomic actions
- **Interference:** disturbing others' assumptions about the state
 - usually, caused by “bad” interleavings
 - particularly with shared variables
- **Race condition:** situation in which correctness depends on the execution order of concurrent activities (“bad interference”)