

2INC0 - Operating Systems

Actions synchronization

Geoffrey Nelissen



Interconnected
Resource-aware
Intelligent Systems

TU/e

Technische Universiteit
Eindhoven
University of Technology

Where innovation starts

- Second homework is available (**deadline on Sunday**)
- If you have **questions about** the practical **assignments**,
contact Olav Bunte

A note on “zombies”

- A zombie is a **child process** that **terminated**, but the **parent does not call wait() or wait_pid()**
- It does not matter what the child does (i.e., whether it calls **exec()** or any other system call or function)



© PopCap studio

```
int main() {  
  
    int i = 0;  
  
    for (int k=0; k< 5; ++k) {  
        fork();  
        execlp("/bin/sh", "/bin/sh", "-c", "ls -l /bin/??", (char *)NULL);  
    }  
  
    i++;  
    return 0;  
}
```

- What is an OS
- Purpose of an OS

- **Introduction to operating systems** (lecture 1)
- **Processes, threads and scheduling** (lectures 2+3)
- **Concurrency and synchronization**
 - atomicity and interference (lecture 4)
 - **actions synchronization (lecture 5)**
 - condition synchronization (lecture 6)
 - deadlock (lecture 7)
- **File systems** (lecture 8)
- **Memory management** (lectures 9+10)
- **Input/output** (lecture 11)

- How to implement concurrency (for efficiency)

- Dangers of concurrency (**race conditions**)
- How to prove properties using **traces**
- Synchronization (**critical sections**)

How can we **check** complex **properties**
without checking all possible traces
of a concurrent program

- **Proving program properties using invariants**
- **Actions synchronization using semaphores**
- **Preventing deadlocks**

How can we **enforce properties** during the
program execution **using semaphores**

Invariant

- An assertion that **holds at all control points** of a program

Topology invariant

- An invariant **derived** directly from the program code
- Examples:
 - number of times an **operation is executed** in comparison to another action.
 - **Value of a variable** based on the number of times a set of actions was executed

```
int x = 1;  
int y = 0;
```

```
Task1 = [[  
  while (true) {  
    < x++; >  
    < y++; >  
  }  
]]
```

||

```
Task2 = [[  
  while (true) {  
    < y--; >  
    < x--; >  
  }  
]]
```

Invariant

- An assertion that **holds at all control points** of a program

Topology invariant

- An invariant **derived** directly from the program code
- Examples:
 - number of times an **operation is executed** in comparison to another action.
 - **Value of a variable** based on the number of times a set of actions was executed

```
int x = 1;  
int y = 0;
```

```
Task1 = [  
  while (true) {  
    X1: < x++; >  
    Y1: < y++; >  
  }  
]
```

||

```
Task2 = [  
  while (true) {  
    Y2: < y--; >  
    X2: < x--; >  
  }  
]
```

cX = number of times
action X executed

$$\text{I1: } 0 \leq cX_1 - cY_1$$

Number of times Y₁ executes is **never more**
than the number of times X₁ executes

Invariant

- An assertion that **holds at all control points** of a program

Topology invariant

- An invariant **derived** directly from the **program code**
- **Examples:**
 - **number of times an operation is executed** in comparison to another action.
 - **Value of a variable** based on the number of times a set of actions was executed

```
int x = 1;  
int y = 0;
```

```
Task1 = ||  
    while (true) {  
        X1: < x++; >  
        Y1: < y++; >  
    }  
||
```

$$I1: 0 \leq cX_1 - cY_1 \leq 1$$

Number of times Y₁ executes is **never more**
than the number of times X₁ executes

||

```
Task2 = ||  
    while (true) {  
        Y2: < y--; >  
        X2: < x--; >  
    }  
||
```

$$I2: 0 \leq cY_2 - cX_2 \leq 1$$

Number of times Y₁ executes is
at most one time less than X₁

cX = number of times
action X executed

Invariant

- An assertion that **holds at all control points** of a program

Topology invariant

- An invariant **derived** directly from the **program** code
- **Examples:**
 - number of times an **operation is executed** in comparison to another action.
 - **Value of a variable** based on the number of times a set of actions was executed

```
int x = 1;  
int y = 0;
```

```
Task1 = [[  
    while (true) {  
        X1: < x++; >  
        Y1: < y++; >  
    }  
]]
```

||

```
Task2 = [[  
    while (true) {  
        Y2: < y--; >  
        X2: < x--; >  
    }  
]]
```

cX = number of times
action X executed

$$I1: 0 \leq cX_1 - cY_1 \leq 1$$

$$I2: 0 \leq cY_2 - cX_2 \leq 1$$

$$I3: y = ?$$

Invariant

- An assertion that **holds at all control points** of a program

Topology invariant

- An invariant **derived** directly from the **program** code
- **Examples:**
 - number of times an **operation is executed** in comparison to another action.
 - **Value of a variable** based on the number of times a set of actions was executed

```
int x = 1;  
int y = 0;
```

```
Task1 = [  
  while (true) {  
    X1: < x++; >  
    Y1: < y++; >  
  }  
]
```

||

```
Task2 = [  
  while (true) {  
    Y2: < y--; >  
    X2: < x--; >  
  }  
]
```

cX = number of times
action X executed

$$I1: 0 \leq cX_1 - cY_1 \leq 1$$

$$I2: 0 \leq cY_2 - cX_2 \leq 1$$

$$I3: y = cY_1 - cY_2$$

$$I4: x = ?$$

Invariant

- An assertion that **holds at all control points** of a program

Topology invariant

- An invariant **derived** directly from the **program** code
- **Examples:**
 - number of times an **operation is executed** in comparison to another action.
 - **Value of a variable** based on the number of times a set of actions was executed

```
int x = 1;  
int y = 0;
```

```
Task1 = [  
  while (true) {  
    X1: < x++; >  
    Y1: < y++; >  
  }  
]
```

||

```
Task2 = [  
  while (true) {  
    Y2: < y--; >  
    X2: < x--; >  
  }  
]
```

cX = number of times
action X executed

$$I1: 0 \leq cX_1 - cY_1 \leq 1$$

$$I2: 0 \leq cY_2 - cX_2 \leq 1$$

$$I3: y = cY_1 - cY_2$$

$$I4: x = 1 + cX_1 - cX_2$$

Using topology invariants to prove program properties

```
int x = 1;  
int y = 0;
```

```
Task1 = [  
  while (true) {  
    X1: < x++; >  
    Y1: < y++; >  
  }  
]
```

||

```
Task2 = [  
  while (true) {  
    Y2: < y--; >  
    X2: < x--; >  
  }  
]
```

cX = number of times
action X executed

$$I_1: 0 \leq cX_1 - cY_1 \leq 1$$

$$I_3: y = cY_1 - cY_2$$

$$I_2: 0 \leq cY_2 - cX_2 \leq 1$$

$$I_4: x = 1 + cX_1 - cX_2$$

Prove that $x > y$

using I_3 and I_4

$$\Leftrightarrow 1 + cX_1 - cX_2 > 0 + cY_1 - cY_2$$

$$\Leftrightarrow 1 + cX_1 + cY_2 > 0 + cY_1 + cX_2$$

by $I_1: cX_1 \geq cY_1$, by $I_2: cY_2 \geq cX_2$, and $1 > 0$

\Rightarrow true

```
int x = 4;  
int y = 0;  
int z = 1;
```

```
Task1 = [  
  while (true) {  
    Z1: < z := z+2; >  
    X1: < x := x+2; >  
  }  
]
```

||

```
Task2 = [  
  while (true) {  
    Y2: < y := y+1; >  
    Z2: < z := z-1; >  
  }  
]
```

||

```
Task3 = [  
  while (true) {  
    X3: < x := x+3 >  
    Y3: < y := y+3 >  
  }  
]
```

Prove that $x \geq y + z$

Solve the problem together with your neighbors.

- You have **10 minutes**.

```
int x = 4;
int y = 0;
int z = 1;
```

```
Task1 = [[
  while (true) {
    Z1: < z := z+2; >
    X1: < x := x+2; >
  }
]]
```

||

```
Task2 = [[
  while (true) {
    Y2: < y := y+1; >
    Z2: < z := z-1; >
  }
]]
```

||


```
Task3 = [[
  while (true) {
    X3: < x := x+3; >
    Y3: < y := y+3; >
  }
]]
```

Prove that $x \geq y + z$

- $x = 4 + 2 * cX1 + 3 * cX3$
 - $y = cY2 + 3 * cY3$
 - $z = 1 + 2 * cZ1 - cZ2$
 - **I1: $cZ1 - cX1 \leq 1$**
 - **I2: $cY2 - cZ2 \leq 1$**
 - **I3: $0 \leq cX3 - cY3$**
- Replacing x, y and z using the topology invariant above, we get that $x \geq y + z$ is equivalent to proving that $4 + 2 * cX1 + 3 * cX3 \geq cY2 + 3 * cY3 + 1 + 2 * cZ1 - cZ2$
- Rewriting: $3 + 3 (cX3 - cY3) \geq (cY2 - cZ2) + 2 (cZ1 - cX1)$
- Using I2, the above is true if: $3 + 3 (cX3 - cY3) \geq 1 + 2 (cZ1 - cX1)$
- Using I3, the above is true if: $3 \geq 1 + 2 (cZ1 - cX1)$
- Using I1, the above is true if: $3 \geq 3$
- true

- Proving program properties using invariants
- **Actions synchronization using semaphores**
- Preventing deadlocks

Example of a synchronization problem



```
int BufferCnt := 0;
```


```
Task Producer =  
[[ while( true ){  
    produce_items();  
  
    BufferCnt := BufferCnt+4;  
    push_items;  
}  
]]
```

||

```
Task Consumer =  
[[ while( true ){  
  
    BufferCnt := BufferCnt-2;  
    pop_items();  
    use_items();  
}  
]]
```

What can go wrong with this program?

Example of a synchronization problem



```
int BufferCnt := 0;
```

```
Task Producer =  
[[ while( true ){  
    produce_items();  
  
    BufferCnt := BufferCnt+4;  
    push_items();  
}  
]]
```

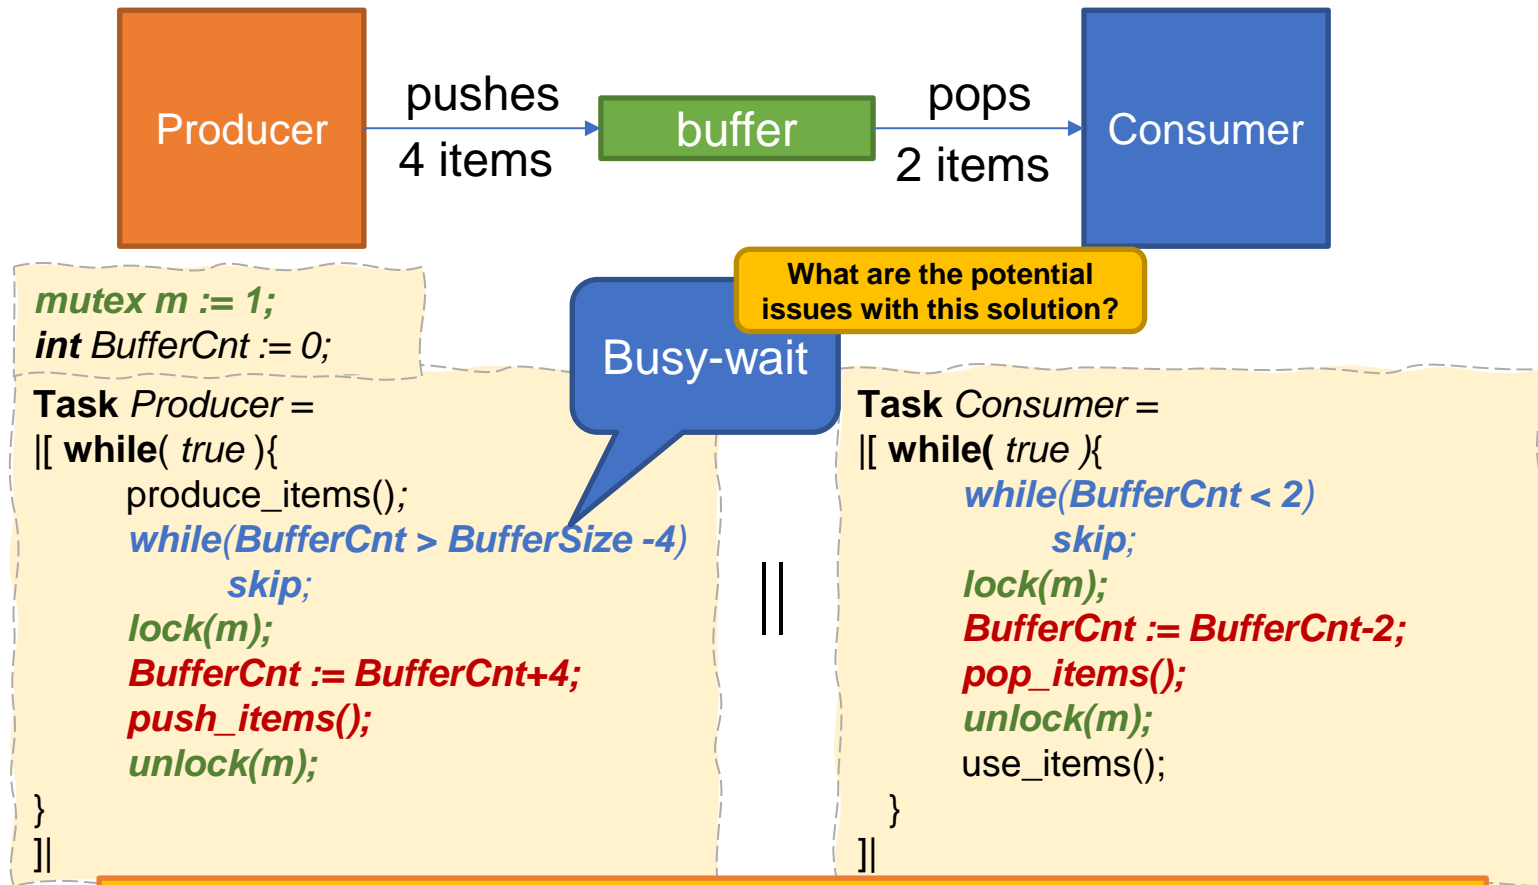
||

```
Task Consumer =  
[[ while( true ){  
  
    BufferCnt := BufferCnt-2;  
    pop_items();  
}]
```

What is the **solution** for the first one?
Use a mutex to protect the **critical sections**

1. **Race condition** when *Producer* and *Consumer* update the value of *BufferCnt* and when they push and pop items in the buffer
2. *Consumer* may try to **read from an empty buffer**
3. *Producer* may try to **write in a buffer that is full** (results in lost items)

Example of a synchronization problem



1. **Race condition** when *Producer* and *Consumer* update the value of *BufferCnt*
2. *Consumer* may try to **read from an empty buffer**
3. *Producer* may try to **write in a buffer that is full** (results in lost items)

Busy-waiting

Repeated testing without going to sleep

- Acceptable only **when**
 - waiting is guaranteed to be **short** in comparison to the cost of context switching or
 - there is **nothing else to do** (e.g., when the task executes on dedicated hardware)
- Busy-waiting using *while(.) skip*; **works only if**
 - the tests in the while and locking the mutex are **executed atomically**

Why?

Alternative solution:
use **semaphores** for **actions synchronization**

- **Non-negative integer** s with initial value s_0 and *atomic* operations $P(s)$ and $V(s)$.

$P(s)$: $\langle \text{await}(s > 0); s := s - 1 \rangle \rightarrow$ the task **sleeps until 's>0'** holds, decrement s

$V(s)$: $\langle s := s + 1 \rangle \rightarrow$ increment s

- From a theoretical perspective, semaphores are always positive


Note1: I use the notations $P(s)$ and $V(s)$, and $lock(s)$ and $unlock(s)$ interchangeably when s is a **mutex** (i.e., a **binary** semaphore).

Note2: the reference book discusses the possibility to implement semaphores that become negative. The goal of such implementation is to record the number of tasks waiting on the semaphore (i.e., if $s = -5$, then five tasks are waiting on semaphore s). It does not have any theoretical implication and may not be supported in all OSs or programming languages. Therefore, we will always assume that a semaphore does not become negative in the lectures.


- Non-negative integer s
- Two atomic actions $P(s)$ and $V(s)$
- A semaphore may be seen as a **bucket containing s tokens**




- Non-negative integer s
- Two atomic actions $P(s)$ and $V(s)$
- A semaphore may be seen as a **bucket containing tokens**
- ***$V(s)$ adds a token in the bucket***



- Non-negative integer s
- Two atomic actions $P(s)$ and $V(s)$
- A semaphore may be seen as a **bucket containing tokens**
- **$V(s)$ adds a token in the bucket**
- **$P(s)$ takes a token from the bucket if there is one in. Otherwise, it waits until one is added and takes it.**



- Non-negative integer s
- Two atomic actions $P(s)$ and $V(s)$
- A semaphore may be seen as a **bucket containing tokens**
- **$V(s)$ adds a token in the bucket**
- **$P(s)$ takes a token from the bucket if there is one in. Otherwise, it waits until one is added and takes it.**
- $V(s)^x$ and $P(s)^x$ places or takes x tokens in or from the bucket



Given:

- a collection of tasks executing a collection of actions A_i and B_j (for $i=0, 1, 2, \dots$ and $j=0, 1, 2, \dots$),
- and a required *synchronization condition* (invariant)

$$\text{Invariant: } \sum_i a_i \times cA_i \leq \sum_j b_j \times cB_j + e$$


for non-negative constants a_i, b_j and e .

Solution:

- introduce semaphore s , with initial value $s_0 = e$
- replace action A_i with $P(s)^{a_i} A_i$
- replace action B_j with $B_j V(s)^{b_j}$

Works because of the
safety property of
semaphores (see video)

Example of a synchronization problem



```
mutex m := 1;  
int BufferCnt := 0;
```


```
Task Producer =  
[[ while( true ) {  
    produce_items();  
  
    lock(m);  
    BufferCnt := BufferCnt+4;  
    push_items();  
    unlock(m);  
}  
]]
```

||

```
Task Consumer =  
[[ while( true ) {  
  
    lock(m);  
    BufferCnt := BufferCnt-2;  
    pop_items();  
    unlock(m);  
    use_items();  
}  
]]
```

1. Race condition when *Producer* and *Consumer* update the value of *BufferCnt* and when they push and pop items in the buffer
2. *Consumer* may try to **read from an empty buffer**
3. *Producer* may try to **write in a buffer that is full** (results in lost items)

Example of a synchronization problem



```
mutex m := 1;  
int BufferCnt := 0;
```

```
Task Producer =  
[[ while( true ) {  
    produce_items();
```

```
    G {  
        lock(m);  
        BufferCnt := BufferCnt+4;  
        push_items();  
        unlock(m);
```

```
    }  
]]
```

||

```
Task Consumer =  
[[ while( true ) {
```

```
    C {  
        lock(m);  
        BufferCnt := BufferCnt-2;  
        pop_items();  
        unlock(m);  
        use_items();
```


```
    }  
]]
```

Equivalent **synchronization conditions**:

1. $BufferCnt \geq 0$
2. $BufferCnt \leq BufferSize$

By **topology invariant**, we know that
 $BufferCnt = ?$

Example of a synchronization problem



```
mutex m := 1;  
int BufferCnt := 0;
```

```
Task Producer =  
[[ while( true ) {  
    produce_items();
```

```
    G {  
        lock(m);  
        BufferCnt := BufferCnt+4;  
        push_items();  
        unlock(m);
```

```
    }
```

```
]]
```

||

```
Task Consumer =  
[[ while( true ) {
```

```
    C {  
        lock(m);  
        BufferCnt := BufferCnt-2;  
        pop_items();  
        unlock(m);  
        use_items();
```

```
    }
```


```
]]
```

Equivalent **synchronization conditions**:

1. $BufferCnt \geq 0$
2. $BufferCnt \leq BufferSize$

By **topology invariant**, we know that
 $BufferCnt = 0 + 4 \times cG - 2 \times cC$

Example of a synchronization problem



```
mutex m := 1;  
int BufferCnt := 0;
```

```
Task Producer =  
[[ while( true ) {  
    produce_items();
```

```
    G {  
        lock(m);  
        BufferCnt := BufferCnt+4;  
        push_items();  
        unlock(m);
```

```
    }
```

```
]]
```

||

```
Task Consumer =  
[[ while( true ) {
```

```
    C {  
        lock(m);  
        BufferCnt := BufferCnt-2;  
        pop_items();  
        unlock(m);  
        use_items();
```


```
    }
```

```
]]
```

Equivalent **synchronization conditions** (injecting the topology invariant):

1. $4 \times cG - 2 \times cC \geq 0$
2. $4 \times cG - 2 \times cC \leq BufferSize$

Example of a synchronization problem



```
mutex m := 1;
int BufferCnt := 0;
```

```
Task Producer =
[[ while( true ){
    produce_items();
```

```

G {
    lock(m);
    BufferCnt := BufferCnt+4;
    push_items();
    unlock(m);
}
]
```

||

```
Task Consumer =
[[ while( true ){
```

```

C {
    lock(m);
    BufferCnt := BufferCnt-2;
    pop_items();
    unlock(m);
    use_items();
}
]
```

Equivalent synchronization conditions


- $2 \times cC \leq 4 \times cG$
- $4 \times cG \leq 2 \times cC + BufferSize$

Given: a synchronization condition: $\sum_i a_i \times cA_i \leq \sum_j b_j \times cB_j + e$
for non-negative constants a_i , b_j , and e .

Solution:

- introduce semaphore s , with initial value $s_0 = e$
- replace action A_i with $P(s)^{a_i} A_i$
- replace action B_j with $B_j V(s)^{b_j}$

Example of a synchronization problem



mutex m := 1;
int BufferCnt := 0;

semaphore s1 := 0;
semaphore s2 := BufferSize;

Introduce one semaphore per synchronization condition

Task Producer =
 || while(true){
 produce_items();
 P(s2)⁴;
 G {
 lock(m);
 BufferCnt := BufferCnt+4;
 push_items();
 unlock(m);
 V(s1)⁴;
 }
 }
 ||

Task Consumer =
 || while(true){
 P(s1)²;
 C {
 lock(m);
 BufferCnt := BufferCnt-2;
 pop_items();
 unlock(m);
 V(s2)²;
 use_items();
 }
 }
 ||

Equivalent synchronization conditions

- $2 \times cC \leq 4 \times cG$
- $4 \times cG \leq 2 \times cC + BufferSize$

Given: a synchronization condition: $\sum_i a_i \times cA_i \leq \sum_j b_j \times cB_j + e$ for non-negative constants a_i , b_j , and e .

Solution:

- introduce semaphore s , with initial value $s_0 = e$
- replace action A_i with $P(s)^{a_i} A_i$
- replace action B_j with $B_j V(s)^{b_j}$

Exercise (typical exam question)

```
int f := 5;  
int w := 0;
```

Given: a synchronization condition: $\sum_i a_i \times cA_i \leq \sum_j b_j \times cB_j + e$
for non-negative constants a_i, b_j , and e .

Solution:

- introduce **semaphore** s , with **initial value** $s_0 = e$
- **replace action** A_i with $P(s)^{a_i} A_i$
- **replace action** B_j with $B_j V(s)^{b_j}$

```
Task T =  
[[ while( true ) {  
    transport();  
    f := f+1;  
    w := w+3;  
}  
]]
```

||

```
Task P =  
[[ while( true ) {  
    w := w-2;  
    produce();  
    f := f+1;  
}  
]]
```

||

```
Task S =  
[[ while( true ) {  
    f := f-2;  
    transport();  
    w := w+3;  
}  
]]
```

Question:

Enforce the following synchronization invariants:

- **I1:** $w \geq 0$
- **I2:** $f \geq 2 \times w$

Prevent **race conditions** on f and w

Functions *produce()* and *transport()* **cannot be in critical sections**

Exercise (typical exam question)

```
int f := 5;  
int w := 0;
```

```
Task T =  
[[ while( true ) {  
    transport();  
    T1: f := f+1;  
    T2: w := w+3;  
}  
]]
```

||

```
Task P =  
[[ while( true ) {  
    P1: w := w-2;  
    produce();  
    P2: f := f+1;  
}  
]]
```

||

```
Task S =  
[[ while( true ) {  
    S1: f := f-2;  
    transport();  
    S2: w := w+3;  
}  
]]
```

$$f = cT1 + cP2 - 2*cS1 + 5$$

$$w = 3*cT2 - 2*cP1 + 3*cS2 + 0$$

Thus, the invariants require

$$I1 : 2*cP1 \leq 3*cT2 + 3*cS2$$

$$I2 : 6*cT2 + 2*cS1 + 6*cS2 \leq cT1 + cP2 + 4*cP1 + 5$$

We need 2 semaphores s1, s2 and two mutexes mf and mw initialized as follows:

s1:=0; s2:=5; mf:=1; mw:=1;

We then replace the actions as follows:

T1 \rightarrow P(mf) T1 V(mf) V(s2)

T2 \rightarrow P⁶(s2) P(mw) T2 V(mw) V³(s1)

P1 \rightarrow P²(s1) P(mw) P1 V(mw) V⁴(s2)

P2 \rightarrow P(mf) P2 V(mf) V(s2)

S1 \rightarrow P²(s2) P(mf) S1 V(mf)

S2 \rightarrow P⁶(s2) P(mw) S2 V(mw) V³(s1)

Exercise (typical exam question)

```
int f := 5;  
int w := 0;
```

```
semaphore s1 := 0;  
semaphore s2 := 5;
```

```
mutex mf := 1;  
mutex mw := 1;
```

```
Task T =  
[[ while( true ){  
    transport();  
  
    P(mf);  
    f := f+1;  
    V(mf);  
    V(s2);  
  
    P(s2)6;  
    P(mw);  
    w := w+3;  
    V(mw);  
    V(s1)3;  
}  
]]
```

||

```
Task P =  
[[ while( true ){  
    P(s1)2;  
    P(mw);  
    w := w-2;  
    V(mw);  
    V(s2)4;  
  
    produce();  
  
    P(mf);  
    f := f+1;  
    V(mf);  
    V(s2);  
}  
]]
```

||

```
Task S =  
[[ while( true ){  
    P(s2)2;  
    P(mf);  
    f := f-2;  
    V(mf);  
  
    transport();  
  
    P(s2)6;  
    P(mw);  
    w := w+3;  
    V(mw);  
    V(s1)3;  
}  
]]
```

- Since we have solved a synchronization problem and introduced blocking we must verify the correctness criteria.
- **Functional correctness** and **minimal waiting** are by construction.
- **Fairness**: the solution is just as fair as the semaphores and mutex.
 - depends on semaphore implementation and order of release from the waiting queue.
- **Absence of deadlock**: see next section

- Proving program properties using invariants
- Actions synchronization using semaphores
- Preventing deadlocks

- A **deadlock state** is a system state in which *a set of tasks* is **blocked indefinitely**
 - each task is blocked on another task in the same set
- We typically **prove** the absence of deadlock **by contradiction**
 - assume a deadlock occurs
 - investigate all task sets that can be blocked at the same time (often: just 1)
 - show a contradiction for all possible combinations of blocking actions of those tasks

```
mutex m := 1;  
int BufferCnt := 0;
```

```
semaphore s1 := 0;  
semaphore s2 := BufferSize;
```

```
Task Producer =  
[[ while( true ){  
    produce_items();  
    P(s2)4  
    lock(m);  
    BufferCnt := BufferCnt+4;  
    push_items();  
    unlock(m);  
    V(s1)4;  
}  
]]
```

Assume we are blocked here indefinitely ...

... then another task must hold the mutex forever.

It can only happen here.

However, it is impossible for Consumer to be blocked forever in its critical section (there is no blocking operation), and Consumer always calls unlock(m) at the end the critical section (i.e., releases the mutex)

```
Task Consumer =  
[[ while( true ){  
    P(s1)2;  
    lock(m);  
    BufferCnt := BufferCnt-2;  
    pop_items();  
    unlock(m);  
    V(s2)2;  
    use_items();  
}  
]]
```

A similar argument must be built for any other point where Producer or Consumer may be blocked

➔ Not a deadlock

When can it go wrong?

Let's swap $P(s1)$ and $lock(m)$ in Consumer

```
mutex m := 1;  
int BufferCnt := 0;
```

```
semaphore s1 := 0;  
semaphore s2 := BufferSize;
```

```
Task Producer =  
[[ while( true ){  
    produce_items();  
    P(s2)4  
    lock(m);  
    BufferCnt := BufferCnt+4;  
    push_items();  
    unlock(m);  
    V(s1)4;  
}  
]]
```

```
Task Consumer =  
[[ while( true ){  
    lock(m);  
    P(s1)2;  
    BufferCnt := BufferCnt-2;  
    pop_items();  
    unlock(m);  
    V(s2)2;  
    use_items();  
}  
]]
```

What can go wrong?

Let's swap $P(s2)$ and $lock(m)$ in Consumer

```
mutex m := 1;  
int BufferCnt := 0;
```

```
semaphore s1 := 0;  
semaphore s2 := BufferSize;
```

```
Task Producer =  
[[ while( true ){  
    produce_items();  
    P(s2)4  
    lock(m);  
    BufferCnt := BufferCnt+4;  
    push_items();  
    unlock(m);  
    V(s1)4;  
}  
]]
```

Assume we are blocked here indefinitely ...

... then another task must hold the mutex forever.

It can only happen here.

Consumer could be blocked here (i.e., buffer is not full enough)

Since *Producer* is blocked, it cannot call $V(s1)$.

Therefore, *Consumer* will remain blocked and will be unable to call $unlock(m)$

→ Possible deadlock

```
Task Consumer =  
[[ while( true ){  
    lock(m);  
    P(s1)2;  
    BufferCnt := BufferCnt-2;  
    pop_items();  
    unlock(m);  
    V(s2)2;  
    use_items();  
}  
]]
```

Note: for a full proof that a deadlock can be reached, a trace should be produced

- Always let critical sections terminate
 - always **call *unlock(m)* after *lock(m)***
- **Avoid cyclic waiting**
 - Try to avoid *P or lock operations* that may block indefinitely *between *lock(m)* and *unlock(m)**

```
mutex m1 := 1;  
mutex m2 := 1;  
mutex m3 := 1;
```

```
Task T1 =  
[[ while( true ){  
    P(m1);  
    P(m2);  
    doSomething();  
    V(m2);  
    V(m1);  
}  
]]
```

```
Task T2 =  
[[ while( true ){  
    P(m2);  
    P(m3);  
    doSomething();  
    V(m3);  
    V(m2);  
}  
]]
```

```
Task T3 =  
[[ while( true ){  
    P(m3);  
    P(m1);  
    doSomething();  
    V(m1);  
    V(m3);  
}  
]]
```

What can go wrong?

```
mutex m1 := 1;  
mutex m2 := 1;  
mutex m3 := 1;
```

```
Task T1 =  
[[ while( true ){  
    P(m1);  
    P(m2);  
    doSomething();  
    V(m2);  
    V(m1);  
}  
]]
```

```
Task T2 =  
[[ while( true ){  
    P(m2);  
    P(m3);  
    doSomething();  
    V(m3);  
    V(m2);  
}  
]]
```

```
Task T3 =  
[[ while( true ){  
    P(m3);  
    P(m1);  
    doSomething();  
    V(m1);  
    V(m3);  
}  
]]
```

(T1.P(m1)) {m1=0} (T2.P(m2)) {m2=0} (T3.P(m3)) {m3=0}

T1 blocked on P(m2), T2 blocked on P(m3),
T3 blocked on P(m1)
→ **deadlock**

Cyclic waiting

solution: lock semaphores in a fixed order

```
mutex m1 := 1;  
mutex m2 := 1;  
mutex m3 := 1;
```

```
Task T1 =  
[[ while( true ){  
    P(m1);  
    P(m2);  
    doSomething();  
    V(m2);  
    V(m1);  
}  
]]
```

```
Task T2 =  
[[ while( true ){  
    P(m2);  
    P(m3);  
    doSomething();  
    V(m3);  
    V(m2);  
}  
]]
```

Swap $P(m1)$ and $P(m3)$

```
Task T3 =  
[[ while( true ){  
    P(m1);  
    P(m3);  
    doSomething();  
    V(m1);  
    V(m3);  
}  
]]
```

Solution: always call $P(m3)$ after $P(m2)$ which is always called after $P(m1)$

Exercise: prove the absence of deadlock for this program (exam 2020-2021)

- Always let critical sections terminate
 - always **call *unlock(m)* after *lock(m)***
- **Avoid cyclic waiting**
 - Try to avoid *P or lock operations* that may block indefinitely *between *lock(m)* and *unlock(m)**
 - Use a **fixed order when calling *P*-operations on semaphores or mutexes**
 - $P(m); P(n); \dots$ in one task may deadlock with $P(n); P(m); \dots$ in another task
 - Look for the “*dining philosopher problem*” for a good example

Assume multiple *Producers*

```
mutex m := 1;  
int BufferCnt := 0;
```

```
semaphore s1 := 0;  
semaphore s2 := 4;
```

Assume
BufferSize = 4

```
Task Producer1 =  
[[ while( true ){  
    produce_items();  
    P(s2)4  
    lock(m);  
    BufferCnt := BufferCnt+4;  
    push_items();  
    unlock(m);  
    V(s1)4;  
}  
]]
```

```
Task Producer2 =  
[[ while( true ){  
    produce_items();  
    P(s2)4  
    lock(m);  
    BufferCnt := BufferCnt+4;  
    push_items();  
    unlock(m);  
    V(s1)4;  
}  
]]
```

What can go wrong?

Assume multiple *Producers*

```
mutex m := 1;  
int BufferCnt := 0;
```

```
semaphore s1 := 0;  
semaphore s2 := 4;
```

Assume
BufferSize = 4

```
Task Producer1 =  
[[ while( true ){  
    produce_items();  
    P(s2)4  
    lock(m);  
    BufferCnt := BufferCnt+4;  
    push_items();  
    unlock(m);  
    V(s1)4;  
}  
]]
```

Assume *Producer1*
executes *P(s2)* two
times...

... then s2 = 2

Now, assume *Producer2*
executes *P(s2)* two times

... then s2 = 0

Both tasks are blocked

→ **deadlock**

```
Task Producer2 =  
[[ while( true ){  
    produce_items();  
    P(s2)4  
    lock(m);  
    BufferCnt := BufferCnt+4;  
    push_items();  
    unlock(m);  
    V(s1)4;  
}  
]]
```


- Always let critical sections terminate
 - always **call *unlock(m)* after *lock(m)***
- **Avoid cyclic waiting**
 - no *P* or *lock* operations that may block indefinitely between *lock(m)* and *unlock(m)*
 - Use a **fixed order when calling *P*-operations on semaphores or mutexes**
 - $P(m); P(n); \dots$ in one task may deadlock with $P(n); P(m); \dots$ in another task
 - Look for the “*dining philosopher problem*” for a good example
- **Avoid greedy consumers**
 - $P(a)^k$ should be **an indivisible atomic operation** when tasks compete for limited resources

Assume multiple *Producers*: solution

```
mutex m := 1;  
int BufferCnt := 0;
```

```
semaphore s1 := 0;  
semaphore s2 := 4;  
mutex mp := 1;
```

New mutex

```
Task Producer1 =  
[[ while( true ){  
    produce_items();  
    lock(mp); P(s2)4; unlock(mp);  
    lock(m);  
    BufferCnt := BufferCnt+4;  
    push_items();  
    unlock(m);  
    V(s1)4;  
}]
```

||

```
Task Producer2 =  
[[ while( true ){  
    produce_items();  
    lock(mp); P(s2)4; unlock(mp);  
    lock(m);  
    BufferCnt := BufferCnt+4;  
    push_items();  
    unlock(m);  
    V(s1)4;  
}]
```

- **Proving program properties** using **topology invariants**
 - **Synchronize tasks** to enforce new properties **using semaphores**
 - How to **prove the existence or absence of deadlock**
 - **Rules of thumb to prevent deadlocks** when using semaphores and mutexes
-
- **We come back to deadlocks in two weeks**
-
- **Additional exercises** will be available on Canvas **to train yourself** at actions synchronization and generating traces
 - **Work on them! They will prepare you for the exam.**
 - **Homework of next week** will be on invariants and action synchronization