

Operating Systems (2INC0)

Scheduling

Dr. Geoffrey Nelissen

Courtesy of Prof. Dr. Johan Lukkien and
Dr. Tanir Ozcelebi



Interconnected
Resource-aware
Intelligent Systems

TU/e

Technische Universiteit
Eindhoven
University of Technology

Where innovation starts

A single process (thread) may not use the system resources efficiently

Only **one task** can use a resource (e.g. processor, memory page, device) **at any given time instant**.

A task **does not use a resource constantly** during its execution

→ We **must schedule** the access to resources


Agenda

- **Resource scheduling**
- **CPU scheduling**
- **Common scheduling algorithms**
- **Priority inversion**

Resource scheduling (allocation)

- Assignment of resources to tasks
 - **Schedule S** is a **function that maps a time and a resource to a task**:
 $S(t,r) = P$ means that task P is allocated resource r at time t
 - Definition of a **task** is **context dependent**
 - e.g. a **process**, a **thread** of execution, **I/O operation** (e.g., the reading of a disk block), servicing an **interrupt**, etc.
 - To define a schedule we must decide
 - **When**: to change the allocation of a resource
 - Usually when there is a **change in the system state**
 - **processor**: e.g. process added into ready queue; end of time slice; process yielding (e.g. blocked, waiting for I/O or for child process to terminate)
 - **memory**: e.g. memory management call; replacement policy by memory subsystem; process termination
 - **How**: what decision *procedures* are used to allocate the resource when in decision mode

Example schedule



- $P0, P1, P2$: tasks (processes, jobs)
- $t0, t1, \dots$: scheduling points
 - system is in decision mode
 - scheduling decisions are taken
 - scheduling points may differ per resource (not shown here)
- processor resource: $PROC$
 - $S([t0..t1], PROC) = P0$,
 $S([t1..t2], PROC) = P1$, etc.
- a memory page frame: $m23$
 - $S([t0..t4], m23) = P0$, etc.
- another memory page frame: $m56$
 - $S([t3..t6], m56) = P2$, etc.

Scheduling policies and mechanisms

- **Scheduling policy** represents the strategy for allocating a resource to a task while in *decision mode*.
 - **policy: an algorithm that decides based on scheduling criteria....**
 - task attributes (deadline, response time, ...)
 - current state (the set of ready processes, available and required resources, ...)
 - **....or based on a pre-computed lookup table**

Agenda


- Resource scheduling
- **CPU scheduling**
 - metrics, task attributes
 - a framework for scheduling
- Common scheduling algorithms
- Priority inversion

Metrics for the quality of CPU scheduling

Scheduling Criteria


Metric: observed property, result of applying scheduling policy

- CPU utilization
 - keep the CPU as busy as possible
- Throughput
 - # of processes that complete their execution per time
- number of deadline misses (real-time scheduling)
- turnaround time
 - time to execute a particular process
- waiting time
 - time a process has been waiting in the ready queue
- response time
 - time it takes from when a request was submitted until the first response is produced → not the entire output (suitable for time-sharing interactive environment)
 - most texts: *response time* is defined as the time elapsing from arrival to completion (in our textbook this is called: *turnaround time*)



➔ Basis for our comparisons of different scheduling schemes in this slide set

Time attributes of a task



- A task has fixed time attributes and ...
 - a name (the j^{th} task)
 - a period, sometimes
 - a (worst case) execution time
 - a relative deadline, sometimes
- ... dynamic time attributes (i^{th} instance or occurrence)
 - an **arrival** time
 - an **absolute deadline**, sometimes (add D_j to arrival time)
 - a **start** time (or beginning time) of execution
 - a departure time, also called **end**, finish or completion time
- (book) Response time: $s_{j,i} - a_{j,i}$
- (book) Turnaround time: $e_{j,i} - a_{j,i}$

τ_j
 T_j
 C_j
 D_j

$a_{j,i}$
 $dl_{j,i}$
 $s_{j,i}$
 $e_{j,i}$

CPU scheduling

- Scheduling framework defines 3 things
 - **When** to schedule?
= *decision mode*
 - **Priority function** (priority scheduling)
 - What task to schedule based on priority?
 - **Arbitration rule**
 - who to schedule **in case of equal priority?**
- Managed by 2 different OS modules (jointly called *the scheduler*)
 - **Process scheduler**
 - policies to determine **which task to execute next**
 - **Process dispatcher**
 - actual **binding of selected task to a processor** (switching context, switching to user mode, jumping to the proper location in the program)

Decision mode: Activating process scheduler

- Decision mode defines the conditions for activating the process scheduler.
 - **Upon activation**
 - Task scheduler selects one ready task for execution
- Scheduler can be invoked (decision mode), e.g.:
 1. **Periodically**
 2. when a **process** switches **from running to waiting state** (e.g. wait for child to terminate)
 3. when a **process** switches **from running to ready state** (e.g. through an interrupt)
 4. when a **process** switches **from waiting to ready** (e.g. i/o completion)
 5. when a **process terminates**
- If scheduling takes place ONLY under 2 & 5 (active process voluntarily yields) → non-preemptive decision mode
otherwise → preemptive decision mode