

Operating Systems (2INC0)

Action Synchronization (06) Checking invariants

Dr. Geoffrey Nelissen

Courtesy of Prof. Dr. Johan Lukkien and
Dr. Tanir Ozcelebi



Interconnected
Resource-aware
Intelligent Systems

TU/e

Technische Universiteit
Eindhoven
University of Technology

Where innovation starts

Synchronization

- **Synchronization** is about **limitation of possible** program **traces** by coordinating execution such that
 - a certain invariant is satisfied
 - i.e., avoid the traces that violate the invariant
 - or the execution has some desired property
 - e.g. a certain assertion holds at a certain control point
- Typically, by blocking execution until an assertion has become true.

Action synchronization

- relies on action counting and invariants on the counting.
- An **invariant** I is an assertion that holds at *all* control points.
 - (Example) I : “mutual exclusion is maintained”
 - (Example) I : $y \leq x$ in the program below
...assuming <atomic> assignments...

Initially: $x=0 \wedge y=0$

```
while true do <x := x+1>; <y := y+1> od
      ||
while true do <y := y-1>; <x := x-1> od
```

Terminology: naming and counting

Naming of actions

Initially: $x=0 \wedge y=0$

while *true* **do** **A:** $\langle x := x+1 \rangle$; **B:** $\langle y := y+1 \rangle$ **od**

||

while *true* **do** **C:** $\langle y := y-1 \rangle$; **D:** $\langle x := x-1 \rangle$ **od**

If A is an action in the program, $\underline{c}A$ denotes the number of completed executions of A . $\underline{c}A$ can be regarded as an auxiliary variable that is initially 0 and is incremented atomically each time A is executed.

$A \rightarrow \langle A; \underline{c}A := \underline{c}A+1 \rangle$

Topology properties

Topology invariants: derived directly from the program text

Example: two actions always occurring one after the other

Initially: $x=0 \wedge y=0$

while *true* **do** **A:** $\langle x := x+1 \rangle$; **B:** $\langle y := y+1 \rangle$ **od**

||

while *true* **do** **C:** $\langle y := y-1 \rangle$; **D:** $\langle x := x-1 \rangle$ **od**

Invariants:

$$I_0: x = \underline{c}A - \underline{c}D$$

$$I_1: y = \underline{c}B - \underline{c}C$$

$$I_2: 0 \leq \underline{c}A - \underline{c}B \leq 1$$

$$I_3: 0 \leq \underline{c}C - \underline{c}D \leq 1$$

Example

We can prove that $I: y \leq x$ holds using topology invariants

$I_0: x = \underline{c}A - \underline{c}D$	$I_2: 0 \leq \underline{c}A - \underline{c}B \leq 1$
$I_1: y = \underline{c}B - \underline{c}C$	$I_3: 0 \leq \underline{c}C - \underline{c}D \leq 1$

$$y \leq x = ?$$

$$= \underline{c}B - \underline{c}C \leq \underline{c}A - \underline{c}D \quad \{ I_0: x = \underline{c}A - \underline{c}D, I_1: y = \underline{c}B - \underline{c}C \}$$

$$= true \quad \{ I_2: \underline{c}B \leq \underline{c}A, I_3: -\underline{c}C \leq -\underline{c}D \}$$

Operating Systems (2INC0)

Action Synchronization (06) Using semaphores

Dr. Geoffrey Nelissen

Courtesy of Prof. Dr. Johan Lukkien and
Dr. Tanir Ozcelebi



Interconnected
Resource-aware
Intelligent Systems



Technische Universiteit
Eindhoven
University of Technology


Where innovation starts

Synchronization conditions

- Action synchronization is specified by inequalities
 - on action counts, or...
 - ...on program variables *directly related to this counting*.
- We refer to such an inequality as
 - a synchronization condition, or
 - a synchronization invariant.

Example: Producer-consumer problem

- A producer process P_Y produces information...
- ...that is consumed by a consumer process P_X .



- Synchronize P_X and P_Y such that the following invariant is maintained.

$$I_0: x \leq y \quad (= \underline{c}A \leq \underline{c}B)$$

- $I_0: x \leq y$ is desired.
 - How do we enforce that invariant?

Recall semaphores (Dijkstra)

- Non-negative integer s with initial value s_0 and *atomic* operations $P(s)$ and $V(s)$.

$P(s)$: $\langle \text{await}(s > 0); s := s - 1 \rangle \rightarrow$ block until ' $s > 0$ ' holds, decrement

$V(s)$: $\langle s := s + 1 \rangle \rightarrow$ increment

- Semaphores can be used to implement mutual exclusion

Semaphore invariants

From the definition, we derive two semaphore properties (invariants):

$$S0: s \geq 0$$

$$S1: s = s_0 + \underline{c}V(s) - \underline{c}P(s)$$

$S0, S1$: functional properties (“safety properties”). Combining the two:

$$S2: \underline{c}P(s) \leq s_0 + \underline{c}V(s)$$

- Hence, semaphores realize a synchronization invariant *by definition*.

In addition, we have a *progress property*:

- **Blocking is allowed only if the safety properties would be violated.**

Solve the producer/consumer problem

$P_X =$

$x := 0;$

while true do

A: $\langle x := x+1 \rangle$

od

\parallel

$P_Y =$

$y := 0;$

while true do

B: $\langle y := y+1 \rangle$

od

Synchronize P_X and P_Y such that the invariant is maintained.

$I_0: x \leq y$

Use the program topology:
hence, I_0 can be rewritten:

$x = \underline{c}A$ and $y = \underline{c}B$
 $I_0: \underline{c}A \leq \underline{c}B$

Solve the producer/consumer problem

Introduce semaphore s ; let A be preceded by $P(s)$ and B be followed by $V(s)$.

```
 $P_X =$   
   $x := 0;$   
  while true do  
     $P(s); A: \langle x := x+1 \rangle;$   
  od
```

||

```
 $P_Y =$   
   $y := 0;$   
  while true do  
     $B: \langle y := y+1 \rangle; V(s);$   
  od
```

From topology:

$I_1: \underline{c}A \leq \underline{c}P(s)$ and $I_2: \underline{c}V(s) \leq \underline{c}B$

Combine with semaphore invariant ($S2: \underline{c}P(s) \leq s_0 + \underline{c}V(s)$)

$\underline{c}A \leq \underline{c}P(s) \leq s_0 + \underline{c}V(s) \leq s_0 + \underline{c}B$

Hence, choosing $s_0 = 0$ does the job. $\rightarrow I_0$ holds.

Action Synchronization Solution (in general)

Given: - collection of tasks/threads executing actions A, B, C, D ;
- a required *synchronization condition (invariant)*

$$\text{SYNC: } a \cdot \underline{c}A + c \cdot \underline{c}C \leq b \cdot \underline{c}B + d \cdot \underline{c}D + e$$

for non-negative constants a, b, c, d, e .

Solution: introduce semaphore s , $s_0 = e$ and replace

$$\begin{aligned} A &\rightarrow P(s)^a; A \\ C &\rightarrow P(s)^c; C \end{aligned}$$

$$\begin{aligned} B &\rightarrow B; V(s)^b \\ D &\rightarrow D; V(s)^d \end{aligned}$$

Exponent = Number of
times $V(s)$ or $P(s)$ is called

Operating Systems (2INC0)

Action Synchronization (06) POSIX implementation

Dr. Geoffrey Nelissen

Courtesy of Prof. Dr. Johan Lukkien and
Dr. Tanir Ozcelebi



Interconnected
Resource-aware
Intelligent Systems

TU/e

Technische Universiteit
Eindhoven
University of Technology

Where innovation starts

Counting semaphores (POSIX 1003.1b)

- Creation and destruction
 - “name” within kernel, persistent until re-boot, like a filename
 - Posix names: for portability
 - or “unnamed” semaphores, for use in shared memory
 - shared memory between processes

```
sem_t *sem;  
sem = sem_open (name, flags, mode, init_val); /* name is system-wide */  
status = sem_close (sem); /* semaphore still reachable */  
status = sem_unlink (name); /* now it is removed */  
  
status = sem_init (sem, pshared, init_val); /* memory space for sem must be defined,  
                                             e.g. through shm or malloc */  
/* pshared flag: whether multiple processes  
   or threads share sem of not */  
  
status = sem_destroy (sem);
```


POSIX semaphore operations

```
status = sem_wait (sem); /* P(sem) locks sem */  
  
status = sem_trywait (sem); /* P(sem) again */  
/* but returns error if sem == 0 */  
  
status = sem_post (sem); /* V(sem) */  
  
status = sem_getvalue (sem, &val); /* current value */  
/* when negative: absolute value = # waiters */
```

sem negative value is interpreted as number of waiters
(length of the waiting queue)

Example: producer – consumer with buffer of depth 4

```
#include <stdio.h>
#include <fcntl.h>
#include <pthread.h>
#include <semaphore.h>
```

```
sem_t *s, *t;
```

```
void Producer ()
{
    int i;
    for (i=0; i<10; i++) {
        sem_wait (t); printf ("Produce "); fflush (stdout);
        sem_post (s); sleep (1);
    }
}
```

```
void Consumer ()
{
    int i;
    for (i=0; i<10; i++) {
        sem_wait (s); printf ("Consume "); fflush (stdout);
        sem_post (t); sleep (2);
    }
}
```

(cnt'd)

```
void main ()
{
    pthread_t thread_id;

    s = sem_open ("Mysem-s", O_CREAT | O_RDWR, 0, 0);
    if (s == SEM_FAILED) { perror ("sem_open"); exit (0); }
    t = sem_open ("Mysem-t", O_CREAT | O_RDWR, 0, 4);
    if (t == SEM_FAILED) { perror ("sem_open"); exit (0); }

    pthread_create (&thread_id, NULL, Producer, NULL);
    Consumer ();
    pthread_join (thread_id, NULL);
    sem_close (s); sem_close (t);
    sem_unlink ("Mysem-s"); sem_unlink ("Mysem-t");
}
```

Output

- Produce Consume Produce Consume Produce Produce Consume
Produce Produce Consume Produce Produce Consume Produce
Consume Produce Consume Consume Consume Consume

(This is one of the many possible outputs.)