

2INC0 - Operating Systems

Condition synchronization

Geoffrey Nelissen



Interconnected
Resource-aware
Intelligent Systems

TU/e

Technische Universiteit
Eindhoven
University of Technology

Where innovation starts

- **Check the additional exercises on Canvas** (in modules)
- New programming **workshop** next Friday (using threads, mutexes, semaphores, ...)

- **Introduction to operating systems** (lecture 1)

- **Processes, threads and scheduling**

- **Concurrency and synchronization**

- atomicity and interference (lecture 4)
- actions synchronization (lecture 5)
- **condition synchronization (lecture 6)**
- deadlock (lecture 7)

- **File Systems** (lecture 8)

- **Memory management** (lectures 9+10)

- **Input/output** (lecture 11)

- Dangers of concurrency (**race conditions**)
- How to prove properties using **traces**
- Synchronization (**critical sections**)

- Check complex properties using **topology invariants**
- Use semaphores to **synchronize actions** and enforce new properties

- Enforce properties that **cannot be checked just by counting** actions

- **Condition variables**
- **Monitors**
 - **Signaling disciplines**

Two complementary tools for condition synchronization

Invariant

- An assertion that **holds at all control points** of a program

Topology invariant

- An invariant **derived** directly from the **program** code
- **Examples:**
 - **number of times an operation is executed** in comparison to another action.
 - **Value of a variable** based on the number of times a set of actions was executed

Actions synchronization

- Uses **semaphores** to **synchronizes actions** in different tasks to **enforce new invariants** (called synchronization invariants or **synchronization conditions**)
- Steps:
 - Name actions that may influence the synchronization condition
 - Write the synchronization condition as a function of actions counts in the form $\sum_i a_i \times cA_i \leq \sum_j b_j \times cB_j + e$ for non-negative constants a_i, b_j and e .
 - introduce **semaphore s** , with **initial value $s_0 = e$**
 - **replace action A_i with $P(s)^{a_i} A_i$**
 - **replace action B_j with $B_j V(s)^{b_j}$**

```
int x := 7;  
int y := 1;  
int z := 6;
```

Go to [menti.com](https://www.menti.com) and use
2438 8084

```
Task1 = [  
  while (true) {  
    X1: x:=x+6;  
    Y1: y:=2*y;  
    if( x < y ) {  
      Z1: z++;  
    }  
  }  
]
```

||

```
Task2 = [  
  while (true) {  
    Y2: y:=y-2;  
    X2: x--;  
    if( x > 8 ) {  
      Z2: z--;  
    }  
  }  
]
```

cX = number of times
action X executed

$x = ?$

```
int x := 7;  
int y := 1;  
int z := 6;
```

```
Task1 = [  
  while (true) {  
    X1: x:=x+6;  
    Y1: y:=2*y;  
    if( x < y ) {  
      Z1: z++;  
    }  
  }  
]
```

||

```
Task2 = [  
  while (true) {  
    Y2: y:=y-2;  
    X2: x--;  
    if( x > 8 ) {  
      Z2: z--;  
    }  
  }  
]
```

cX = number of times
action X executed

~~$$x = 7 + 6 \times cX_1 - cX_2$$~~

Is it really correct?

No, because X1 and X2 are **not atomic**.
We must first protect them against **race conditions**

```
int x := 7;  
int y := 1;  
int z := 6;
```

```
Task1 = [[  
  while (true) {  
    X1: <x:=x+6>;  
    Y1: <y:=2*y>;  
    if( x < y ){  
      Z1: <z++>;  
    }  
  }  
]]
```

||

```
Task2 = [[  
  while (true) {  
    Y2: <y:=y-2>;  
    X2: <x-->;  
    if( x > 8 ) {  
      Z2: <z-->;  
    }  
  }  
]]
```

cX = number of times
action X executed

Now assume all actions are atomic:

$$x = 7 + 6 \times cX_1 - cX_2$$

$$z = ?$$


```
int x := 7;  
int y := 1;  
int z := 6;
```

```
Task1 = [  
  while (true) {  
    X1: <x:=x+6>;  
    Y1: <y:=2*y>;  
    if( x < y ){  
      Z1: <z++>;  
    }  
  }  
]
```

||

```
Task2 = [  
  while (true) {  
    Y2: <y:=y-2>;  
    X2: <x-->;  
    if( x > 8 ) {  
      Z2: <z-->;  
    }  
  }  
]
```

cX = number of times
action X executed

Now assume all actions are atomic:

$$x = 7 + 6 \times cX_1 - cX_2$$

$$z = 6 + cZ_1 - cZ_2$$

Actions synchronization: what can we enforce?

```
int x := 7;  
int y := 1;  
int z := 6;
```

Go to [menti.com](https://www.menti.com) and use
2438 8084

```
Task1 = [[  
  while (true) {  
    X1: x:=x+6;  
    Y1: y:=2*y;  
    if( x < y ) {  
      Z1: z++;  
    }  
  }  
]]
```

||

```
Task2 = [[  
  while (true) {  
    Y2: y:=y-2;  
    X2: x--;  
    if( x > 8 ) {  
      Z2: z--;  
    }  
  }  
]]
```

cX = number of times
action X executed

Which of the **synchronization invariants** presented to you
can be enforced **with actions synchronization**?

Value of y cannot be derived from the number of times we execute $Y1$ and $Y2$ because of the multiplication in $Y1$

- $x \geq 4$
- $y > 0$
- The execution repeats the exact sequence $X1 Y2 X2 Y1$
- $z < x$
- $x \geq 4$ and $z < x$
- $y < x$
- $Y1$ executes at least twice as often as $Z2$
- $x \geq 4$ or $z < x$

Using action synchronization can only check and enforce one of the two conditions, and will thus not allow for minimal waiting

Green = can be enforced with actions synchronization
Red = cannot be enforced with actions synchronization

- **Condition synchronization**
 - **explicit communication (signalling)** between tasks
 - when just **counting is not enough** to solve a synchronization problem
 - or to **simplify** otherwise **complex sequences of semaphore operations**,
e.g., $P(a); P(a); P(b); P(m); \dots V(a); V(b);$

Two principles:

Where a condition **may be violated**: **check** and potentially **wait**

Where a condition **may have become true**: **signal** waiters

Condition synchronization: building blocks

- *ConditionVar* *cv* ; (often **associated with** a condition: a **boolean expression** *B* in terms of program variables;)
- 4 basic operations on variable *cv*:
 - *Wait(..., cv)* **suspend** execution of caller
 - *Signal(cv)* **free one task** suspended on *Wait(cv)*
(void if there is none)
 - *Sigall(cv)* **free all tasks** suspended on *Wait(cv)*
 - *Empty(cv)* Check if “there is no task suspended on *Wait(cv)*” (returns true or false)

Enforce the **synchronization invariant** $x \geq 0$

```
int x := 0;
```

```
Task T1
[[
  while( true ){
    z:=z+1;
    x := x + 50;
    x := 2*x;
  }
]]
```

||

```
Task T2
[[
  while( true ){
    x := x-10;
    y := x+3;
  }
]]
```

Step 1: identify critical sections and protect them

```
mutex m := 1;  
int x := 0;
```

Enforce the **synchronization invariant** $x \geq 0$

```
Task T1  
[[  
  while( true ){  
    z := z + 1;  
    lock(m);  
    x := x + 50;  
    x := 2 * x;  
    unlock(m);  
  }  
]]
```

||

```
Task T2  
[[  
  while( true ){  
    lock(m);  
    x := x - 10;  
    unlock(m);  
    y := x + 3;  
  }  
]]
```

Step 1: identify critical sections and protect them

```
mutex m := 1;  
int x := 0;
```

Enforce the **synchronization**
invariant $x \geq 0$

```
Task T1  
[[  
  while( true ){  
    z := z + 1;  
    lock(m);  
    x := x + 50;  
    x := 2 * x;  
    unlock(m);  
  }  
]]
```

||

```
Task T2  
[[  
  while( true ){  
    lock(m);  
    x := x - 10;  
    unlock(m);  
    y := x + 3;  
  }  
]]
```

Step 2: identify where a condition **may become false**
and add a guard


```
ConditionVar cv;  
mutex m := 1;  
int x := 0;
```

Enforce the **synchronization**
invariant $x \geq 0$

```
Task T1  
[[  
  while( true ){  
    z:=z+1;  
    lock(m);  
    x := x + 50;  
    x := 2*x;  
    unlock(m);  
  }  
]]
```

||

```
Task T2  
[[  
  while( true ){  
    lock(m);  
    while( x<10 ){  
      Wait( m, cv);  
    }  
    x := x-10;  
    unlock(m);  
    y := x+3;  
  }  
]]
```

Step 2: identify where a condition **may become false**
and add a guard

```
ConditionVar cv;  
mutex m := 1;  
int x := 0;
```

Enforce the **synchronization**
invariant $x \geq 0$

```
Task T1  
[[  
  while( true ){  
    z:=z+1;  
    lock(m);  
    x := x + 50;  
    x := 2*x;  
    unlock(m);  
  }  
]]
```

||

```
Task T2  
[[  
  while( true ){  
    lock(m);  
    while( x<10 ){  
      Wait( m, cv);  
    }  
    {x>=10}  
    x := x-10;  
    unlock(m);  
    y := x+3;  
  }  
]]
```

Step 3: identify where a condition **may become true**
and signal

```
ConditionVar cv;  
mutex m := 1;  
int x := 0;
```

Enforce the **synchronization invariant** $x \geq 0$

```
Task T1  
[[  
  while( true ){  
    z:=z+1;  
    lock(m);  
    x := x + 50;  
    x := 2*x;  
    Sigall(cv);  
    unlock(m);  
  }  
]]
```

||

```
Task T2  
[[  
  while( true ){  
    lock(m);  
    while( x < 10 ){  
      Wait( m, cv );  
    }  
    {x >= 10}  
    x := x - 10;  
    unlock(m);  
    y := x + 3;  
  }  
]]
```

Wait(m,cv) is the short-hand notation for
<unlock(m); wait(cv);> lock(m);

Must be **atomic**!
Why?

```
ConditionVar cv;
mutex m := 1;
int x := 0;
```

Enforce the **synchronization invariant** $x \geq 0$

```
Task T1
[[
  while( true ){
    z:=z+1;
    lock(m);
    x := x + 50;
    x := 2*x;
    Sigall(cv);
    unlock(m);
  }
]]
```

Why do we use **Sigall(.)** instead of **Signal(.)**?

Wait(m,cv) is the short-hand notation for
~~unlock(m); wait(cv);~~ lock(m);

Must be **atomic**!
 Otherwise, another task may **signal before** T2 starts to wait and T2 will never wakeup

```
Task T2
[[
  while( true ){
    lock(m);
    while( x<10 ){
      Wait( m, cv);
    }
    {x>=10}
    x := x-10;
    unlock(m);
    y := x+3;
  }
]]
```

```
ConditionVar cv;
mutex m := 1;
int x := 0;
```

Enforce the **synchronization invariant** $x \geq 0$

Task T1

```
[[
while( true ){
  z:=z+1;
  lock(m);
  x := x + 50;
  x := 2*x;
  Sigall(cv);
  unlock(m);
}
```

Why do we use **Sigall(.)** instead of **Signal(.)**?
More than one instance of T2 could execute their **critical section** for each execution of T1's critical section

How do we decide what to write as a guard condition?

Wait(m,cv) is the short-hand notation for
`<unlock(m); wait(cv);> lock(m);`

Must be **atomic**!
 Otherwise, another task may **signal before** T2 starts to wait and T2 will never wakeup

Task T2

```
[[
while( true ){
  lock(m);
  while( x<10 ){
    Wait( m, cv);
  }
  {x>=10}
  x := x-10;
  unlock(m);
  y := x+3;
}
```

What to write as a guard condition?

1. Take the **synchronization invariant** to maintain, i.e.,
 $x \geq 0$
2. Compute **preconditions** for the **statements that might disturb the invariant**, i.e.,
 $x := x - 10$
→ Substitute disturbing statement in the synchronization invariant
 $x - 10 \geq 0 \Leftrightarrow x \geq 10$
3. **Negate** the resulting condition, i.e.,
 $!(x \geq 10)$
or equivalently
 $(x < 10)$

```
Task T2
[[
    while( true ){
        lock(m);
        while( x < 10 )
            Wait( m, cv);
        }
        x := x - 10;
        unlock(m);
        y := x + 3;
    }
]]
```

Exercise (typical exam question)

```
int f := 5;  
int w := 0;
```

```
Task T =  
[[ while( true ) {  
    transport();  
    f := f+1;  
    w := w+3;  
}  
]]
```

||

```
Task P =  
[[ while( true ) {  
    w := w-2;  
    produce();  
    f := f+1;  
}  
]]
```

||

```
Task S =  
[[ while( true ) {  
    f := f-2;  
    transport();  
    w := w+3;  
}  
]]
```

Question:

Enforce the following **synchronization invariants** using **condition variables**

- **I1:** $w \geq 0$
- **I2:** $f \geq 2 * w$

Prevent **race conditions** on f and w

Functions *produce()* and *transport()* cannot be in critical sections

Exercise (typical exam question)

Go to menti.com and use
2438 8084

```
int f := 5;  
int w := 0;  
mutex m := 1;  
ConditionVar cv1, cv2;
```

```
Task T =  
[[ while( true ){  
    transport();  
  
    lock(m);  
    f := f+1;  
    sigall(cv2);  
    unlock(m);  
  
    lock(m);  
    while(...)  
        wait(m, cv2);  
    w := w+3;  
    sigall(cv1);  
    unlock(m);  
}  
]]
```

||

```
Task P =  
[[ while( true ){  
    lock(m);  
    while(...)  
        wait(m, cv1);  
    w := w-2;  
    sigall(cv2);  
    unlock(m);  
  
    produce();  
  
    lock(m);  
    f := f+1;  
    sigall(cv2);  
    unlock(m);  
}  
]]
```

What should
be the guard
condition?

||

```
Task S =  
while( true ){  
    lock(m);  
    while(...)  
        wait(m, cv2);  
    f := f-2;  
    unlock(m);  
  
    transport();  
  
    lock(m);  
    while(...)  
        wait(m, cv2);  
    w := w+3;  
    sigall(cv1);  
    unlock(m);  
}
```

Question:

Enforce the following **synchronization invariants** using **condition variables**

- **I1:** $w \geq 0$
- **I2:** $f \geq 2 * w$

Prevent **race conditions** on f and w

Functions *produce()* and *transport()* cannot be in critical sections

Exercise (typical exam question)

Go to [menti.com](https://www.menti.com) and use
2438 8084

```
int f := 5;  
int w := 0;  
mutex m := 1;  
ConditionVar cv1, cv2;
```

```
Task T =  
[[ while( true ){  
    transport();
```

```
    lock(m);  
    f := f+1;  
    sigall(cv2);  
    unlock(m);
```

```
    lock(m);  
    while(...)  
        wait(m, cv2);  
    w := w+3;  
    sigall(cv1);  
    unlock(m);
```

```
}  
]]
```

What should
be the guard
condition?

```
Task P =  
[[ while( true ){  
    lock(m);  
    while(w<2)  
        wait(m, cv1);  
    w := w-2;  
    sigall(cv2);  
    unlock(m);
```

```
    produce();
```

```
    lock(m);  
    f := f+1;  
    sigall(cv2);  
    unlock(m);
```

```
}  
]]
```

||

```
Task S =  
[[ while( true ){  
    lock(m);  
    while(...)  
        wait(m, cv2);  
    f := f-2;  
    unlock(m);  
  
    transport();  
  
    lock(m);  
    while(...)  
        wait(m, cv2);  
    w := w+3;  
    sigall(cv1);  
    unlock(m);
```

```
}  
]]
```

Question:

Enforce the following **synchronization invariants** using **condition variables**

- **I1:** $w \geq 0$
- **I2:** $f \geq 2 * w$

Prevent **race conditions** on f and w

Functions *produce()* and *transport()* cannot be in critical sections

Exercise (typical exam question)

```
int f := 5;  
int w := 0;  
mutex m := 1;  
ConditionVar cv1, cv2;
```

```
Task T =  
[[ while( true ){  
    transport();  
  
    lock(m);  
    f := f+1;  
    sigall(cv2);  
    unlock(m);  
  
    lock(m);  
    while( w > (f-6)/2 )  
        wait(m, cv2);  
    w := w+3;  
    sigall(cv1);  
    unlock(m);  
}  
]]
```

||

```
Task P =  
[[ while( true ){  
    lock(m);  
    while(w < 2)  
        wait(m, cv1);  
    w := w-2;  
    sigall(cv2);  
    unlock(m);  
  
    produce();  
  
    lock(m);  
    f := f+1;  
    sigall(cv2);  
    unlock(m);  
}  
]]
```

||

```
Task S =  
[[ while( true ){  
    lock(m);  
    while(...)  
        wait(m, cv2);  
    f := f-2;  
    unlock(m);  
  
    transport();  
  
    lock(m);  
    while(...)  
        wait(m, cv2);  
    w := w+3;  
    sigall(cv1);  
    unlock(m);  
}  
]]
```

Question:

Enforce the following **synchronization invariants** using **condition variables**

- **I1:** $w \geq 0$
- **I2:** $f \geq 2 * w$

Prevent **race conditions** on f and w

Functions *produce()* and *transport()* cannot be in critical sections

Exercise (typical exam question)

```
int f := 5;  
int w := 0;  
mutex m := 1;  
ConditionVar cv1, cv2;
```

```
Task T =  
[[ while( true ){  
    transport();  
  
    lock(m);  
    f := f+1;  
    sigall(cv2);  
    unlock(m);  
  
    lock(m);  
    while( w > (f-6)/2 )  
        wait(m, cv2);  
    w := w+3;  
    sigall(cv1);  
    unlock(m);  
}  
]]
```

||

```
Task P =  
[[ while( true ){  
    lock(m);  
    while(w < 2)  
        wait(m, cv1);  
    w := w-2;  
    sigall(cv2);  
    unlock(m);  
  
    produce();  
  
    lock(m);  
    f := f+1;  
    sigall(cv2);  
    unlock(m);  
}  
]]
```

||

```
Task S =  
[[ while( true ){  
    lock(m);  
    while(f < 2*w+2)  
        wait(m, cv2);  
    f := f-2;  
    unlock(m);  
  
    transport();  
  
    lock(m);  
    while(w > (f-6)/2)  
        wait(m, cv2);  
    w := w+3;  
    sigall(cv1);  
    unlock(m);  
}  
]]
```

Question:

Enforce the following **synchronization invariants** using **condition variables**

- **I1:** $w \geq 0$
- **I2:** $f \geq 2 * w$

Prevent **race conditions** on f and w

Functions *produce()* and *transport()* cannot be in critical sections

- Two signalling strategies
 - *Signal* wakes up **one** waiter on cv;
 - *Sigall* wakes up **all** waiters on cv.
- Which one to choose: *Signal* or *Sigall* ?
 - Requires a careful analysis to ensure both correctness and efficiency.
 - *Sigall* is **always correct**, but **often inefficient**.
- ***Wait(m, cv)*** can also be **extended with** a **timeout** mechanism
 - To recheck a condition even when there is no signal
 - **May increase robustness** against (programming) errors
(e.g., signal **not sent**, or use of ***signal()*** instead of ***sigall()*** in the wrong situation, normally or abnormally **terminated signaller**, ...)

```
pthread_cond_t cond = PTHREAD_COND_INITIALIZER;  
status = pthread_cond_init (&cond, attr);          /* should return 0      */  
status = pthread_cond_destroy (&cond);              /* idem                */  
status = pthread_cond_wait (&cond, m);              /* semaphore m is associated with  
                                                    * all critical sections */  
status = pthread_cond_timedwait (&cond, m, exp); /* exp: max. waiting time; returns  
                                                    * ETIMEDOUT after exp.*/  
status = pthread_cond_signal (&cond);               /* signal one waiter    */  
status = pthread_cond_broadcast (&cond);            /* signal all waiters   */
```

```
m = threading.Lock()
cv = threading.Condition(lock=m)
```

```
cv.acquire()
while not condition_respected():
    cv.wait()
execute_critical_section()
cv.release()
```

```
cv.acquire()
do_something()
cv.notify()
cv.release()
```

```
cv = threading.Condition()
```

with cv:

```
cv.wait_for(condition_respected)  
execute_critical_section()
```

with cv:

```
do_something()  
cv.notify()
```

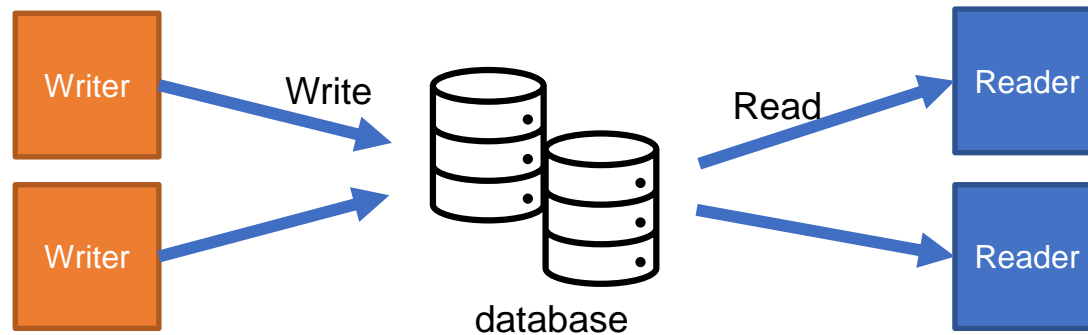
- Condition variables
- **Monitors**
 - Signaling disciplines

- A **monitor** is an object that **contains** a combination of **data** and **operations on those data**
 - **Data** record the **state of the monitor**
 - **Procedures** of a monitor are executed **under mutual exclusion**
 - **At most one task** can be “**inside the monitor**”
 - **Synchronization** is implemented **using condition variables**
- Often used as a **programming pattern** rather than natively supported by OS and programming languages

- Implemented in **Java** using the keywords *synchronize*, *wait*, and *notify*
- Supported by **Ada** with ***protected objects***

Programming language for **safety-critical systems** (used in **avionics** and **space**)

Example: readers-writers problem



Why using **mutexes** or semaphores is **not a good solution**?

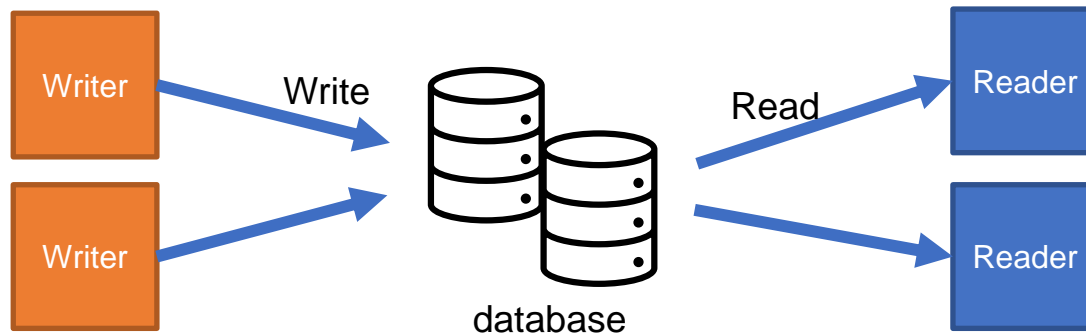
At any given time, we can have **several** readers but no writer, or **one** writer and no readers

```
Proc Writer =  
[[  
  while( true ){  
    Write_action();  
  }  
]]
```

```
Proc Reader =  
[[  
  while( true ){  
    Read_action();  
  }  
]]
```

Example: readers-writers problem

Synchronization invariant to satisfy
 $(nr = 0 \wedge nw \leq 1) \vee nw = 0$



At any given time, we can have **several** readers but no writer, or **one** writer and no readers

We use a **monitor** to **synchronize accesses** to the read and write actions

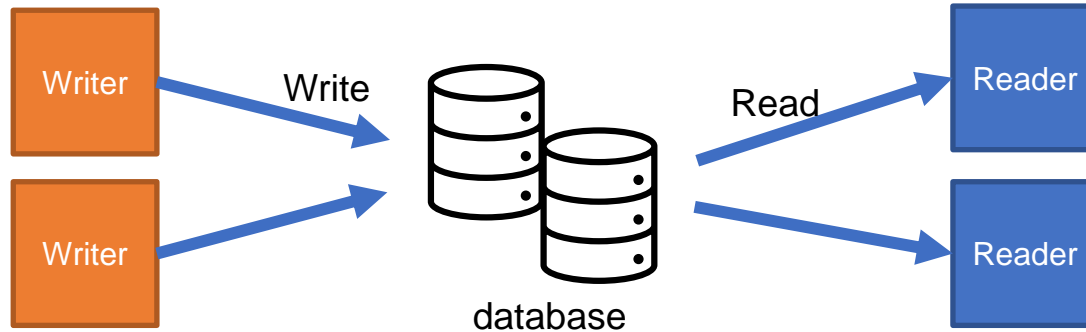
```
Proc Writer =  
[[  
  while( true ){  
    Write_action();  
  }  
]]
```

```
Monitor MonReadWrite =  
[[  
  int nr := 0; // number of readers reading  
  int nw := 0; // number of writers writing  
]]
```

```
Proc Reader =  
[[  
  while( true ){  
    Read_action();  
  }  
]]
```

Example: readers-writers problem

Synchronization invariant to satisfy
 $(nr = 0 \wedge nw \leq 1) \vee nw = 0$



MonReadWrite *RW*;

```
Proc Writer =  
[[  
  while( true ){  
    RW.writeEntry();  
    Write_action();  
    RW.writeExit()  
  }  
]]
```

```
Proc Reader =  
[[  
  while( true ){  
    RW.readEntry();  
    Read_action();  
    RW.readExit();  
  }  
]]
```

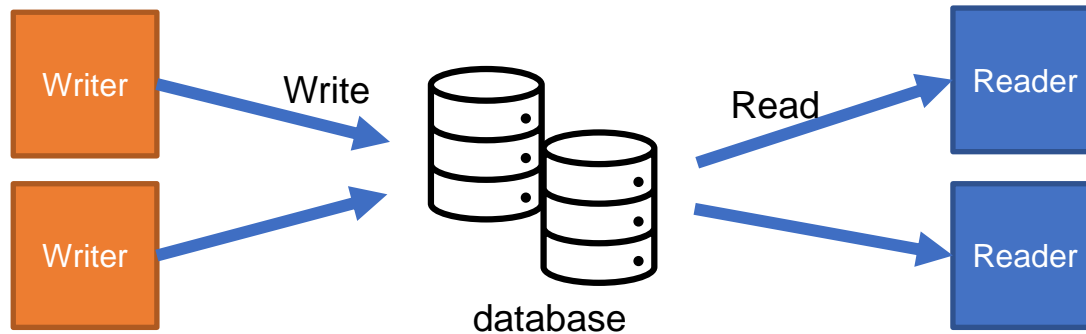
Monitor *MonReadWrite* =

```
[[  
  int nr := 0; // number of readers reading  
  int nw := 0; // number of writers writing  
  Proc writeEntry() { nw++; }  
  Proc writeExit() { nw--; }  
  Proc readEntry() { nr++; }  
  Proc readExit() { nr--; }  
]]
```

Add **entry and exit protocols**
around the read and write actions

Example: readers-writers problem

Synchronization invariant to satisfy
 $(nr = 0 \wedge nw \leq 1) \vee nw = 0$



MonReadWrite RW;

```
Proc Writer =  
[[  
  while( true ){  
    RW.writeEntry();  
    Write_action();  
    RW.writeExit()  
  }  
]]
```

Note 2: we use
signal for write but
sigall for read

Why?

```
Proc Reader =  
[[  
  while( true ){  
    RW.readEntry();  
    Read_action();  
    RW.readExit();  
  }  
]]
```

Note 1: we do not
use a mutex because
monitors already
enforce mutual
exclusion between
their procedures

Monitor MonReadWrite =

```
[[  
  ConditionVar readCond, writeCond;  
  int nr := 0; // number of readers reading  
  int nw := 0; // number of writers writing  
  Proc writeEntry() {  
    while(nw>0 or nr>0) wait(writeCond);  
    nw++;  
  }  
  Proc writeExit() {  
    nw--; signal(writeCond); sigall(readCond);  
  }  
  Proc readEntry() {  
    while(nw>0) wait(readCond);  
    r++;  
  }  
  Proc readExit() {  
    nr--; signal(writeCond);  
  }  
]]
```

Example: readers-writers problem

Synchronization invariant to satisfy
 $(nr = 0 \wedge nw \leq 1) \vee nw = 0$

Why did we **not** implement the *Write_action()* and *Read_action()* procedures in the monitor, instead of using *Entry* and *Exit* procedures?

Would limit the number of readers reading concurrently to one because the monitor procedures are mutually exclusive.

MonReadWrite RW;

```
Proc Writer =  
[[  
  while( true ){  
    RW.writeEntry();  
    Write_action();  
    RW.writeExit()  
  }  
]]
```

```
Proc Reader =  
[[  
  while( true ){  
    RW.readEntry();  
    Read_action();  
    RW.readExit();  
  }  
]]
```

```
Monitor MonReadWrite =  
[[  
  ConditionVar readCond, writeCond;  
  int nr := 0; // number of readers reading  
  int nw := 0; // number of writers writing  
  Proc writeEntry() {  
    while(nw>0 or nr>0) wait(writeCond);  
    nw++;  
  }  
  Proc writeExit() {  
    nw--; signal(writeCond); sigall(readCond);  
  }  
  Proc readEntry() {  
    while(nw>0) wait(readCond);  
    r++;  
  }  
  Proc readExit() {  
    nr--; signal(writeCond);  
  }  
]]
```

Implementing a monitor in Java

Java has a **single** implicit **condition variable** associated with every object

```
class MonReadWrite {
    private int nr = 0; // number of readers reading
    private int nw = 0; // number of writers writing

    public synchronized void writeEntry() {
        while(nw>0 or nr>0) {
            try{ wait(); }
            catch(InterruptedException e){}
            finally{}
        }
        nw++;
    }

    public synchronized void writeExit() {
        nw--; notifyAll();
    }

    public synchronized void readEntry() {
        while(nw>0) {
            try{ wait(); }
            catch(InterruptedException e){}
            finally{}
        }
        nr++;
    }

    public synchronized void readExit() {
        nr--; notifyAll();
    }
}
```

- Condition variables
- Monitors
 - **Signaling disciplines**

- A monitor:
 - **encapsulates** relevant **shared variables**
 - provides well-defined **operations on the shared variables**
- A monitor **provides exclusion**.
 - **At most one task may be inside** the monitor at any given time.
- Question: *Which task is inside the monitor right after the signal?*
 - The answer depends on the *signalling discipline (scheduling discipline)*.

- The discipline defines what happens to the monitor upon a *signal*.
- Four disciplines are used in various implementations
 - **signal & exit**
 - **signal & continue**
 - **signal & wait**
 - **signal & urgent wait**
- *Correctness of solution depends on signalling discipline!*

- **The task executing a signal exits** the monitor immediately.
 - Monitor access is given to a waiter on that variable, if any.
 - ***Sigall* is not supported**
- **Signalling strategy:**
 - upon leaving the monitor, a task performs **at most one signal** on a condition variable cv that is non-empty...
 - ... and for which the corresponding condition $B(cv)$ holds.
- **Property:** **The condition $B(cv)$ certainly holds after *Wait()*.**

- **The task executing a signal continues to use the monitor**, until a *Wait* or until the end of the routine.
 - Any **task released** by this signal **compete** with the other tasks to obtain monitor access again.
- **Property: The condition may or may not hold after *Wait*.**
- **Signalling strategy:**
 - during execution of a monitor routine **all relevant conditions are signalled**.
 - **after each wait()**, the condition is **evaluated** again.

- **The task executing a signal is stopped in favour of the signalled process.**
 - The **signaller waits** to access the monitor again after the *Signal*, ...
 - ... **and competes** with all other tasks **again**.
- **Signalling strategy:** similar to *Signal & Exit*.
- **Property:** **The condition holds after *Wait*.**

- The task executing a signal is stopped in favour of the signalled process.
 - The **signaller** waits to access the monitor again after the *Signal*, ...
 - ... but **does not compete** with all other tasks again.
 - **Signaller has priority** over other tasks.
- Signalling strategy: similar to *Signal & Exit*.
- Property: The condition holds after *Wait*.

- **Limitations of actions synchronization** using semaphores
 - Can only enforce invariants based on **action counts**
- How to **synchronize** tasks **using condition variables**
 - Using **guards** and **signals**
- Presented the notion of **monitors**
 - **Object** made of procedures and private data
 - **Mutual exclusion** on the execution of **procedures** (only one task in the monitor)
 - **Condition variables** can be used to enforce invariants on states reachable by the monitor
- **Next week**, we discuss how to **detect, prevent** and **recover from deadlocks** in more details
- I **released additional exercises** on condition synchronization to train yourself