

# Operating Systems (2INC0)

## Memory Management (09) Reminders

**Dr. Geoffrey Nelissen**

Courtesy of Dr. Tanir Ozcelebi,  
Dr. I. Radovanovic, and Dr. R. Mak  
(figures from Bic & Shaw)



Interconnected  
Resource-aware  
Intelligent Systems

**TU/e**

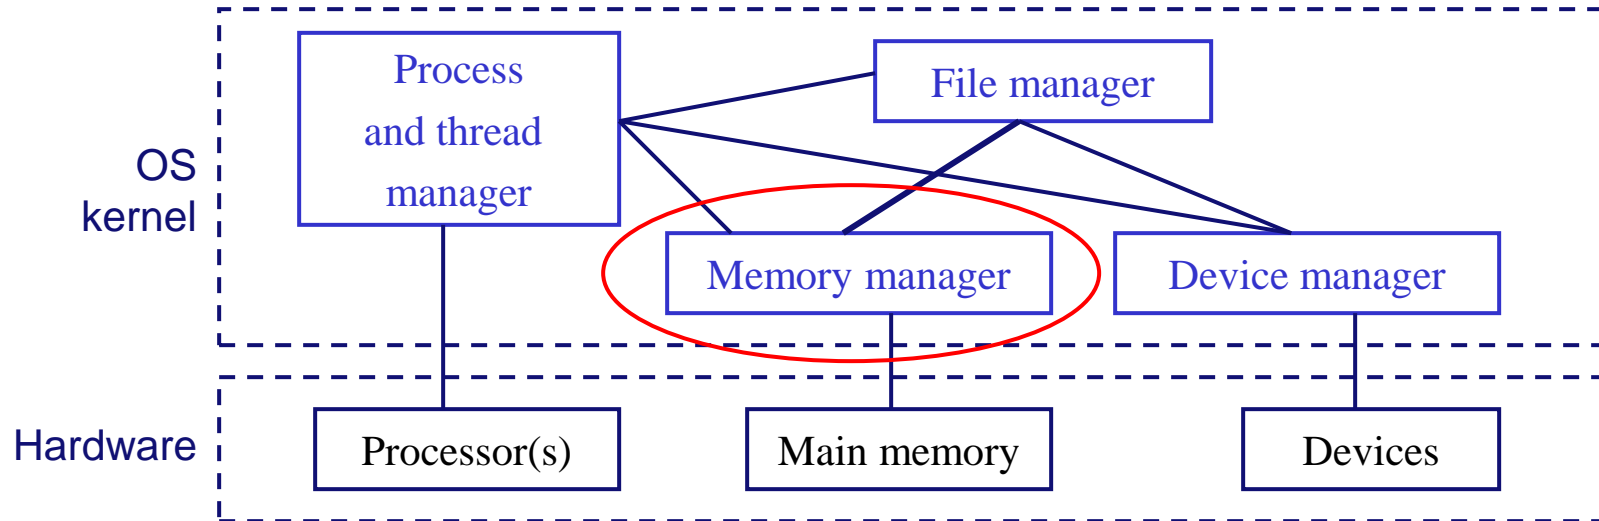
Technische Universiteit  
**Eindhoven**  
University of Technology

Where innovation starts

# OS as resource manager

- Resource: Anything that is needed for a process to run
  - Main Memory (MM)
  - Space on a disk
  - CPU
- **OS roles**:
  - create resource **abstractions** (hide HW details)
    - files for disk space, processes for memory use, threads for CPU use, ...
  - **manage resource sharing**
    - space-sharing vs time-sharing
    - provides a mechanism to isolate (protect) resources and programs.

# Logical OS Organization



- All modules interact to coordinate their activities.
- Examples:
  - **For virtual memory:** Process Man. – Memory Man.
    - Coordinate scheduling activity with memory allocation
  - **For performance improvement:** File Man. – Memory Man.
    - Information from storage device read prior to request by a thread

# Process execution

- Requirements:
  - A running program must be brought (from disk) into MM.
  - Each process must have a distinct address space (protection).
- ➔ Main memory is shared in time and space between concurrent programs

# User/programmer perspective

A good memory is one that...


- ...is non-volatile
- ...is private for a program (protected)
- ...has infinite capacity
- ...has zero access time
- ...is simple to use
- ...is cheap

Such a memory clearly does not exist, ...

- ... but a hierarchical hardware organization combined with virtualization can help!
  - We discuss virtual memory in the next lecture...

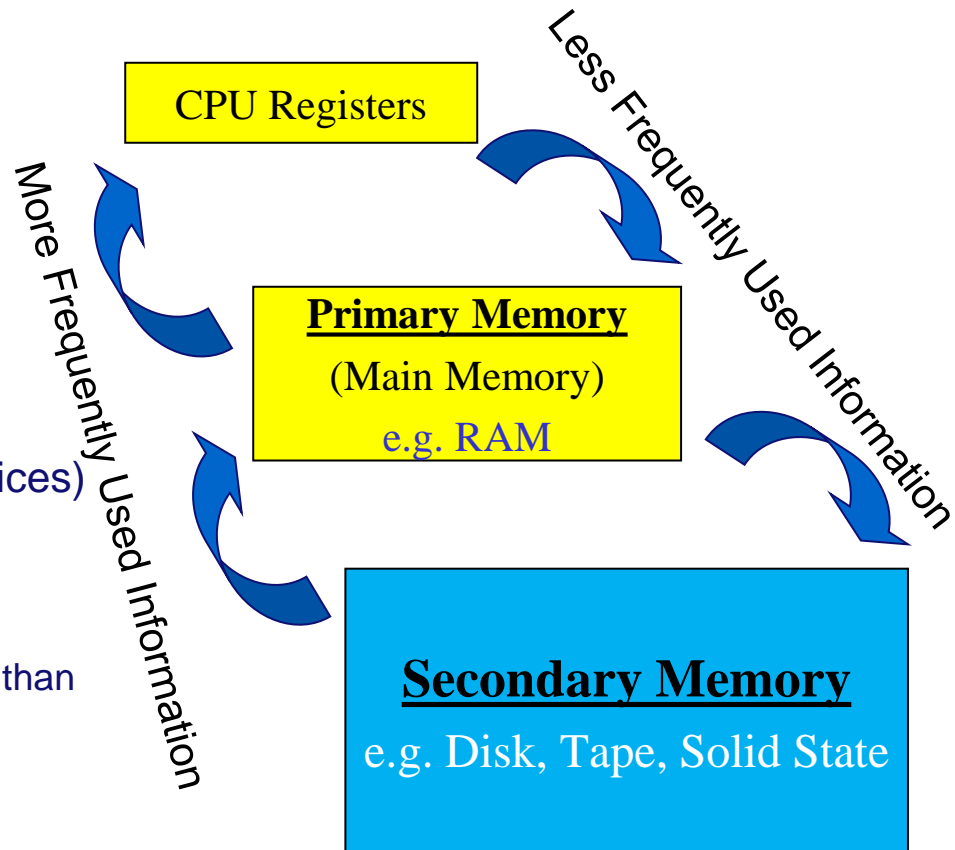
# Analogy: Office storage hierarchy

- Office storage hierarchy
  - Desktop
  - File folder
  - File cabinet
  - Warehouse
- Upper-layers:
  - info more frequently used
- Lower layers:
  - info less frequently used

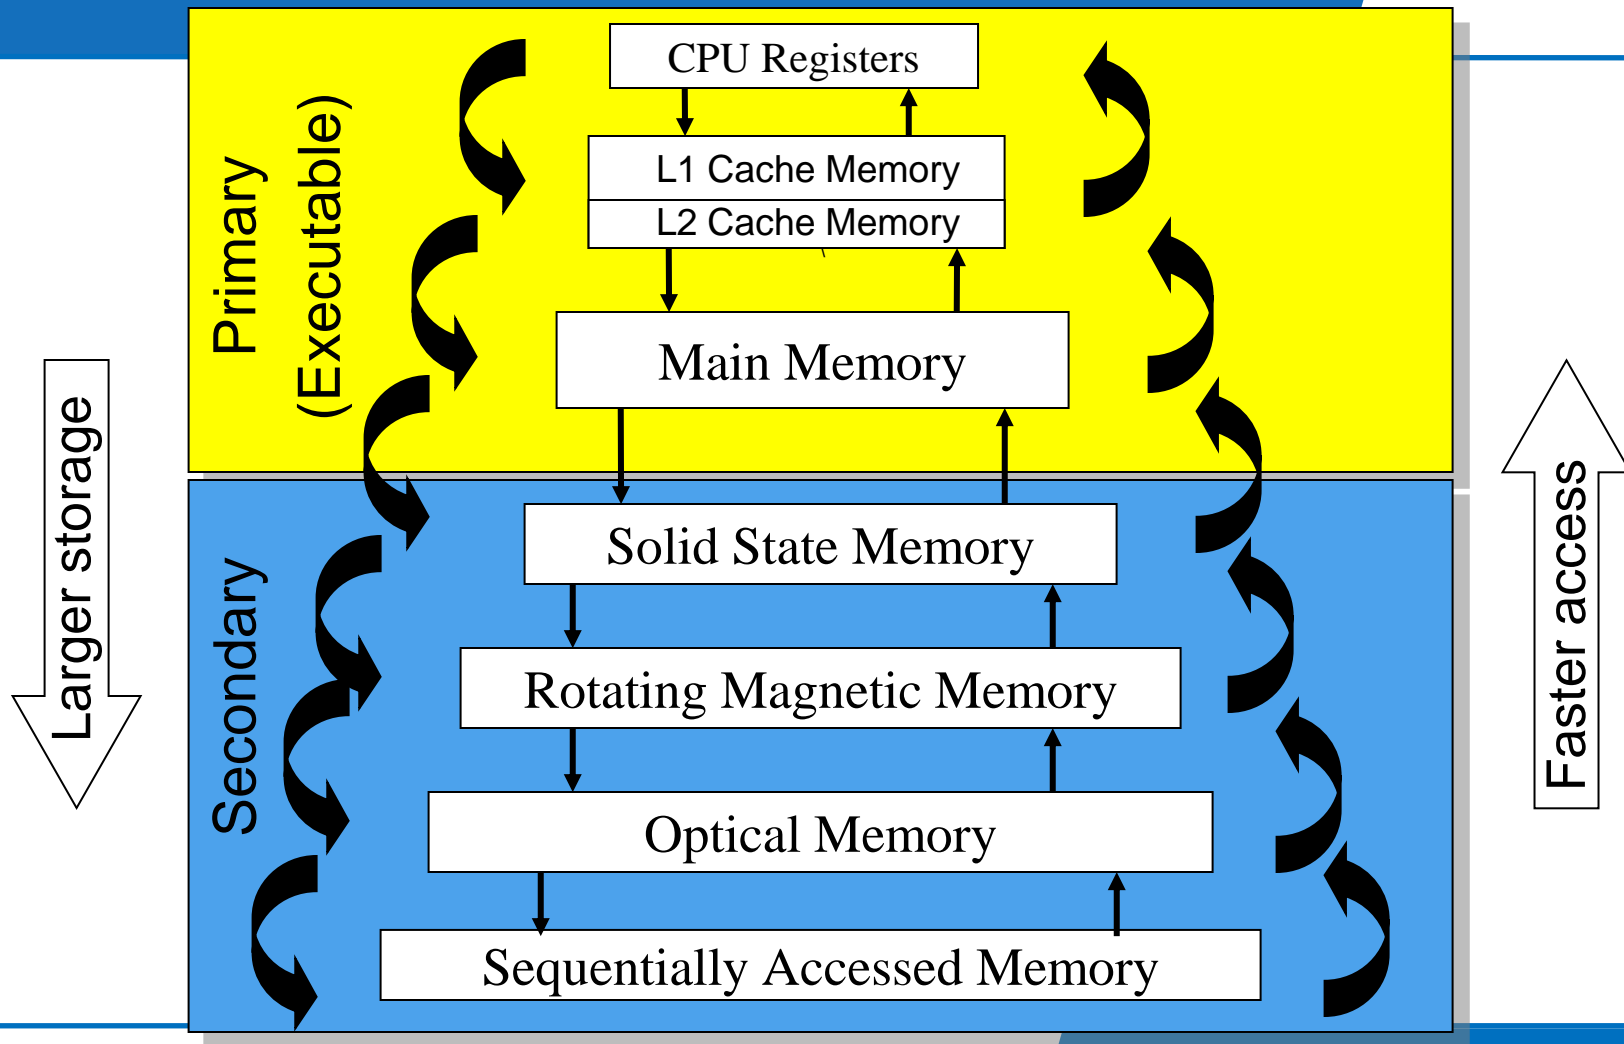


# Memory hierarchy

- **CPU registers**
  - **fast access** (1 clk cycle)
  - **limited in size**
- **Primary (executable) memory**
  - **direct access** (~a 100 clk cycles),
  - **relatively larger**
- **Secondary memory** (storage devices)
  - **Accessed through I/O operations**,
  - **very cheap / very large**
  - data stored for a long period of time
  - **slow** (3 orders of magnitude slower than CPU register memory)



# Contemporary memory hierarchy & dynamic loading





# OS exploiting the hierarchy

- **Place**
  - frequently-used info high in the hierarchy
  - infrequently-used info low in the hierarchy

# OS exploiting the hierarchy

- **Place**
  - frequently-used info high in the hierarchy
  - infrequently-used info low in the hierarchy
- Updates are first applied to upper memory.
  - Upward moves are (usually) copy operations
    - Image exists in both higher & lower memories
  - Downward moves are (usually) destructive
    - Usually to save space in upper memory.
    - Destroy image in upper memory
    - Update image in lower memory
      - avoid when the copy in lower memory is still consistent!

# Memory manager functional requirements

- Classical memory managers provide
  - **Abstraction**
    - MM (virtually) **appears to be larger** than the physical machine mem.
    - Memory is presented as an **array of contiguously addressed bytes**
    - Abstract set of **logical addresses** to reference physical memory
  - **Allocation**
    - Allocate primary memory to processes as requested
  - **Isolation**
    - Enable **mutually exclusive use of memory** by processes
  - **Memory Sharing**
    - **Enable memory to be shared by processes**

# Operating Systems (2INC0)

## Memory Management (09) Using partitions

**Dr. Geoffrey Nelissen**

Courtesy of Dr. Tanir Ozcelebi,  
Dr. I. Radovanovic, and Dr. R. Mak  
(figures from Bic & Shaw)



Interconnected  
Resource-aware  
Intelligent Systems

**TU/e**

Technische Universiteit  
**Eindhoven**  
University of Technology


Where innovation starts

# Early systems

- **Single-programmed solution**
- Multi-programmed solutions
  - Fixed partitions
  - Dynamic partitions
  - Relocatable dynamic partitions

# Single-programmed solution

- **Program loaded in its entirety into MM** and allocated as much contiguous MM space as needed
- Once in MM, program **stays there until execution ends.**
- **Problems:**
  - **No multiprocessing support.**
  - **Large context-switch overhead**
  - If a program does not fit in the memory, **it cannot be executed.**
- **Advantages:**
  - **simple hardware** requirements (base address and limit registers)
  - **Easy protection** between OS and processes and between processes



# Early systems

- Single-programmed solution
- Multi-programmed solution
  - **Fixed partitions**
  - Dynamic partitions

# Fixed partitions

- **Divide the memory into fixed partitions**
  - **one process per partition**
  - partitions may have different sizes to accommodate various types of programs (fixed at OS initialization)
  - a process can only access its own partition (protection)
- **Partition allocation scheme** needed (next slide uses first-fit)
- **Disadvantage:**
  - Entire program must be stored in MM
  - **Internal fragmentation**
    - Solution: **Choosing the right partition size**
      - Too small – long turnaround time (doesn't fit in any partition in the worst case)
      - Too big – wasted memory



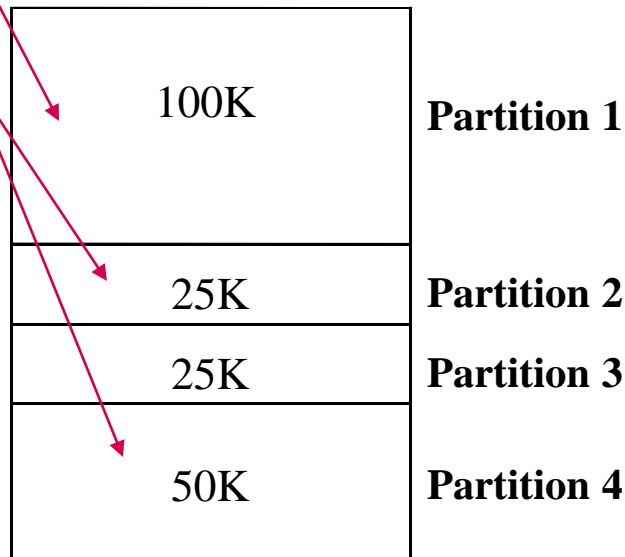
# Fixed partition Example

## Process List :

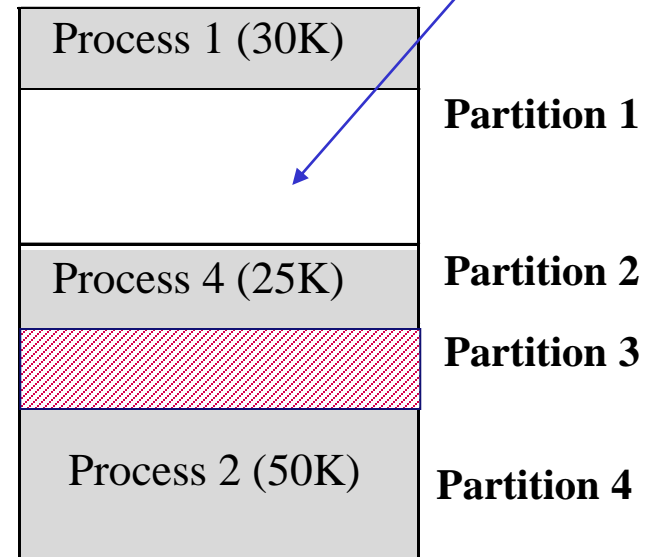
P1 30K  
P2 50K  
P3 30K  
P4 25K

**Process 3 must wait** even though **70K of contiguous free space** is available (and **70K+25K** available in total).

### Original State



### After Process Entry



# Early systems

- Single-programmed configuration
- Multi-programmed configurations
  - Fixed partitions
  - **Dynamic partitions**
    - Allocation
    - De-allocation

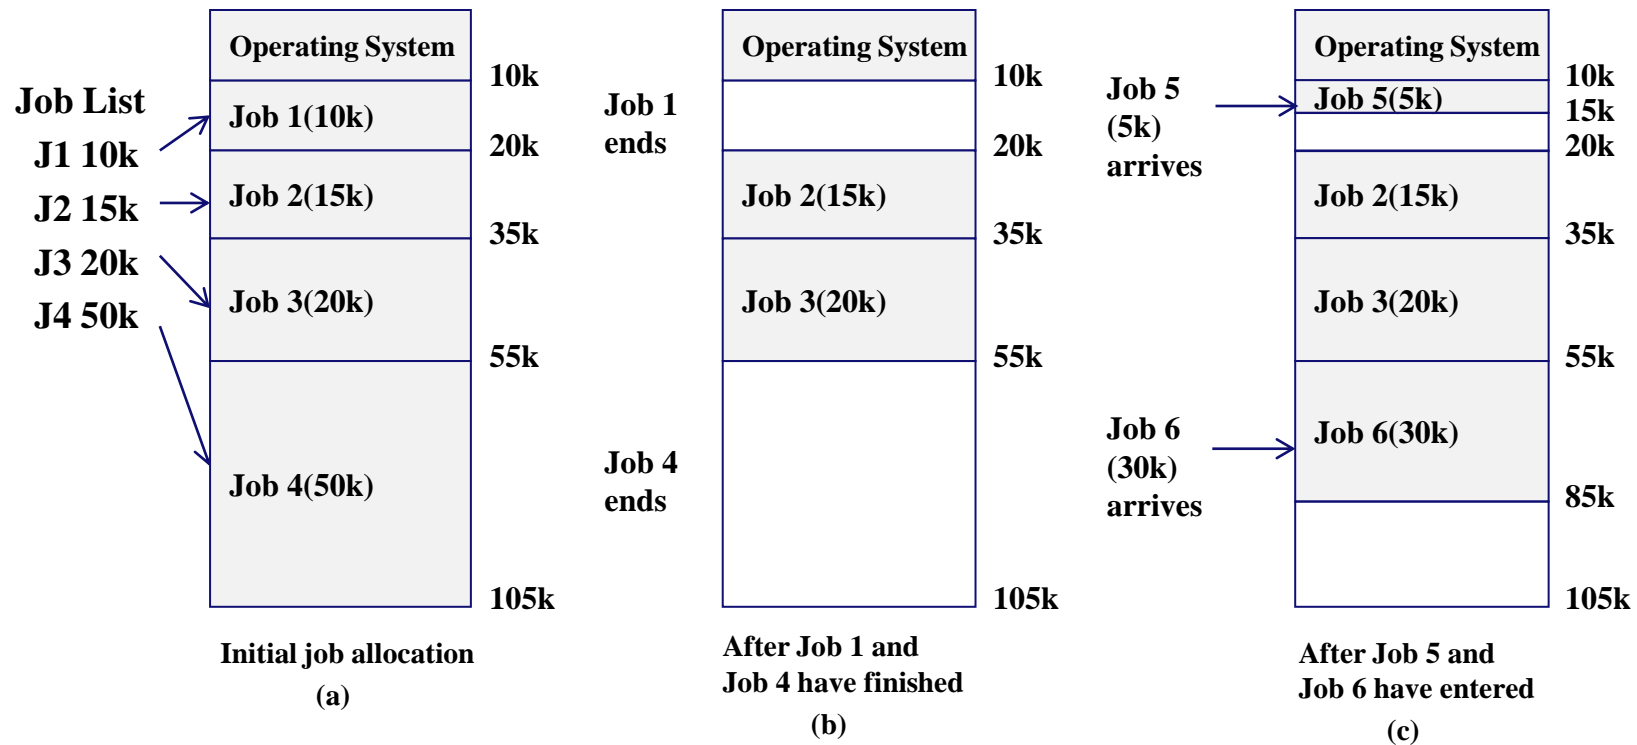
# Dynamic partitions

- **Partition size** not fixed but determined **based on the process size**.
- **Advantage:**
  - **no** memory waste due to internal fragmentation
- **Disadvantages:**
  - **External fragmentation**
    - Memory utilization deteriorates as new processes enter and old processes exit the system
  - **The complete program must still be stored contiguously.**

# Dynamic partitions

## Example

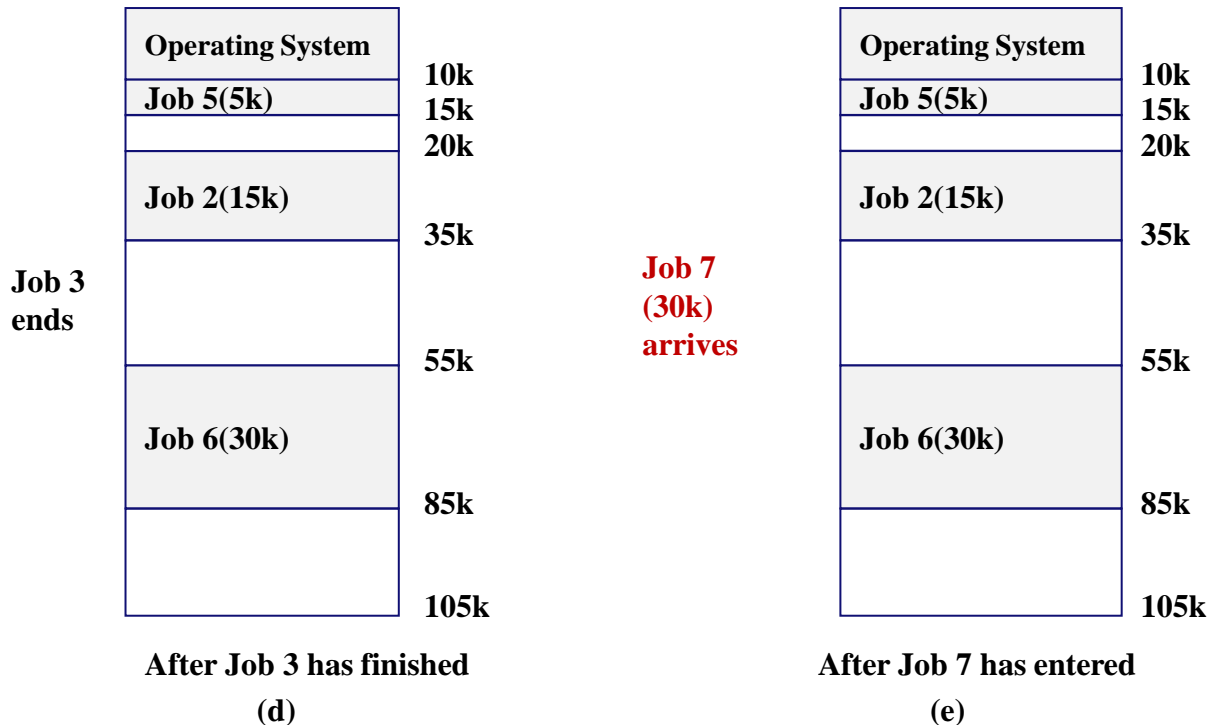
### First-fit allocation scheme



# Dynamic partitions

## Example

First-fit allocation scheme (cont'd)



**Job 7 must wait** even though there is **45K of free memory** space available

# Dynamic partition: Allocation schemes

- Goals: decrease allocation time, increase memory utilization
- First-fit
  - Allocate the *first* partition that is big enough
  - Advantage:
    - Faster in making the allocation
  - Disadvantage:
    - Waste more memory space
- Best-fit
  - Allocate the smallest partition fitting the requirements
  - Advantage:
    - Produces the smallest leftover partition
    - Makes best use of memory *when it is possible to allocate a new job*
  - Disadvantage:
    - Takes more time

# Memory allocation

## Keep track of holes using a list

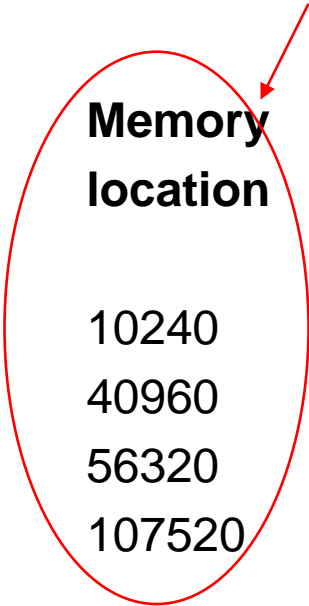
Memory location	Memory block size
10240	30K
40960	15K
56320	50K
107520	20K

Total Available: 115K

# Memory allocation

## First-fit

**Order the list** by  
memory location



Memory location	Memory block size
10240	30K
40960	15K
56320	50K
107520	20K

Total Available: 115K



# Memory allocation

## Best-fit

**Order the list** by  
block sizes

Memory location	Memory block size
40960	15K
107520	20K
10240	30K
56230	50K

Total Available: 115K

# Exercise

- The table of free memory spaces is given below. Show how this table will look after allocating a process that occupies 200 spaces using the Best-Fit algorithm.

<u>Before request</u>		<u>After request</u>	
Beginning address	Memory block size	Beginning address	Memory block size
4075	105	4075	105
5225	5	5225	5
6785	600	6785	600
7560	20	7560	20
7600	205	*7800	5
10250	4050	10250	4050
15125	230	15125	230
24500	1000	24500	1000

Best-Fit is **not** the best for dynamic partitions:

Excessive memory fragmentation due to smaller free space

# Allocation schemes

- First-Fit: clusters small holes at the top
- Next-Fit (rotating First-Fit):
  - start each search at the point where the previous search stopped  
(faster search, slightly worse memory utilization than First-Fit)
- Best-Fit: worst performance due to excessive fragmentation
- Worst-Fit: always uses the largest available partition
  - similar performance to first-fit and next-fit
- Optimized schemes
  - statistics on average process size etc. needed

# Early systems

- Single-programmed configuration
- Multi-programmed configurations
  - Fixed partitions
  - **Dynamic partitions**
    - Allocation
    - **De-allocation**

# Release of memory space

- Occurs **when** process terminates or suspends
- For **fixed partitions**:
  - Memory Manager **resets status of memory block to “free”**.
- For **dynamic partitions** tries to **combine free areas** of memory
  - Is the block **adjacent** to another free block?
    - If yes, combine the 2 blocks
  - Is the block **between** 2 free blocks?
    - If yes, combine those 3 blocks
  - Is the block **isolated** from other free blocks?
    - If yes, make a new table entry



# De-allocation Example

## *Before Deallocation*

Beginning address	Memory block size	Status
4075	105	Free
5225	5	Free
6785	600	Free
7560	40	Free
(7600)	(200)	(Busy) <sup>1</sup>
*7800	5	Free
10250	4050	Free
15125	230	Free
24500	1000	Free

De-allocate →

# Important note

- **De-allocate only processes that are really idle**
- **Example: not waiting for an I/O device** to write in an I/O buffer in the process address space  
→ otherwise, the **I/O manager** may write in the address space of the **wrong process**

# Operating Systems (2INC0)

## Memory Management (09) Relocatable partitions

**Dr. Geoffrey Nelissen**

Courtesy of Dr. Tanir Ozcelebi,  
Dr. I. Radovanovic, and Dr. R. Mak  
(figures from Bic & Shaw)



Interconnected  
Resource-aware  
Intelligent Systems

**TU/e**

Technische Universiteit  
**Eindhoven**  
University of Technology

Where innovation starts




# Relocatable dynamic partitions

- *Memory manager* **relocates programs to gather all empty blocks** and generate one big free memory block.
- **Solves both internal and external fragmentation problems**
  - relocation and compaction **avoids memory waste**
  - “insufficient memory” problems less frequent

# Compaction steps

- **Relocate** every program in memory so that they're **contiguous**.
- **Adjust every address**, and every reference to an address, within each program to account for programs' new locations in memory.
  - Other values within the program are kept unchanged (e.g., data values).
- **Question: When should compaction be done?**
  - When a certain percent of memory becomes busy (e.g. 75%),
  - When there are processes pending, or
  - After a prescribed amount of time
- **Note: Compaction cannot be done if address binding is static.**

# Compaction Example




# Memory compaction (strategies)

**complete**  
compaction


**partial**  
compaction

**minimal**  
data movement


We need  
a block of  
size 10




(a)



(b)



(c)



(d)

# Memory compaction (remarks)

- **Disadvantage: very costly overhead**
  - requires each word to be read from and written into memory
    - May take several seconds
- Contemporary systems
  - Virtual Memory techniques (paging, segmentation) instead of compaction

# Relocation

Q: How to keep track of where each process is with respect to its original location?

A:

- Special-purpose registers
  - **Relocation register**
    - Contains value to be added to each address referenced in the program
  - **Limit register**
    - Stores the size of the memory space accessible by each program

# Swapping

- When?
  - a program must be loaded into MM and there is not enough room
  - move something else to secondary storage to make room.
- How?
  - OS selects a process to swap-out and replaces it with another process from the secondary storage
  - swap-out only those parts of process space that are not already on disk
    - typically code is not modified
    - benefits from hardware support that registers (keeps track of) modification
- Where to (on secondary memory)?
  - either to an arbitrary file (introduces file management overhead)
  - or to a special partition on disk, the *swap space*
    - more efficient (low-level I/O)