

2INC0 - Operating Systems

Virtual Memory Part 1

Geoffrey Nelissen



Interconnected
Resource-aware
Intelligent Systems

TU/e

Technische Universiteit
Eindhoven
University of Technology

Where innovation starts

- **Introduction to operating systems** (lecture 1)
- **Processes, threads and scheduling** (lectures 2+3)
- **Concurrency and synchronization**
 - atomicity and interference (lecture 4)
 - actions synchronization (lecture 5)
 - condition synchronization (lecture 6)
 - deadlock (lecture 7)
- **File systems** (lecture 8)
- **Memory management** (lectures 9+10)
- **Input/output** (lecture 11)

One of the most important subsystems in an OS. It impacts:

- **speed** of execution
- **Maximum size** of executable programs
- **how many** processes can be executed concurrently
- how data and code can be shared
- ...

Program₁:

```
int x=0;
void main(){
  int x=0;
  while (true) {
    x--;
    if( x > 8 ) {
      something(x);
    } else {
      func();
    }
  }
}
```

compilation



Code	Mem address
... sll \$t1, \$s3, 2	0x1000A800
add \$t1, \$t1, \$s6	0x1000A804
... lw \$a0, x	0x1000AFA0
jal Proc	0x1000AFA4
addi \$s3, \$s3, 1	0x1000AFA8
...	...
...	...
...	...
Data	
x	0x1000FF44

The compiler/linker associate a **(logical)** address to every instruction and data

Processes creation



Process 1

Code
... sll \$t1, \$s3, 2
add \$t1, \$t1, \$s6
... lw \$a0, x
jal Proc
addi \$s3, \$s3, 1
...
...
...
Data
x 0x1000FF44

Process 2

Code
... sll \$t1, \$s3, 2
add \$t1, \$t1, \$s6
... lw \$a0, x
jal Proc
addi \$s3, \$s3, 1
...
...
...
Data
x 0x1000FF44

How do we differentiate between x in Process 1 and Process 2 if they have the same address?

Program₁:

```
int x=0;
void main(){
  int x=0;
  while (true) {
    x--;
    if( x > 8 ) {
      something(x);
    } else {
      func();
    }
  }
}
```

compilation

Code	Mem address
... sll \$t1, \$s3, 2	0x1000A800
add \$t1, \$t1, \$s6	0x1000A804
... lw \$a0, x	0x1000AFA0
jal Proc	0x1000AFA4
addi \$s3, \$s3, 1	0x1000AFA8
...	...
...	...
Data	
x	0x0000FF44

The compiler/linker associate a **(logical)** address to every instruction and data

Physical memory

Processes creation

Process 1

Code
... sll \$t1, \$s3, 2
add \$t1, \$t1, \$s6
... lw \$a0, x
jal Proc
addi \$s3, \$s3, 1
...
...
Data
x 0x0000FF44

Process 2

Code
... sll \$t1, \$s3, 2
add \$t1, \$t1, \$s6
... lw \$a0, x
jal Proc
addi \$s3, \$s3, 1
...
...
Data
x 0x0000FF44

Process loading in main memory

The **physical** address of x in main memory is different

...	0
sll \$t1, \$s3, 2	
add \$t1, \$t1, \$s6	
...	
lw \$a0, x	
jal Proc	
addi \$s3, \$s3, 1	
...	
...	
x 0x0000FF44	300K
...	
sll \$t1, \$s3, 2	388K
add \$t1, \$t1, \$s6	
...	
lw \$a0, x	
jal Proc	
addi \$s3, \$s3, 1	
...	
...	
x 0x000070F44	688K

Program₁:

```
int x=0;
void main(){
  int x=0;
  while (true) {
    x--;
    if( x > 8 ) {
      something(x);
    } else {
      func();
    }
  }
}
```

compilation



Code	Mem address
... sll \$t1, \$s3, 2	0x1000A800
add \$t1, \$t1, \$s6	0x1000A804
... lw \$a0, x	0x1000AFA0
jal Proc	0x1000AFA4
addi \$s3, \$s3, 1	0x1000AFA8
...	...
...	...
Data	
x	0x0000FF44

The compiler/linker associate a **(logical)** address to every instruction and data

Physical memory

Processes creation



Process 1

Code
... sll \$t1, \$s3, 2
add \$t1, \$t1, \$s6
... lw \$a0, x
jal Proc
addi \$s3, \$s3, 1
...
...
Data
x 0x0000FF44

Process 2

Code
... sll \$t1, \$s3, 2
add \$t1, \$t1, \$s6
... lw \$a0, x
jal Proc
addi \$s3, \$s3, 1
...
...
Data
x 0x0000FF44

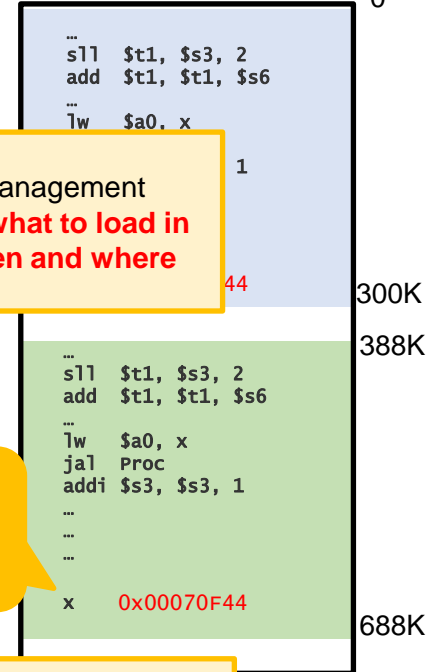
→ the memory management subsystem decides **what to load in main memory, when and where**

Process loading in main memory



The **physical** address of x in main memory is different

→ the memory management subsystem **translates logical addresses to physical addresses** for each process in main memory



- **Reminder**
- **Memory paging**
- **Conclusion**

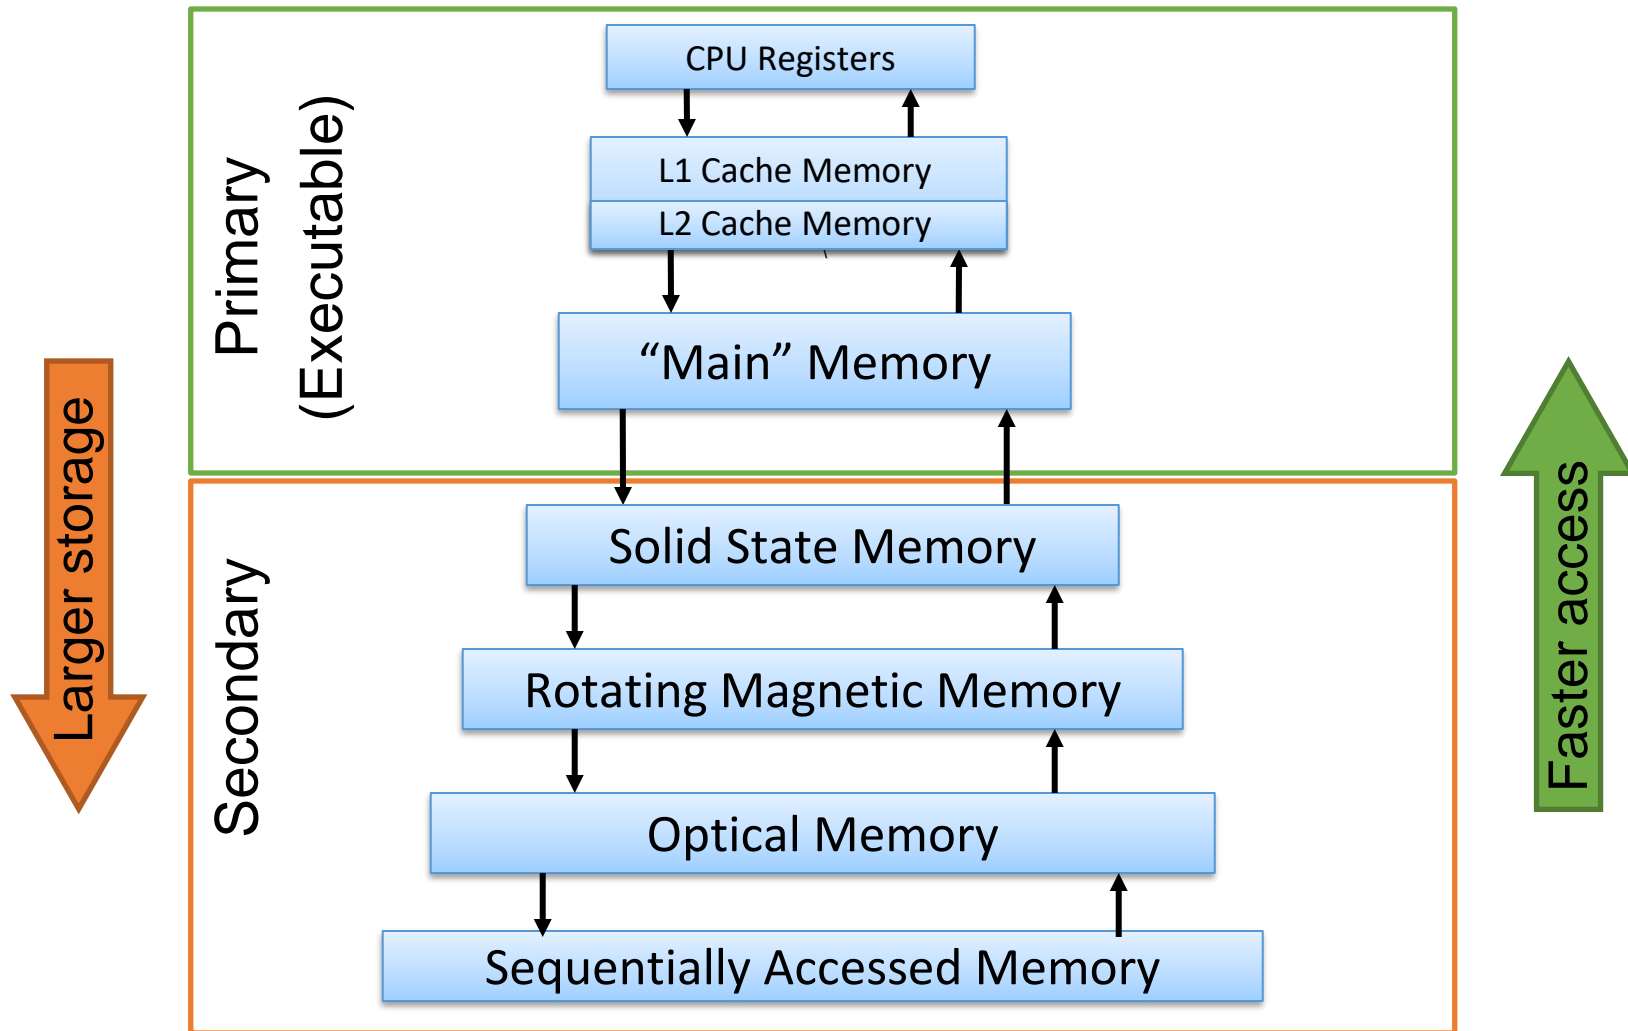
What would we expect from an **ideal memory**?

- **Simple** to use
 - The programmer should not have to know where data and code are loaded
- **Private** (isolation)
 - Tasks should be able to share only what they want to
- **Non-volatile / permanent**
 - Data and code should always be available in memory
- **Fast** access
 - For better performance
- **Huge** (unlimited) capacity
 - To be able to execute large programs and treat lots of data
- **Cheap**
 - To be cost-effective

Conflicting requirements!

We use a **memory hierarchy** and **virtualization** to find a good **tradeoff**.

Memory Hierarchy



Memory management in early systems

- Used to use **memory partitioning**
- Every active process resides **in its entirety in main memory**
 - **Large programs cannot execute**
- Every active process is **allocated a contiguous part of main memory**.

See preparatory material

Memory management in early systems

- Every active process resides **in its entirety in main memory**
 - **Large programs cannot execute**
- Every active process is **allocated a contiguous part of main memory**.
- Three **partitioning schemes**

See preparatory material

Fixed

- **Partitions** have **fixed size**
- Processes are **assigned to a free partition** that is big enough

Dynamic

Relocatable

Fixed partitioning: example

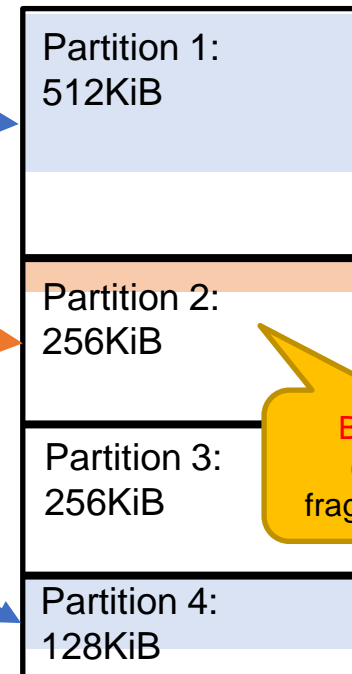
Assume a main memory partitioned as shown below, and a process to partition mapping using **Best Fit**.

Consider the following **sequence of process arrivals**, and the size of their memory address space.

- P1: 300KiB
- P2: 78KiB
- P3: 512KiB
- P4: 32KiB
- P5: 600KiB

P3 cannot be loaded until P1 is swapped out,
even though there is enough free space

P5 will never be able to execute
because it is too big



Memory management in early systems

- Every active process resides **in its entirety in main memory**
 - **Large programs cannot execute**
- Every active process is **allocated a contiguous part of main memory**.

See preparatory material

- Three **partitioning schemes**

Fixed

- **Partitions** have **fixed size**
- Processes are **assigned to a free partition** that is big enough
- **Number of active processes** limited by **number of partitions**
- **Size of processes** limited by the **size of** the largest **partition**
- Wastes memory due to **internal fragmentation**

Dynamic


- Processes are **dynamically allocated** space in main memory

Relocatable

Assume a main memory of 1MB as shown below, and a process to partition mapping using **Best Fit**.

Consider the following **sequence of process arrivals and leavings**, and the size of their memory address space.

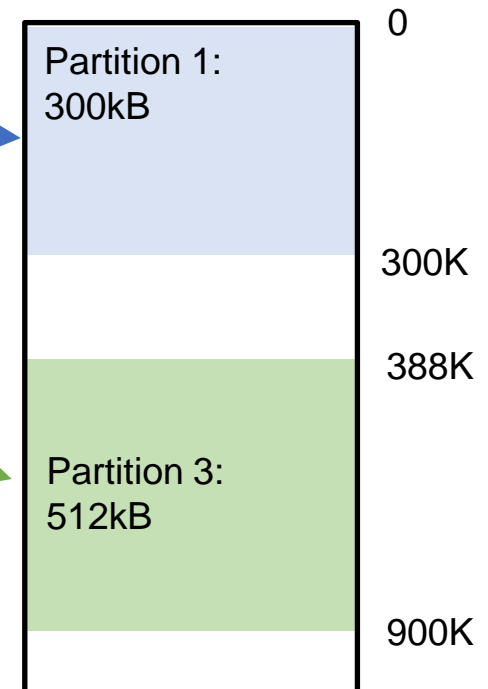
- P1: 300kB
- P2: 88kB
- P3: 512kB
- P2 completes and releases memory
- P4: 150kB
- P5: 1MB
- P6: 2MB



Assume a main memory of 1MB as shown below, and a process to partition mapping using **Best Fit**.

Consider the following **sequence of process arrivals and leavings**, and the size of their memory address space.

- P1: 300kB
- P2: 88kB
- P3: 512kB
- P2 completes and release memory
- P4: 150kB P4 cannot be loaded even though there is enough free space (**external fragmentation**)
- P5: 1MB **P5 requires exclusive access to the main memory**
- P6: 2MB P6 will never be able to execute because it is too big



Memory management in early systems

- Every active process resides **in its entirety in main memory**
 - **Large programs cannot execute**
- Every active process is **allocated a contiguous part of main memory**.

See preparatory material

- Three **partitioning schemes**

Fixed

- **Partitions** have **fixed size**
- Processes are **assigned to a free partition** that is big enough
- **Number of active processes** limited by **number of partitions**
- **Size of processes** limited by the **size of the largest partition**
- Wastes memory due to **internal fragmentation**

Dynamic

- Processes are **dynamically allocated** space in main memory
- **Number of active processes** is only limited by the size of the memory
- **Size of a process** is limited by the **size of the memory**
- Wastes memory due to **external fragmentation**


Relocatable

- Same as dynamic, but **can compact memory** to make space for more processes

Assume a main memory of 1MB as shown below, and a process to partition mapping using **Best Fit**.

Consider the following **sequence of process arrivals and leavings**, and the size of their memory address space.

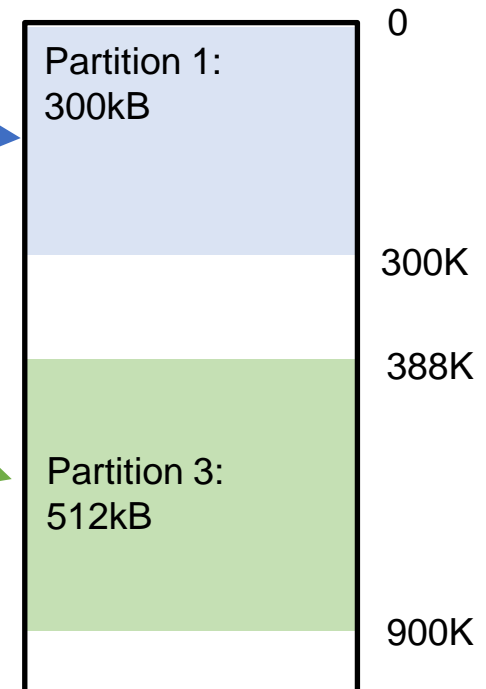
- P1: 300kB
- P2: 88kB
- P3: 512kB
- P2 leaves
- P4: 150kB



Assume a main memory of 1MB as shown below, and a process to partition mapping using **Best Fit**.

Consider the following **sequence of process arrivals and leavings**, and the size of their memory address space.

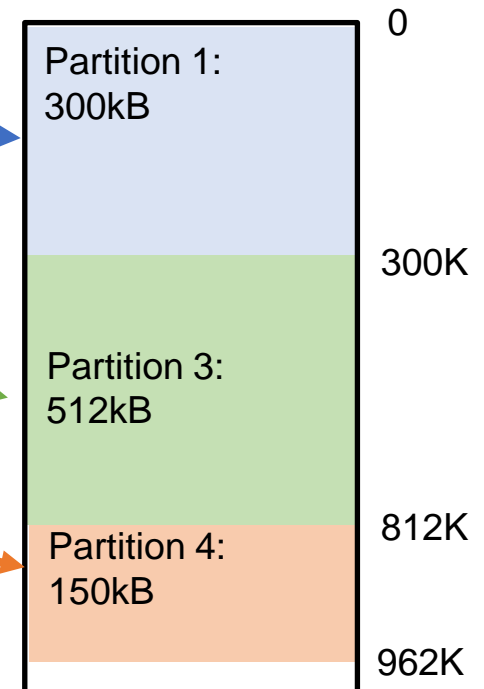
- P1: 300kB
- P2: 88kB
- P3: 512kB
- P2 leaves
- P4: 150kB



Assume a main memory of 1MB as shown below, and a process to partition mapping using **Best Fit**.

Consider the following **sequence of process arrivals and leavings**, and the size of their memory address space.

- P1: 300kB
- P2: 88kB
- P3: 512kB
- P2 leaves
- **compaction**
- P4: 150kB



Memory management in early systems

- Every active process resides **in its entirety in main memory**
 - **Large programs cannot execute**
- Every active process is **allocated a contiguous part of main memory**.
- Three **partitioning schemes**

See preparatory material

Fixed

- **Partitions** have **fixed size**
- Processes are **assigned to a free partition** that is big enough
- **Number of active processes** limited by **number of partitions**
- **Size of processes** limited by the **size of the largest partition**
- Wastes memory due to **internal fragmentation**

Dynamic

- Processes are **dynamically allocated** space in main memory
- **Number of active processes** is only limited by the size of the memory
- **Size of a process** is limited by the **size of the memory**
- Wastes memory due to **external fragmentation**

Relocatable

- Same as dynamic, but **can compact memory** to make space for more processes
- Compaction is a **costly operation**
- Requires **dynamic address binding**
- **Different compaction algorithms exist. They're explained in the preparatory material**

Preparatory material is part of the exam!

Translating virtual addresses in physical addresses (early systems) TU/e

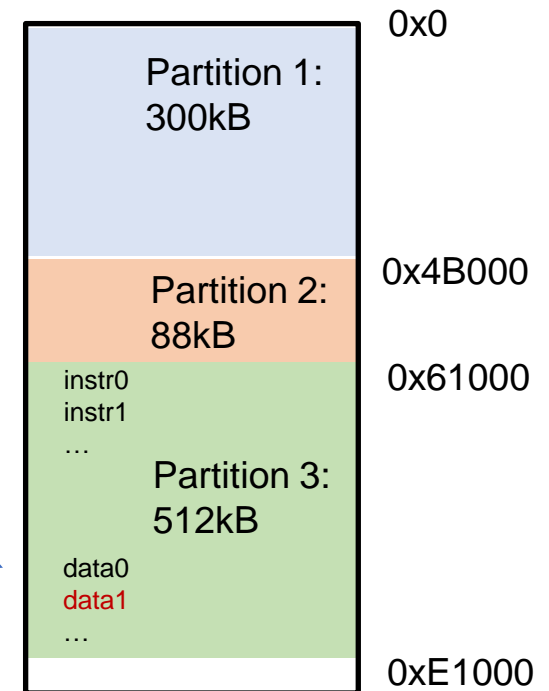
Process X

Code/data **Logical** address

instr0	0x0
instr1	0x1
...	...
data0	0xA10
data1	0xA11
...	...

What is the address of data1 in main memory (i.e., the physical address of data1)?

start address of process in physical memory + virtual address of data1 =
 $0x61000 + 0xA11$
 $= 0x61A11$




Do we have the ideal memory properties so far?

Simple	Yes	
Private	Yes/No	We have isolation but no sharing
Permanent	No	Unless the programmer enforces a process stays in memory
Fast	No	Process-based swapping is expensive (needs to move a complete process in and out) Compaction is expensive
Huge	No	Process size cannot exceed physical size of the main memory
Cost-effective	Yes	Using memory hierarchy


- Reminder
- **Memory paging**
- **Conclusion**

Let's use an analogy

Warehouse (secondary storage)




Office (primary storage)



Let's use an analogy

Warehouse (secondary storage)



A process moved
to the office



Office (primary storage)




Old (memory) management systems, **move a full pallet (process) in the office** whenever we need **to access a single a file (data/instruction)** from one of the boxes


- Long loading time
- Waste of office space
- Few pallets may be accessed at the same time in the office
➔ must often move a pallet back in the warehouse to access a new pallet (i.e., run another process)

Let's use an analogy

Warehouse (secondary storage)




Office (primary storage)



Let's use an analogy

Why are we not moving files instead of boxes?

Warehouse (secondary storage)



Office (primary storage)



Move one box
to the office



New (memory) management systems, **move one** (or a set of) **box** at a time **in the office** whenever we need **to access a file** in that box.

Store boxes on shelves of the same size than a box.

- **Short loading** time (i.e., quick to start executing a new task)
- Efficient **space usage**
- **Many more concurrent “tasks”** running in the office
- **Quick to swap** boxes

Divide the process address space in **segments of identical sizes called pages**

Process 1	
1 st 100 lines	
2 nd 100 lines	
3 rd 100 lines	
Remaining 50 lines	
Wasted space	

Page 0

Page 1

Page 2

Page 3

Main Memory	
Operating system	
Process 1 – Page 2	
Process 1 – Page 0	
Process 1 – Page 1	
Process 1 – Page 3	

Page
frame #

0
1
2
3
4
5
6
7
8
9
10
11
12

Divide the main memory in **segments of identical sizes called pages frames**

Advantage:

Every free frame is always usable
→ No external fragmentation

The size of page = the size of a page frame

A page can be loaded in any free page frame (one page per frame)

- programs do not have to be stored contiguously in main memory
- No external fragmentation
- Internal fragmentation is limited to the size of one page
- We only need to load the pages we actually use instead of the whole process

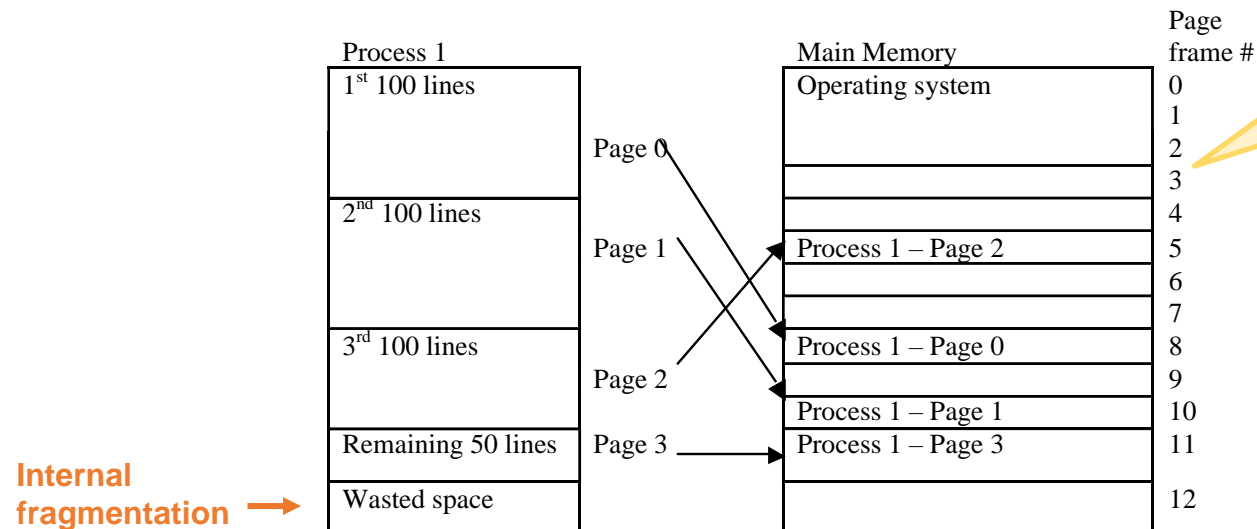
Memory paging: example

Assume that the **main memory** has a **size** of **1300 lines**. Each **page** and each **frame** has a **size** of **100 lines**.

The **OS** memory space **requires 280 lines**, and a loaded **process P1** has a **size** of **350 lines**

How many free frames are left in main memory?

There are 13 frames in main memory. OS needs 3 frames, P1 needs 4 => 6 frames left




Remaining problems:


1. We must **keep track of where pages are stored**
2. We must provide a mechanism for **address binding** (i.e., translate logical/virtual addresses into physical addresses)
3. How to decide **which page to load** in main memory and when?

Problem 1: keep track of pages (boxes) location

Warehouse (secondary storage)



Office (primary storage)



What solution would you propose to remember which box of which pallet is on which shelf in the office?

Solution 1:


For each pallet in the warehouse, keep track in a ledger **on which shelf each box is** located

Solution 2:

For each shelf in the office, keep track in a ledger **which box of which pallet** is stored there

Problem 1: keep track of pages (boxes) location

Warehouse (secondary storage)



Office (primary storage)



What solution would you propose to remember which box of which pallet is on which shelf in the office?

Solution 1:

For each pallet in the warehouse, keep track in a ledger **on which shelf each box is** located

Called a **page tables** in an OS

Solution 2:

For each shelf in the office, keep track in a ledger **which box of which pallet** is stored there

Called a **frame table** in an OS

Problem 1: keep track of pages location

Solution 1

- The OS maintains a **page table** for each process
- The address at which the page table is saved is recorded in the processes' **PCBs**

Process 0

Page ID	Frame ID
0	2
1	7
2	-
3	5

Process 1

Page ID	Frame ID
0	1
1	8
2	-
3	14
4	4
5	-
6	15

Process n

Page ID	Frame ID
0	-
1	0
2	-

...

Some pages may not be in main memory yet
→ no frame ID recorded

Records the **frame used by each page** of that process

Advantage:

- **Easy to find a page** in main memory (requires a **single look-up in the page table**)

Disadvantage:

- **Consumes a lot of memory** if each process is made of tens of thousands of pages

Problem 1: keep track of pages location

- Alternative to page tables (Solution 2):
the OS maintains a single **inverted page table** (also called **frame table**)

Frame ID	Process ID	Page ID
0	3	1
1	2	8
2	-	-
3	1	14
4	0	4
5	2	6
6	3	15
7	-	-

Some frames may not be used yet → no page ID recorded

Records **which page of which process** is loaded **in which frame**

Advantage:

- Memory consumption** is constant and independent from the number of processes or their size

Disadvantage:

- Time consuming to find the location of a page** in main memory
(requires **up to as many look-ups as there are frames** in the main memory)

Consider mapping the virtual memory of a process addressing 1GiB onto a physical memory organized in 256 page frames of 8KiB.

What is the physical memory size?

The virtual memory is made of how many pages?
(assume all virtual memory addresses are consecutive)

How many entries is there in the page table?

What is the size (in bytes) of the page table?

On how many pages is the page table encoded?

Page table: records the **frame used by each page** of the process

Exercise

Consider mapping a virtual memory of a process addressing 1GiB onto a physical memory organized in 256 page frames of 8KiB.

What is the physical memory size?

256 frames of 8KiB \rightarrow physical memory is $256 \cdot 8\text{KiB} = 2\text{MiB}$

The virtual memory is made of how many pages? (assume all virtual memory addresses are consecutive)

1GiB of virtual memory divided in pages of 8KiB $\rightarrow 1\text{GiB}/8\text{KiB} = 2^{17}$ pages

How many entries is there in the page table?

As many as the number of pages $\rightarrow 2^{17}$ entries

What is the size (in bytes) of the page table?

Each entry in the page table encodes the page frame in which the corresponding page is loaded. Since we have 256 page frames, we need 8 bits to identify one page frame \rightarrow each entry in the page table contains 8 bits \rightarrow the page table is $2^{17} \text{ B} = 128\text{KiB}$

On how many pages is the page table encoded?

$128\text{KiB} / 8\text{KiB} = 16$ pages

Consider every process addresses a virtual memory of 1GiB (per process). Consider a physical memory organized in 256 page frames of 8KiB. Assume that there are a maximum of 128 processes in the system.

How many entries is there in the frame table (inverted page table)?

What is the size (in bytes) of the frame table?

On how many pages is the frame table encoded?

Frame table: records **which page of which process** is loaded **in which frame**

Exercise

Consider every process addresses a virtual memory of 1GiB (per process). Consider a physical memory organized in 256 page frames of 8KiB. Assume that there are a maximum of 128 processes in the system.

Physical memory size = 2MiB

Virtual memory of each process is organized in 2^{17} pages

How many entries is there in the frame table?

As many as the number of page frames → 256

What is the size (in bytes) of the frame table?


Each entry in the frame table encodes which page from which process is loaded in the corresponding page frame. Since we have 128 processes, we need 7 bits to identify the process. Since we have 2^{17} pages in each process, we need 17 bits to identify the page loaded in the frame → each entry needs $7+17=24\text{bits}=3\text{ bytes}$.
→ the frame table is $256*3\text{bytes} = 768\text{B}$

On how many pages is the page table encoded?


Since a page has a size of 8KiB, we need a single page to encode the frame table

Problem 2: address binding


Warehouse (secondary storage)




Office (primary storage)



Say I have a ledger of where boxes are stored (page or frame table).
How can I find file 56 in box 23 of pallet 48 in the office?



physical address of the
file in the office




Problem 2: address binding

Solution

- Use **virtual addresses** in the program code
- **Translate** virtual addresses in **physical addresses** during the execution of the code

Virtual address

- Can refer to $2^{|p|}$ pages
- Page size is $2^{|w|}$ words/page
- **Virtual memory size: $2^{|p|+|w|}$ words**




Physical address

- Can refer to $2^{|f|}$ frames
- Frame size is $2^{|w|}$ words/frame
- **Physical memory size: $2^{|f|+|w|}$ words**

Virtual memory size seen by each process can be **larger or smaller than** the physical memory size


Frame number in which the page we want to access is loaded



Same as the offset in the virtual address

Problem 2: address binding using page tables


Virtual address:



Pointer to the page table

PT

Page table of the running process



```
address_translation( $p, w$ ) {  
    return  $pa = *(PT+p) + w$ ;  
}
```

PT = pointer to the page table

Frame in which the page is loaded


Word that is being accessed

Assuming hardware support: a **Page Table Base Register (PTBR)** is used to hold the value of PT for fast access to the page table




Reading/writing from/to memory requires **two memory accesses**: one for **accessing the page table**, and one for **accessing actual data**

Problem 2: address binding using the frame table

Virtual address:



Frame table
(inverted page
table)



```

address_translation( $id, p, w$ ) {
     $pa = UNDEFINED$ ;
    for ( $f = 0; f < F; f++$ ) {
        if ( $FT[f].pid == id \ \&\& \ FT[f].page == p$ )
             $pa = f + w$ ;
    }
    return  $pa$ ;
}
    
```

FT = pointer to
the frame table

Frame in which the
page is loaded

Word that is
being accessed

Reading/writing from/to memory requires up to **as many memory accesses as there are entries in the frame table**

Exercise

Consider a **page table** with 8 entries as shown below. Each page has a size of 256 bytes. The smallest addressable unit is a byte.

	Page table
0	0x5AC0
1	-
2	-
3	0xFFFF
4	0x0012
5	-
6	0x1234
7	-

What is the size of the virtual memory of the process?

What is the size of the physical memory?

On how many bits is the offset of a word encoded?

What is the physical address for the four following virtual addresses?

0x056

0x434

0x500

0x912

Page table: records the page **frame used by**
each page of the process

Exercise

Consider a **page table** with 8 entries as shown below. Each page has a size of 256 bytes. The smallest addressable unit is a byte.

Page table	
0	0x5AC0
1	-
2	-
3	0xFFFF
4	0x0012
5	-
6	0x1234
7	-

What is the size of the virtual memory of the process?

8 pages of 256B $\rightarrow 8 \cdot 256B = 2KiB$

What is the size of the physical memory?

We see that there is a page frame with ID 0xFFFF

\rightarrow There are at least 2^{16} frames, each of size 256B

\rightarrow The physical memory is **at least** $2^{16} \cdot 256B = 2^{24}B = 16MiB$

On how many bits is the offset of a word encoded?

Pages are 256B and we address bytes \rightarrow 256 addresses \rightarrow we use 8 bits for the address of a word in a page

What is the physical address for the four following virtual addresses?

0x056 Physical address: 0x5AC056

0x434 Physical address: 0x001234

0x500 Undefined. The page must first be loaded in main memory

0x912 Error. Virtual address is outside the address space of the process

Exercise

Consider a **frame table** with 8 entries as shown below. Each page has a size of 4KiB. The smallest addressable unit is a byte.

	Process ID	Page ID
0	5	0x90
1	-	
2	2	0x22
3	2	0x21
4	5	0x80
5	-	
6	1	0x00
7	2	0x04

What is the size of the physical memory?

On how many bits is the offset of a word encoded?

What is the physical address for the four following virtual addresses **of process 2**?

0x21650

0x80123


0x4341

0x2022

Frame table: records **which page of which process** is loaded **in which frame**

Problem 2: address binding (acceleration)

- A **Translation Look-aside Buffer (TLB)** is a small cache memory in the processor that **keeps track of the locations of the most recently used pages** in main memory
- It **accelerates address translation** and thus memory accesses
- It **does not contain data or instructions**, only the frame id in which the most recently accessed pages are loaded



Source: Fig 9.12, Silberschatz, Galvin, Gagne: Operating System Concepts, 10th Edition, Global edition

- **Whenever a page is accessed**
 - First, check if it is in the TLB
 - if the page location **cannot be found in the TLB**, use the **page table**
 - if the page location still cannot be found a **page fault** is generated (i.e., the page is not in physical memory)

Problem 3: How to decide which page to load and when?

Use **demand paging**:

- Bring a **page** into main memory **only when it is accessed for the first time**
 - No need to have the entire process stored in memory

Takes advantage of the fact that **not all pages are necessary at once**

Examples:

User-written **error handling code**

Mutually exclusive modules/segments of code

Only **fractions of large tables** are actually used

More problems:

- **For how long** should a page stay in main memory
- **How many pages** of each process should be kept in main memory
- What to do when there is no more free frames?

See next lecture.

- The **memory management subsystem** **decides what to load, where and when in main memory**
- **Early systems** required to **load entire processes in contiguous memory regions**
 - Slow
 - Limited concurrency level
 - Fragmentation (internal and external)
- **Memory paging** **eliminates external fragmentation** and **strongly limits internal fragmentation**
- **Page or Frame tables** allow to **dynamically bind virtual addresses** in the process address space **to physical addresses** in main memory

- **Load and replacement strategies** for virtual memory
- A new homework will be released during the weekend, deadline after the vacations