

2INC0 - Operating Systems

I/O management

Geoffrey Nelissen



Interconnected
Resource-aware
Intelligent Systems

TU/e

Technische Universiteit
Eindhoven
University of Technology

Where innovation starts

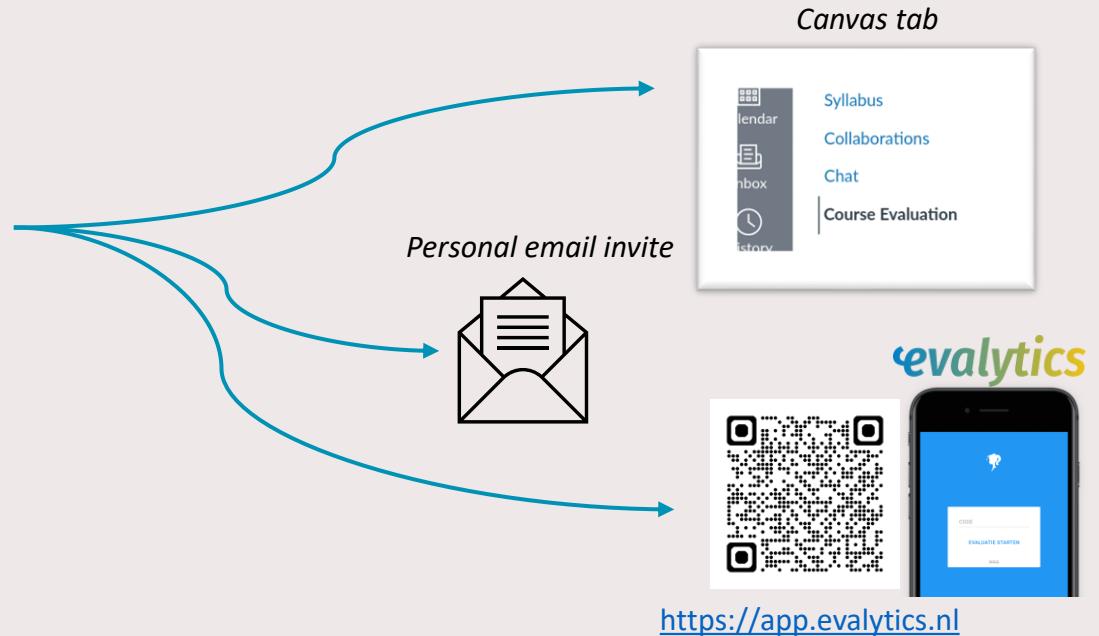
Announcement:

The course evaluation surveys are open

Where can you give feedback?

Feedback on the final exam/assignment

The course evaluation surveys will be closed one day before the final exam; if you want to give feedback about the final exam/assignment, you can use [this link](#) or QR code.



Questions about anything regarding quality assurance?
Mail us at mcs.quality.assurance@tue.nl

You can use these tips to provide more impactful feedback



Be **specific** and **focused** on your feedback. Use examples and **suggestions** to avoid vague statements.



Always be **respectful** when giving feedback.



Give positive feedback as well as areas for improvement. Stay **solution-oriented**.



When giving feedback, focus on the **1 or 2** most important points that apply to you.

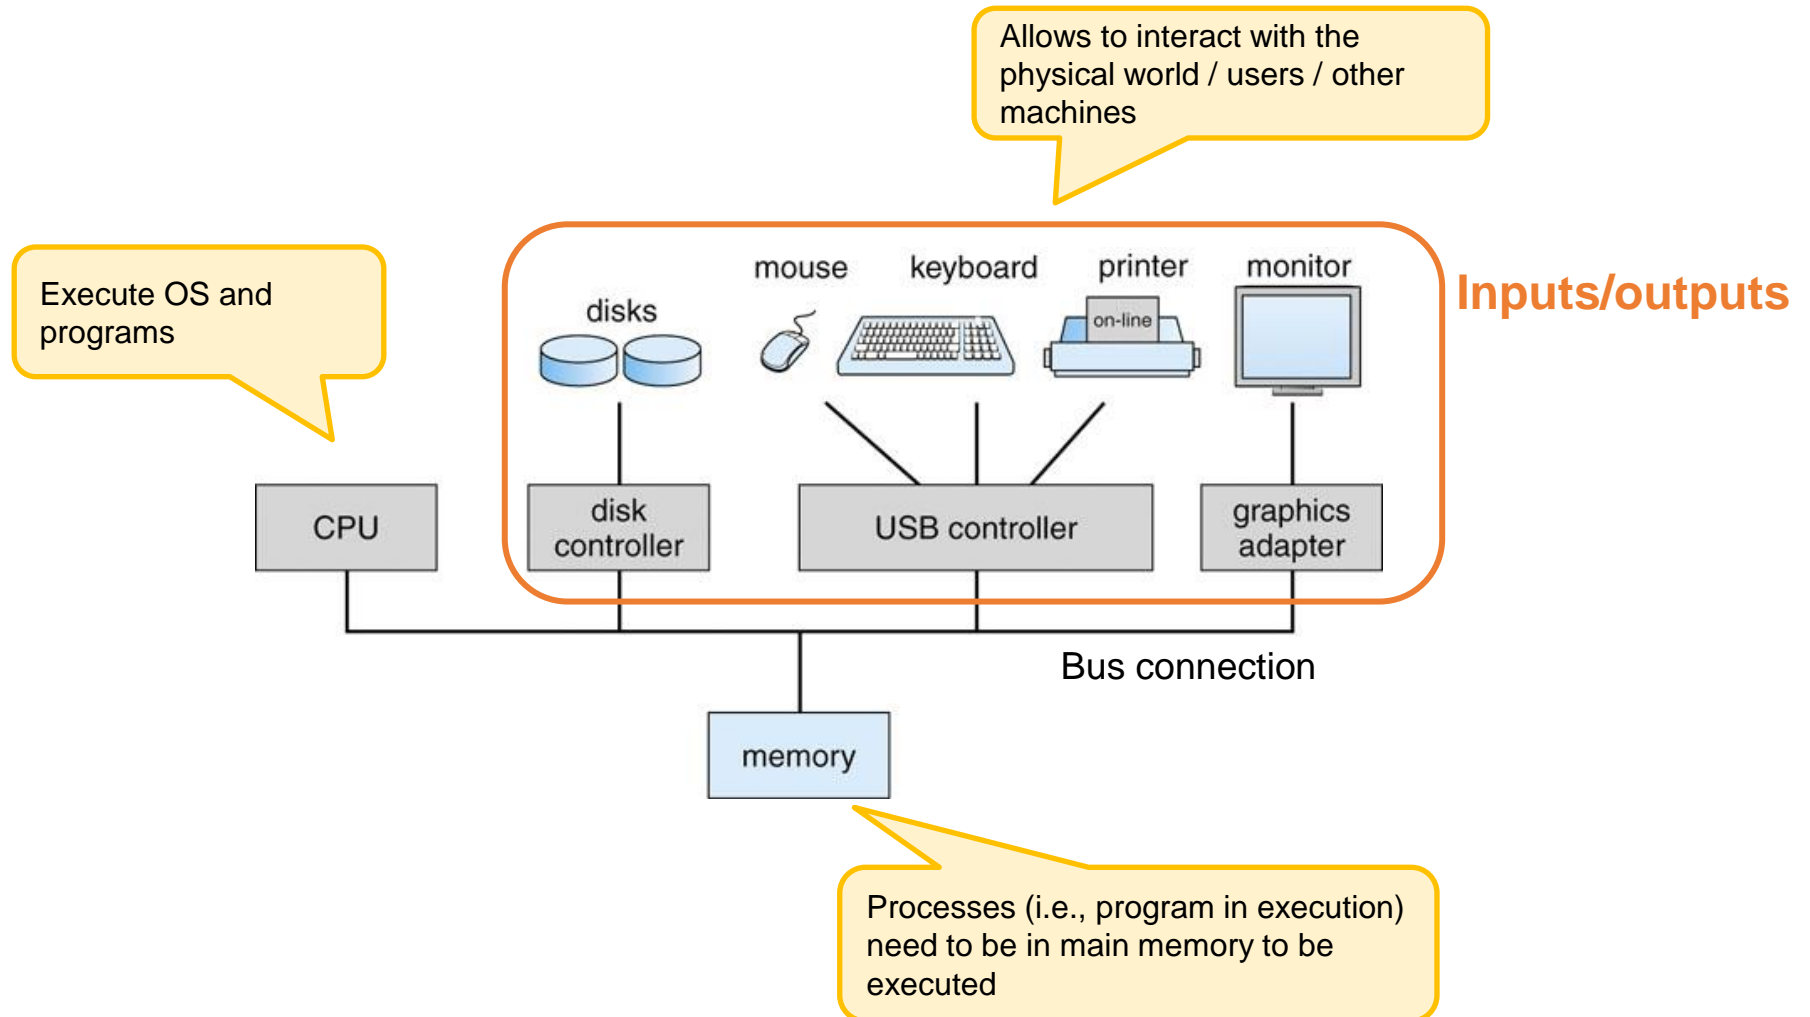
What happens with the results of the surveys?



The responsible teacher reads the **anonymous results** and reflects on their course.



Results and reflection are discussed during committee meetings to **improve the courses and programs!**



Purpose of an operating system

Structure of running programs.
How to **implement concurrency**.
How to decide **what to execute and when**.

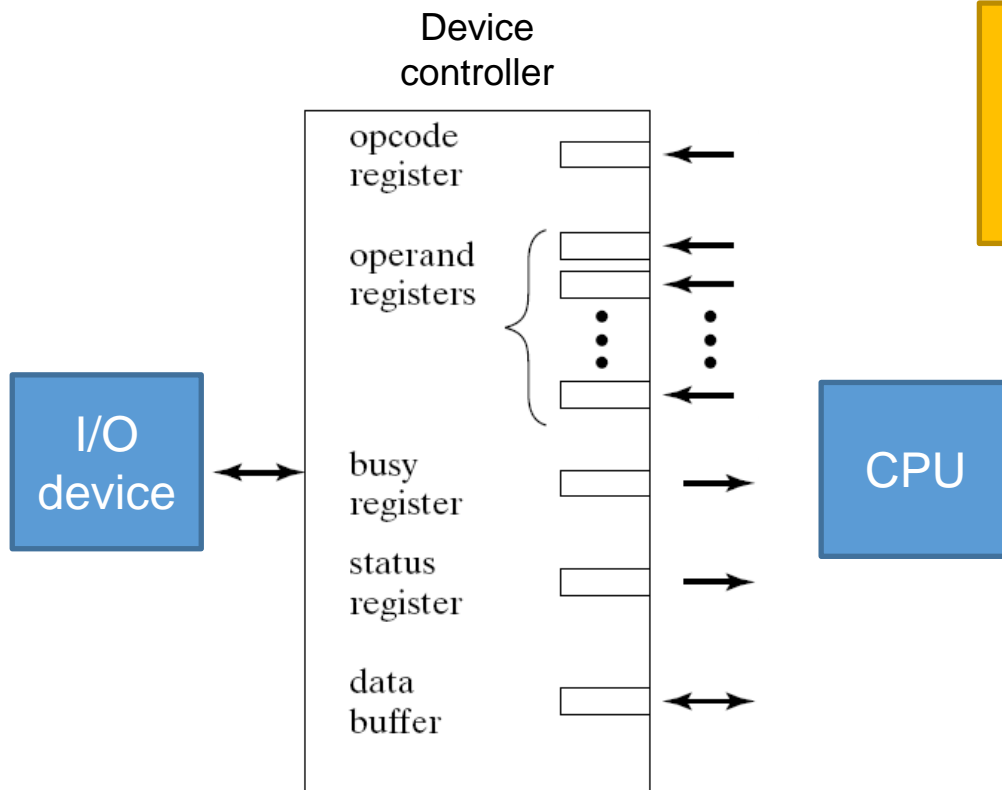
- **Introduction to operating systems** (lecture 1)
- **Processes, threads and scheduling** (lectures 2+3)
- **Concurrency and synchronization**
 - atomicity and interference (lecture 4)
 - actions synchronization (lecture 5)
 - condition synchronization (lecture 6)
 - deadlock (lecture 7)
- **File systems** (lecture 8)
- **Memory management** (lectures 9+10)
- **Input/output** (lecture 12)

Problems associated to **concurrent executions**.
How to prove program properties (**topology invariants, traces**)
How to protect **critical sections**.
How to **synchronize the execution of** programs to enforce new properties
How to analyze **deadlocks**, prevent them and detect them.

How files are **organized (virtually)**
How can they be **accessed (physically on hard drive)**

How to **efficiently load** processes in main memory.
How to efficiently manage **limited physical memory space**.
How to **share memory space** between concurrent processes.

- **I/O device controllers**
- **I/O subsystem**
- **I/O buffering**
- **Disk scheduling**



An I/O device **communicates** with the CPU through **hardware registers** available in a **device controller**

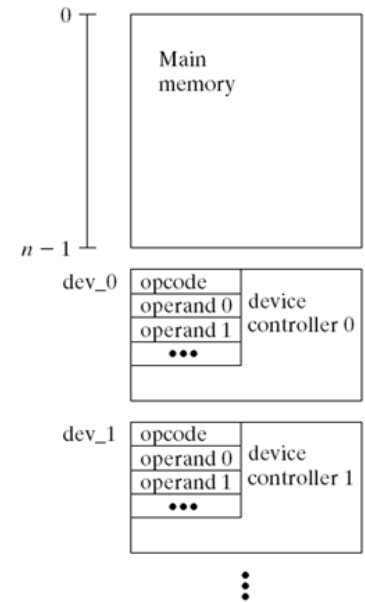
- Opcode register: code of the operation to be performed, e.g. read, write
- Operand registers: parameters associated with the operation
 - e.g. addresses to be read from or written at, DMA parameters, etc.
- Busy and status registers: provide information about availability, readiness, errors
- Data buffer registers: bytes transferred to or from the I/O device
 - e.g. value typed in on the keyboard, or text to be printed

Two methods for accessing the device controller

Method 1

The **CPU instruction set is extended** with special I/O instructions

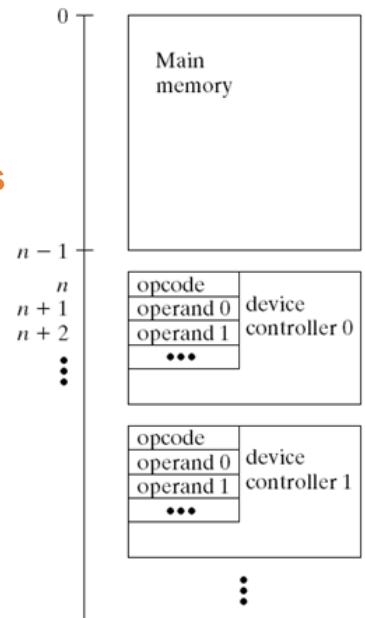
- **Example** of assembly instructions:
`io_store cpu_reg, dev_no, dev_reg`
`io_load dev_no, dev_reg, cpu_reg`
- **No physical address** associated to each register
 - ➔ **Advantage:** **no interference with virtual memory**
 - ➔ **Disadvantage:** **not possible to map device in user space**
 - ➔ user process cannot directly access the device as a normal data structure.



Method 2

Physical address space is extended to directly refer to device registers

- Each device controller **register is given a physical address**
 - ➔ **Advantage:** the **device may be mapped to the user space** (i.e., we associate a virtual address to the physical address of each I/O register).
 - ➔ **Disadvantage:** **complexify virtual memory management**



Method 1

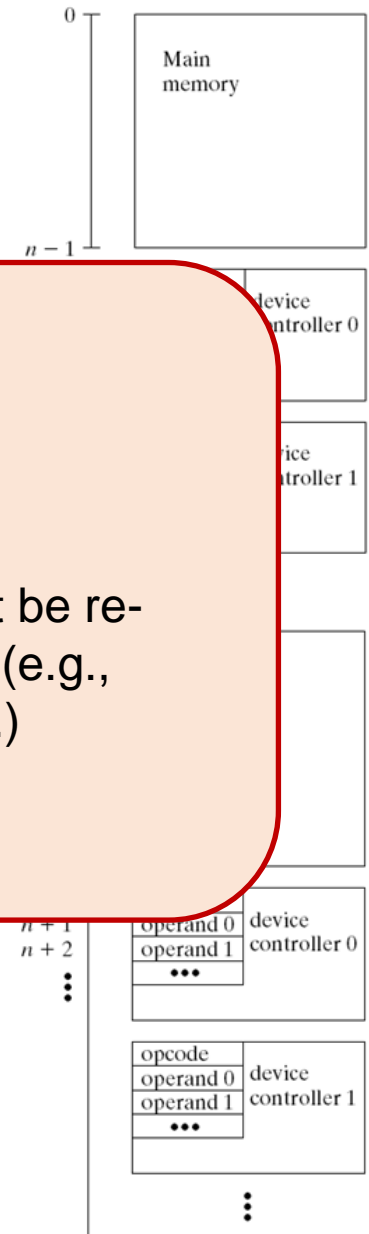
The **CPU instruction set is extended** with special I/O instructions

- **Example** of assembly instructions:
`io_store cpu_reg, dev_no, dev_reg`
`io_load dev_no, dev_reg, cpu_reg`

Issue:

- **Each I/O device controller** may have a **different** set of registers/opcodes/operands
 - Code is written for a specific I/O controller and must be re-written if we change the type or brand of I/O device (e.g., move from HDD to SSD, change brand of SSD, etc.)
 - **Limits portability**
 - **Increases work, risk of bugs, etc.**

→ Disadvantage: **complexify virtual memory management**



- I/O device controllers
- I/O subsystem
- I/O buffering
- Disk scheduling

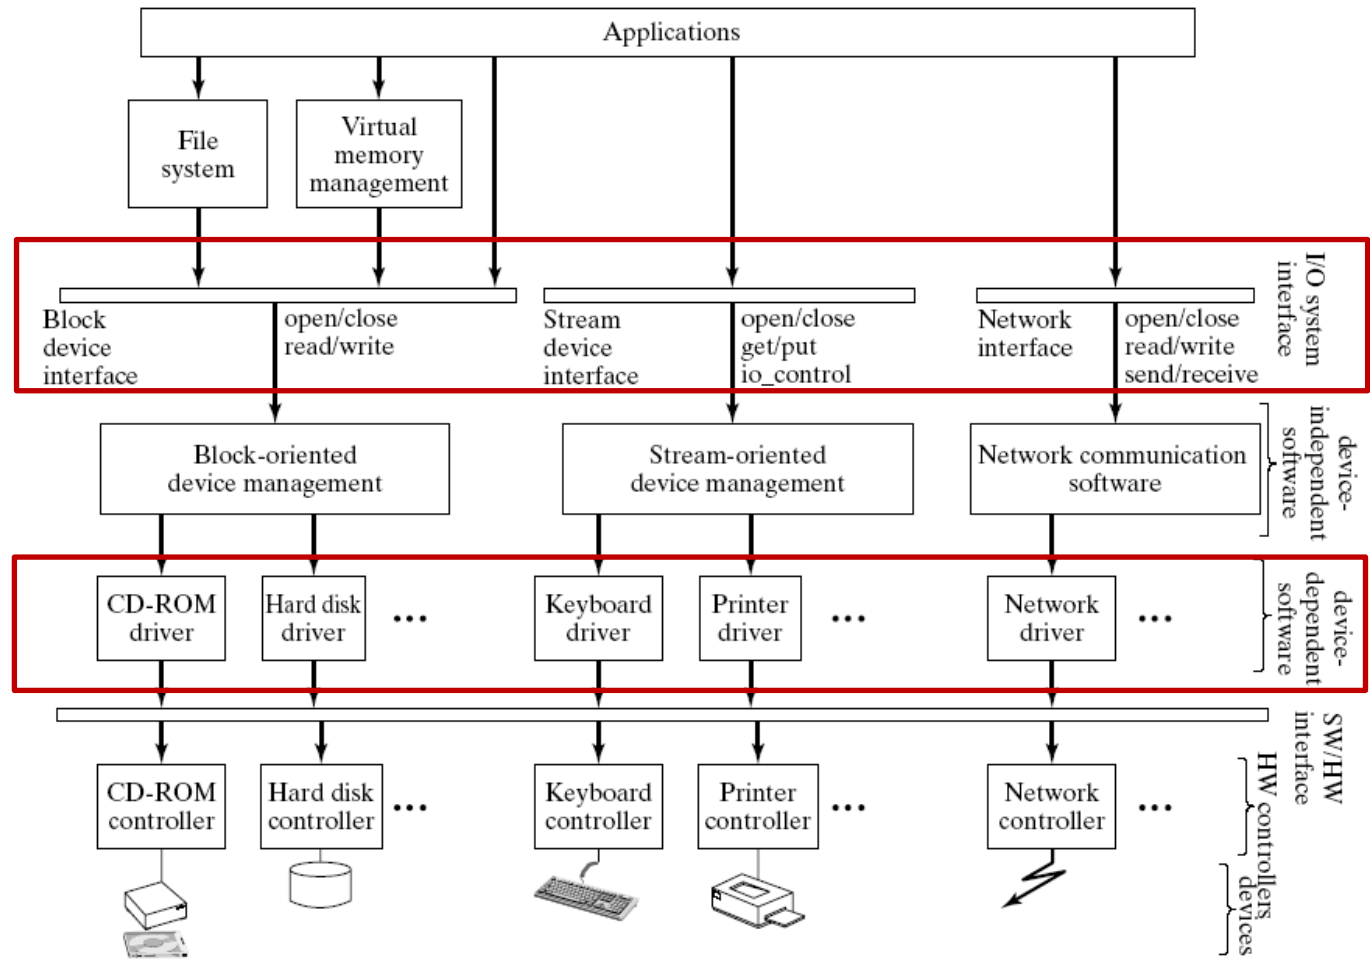
- Part of the OS that manages input and output devices
- **Goals:**
 - to present a **logical/abstract view** of I/O devices
 - to **facilitate sharing** of devices
 - to provide **efficiency** and **optimize performance**
 - Examples:
 - ensures the CPU and multiple I/O devices run in parallel
 - re-ordering I/O requests to improve throughput
 - buffering to hide latency
 - ...

Design concerns:

- large **variety** of I/O device types
 - keyboard, display, printer, disk, temperature sensor, network cards
- **different speeds** and **(brand-)specific approaches**
- ensure we can **add new devices** after the OS development and installation

High level I/O abstraction
(for the user)

Low level I/O abstraction
(for the OS)



Reference: Bic, Lubomir, and Alan C. Shaw. *Operating systems principles*. Prentice Hall, 2003.

User level abstraction:
standardized APIs defined
per device class

**High level I/O
abstraction**

**Provide generic code for handling
classes/families of I/O devices**

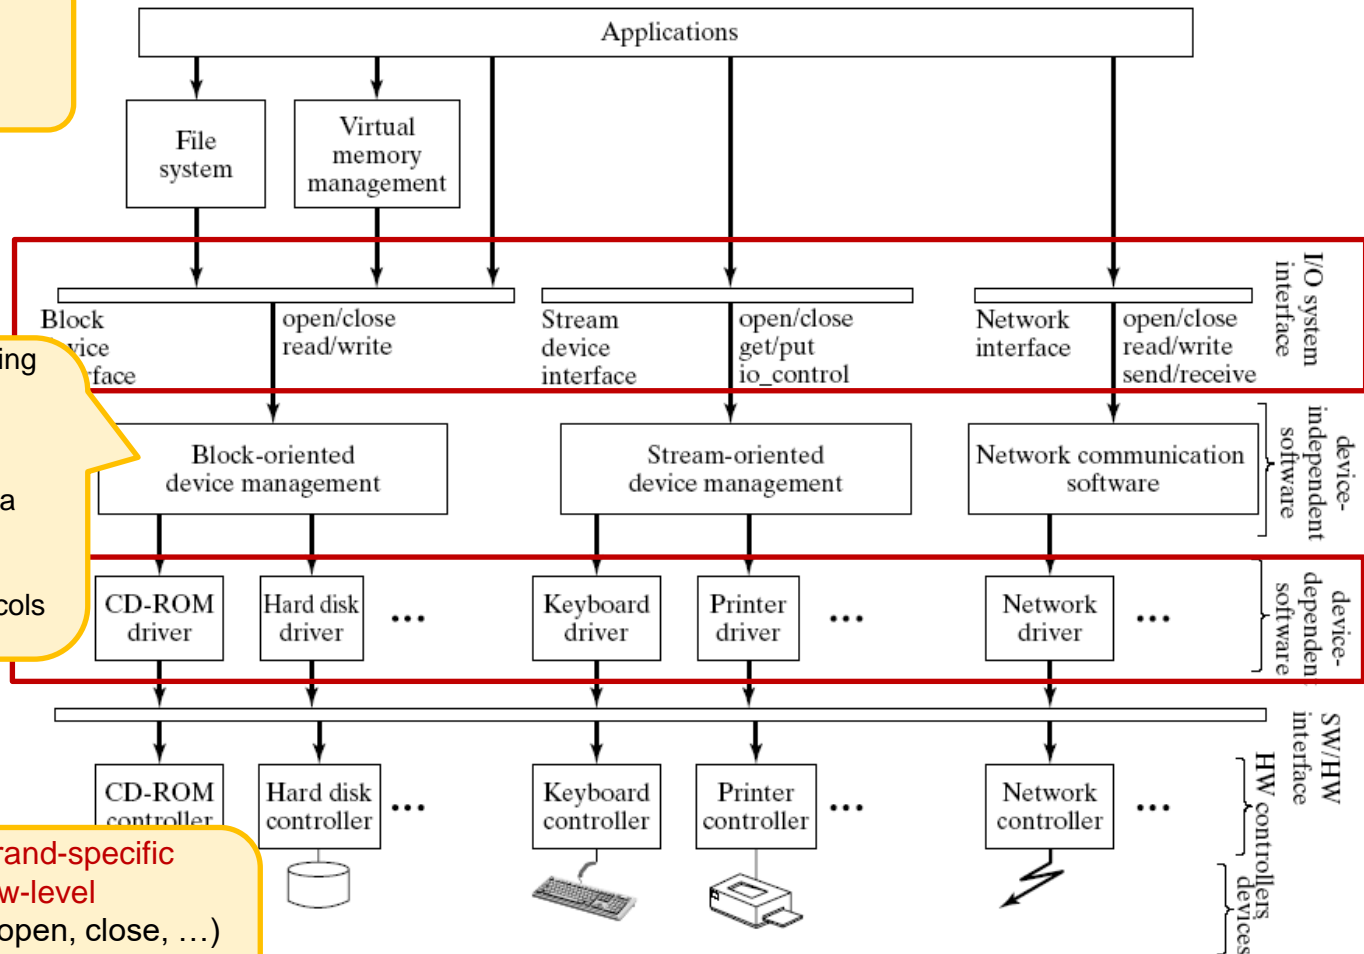
Examples:

- Sequence of high-level operations to read or write data on disk/cd-rom/...
- implementation of specific network communication protocols (e.g., TCP)

**Low level I/O
abstraction**

OS level: encapsulation of brand-specific
issues behind well-defined low-level
operations (e.g., read, write, open, close, ...)

For each specific device, a **device driver**
implements those low-level operations



Source: Bick, Lubomir, and Alan C. Shaw. *Operating systems principles*. Prentice Hall, 2003.

- I/O device controllers er
- I/O subsystem
 - High level abstraction
 - Low-level abstraction
- I/O buffering
- Disk scheduling


- **Block-oriented** or storage devices

- Read and write blocks of **data in arbitrary order**

i.e. readers/writers model

→ Each block of data **has an address that can be accessed at any time**

- **Sequence of actions that must be performed to access data depends on** the operation (e.g., read/write/copy), **data location(s)**, and **last** (set of) **data accessed before**
- examples: hard drive, ssd, CD reader, magnetic tape




- **Stream-oriented** devices

- input and output **data in a sequential way**



i.e. producer/consumer model

- **Optimized for single byte access**
- examples: keyboard, sensors, actuators, mouse



- **Network communication** devices

- send and receive packets on a network (socket) interface
- Provide connection and communication protocols

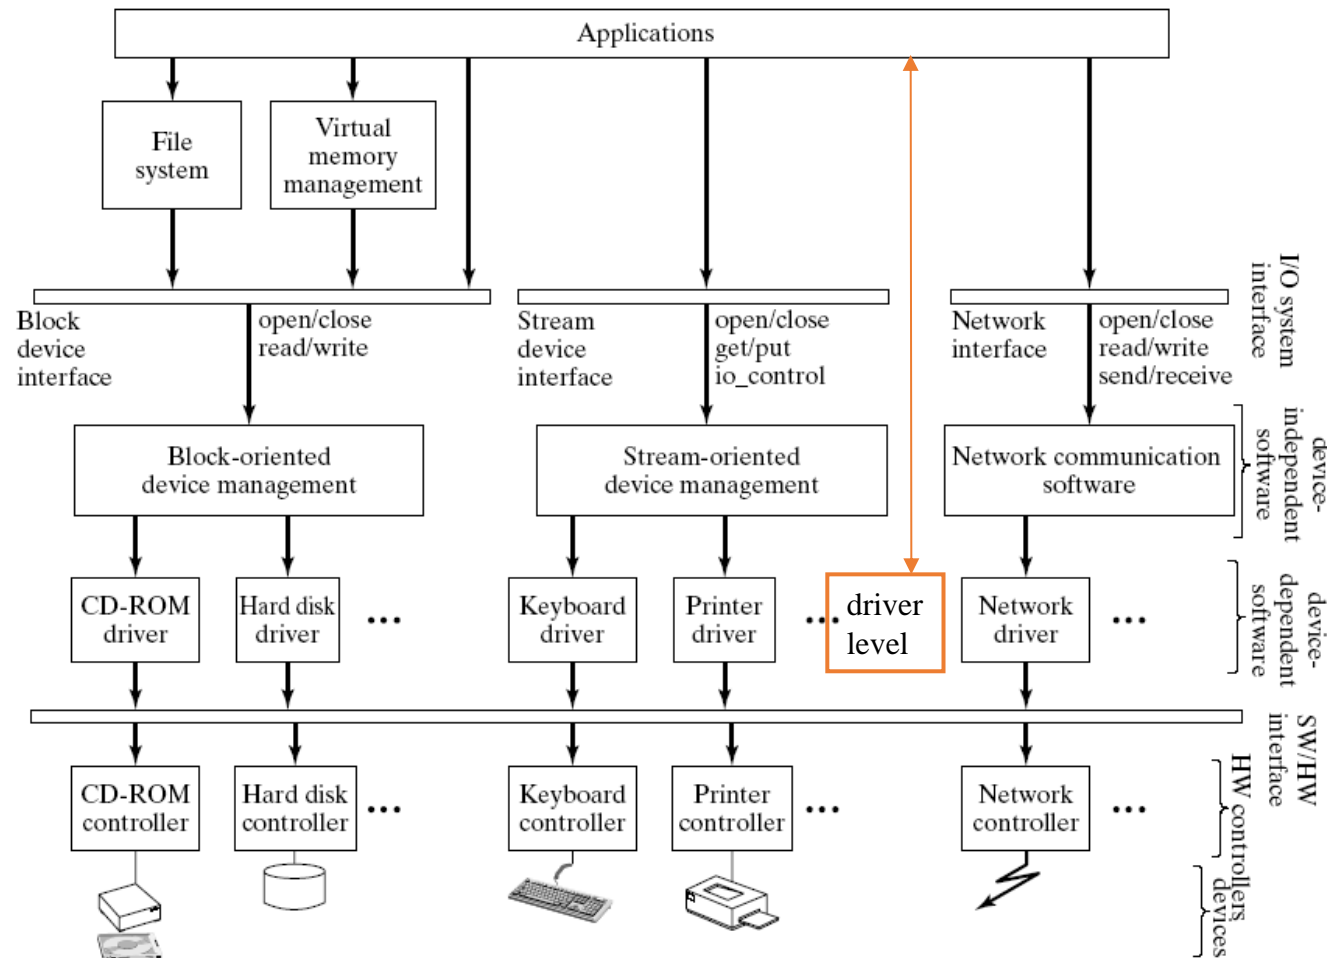


Note: the above ‘**classical classification**’ **may not be enough**

- **Example:** accessing graphics hardware not really covered by the standard interface: read() / write()
→ very low performance with classical I/O system interfaces

Two main solutions when I/O devices do not fit in any class

- Solution 1:** let the user application directly access the driver




Two main solutions when I/O devices do not fit in any class

- **Solution 1:** let the user application directly access the driver
 - **Drawbacks:**
 - Application code becomes *device specific*
 - Application code *operates at a low level of abstraction*
- **Solution 2:** extend the capabilities of the OS with **new APIs via external libraries** to
 - provide an *abstraction* from vendor-specific issues
 - support **domain-specific concepts** needed for device-independent application development

Examples: audio sink, filter operations, graphics scene, ...

 - provide *optimized performance*
 - **Examples:** Microsoft DirectX or OpenGL for graphic or CUDA for GPGPU management



- I/O device controllers
- I/O subsystem
 - High level abstraction
 - Low-level abstraction
- I/O buffering
- Disk scheduling

High level I/O abstraction

Provide generic code for handling classes/families of I/O devices

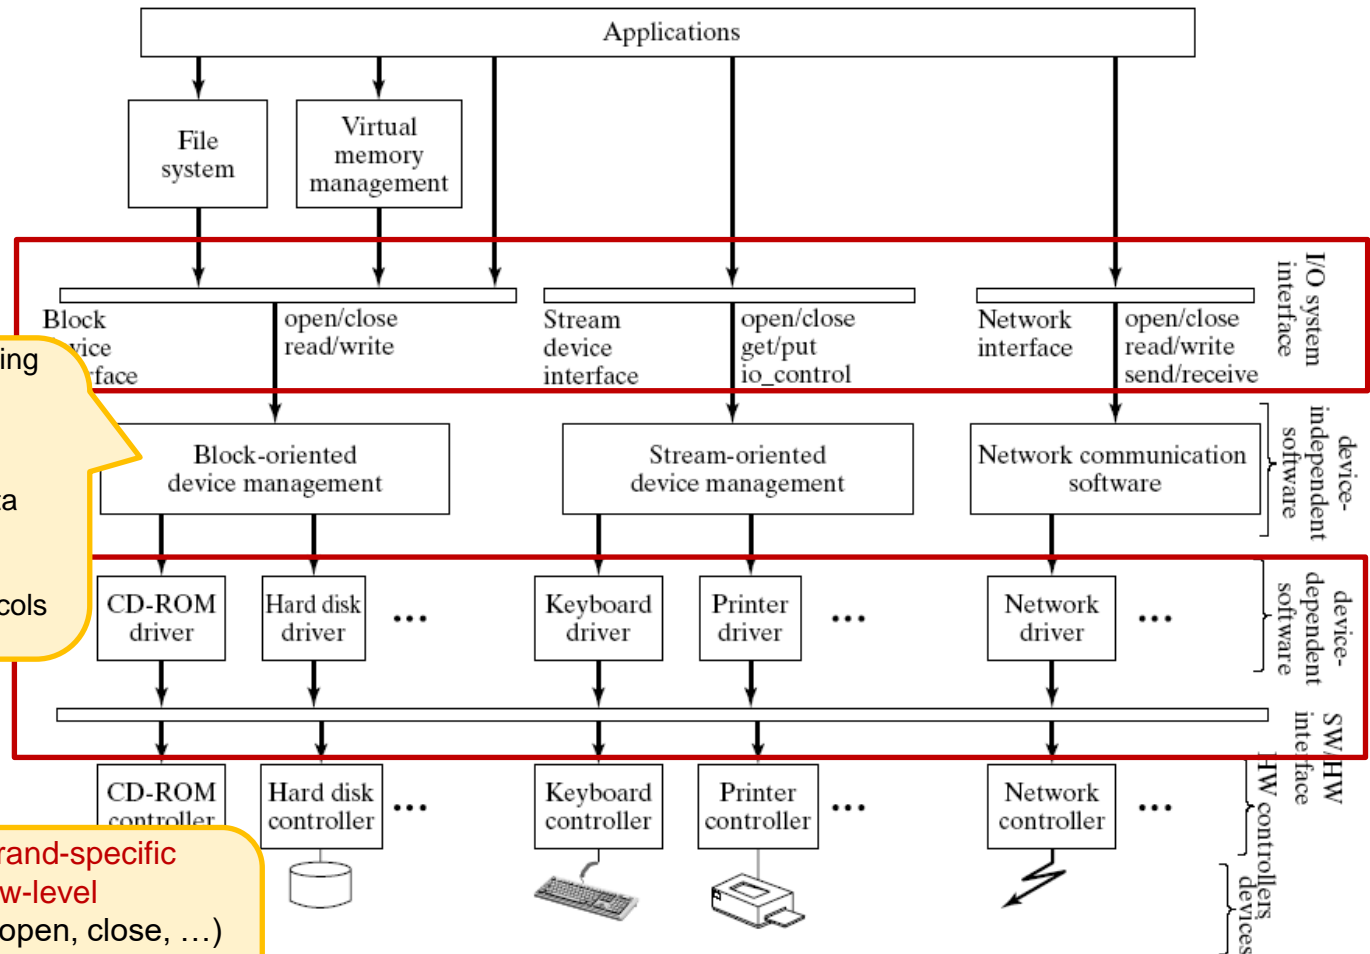
Examples:

- Sequence of high-level operations to read or write data on disk/cd-rom/...
- implementation of specific network communication protocols (e.g., TCP)

Low level I/O abstraction

OS level: encapsulation of brand-specific issues behind well-defined low-level operations (e.g., read, write, open, close, ...)

For each specific device, a **device driver** implements those low-level operations



Source: Bick, Lubomir, and Alan C. Shaw. *Operating systems principles*. Prentice Hall, 2003.

- The device driver implements a **collection of standard operations** (functions)
 - in Linux: the functions (say, for device xxx) are registered in a data structure

```
struct file_operations xxx_fops = {  
    NULL,          /* lseek()    */  
    xxx_read,      /* read()    */  
    xxx_write,     /* write()   */  
    NULL,          /* readdir() */  
    NULL,          /* select()  */  
    xxx_ioctl,     /* ioctl()   */  
    NULL,          /* mmap()    */  
    xxx_open,      /* open()    */  
    xxx_close      /* close()   */  
};
```

- *read*: function to read data
- *write*: function to write data
- *lseek*: move the read/write pointer
- *ioctl*: i/o control to modify device/driver parameters
- *select*: notify when the i/o device is ready to perform a specific operation
- ...

- **Polling**

- Driver initiates I/O operations in the device controller and observes completion (i.e., driver busy-waits) or periodically wakes-up to check completion
- The driver polls the device controller repeatedly and tests
- **Wastes CPU time**

When is it acceptable to busy-wait?

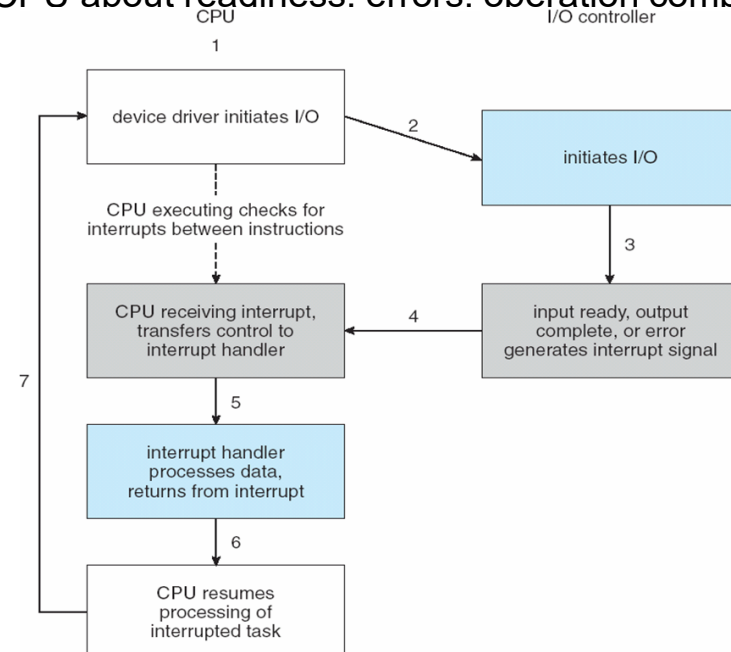
- Only **when I/O operations are** known to be **fast in comparison to** the overhead of **context switches**

- **Polling**

- Driver initiates I/O operations in the device controller and observes completion (i.e., driver busy-waits) or periodically wakes-up to check completion
- The driver polls the device controller repeatedly and tests
- **Wastes CPU time**

- **Interrupts**

- Driver initiates I/O operations in the device controller and then **yields the CPU**
- Device controller uses interrupts to inform the CPU about readiness, errors, operation completion, ...



- **Polling**

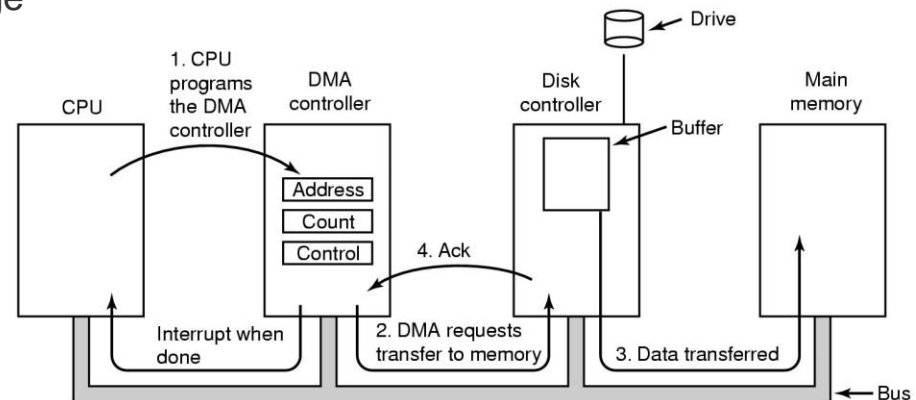
- Driver initiates I/O operations in the device controller and observes completion (i.e., driver busy-waits)
- The driver polls the device controller repeatedly and tests
- **Wastes CPU time**

- **Interrupts**

- Driver initiates I/O operations in the device controller and then **yields the CPU**
- Device controller uses interrupts to inform the CPU about readiness, errors, operation completion, ...

- **DMA**

- DMA can be implemented inside the device controller or as a separate hardware component
- After initialization, the DMA independently **moves groups of data** between the device controller and memory
- **less overheads for the CPU** only one interrupt handling when the whole transfer is completed instead of one interrupt per word/message



Reference: Tanenbaum, Andrew. Modern operating systems. Pearson Education, Inc., 2009.

- I/O device controllers
- I/O subsystem
 - High level abstraction
 - Low-level abstraction
- I/O buffering
- Disk scheduling


- As an example, consider a process reading or writing in a file located on a hard drive
 - It issues the command
 - then the process is either waiting or suspended until an interrupt
- **Potential problems**
 - **Speed / latency mismatch:**
Process must wait for the relatively slow I/O to complete before it can send new data to write on the disk
 - **Data granularity mismatch** (byte vs line vs block):
Application may expect to receive data in smaller pieces than a block (and vice versa)
 - **Conflict with the swapping decisions made by the OS:**
Pages containing the virtual address range must remain in physical memory until I/O completes (otherwise, the driver/DMA may write the data at the wrong physical address or corrupt the address space of another process)

Buffer

dedicated (kernel) memory space or hardware registers that holds data of a producer until its consumer is ready to consume

Main task: **decouple producer and consumer**

- resolve **speed difference** and latency problems
- resolve **granularity differences**
 - driver returns characters, application needs lines
 - driver returns blocks, application needs bytes
- resolve **swapping issues**
 - The buffer remains permanently in physical memory even when the process is swapped out
- **improve efficiency**
 - Perform input **transfers before request**: try to **predict what will be needed in the future**
 - **Delay outputs** on purpose: wait for the right time to perform output (e.g., more data can be transferred at once, or optimize seek time (**see disk scheduling**))
 - **Caching** data



picture from Stallings,
Operating Systems – Internals and Design Principles

Buffering inputs:

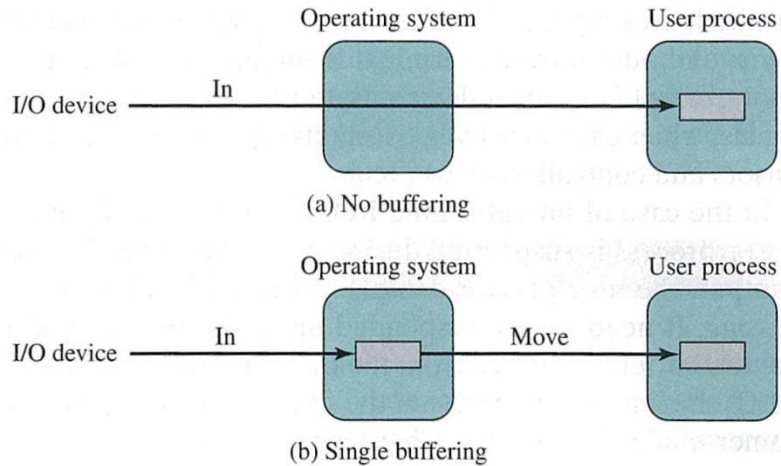
The data is written into the buffer first (e.g., from disk to kernel memory address space).

Then, the data is moved into the user process address space.

Buffering outputs:

The data is moved to the buffer first (e.g., from user space to kernel space).

Then, data is output (e.g., written to a disk) directly from the buffer.



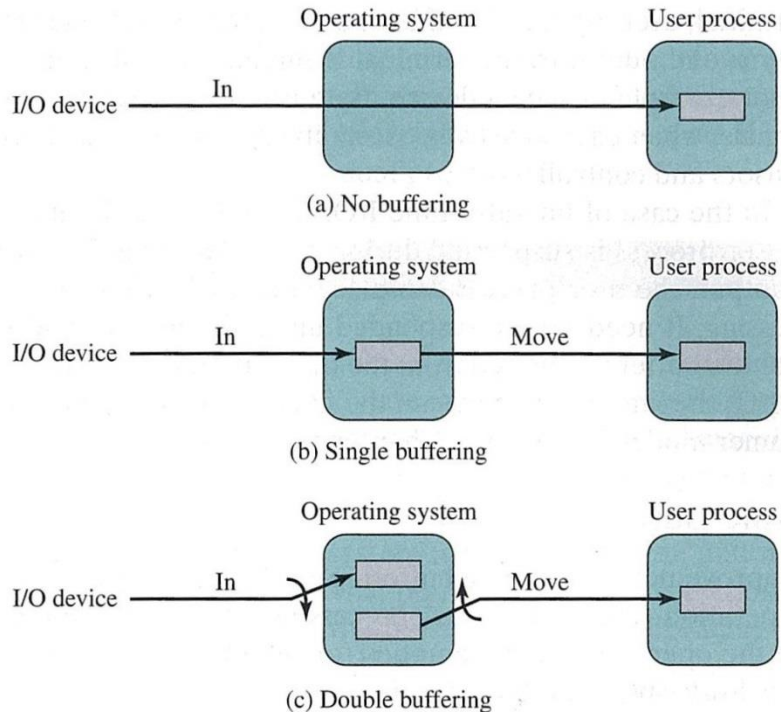
picture from Stallings,
Operating Systems – Internals and Design Principles

- **Single buffering**





- **enables asynchronous transfer**

- The input device controller can **produce data without waiting** for the process to be ready to consume it.
→ releases I/O device to perform another operation
 - The process can **buffer output data even if the device controller is not yet ready** to transfer it.
→ allows process to continue doing something else

- **allows swapping the process out of main memory**



- **Double buffering** (buffer swapping)
 - reduces idle time
 - OS can move the content of one buffer from kernel to user space while the device controller is filling the other buffer
 - still poor in handling bursts of data



- **Double buffering** (buffer swapping)
 - reduces idle time
 - OS can move the content of one buffer from kernel to user space while the device controller is filling the other buffer
 - still poor in handling bursts of data
- **Circular buffering**
 - to handle bursts of data

picture from Stallings,
Operating Systems – Internals and Design Principles

Throughput

- How many data can be transferred per second?
(we limit ourselves to analyzing the input scenario)

T : time to **transfer** one data in main memory

M : time for **moving** one data from kernel address space to user address space

C : time the process need to operate on one data (**computation time** on received data)

D : Minimum time until the next data may become available in user space

- **Throughput** = $1/D$


- **no buffer**: $D = T + C$

(because $M=0$ in this case, and we cannot transfer a new data until the previous data is processed)

- **single buffer**: $D = \max(C, T) + M$

(because we can move a data between kernel and user space only when a new data is transferred in the buffer and the previous data has been processed, but the computation may happen in parallel to the data transfer in the buffer)

Because T and C can happen in parallel



T : time to **transfer** one data in main memory

M : time for **moving** one data from kernel address space to user address space

C : time the process need to operate on one data (**computation time** on received data)

D : total time until a data is available in user space

- **Throughput** = $1/D$

- **no buffer**: $D = T + C$


(because $M=0$ in this case, and we cannot transfer a new data until the previous data is processed)

- **single buffer**: $D = \max(C, T) + M$

(because we can move a data between kernel and user space only when a new data is transferred in the buffer and the previous data has been processed, but the computation may happen in parallel to the data transfer in the buffer)



- **double buffer**: $D = \max(M+C, T)$

(because the move cannot start before the computation completes, but the move and computation happen in parallel to the data transfer)



- I/O device controllers
- I/O subsystem
 - High level abstraction
 - Low-level abstraction
- I/O buffering
- **Disk scheduling**

- Several **platters**, divided in **tracks**, divided in **sectors**
- **Cylinder**: set of tracks that are at the same arm position (same radius).



- Average time to read/write one block on the disk

$$T_{read} = T_{seek\ cylinder} + T_{move\ to\ sector} + T_{read\ block}$$



$$= T_{seek\ cylinder} + \frac{1}{2r} + \frac{B}{rN}$$

r = rotations per second

N = #bytes per track

B = #bytes per block/sector

- Reading/writing a block of data in a sector:
 - **head-positioning time on correct cylinder (= seek time)**
 - On average: 9-12 msec
 - ~1-2 msec for neighbouring tracks
 - **read/write within a track** (depending on rpm): ~3-5ms





→ The OS must optimize the disk head movement

- Requests for blocks are buffered.
- The block requests in the buffer are treated according to a scheduling policy:
 - **FIFO**
 - treats requests in the same **order they are submitted**
 - **Shortest Seek Time First**
 - treats the buffered request that requires the **smallest head movement**
→ possible starvation, unpredictable
 - **Elevator Scan**
 - completes **full swing** of the head in either direction
- Many other versions investigated in the literature

Comparison on an example

Assume we just moved from cylinder 0 to cylinder 5, and that the requests in the buffer are requesting accesses to cylinders 12, 4 and 7.




Notes:


- SSTF may lead to **starvation**

Comparison on an example


Assume we just moved from cylinder 0 to cylinder 5, and that the requests in the buffer are requesting accesses to cylinders 12, 4 and 7.



(a)



(b)



- **The I/O subsystem adopts a layered approach** based on **two levels of abstraction**
 - User interface
 - Driver interface
- **I/O buffering** decouples producers from consumers
 - Addresses speed, granularity, process swapping and efficiency issues
- **Disk scheduling**: example of how I/O buffering can be used to **reorder I/O transfer requests and improve efficiency**
- **Deadline of Homework 5 at the end of the week**
- **One more homework on file systems will be released**
- **Q&A session next Friday at 9:45**
- **Send me your questions before the session for more complete answers**
- **Course survey is open**